

AWS TerraOps - Infrastructure Automation and Scalability

Manual

By

Vivek Vashisht



Date: 13 November 2024

Contact:-

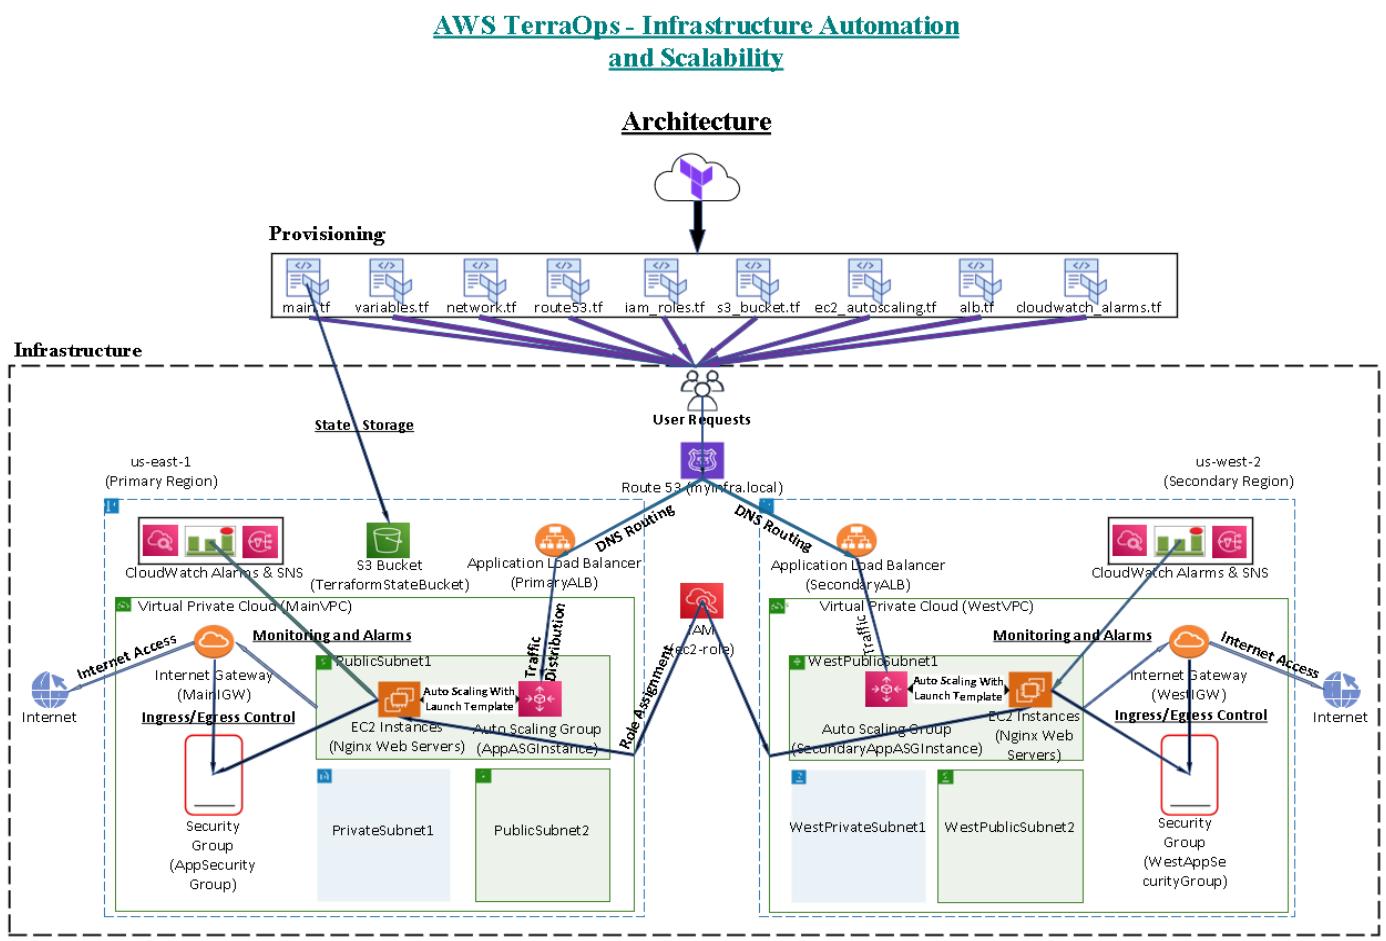
LinkedIn: <https://www.linkedin.com/in/vivek-vashisht04/>

GitHub: <https://github.com/vivekvashisht04/AWS-TerraOps>

Table of Contents

- Project Architecture Diagram
- Introduction
- Initial Setup
- VPC Creation and Networking
- EC2 Instances and Auto Scaling
- IAM Roles and S3 Integration
- Multi-Region Deployment and Failover
- Monitoring and Alerts with CloudWatch
- Parameterization and Validation
- Testing
- Cleanup
- Conclusion
- Lessons Learned

Project Architecture Diagram



Introduction

Project Overview

The AWS TerraOps: Infrastructure Automation and Scalability project is a strategic initiative designed to automate infrastructure deployment and achieve scalable cloud solutions using AWS services. This project leverages AWS's suite of tools to create a resilient, automated, and highly available environment capable of adapting to fluctuating workloads. It emphasizes seamless scaling, effective resource management, and failover mechanisms to ensure reliability and optimal performance.

Throughout the project, I utilized various AWS services, including EC2, Auto Scaling Groups, Application Load Balancers (ALBs), Route 53, and S3 Buckets. The project demonstrates the complete lifecycle, from infrastructure provisioning and configuration management to monitoring and failover testing. Each component was meticulously implemented and documented, capturing all critical aspects of infrastructure automation and scalability within AWS.

Purpose of this Documentation

This document provides a comprehensive, step-by-step guide for replicating the AWS TerraOps project, complete with detailed instructions and screenshots. It serves as a practical resource for cloud enthusiasts, professionals, and employers interested in learning infrastructure automation and scalability practices in AWS. The primary goal is to document the project thoroughly, showcasing my technical expertise in AWS and Terraform, while providing a valuable reference for potential employers and the broader cloud community.

Prerequisites

Before starting this project, it's essential to have a foundational understanding of cloud infrastructure, networking, and automation. Additionally, the following prerequisites are recommended:

1. **AWS Account:** A valid AWS account is necessary to deploy and manage cloud resources.
2. **Basic Networking Knowledge:** Familiarity with IP addressing, subnets, and security configurations.
3. **Terraform Installed:** Knowledge of Infrastructure as Code (IaC) concepts and experience with Terraform for resource provisioning.
4. **AWS CLI (Optional):** Familiarity with AWS CLI can be beneficial for automation and advanced configurations.
5. **Text Editor (Optional):** A text editor, such as Visual Studio Code, may be useful for managing configuration files and scripts.

This documentation aims to showcase the steps I took to complete the project, providing an informative guide to building scalable and automated infrastructure in AWS.

Initial Setup

a) Terraform and AWS CLI Installation

To start, I ensured that **terraform** was installed on my local machine and verified its installation by running the commands:

- `terraform`
- `terraform -version`

Next, I created a folder called **Binaries** in the C drive of my local machine, moved the Terraform executable file into this folder, and set up an environment variable. To do this, I accessed the **System Properties** window, selected **Environment Variables**, chose the **PATH** variable, clicked **Edit**, and added a path to the **Binaries** folder. This setup allowed Terraform to be accessible from any command line interface on my machine.

Then, I proceeded with the **AWS CLI** installation. I verified the installation by running the command:

- `aws --version`

Afterward, I configured AWS credentials. I logged into my AWS console and ensured that my account had the **AdministratorAccess** permission. In the **IAM** section, I created an access key. Once the **Access Key ID** and **Secret Access Key** were generated, I ran the following command in CMD:

- `aws configure`

I entered the credentials, completing the AWS configuration.

b) Version Control Setup

For version control, I ensured that **Git** was installed on my local machine. I visited [Git's official download page](#) to download and install Git.

To initialize a Git repository, I created a folder named **Git**, navigated to this directory in CMD, and ran:

- `git init`

Within the **Git** folder, I created a **.gitignore** file to exclude specific files and folders, adding the following entries:

1. `.terraform/`
2. `terraform.tfstate`

3. `terraform.tfstate.backup`

I then staged and committed the files with:

- `git add .`
- `git commit -m "Initial commit with Terraform setup"`

After committing, I created a new branch named **dev** with:

- `git checkout -b dev`

To set up a remote repository, I added the GitHub repository as a remote by running:

- `git remote add origin https://github.com/VivekVashisht2004/AWS-TerraOps.git`

I then pushed the changes to the **dev** branch and authorized Git credentials:

- `git push -u origin dev`

c) Creating the Main Terraform Configuration File (**main.tf**)

Next, I created a **main.tf** file containing the basic Terraform configuration to initialize the provider and create an S3 bucket to store the `terraform.tfstate` file as a remote backend.

Using references from the **Terraform Registry**, **AWS Documentation**, and AI validation, I configured the **main.tf** file as follows:

```
provider "aws" {
    region = "us-east-1"
}

# S3 bucket
resource "aws_s3_bucket" "terraform_state" {
    bucket = "aws-tf-backend-vv-bucket"

    tags = {
        Name      = "TerraformStateBucket"
        Environment = "Dev"
    }
}

# Public access block for the S3 bucket
resource "aws_s3_bucket_public_access_block" "terraform_state_block" {
    bucket = aws_s3_bucket.terraform_state.id

    block_public_acls      = true
    block_public_policy     = true
    restrict_public_buckets = true
    ignore_public_acls      = true
}
```

After saving the **main.tf** file, I navigated to the **Terraform** folder and ran the following commands to initialize, plan, and apply the configuration:

- `terraform init`
- `terraform plan`
- `terraform apply -auto-approve`

Following successful initialization, planning, and application, I confirmed the resource creation by checking the **aws-tf-backend-vv-bucket** in my AWS Console.

d) Remote Backend Configuration

To complete the remote backend configuration, I updated the **main.tf** file with the following block:

```
terraform {  
  backend "s3" {  
    bucket = "aws-tf-backend-vv-bucket"  
    key    = "terraform/dev/terraform.tfstate"  
    region = "us-east-1"  
  }  
}
```

After saving the **main.tf** file with the new backend configuration, I re-ran:

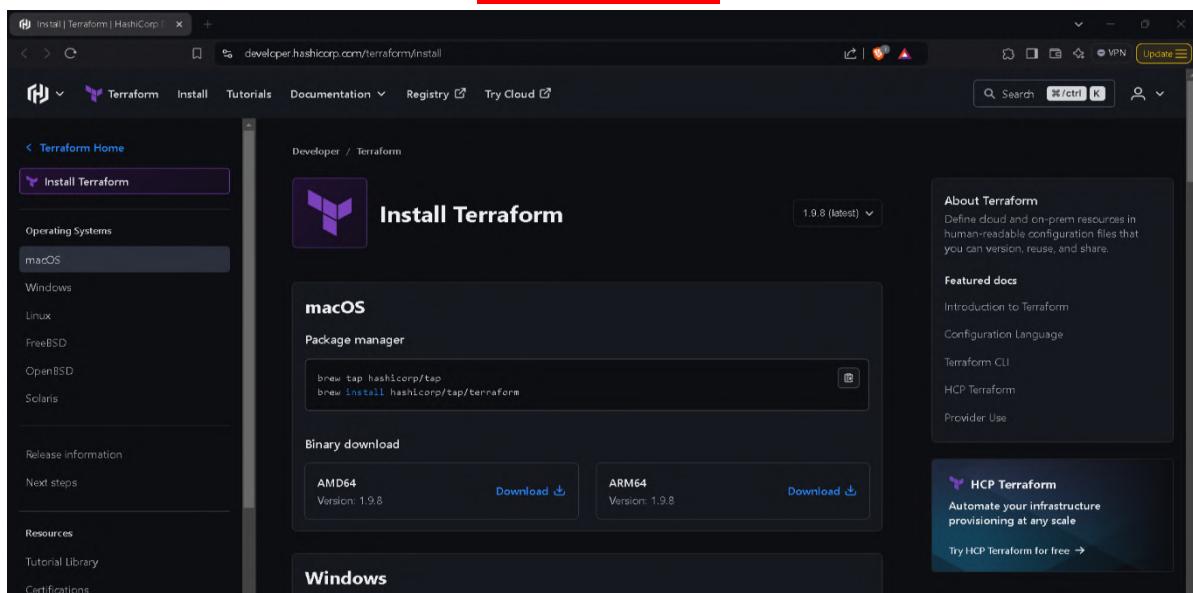
- `terraform init`

The command prompted to copy the existing state to the new backend. I confirmed, and with successful initialization, the S3 backend configuration was completed.

To verify, I checked the **aws-tf-backend-vv-bucket** in the AWS Console and found the **terraform.tfstate** file at the specified path **terraform/dev/terraform.tfstate**.

With this, the **Initial Setup** phase of the project was successfully completed.

Screenshots



Windows

Binary download

- 32-bit Version: 1.9.8 [Download](#)
- AMD64 Version: 1.9.8 [Download](#)

Linux

Package manager

- Ubuntu/Debian
- CentOS/RHEL
- Fedora
- Amazon Linux
- Homebrew

```
 wget -O- https://apt.releases.hashicorp.com/gpg | sudo gpg --dearmor -o /usr/share/keyrings/hashicorp-archive-keyring.gpg
 echo "deb [signed-by=/usr/share/keyrings/hashicorp-archive-keyring.gpg] https://apt.releases.hashicorp.com $(lsb_release -cs) main" | sudo tee /etc/apt/sources.list.d/hashicorp.list
 curl -fsSLO https://releases.hashicorp.com/1.9.8/terraform_1.9.8_windows_amd64.zip
```

```
C:\Users\hp>Terraform
Usage: terraform [global options] <subcommand> [args]

The available commands for execution are listed below.
The primary workflow commands are given first, followed by
less common or more advanced commands.

Main commands:
  init      Prepare your working directory for other commands
  validate   Check whether the configuration is valid
  plan      Show changes required by the current configuration
  apply     Create or update infrastructure
  destroy    Destroy previously-created infrastructure

All other commands:
  console    Try Terraform expressions at an interactive command prompt
  fmt        Reformat your configuration in the standard style
  force-unlock Release a stuck lock on the current workspace
  get        Install or upgrade remote Terraform modules
  graph     Generate a graph of the steps in an operation
  import    Associate existing infrastructure with a Terraform resource
  login     Obtain and save credentials for a remote host
  logout    Remove locally-stored credentials for a remote host
  metadata   Metadata related commands
  output    Show output values from your root module
  providers Show the providers required for this configuration
  refresh   Update the state of each remote system
  show      Show the current state of a saved plan
  state     Advanced state management
  taint     Mark a resource instance as not fully functional
  test      Execute integration tests for Terraform modules
  untaint  Remove the 'tainted' state from a resource instance
  version   Show the current Terraform version
  workspace Workspace management

Global options (use these before the subcommand, if any):
  -chdir=DIR  Switch to a different working directory before executing the
             given subcommand.
  -help       Show this help output, or the help for a specified subcommand.
  -version    An alias for the "version" subcommand.

C:\Users\hp>Terraform --version
Terraform v1.9.2
on windows_amd64

Your version of Terraform is out of date! The latest version
is 1.9.8. You can update by downloading from https://www.terraform.io/downloads.html
C:\Users\hp>
```

AWS Command Line Interface

Windows

Install and update requirements

- We support the AWS CLI on Microsoft-supported versions of 64-bit Windows.
- Admin rights to install software

Install or update the AWS CLI

To update your current installation of AWS CLI on Windows, download a new installer each time you update to overwrite previous versions. AWS CLI is updated regularly. To see when the latest version was released, see the [AWS CLI version 2 Changelog](#) on GitHub.

- Download and run the AWS CLI MSI installer for Windows (64-bit):
<https://awscli.amazonaws.com/AWSCLIV2.msi>

Alternatively, you can run the `msiexec` command to run the MSI installer.

```
C:\> msieexec.exe /i https://awscli.amazonaws.com/AWSCLIV2.msi
```

For various parameters that can be used with `msiexec`, see [msiexec](#) on the Microsoft Docs website. For example, you can use the `/qn` flag for a silent installation.

```
C:\> msieexec.exe /i https://awscli.amazonaws.com/AWSCLIV2.msi /qn
```

The screenshot shows the AWS Command Line Interface User Guide for Version 2. The left sidebar has sections like 'About the AWS CLI', 'Get started', 'Prerequisites', 'Install/Update', 'Configure the AWS CLI', 'Authentication and access credentials', 'Using the AWS CLI', 'Code examples', 'Security', 'Migration guide', 'Uninstall', and 'Document History'. The main content area is titled 'AWS Command Line Interface v2 Setup' and shows the 'Completed the AWS Command Line Interface v2 Setup Wizard' step. It includes instructions to click 'Finish' and a note about updating each time a new version is released. Below this is a 'Install or update' section with a link to download the MSI file from <https://awscli.amazonaws.com/AWSCLIV2.msi>. There are also alternative download links via msixexec.exe.

```
C:\Users\hp>aws --version
aws-cli/2.18.12 Python/3.12.6 Windows/11 exe/AMD64
C:\Users\hp>
```

The screenshot shows the AWS Cloud Console Home page. On the left, under 'Recently visited', there are links to IAM, Billing and Cost Management, and CloudWatch. In the center, the 'Applications' section shows '0' applications with a 'Create application' button. At the bottom, there are sections for 'Welcome to AWS' (Getting started with AWS), 'AWS Health' (Open issues), and 'Cost and usage' (Current month costs \$0.00, Cost breakdown, No cost data available).

The screenshot shows the AWS IAM User Details page for the user 'vivekvashisht'. The 'Summary' section displays the ARN (arn:aws:iam::134575801911:user/vivekvashisht), which is enabled without MFA. It was created on September 06, 2023, at 22:50 UTC-06:00, and the last console sign-in was today. The 'Permissions' tab is selected, showing two attached policies. The 'Access key 1' section indicates one access key has been created. The left sidebar shows navigation options like Dashboard, Access management, and Access reports.

The screenshot shows the 'Create access key' wizard, Step 3: Retrieve access keys. It displays the 'Access key' and 'Secret access key' fields, both of which are redacted. Below this, the 'Access key best practices' section provides several recommendations. At the bottom right, there are 'Download .csv file' and 'Done' buttons.

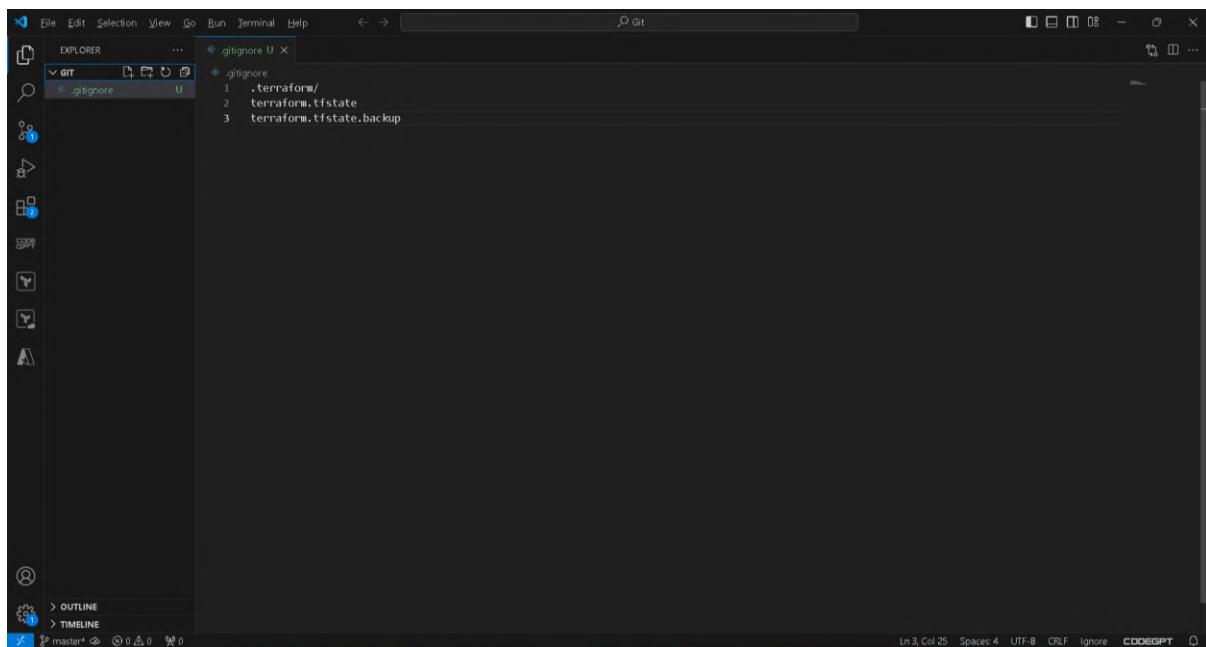
```
Microsoft Windows [Version 10.0.22631.4317]
(c) Microsoft Corporation. All rights reserved.

C:\Users\hp>aws --version
aws-cli/2.18.12 Python/3.12.6 Windows/11 exe/AMD64

C:\Users\hp>aws configure
AWS Access Key ID [None]: AKIAR6VK[REDACTED]
AWS Secret Access Key [None]: 7eyS8FLB02tJ0zm/pczbPB[REDACTED]
Default region name [None]: us-east-1
Default output format [None]: json

C:\Users\hp>
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git init
Initialized empty Git repository in C:/Users/hp/OneDrive/Desktop/LINKEDIN/Projects/6) AWS TerraOps/Screenshots/Step 1 Initial Setup/Git/.git/
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>
```



```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git init
Initialized empty Git repository in C:/Users/hp/OneDrive/Desktop/LINKEDIN/Projects/6) AWS TerraOps/Screenshots/Step 1 Initial Setup/Git/.git/
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git add .
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git commit -m "Initial commit with Terraform setup"
[master (root-commit) f4ce857] Initial commit with Terraform setup
 1 file changed, 3 insertions(+)
 create mode 100644 .gitignore
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>
```

```
Command Prompt x + v

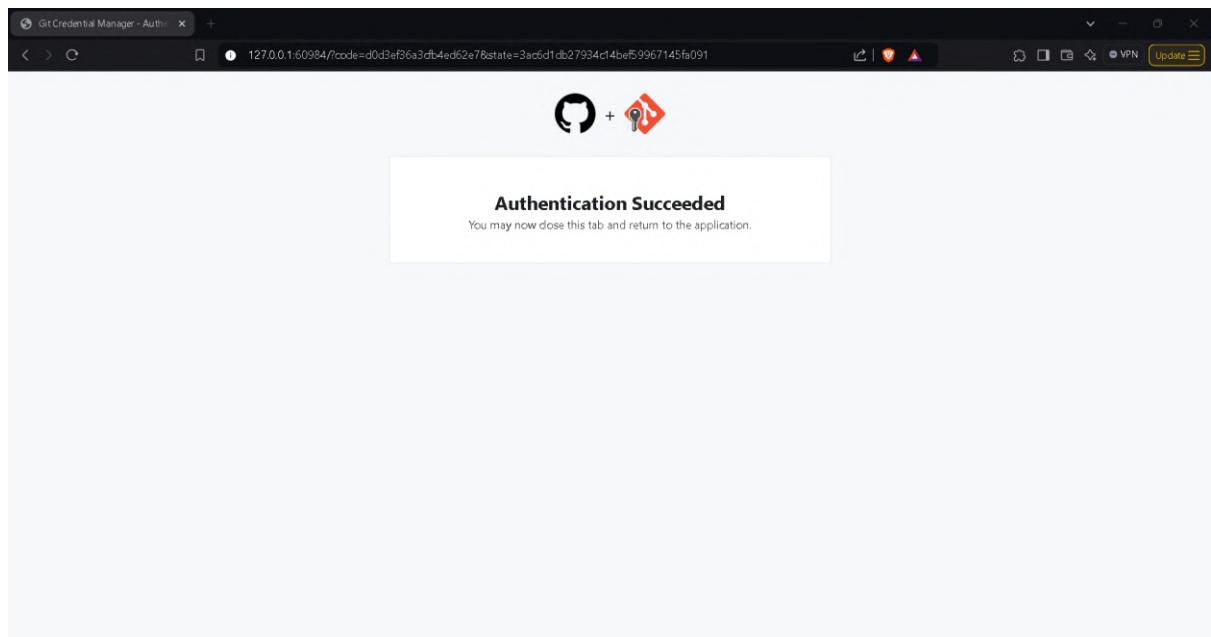
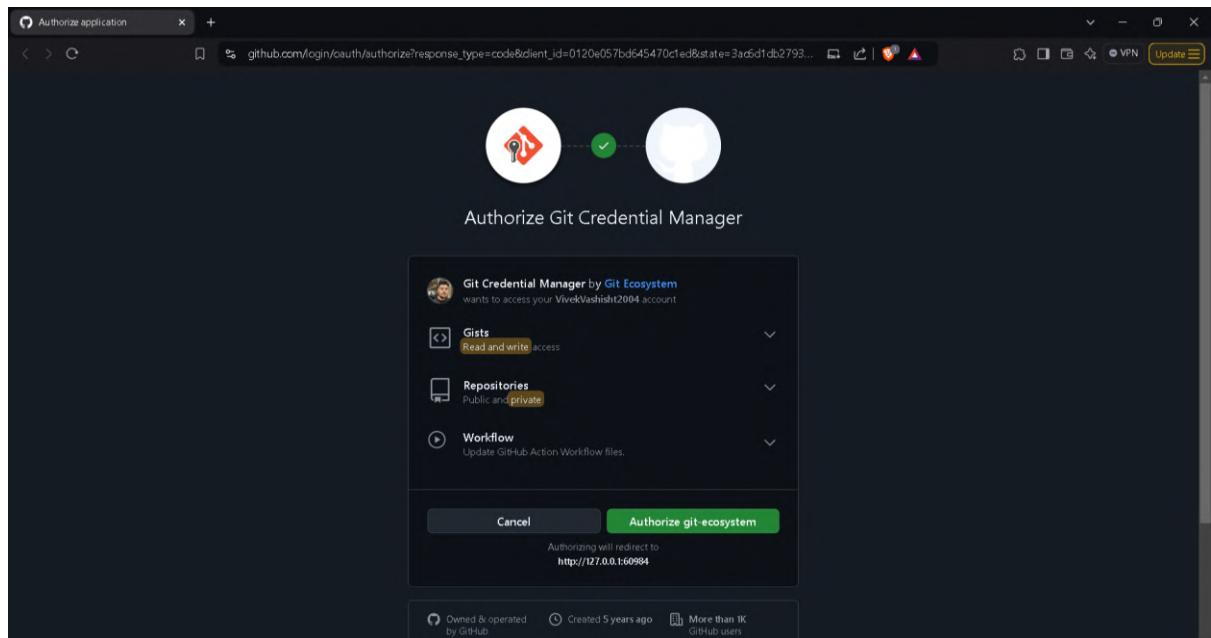
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git init
Initialized empty Git repository in C:/Users/hp/OneDrive/Desktop/LINKEDIN/Projects/6) AWS TerraOps/Screenshots/Step 1 Initial Setup/Git/.git/
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git add .
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git commit -m "Initial commit with Terraform setup"
[master (root-commit) f4ce857] Initial commit with Terraform setup
 1 file changed, 3 insertions(+)
 create mode 1006404 .gitignore

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git checkout -b dev
Switched to a new branch 'dev'
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git remote add origin https://github.com/VivekVashisht2004/AWS-TerraOps.git

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git push -u origin dev
info: please complete authentication in your browser...
```



```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>git push -u origin dev
info: please complete authentication in your browser...
Enumerating objects: 3, done.
Counting objects: 100% (3/3), done.
Delta compression using up to 12 threads
Compressing objects: 100% (2/2), done.
Writing objects: 100% (3/3), 260 bytes | 260.00 KiB/s, done.
Total 3 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
remote:
remote: Create a pull request for 'dev' on GitHub by visiting:
remote:   https://github.com/VivekVashisht2004/AWS-TerraOps/pull/new/dev
remote:
To https://github.com/VivekVashisht2004/AWS-TerraOps.git
 * [new branch]      dev -> dev
branch 'dev' set up to track 'origin/dev'.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Git>
```

```
main.tf
1 provider "aws" {
2   region = "us-east-1"
3 }
4
5 # S3 bucket
6 resource "aws_s3_bucket" "terraform_state" {
7   bucket = "aws-tf-backend-vv-bucket"
8
9   tags = [
10     { Name = "TerraformStateBucket" }
11     { Environment = "Dev" }
12   ]
13 }
14
15 # Public access block for the s3 bucket
16 resource "aws_s3_bucket_public_access_block" "terraform_state_block" [
17   bucket = aws_s3_bucket.terraform_state.id
18
19   block_public_acls = true
20   block_public_policy = true
21   restrict_public_buckets = true
22   ignore_public_acls = true
23 ]
24 }
```

```
Command Prompt
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
run this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_s3_bucket.terraform_state will be created
+ resource "aws_s3_bucket" "terraform_state" {
  + acceleration_status          = (known after apply)
  + cors_rules                   = (known after apply)
  + id                           = (known after apply)
  + arn                          = "aws-tf-backend-vv-bucket"
  + bucket                       = (known after apply)
  + bucket_domain_name           = (known after apply)
  + bucket_prefix                = (known after apply)
  + bucketRegionalDomainName    = (known after apply)
  + force_destroy               = false
  + hostedZoneId                = (known after apply)
  + id                           = (known after apply)
  + objectLockEnabled            = (known after apply)
  + policy                        = (known after apply)
  + region                        = (known after apply)
  + requestPayer                 = (known after apply)
  + tags                          = {
      + "Environment" = "Dev"
      + "Name"        = "TerraformStateBucket"
    }
  + tags_all                     = {
      + "Environment" = "Dev"
      + "Name"        = "TerraformStateBucket"
    }
  + website_domain               = (known after apply)
  + website_endpoint              = (known after apply)
}
```

```

Command Prompt x + v
+ cors_rule (known after apply)
+ grant (known after apply)
+ lifecycle_rule (known after apply)
+ logging (known after apply)
+ object_lock_configuration (known after apply)
+ replication_configuration (known after apply)
+ server_side_encryption_configuration (known after apply)
+ versioning (known after apply)
+ website (known after apply)
}

# aws_s3_bucket_public_access_block.terraform_state_block will be created
resource "aws_s3_bucket_public_access_block" "terraform_state_block" {
  block_public_acls      = true
  block_public_policy     = true
  bucket                 = (known after apply)
  id                     = (known after apply)
  ignore_public_acls     = true
  restrict_public_buckets = true
}

Plan: 2 to add, 0 to change, 0 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.
C:\Users\hp\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform apply -auto-approve
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create
Terraform will perform the following actions:

# aws_s3_bucket.terraform_state will be created
resource "aws_s3_bucket" "terraform_state" {
  acceleration_status      = (known after apply)
  acl                      = (known after apply)
  arn                      = (known after apply)
  bucket                   = "aws-tf-backend-vv-bucket"
  bucket_domain_name       = (known after apply)
}

```

```

Command Prompt x + v
+ request_payer           = (known after apply)
+ tags                     = {
    "Environment" = "Dev"
    "Name"        = "TerraformStateBucket"
}
+ tags_all                = {
    "Environment" = "Dev"
    "Name"        = "TerraformStateBucket"
}
+ website_domain          = (known after apply)
+ website_endpoint         = (known after apply)

+ cors_rule (known after apply)
+ grant (known after apply)
+ lifecycle_rule (known after apply)
+ object_lock_configuration (known after apply)
+ replication_configuration (known after apply)
+ server_side_encryption_configuration (known after apply)
+ versioning (known after apply)
+ website (known after apply)
}

# aws_s3_bucket_public_access_block.terraform_state_block will be created
resource "aws_s3_bucket_public_access_block" "terraform_state_block" {
  block_public_acls      = true
  block_public_policy     = true
  bucket                 = (known after apply)
  id                     = (known after apply)
  ignore_public_acls     = true
  restrict_public_buckets = true
}

Plan: 2 to add, 0 to change, 0 to destroy...
aws_s3_bucket.terraform_state: Creating...
aws_s3_bucket.terraform_state: Creation complete after 2s [id=aws-tf-backend-vv-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Creating...
aws_s3_bucket_public_access_block.terraform_state_block: Creation complete after 0s [id=aws-tf-backend-vv-bucket]

Terraform applied 2 changes to your infrastructure.

Resources: 2 added, 0 changed, 0 destroyed.
C:\Users\hp\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

```

aws-tf-backend-vv-bucket - S3 x +

us-east-1.console.aws.amazon.com/s3/buckets/aws-tf-backend-vv-bucket?region=us-east-1&bucketType=gen... [Alt+S] N. Virginia vivekavashisth@vv3281

Amazon S3 > Buckets > aws-tf-backend-vv-bucket

aws-tf-backend-vv-bucket Info

Objects (0) Info

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name	Type	Last modified	Size	Storage class
No objects				
You don't have any objects in this bucket.				
Upload				

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Do you want to copy existing state to the new backend?
Preexisting state was found while migrating the previous "local" backend to the
newly configured "s3" backend. No existing state was found in the newly
configured "s3" backend. Do you want to copy this state to the new "s3"
backend? Enter "yes" to copy and "no" to start with an empty state.

Enter a value: yes

Successfully configured the backend "s3". Terraform will automatically
use this backend unless the backend configuration changes.
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

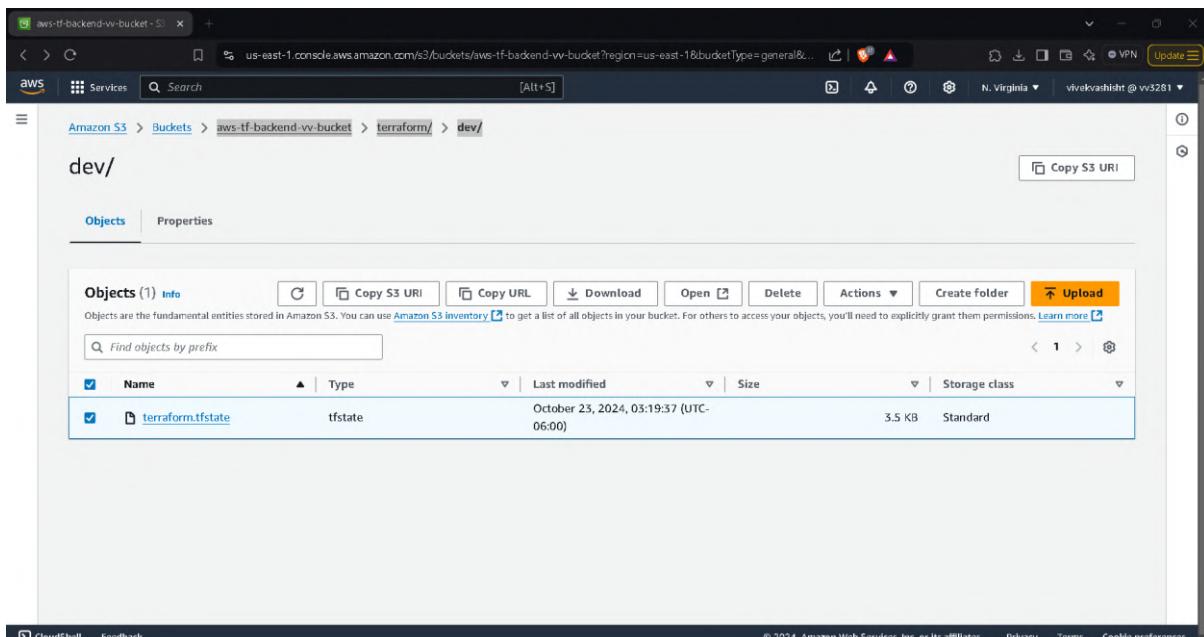
If you ever set or change modules or backend configuration for Terraform,
re-run this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]

No changes. Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes are needed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```



VPC Creation and Networking

a) Overview

In this phase of the project, I focused on creating a Virtual Private Cloud (VPC) along with its essential networking components, including subnets, an Internet Gateway, and route tables. These resources were defined in a Terraform file named `network.tf`, located in the Terraform project directory. By referencing both the Terraform Registry and AWS Documentation, I ensured each resource was configured correctly to support public and private networking.

b) VPC Setup

The first step was to define the VPC itself. I specified a CIDR block with a `/16` range, allowing for sufficient IP addresses across subnets. The VPC resource block in `network.tf` was as follows:

```
# VPC
resource "aws_vpc" "main_vpc" {
  cidr_block = "10.0.0.0/16"  # VPC with a /16 subnet range

  tags = {
    Name = "MainVPC"
  }
}
```

This created a VPC named `MainVPC` with a large address space, allowing for both public and private subnets to be defined within it.

c) Public and Private Subnet Configuration

Next, I defined both a public and a private subnet within the VPC. The public subnet was set to assign public IPs to instances, making them accessible from the Internet. The private subnet, on the other hand, was configured without public IPs, limiting access to internal networking.

```
# Public Subnet
resource "aws_subnet" "public_subnet_1" {
  vpc_id          = aws_vpc.main_vpc.id
  cidr_block      = "10.0.1.0/24"  # Public subnet's CIDR block
  availability_zone = "us-east-1a"  # Availability zone
  map_public_ip_on_launch = true    # Ensures EC2 instances in this subnet
  get a public IP

  tags = {
    Name = "PublicSubnet1"
  }
}

# Private Subnet
```

```

resource "aws_subnet" "private_subnet_1" {
  vpc_id          = aws_vpc.main_vpc.id
  cidr_block      = "10.0.2.0/24" # Private subnet's CIDR block
  availability_zone = "us-east-1a" # Same AZ as the public subnet
  map_public_ip_on_launch = false # Instances in this subnet will not
  have public IPs

  tags = {
    Name = "PrivateSubnet1"
  }
}

```

The public subnet was assigned 10.0.1.0/24, and the private subnet 10.0.2.0/24, both within the us-east-1a availability zone. This segmentation allows for structured and secure network architecture, where only certain resources can be accessed externally.

d) Internet Gateway (IGW)

To enable internet connectivity for resources in the public subnet, I created an Internet Gateway (IGW) and attached it to the VPC.

```

# Internet Gateway (IGW)
resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main_vpc.id

  tags = {
    Name = "MainIGW"
  }
}

```

The IGW, named MainIGW, allows traffic from the internet to reach instances within the public subnet when properly routed.

e) Public Route Table and Routing

To direct internet-bound traffic from the public subnet, I created a public route table and associated it with the subnet. The route table includes a route to forward all outbound traffic (0.0.0.0/0) to the Internet Gateway.

```

# Public Route Table and Route
resource "aws_route_table" "public_rt" {
  vpc_id = aws_vpc.main_vpc.id

  tags = {
    Name = "PublicRouteTable"
  }
}

resource "aws_route" "internet_access" {
  route_table_id      = aws_route_table.public_rt.id
  destination_cidr_block = "0.0.0.0/0" # Sends traffic to the internet
  gateway_id          = aws_internet_gateway.igw.id
}

resource "aws_route_table_association" "public_subnet_assoc" {
  subnet_id      = aws_subnet.private_subnet_1.id
}

```

```

    route_table_id = aws_route_table.public_rt.id
}

```

This setup enables instances within `PublicSubnet1` to send and receive internet traffic through the IGW, as specified in the `PublicRouteTable`.

f) network.tf Summary

The complete `network.tf` file, with all configurations, looked as follows:

```

# VPC
resource "aws_vpc" "main_vpc" {
  cidr_block = "10.0.0.0/16" # VPC with a /16 subnet range

  tags = {
    Name = "MainVPC"
  }
}

# Public Subnet
resource "aws_subnet" "public_subnet_1" {
  vpc_id          = aws_vpc.main_vpc.id
  cidr_block      = "10.0.1.0/24"
  availability_zone = "us-east-1a"
  map_public_ip_on_launch = true

  tags = {
    Name = "PublicSubnet1"
  }
}

# Private Subnet
resource "aws_subnet" "private_subnet_1" {
  vpc_id          = aws_vpc.main_vpc.id
  cidr_block      = "10.0.2.0/24"
  availability_zone = "us-east-1a"
  map_public_ip_on_launch = false

  tags = {
    Name = "PrivateSubnet1"
  }
}

# Internet Gateway (IGW)
resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main_vpc.id

  tags = {
    Name = "MainIGW"
  }
}

# Public Route Table and Route
resource "aws_route_table" "public_rt" {
  vpc_id = aws_vpc.main_vpc.id

  tags = {
    Name = "PublicRouteTable"
  }
}

```

```

}

resource "aws_route" "internet_access" {
  route_table_id      = aws_route_table.public_rt.id
  destination_cidr_block = "0.0.0.0/0"
  gateway_id          = aws_internet_gateway.igw.id
}

resource "aws_route_table_association" "public_subnet_assoc" {
  subnet_id      = aws_subnet.public_subnet_1.id
  route_table_id = aws_route_table.public_rt.id
}

```

g) Execution and Verification

After saving `network.tf`, I proceeded with the following commands in CMD to initialize, plan, and apply the configuration:

```

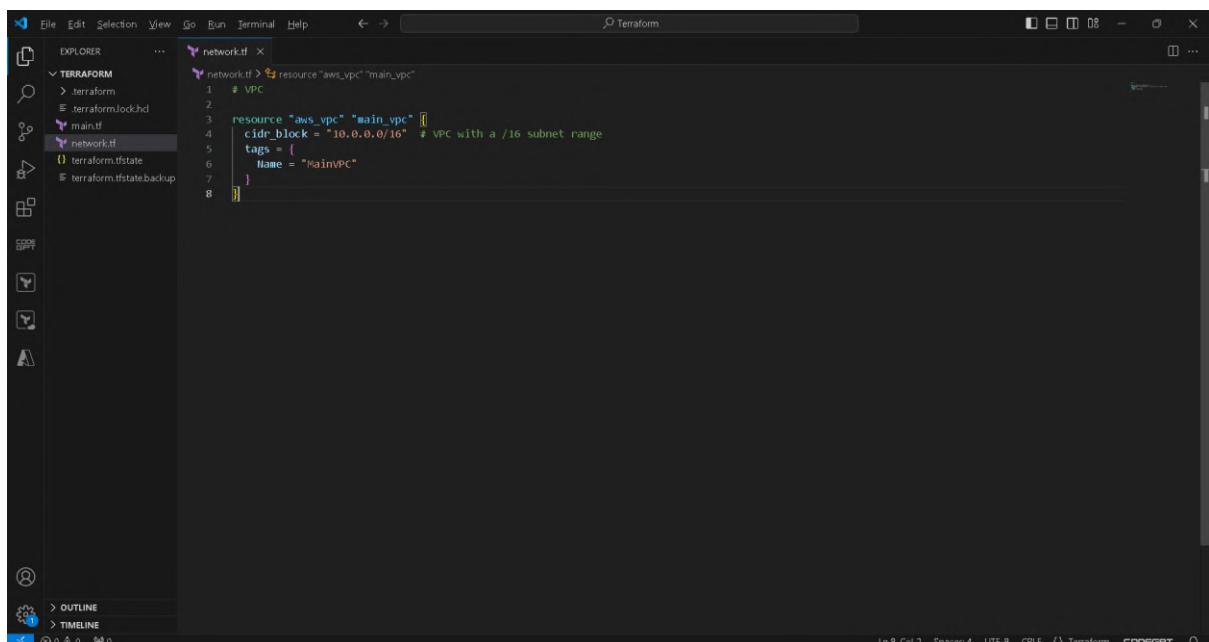
terraform init
terraform plan
terraform apply -auto-approve

```

Upon confirmation and successful application, I verified the resources in the AWS Console. I checked that all resources, including MainVPC, PublicSubnet1, PrivateSubnet1, MainIGW, PublicRouteTable, and the associated routes and subnets, were created as defined in the `network.tf` file.

This phase successfully established the foundational networking infrastructure for the project, allowing public and private instances to operate securely within a structured network environment.

Screenshots



```
File Edit Selection View Go Run Terminal Help < > ⚡ Terraform EXPLORER ... networktf > TERRAFORM > .terraform > main.tf > networktf > network.tf > terraform.tfstate & terraform.tfstate.backup CODEGPT networkif > resource "aws_vpc" "main_vpc" { 1 # VPC 2 3 cidr_block = "10.0.0.0/16" # VPC with a /16 subnet range 4 tags = [ 5 Name = "MainVPC" 6 ] 7 } 8 } 9 10 # Public Subnet 11 12 resource "aws_subnet" "public_subnet_1" { 13 vpc_id = aws_vpc.main_vpc.id 14 cidr_block = "10.0.1.0/24" # Public subnet's CIDR block 15 availability_zone = "us-east-1a" # Availability zone 16 map_public_ip_on_launch = true # Ensures EC2 instances in this subnet get a public IP 17 tags = [ 18 Name = "PublicSubnet1" 19 ] 20 }
```

The screenshot shows the CodeGPT interface with the 'TERRAFORM' section selected in the sidebar. The main area displays the first part of a Terraform configuration for a VPC. It includes a single VPC resource and one public subnet resource. The configuration uses AWS provider syntax with variables like `cidr_block` and `availability_zone`.

```
File Edit Selection View Go Run Terminal Help < > ⚡ Terraform EXPLORER ... networktf > TERRAFORM > .terraform > main.tf > networktf > network.tf > terraform.tfstate & terraform.tfstate.backup CODEGPT networkif > resource "aws_vpc" "main_vpc" { 2 3 cidr_block = "10.0.0.0/16" # VPC with a /16 subnet range 4 tags = [ 5 Name = "MainVPC" 6 ] 7 } 8 } 9 10 # Public Subnet 11 12 resource "aws_subnet" "public_subnet_1" { 13 vpc_id = aws_vpc.main_vpc.id 14 cidr_block = "10.0.1.0/24" # Public subnet's CIDR block 15 availability_zone = "us-east-1a" # Availability zone 16 map_public_ip_on_launch = true # Ensures EC2 instances in this subnet get a public IP 17 tags = [ 18 Name = "PublicSubnet1" 19 ] 20 } 21 22 # Private Subnet 23 24 resource "aws_subnet" "private_subnet_1" { 25 vpc_id = aws_vpc.main_vpc.id 26 cidr_block = "10.0.2.0/24" # Private subnet's CIDR block 27 availability_zone = "us-east-1a" # Same AZ as the public subnet 28 map_public_ip_on_launch = false # Instances in this subnet will not have public IPs 29 tags = [ 30 Name = "PrivateSubnet1" 31 ] 32 }
```

This screenshot continues the Terraform configuration from the previous one. It adds a second subnet, labeled 'Private Subnet', which has the same CIDR range and availability zone as the public subnet but with `map_public_ip_on_launch` set to `false`, meaning instances will not have public IPs.

```
File Edit Selection View Go Run Terminal Help < > ⚡ Terraform EXPLORER ... networktf > TERRAFORM > .terraform > main.tf > networktf > network.tf > terraform.tfstate & terraform.tfstate.backup CODEGPT networkif > resource "aws_vpc" "main_vpc" { 12 cidr_block = "10.0.0.0/16" # VPC with a /16 subnet range 13 tags = [ 14 Name = "MainVPC" 15 ] 16 } 17 } 18 19 # Private Subnet 20 21 resource "aws_subnet" "private_subnet_1" { 22 vpc_id = aws_vpc.main_vpc.id 23 cidr_block = "10.0.2.0/24" # Private subnet's CIDR block 24 availability_zone = "us-east-1a" # Same AZ as the public subnet 25 map_public_ip_on_launch = false # Instances in this subnet will not have public IPs 26 tags = [ 27 Name = "PrivateSubnet1" 28 ] 29 } 30 31 32 # Internet Gateway (IGW) 33 34 resource "aws_internet_gateway" "igw" { 35 vpc_id = aws_vpc.main_vpc.id 36 tags = [ 37 Name = "MainIGW" 38 ] 39 }
```

This final screenshot completes the Terraform configuration. It adds an Internet Gateway (IGW) resource, which is associated with the VPC defined earlier. The IGW is named 'MainIGW'. The configuration ensures that the VPC has a /16 CIDR range, two subnets (one public and one private), and an IGW for internet connectivity.

```

EXPLORER          network.tf
TERRAFORM         > .terraform
> .terraform
> .terraform.lock.hcl
> main.tf
> network.tf
> terraform.tfstate
> terraform.tfstate.backup

resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main_vpc.id
  tags = [
    { Name = "MainIGW" }
  ]
}

# Public Route Table and Route
resource "aws_route_table" "public_rt" {
  vpc_id = aws_vpc.main_vpc.id
  tags = [
    { Name = "PublicRouteTable" }
  ]
}

resource "aws_route" "internet_access" {
  route_table_id      = aws_route_table.public_rt.id
  destination_cidr_block = "0.0.0.0/0" # sends traffic to the internet
  gateway_id          = aws_internet_gateway.igw.id
}

resource "aws_route_table_association" "public_subnet_assoc" {
  subnet_id           = aws_subnet.public_subnet_1.id
  route_table_id      = aws_route_table.public_rt.id
}

```

```

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_internet_gateway.igw will be created
+ resource "aws_internet_gateway" "igw" {
  + arn      = (known after apply)
  + id       = (known after apply)
  + owner_id = (known after apply)
  + tags     = {
      + "Name" = "MainIGW"
    }
  + tags_all = {
      + "Name" = "MainIGW"
    }
  + vpc_id   = (known after apply)
}

# aws_route.internet_access will be created

```

```

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
+ resource "aws_route" "internet_access" {
  + destination_cidr_block = "0.0.0.0/0"
  + gateway_id              = (known after apply)
  + id                      = (known after apply)
  + instance_id              = (known after apply)
  + instance_owner_id        = (known after apply)
  + network_interface_id     = (known after apply)
  + origin                  = (known after apply)
  + route_table_id           = (known after apply)
  + state                   = (known after apply)
}

# aws_route_table.public_rt will be created
+ resource "aws_route_table" "public_rt" {
  + arn      = (known after apply)
  + id       = (known after apply)
  + owner_id = (known after apply)
  + propagating_vgwss = (known after apply)
  + route     = (known after apply)
  + tags     = {
      + "Name" = "PublicRouteTable"
    }
  + tags_all = {
      + "Name" = "PublicRouteTable"
    }
  + vpc_id   = (known after apply)
}

# aws_route_table_association.public_subnet_assoc will be created
+ resource "aws_route_table_association" "public_subnet_assoc" {
  + id           = (known after apply)
  + route_table_id = (known after apply)
  + subnet_id    = (known after apply)
}

# aws_subnet.private_subnet_1 will be created
+ resource "aws_subnet" "private_subnet_1" {
  + arn      = (known after apply)
  + assign_ipv6_address_on_creation = false
  + availability_zone                = "us-east-1a"
  + availability_zone_id             = (known after apply)
}
```

```

Command Prompt x + v
+ cidr_block = "10.0.2.0/24"
+ enable_dns64 = false
+ enable_resource_name_dns_a_record_on_launch = false
+ enable_resource_name_dns_aaaa_record_on_launch = false
+ id = (Known after apply)
+ ipv6_cidr_block_association_id = (Known after apply)
+ ipv6_native = false
+ map_public_ip_on_launch = false
+ owner_id = (Known after apply)
+ private_dns_hostname_type_on_launch = (Known after apply)
+ tags = {
    + "Name" = "PrivateSubnet1"
}
+ tags_all = {
    + "Name" = "PrivateSubnet1"
}
+ vpc_id = (known after apply)

# aws_subnet_public_subnet_1 will be created
resource "aws_subnet" "public_subnet_1" {
+ arn = (Known after apply)
+ assign_ipv6_address_on_creation = false
+ availability_zone = "us-east-1a"
+ availability_zone_id = (Known after apply)
+ cidr_block = "10.0.1.0/24"
+ enable_dns64 = false
+ enable_resource_name_dns_a_record_on_launch = false
+ enable_resource_name_dns_aaaa_record_on_launch = false
+ id = (Known after apply)
+ ipv6_cidr_block_association_id = (Known after apply)
+ ipv6_native = false
+ map_public_ip_on_launch = true
+ owner_id = (Known after apply)
+ private_dns_hostname_type_on_launch = (Known after apply)
+ tags = {
    + "Name" = "PublicSubnet1"
}
+ tags_all = {
    + "Name" = "PublicSubnet1"
}
}

```

```

Command Prompt x + v
+ tags = {
    + "Name" = "PublicSubnet1"
}
+ tags_all = {
    + "Name" = "PublicSubnet1"
}
+ vpc_id = (known after apply)

# aws_vpc.main_vpc will be created
resource "aws_vpc" "main_vpc" {
+ arn = (known after apply)
+ cidr_block = "10.0.0.0/16"
+ default_network_acl_id = (known after apply)
+ default_route_table_id = (known after apply)
+ default_security_group_id = (known after apply)
+ dhcp_options_id = (known after apply)
+ enable_dns_hostnames = (known after apply)
+ enable_dns_support = true
+ enable_network_address_usage_metrics = (known after apply)
+ id = (known after apply)
+ instance_tenancy = "default"
+ ipv6_association_id = (known after apply)
+ ipv6_cidr_block = (known after apply)
+ ipv6_cidr_block_network_border_group = (known after apply)
+ main_route_table_id = (known after apply)
+ owner_id = (known after apply)
+ tags = {
    + "Name" = "MainVPC"
}
+ tags_all = {
    + "Name" = "MainVPC"
}
}

```

Plan: 7 to add, 0 to change, 0 to destroy.

Note: You didn't use the `-out` option to save this plan, so Terraform can't guarantee to take exactly these actions if you run `"terraform apply"` now.

```

Command Prompt x + v
Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup>terraform apply
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_internet_gateway.igw will be created
resource "aws_internet_gateway" "igw" {
+ arn = (known after apply)
+ id = (known after apply)
+ owner_id = (known after apply)
+ tags = {
    + "Name" = "MainIGW"
}
+ tags_all = {
    + "Name" = "MainIGW"
}
+ vpc_id = (known after apply)
}

# aws_route.internet_access will be created
resource "aws_route" "internet_access" {
+ destination_cidr_block = "0.0.0.0/0"
+ gateway_id = (known after apply)
+ id = (known after apply)
+ instance_id = (known after apply)
+ instance_owner_id = (known after apply)
+ network_interface_id = (known after apply)
+ origin = (known after apply)
+ route_table_id = (known after apply)
+ state = (known after apply)
}

# aws_route_table.public_rt will be created
resource "aws_route_table" "public_rt" {

```

```
Command Prompt x + v
# aws_route_table.public_rt will be created
+ resource "aws_route_table" "public_rt" {
+   arn          = (known after apply)
+   id           = (known after apply)
+   owner_id     = (known after apply)
+   propagating_vgws = (known after apply)
+   route        = (known after apply)
+   tags         = {
+     + "Name" = "PublicRouteTable"
+   }
+   tags_all     = {
+     + "Name" = "PublicRouteTable"
+   }
+   vpc_id       = (known after apply)
}

# aws_route_table_association.public_subnet_assoc will be created
+ resource "aws_route_table_association" "public_subnet_assoc" {
+   id           = (known after apply)
+   route_table_id = (known after apply)
+   subnet_id    = (known after apply)
}

# aws_subnet.private_subnet_1 will be created
+ resource "aws_subnet" "private_subnet_1" {
+   arn          = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone      = "us-east-1a"
+   availability_zone_id   = (known after apply)
+   cidr_block            = "10.0.2.0/24"
+   enable_dns64           = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id                   = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native           = false
+   map_public_ip_on_launch = false
+   owner_id              = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags                 = {
+     + "Name" = "PrivateSubnet1"
+   }
}
```

```
Command Prompt x + v
# aws_subnet.public_subnet_1 will be created
+ resource "aws_subnet" "public_subnet_1" {
+   arn          = (known after apply)
+   assign_ipv6_address_on_creation = false
+   availability_zone      = "us-east-1a"
+   availability_zone_id   = (known after apply)
+   cidr_block            = "10.0.1.0/24"
+   enable_dns64           = false
+   enable_resource_name_dns_a_record_on_launch = false
+   enable_resource_name_dns_aaaa_record_on_launch = false
+   id                   = (known after apply)
+   ipv6_cidr_block_association_id = (known after apply)
+   ipv6_native           = false
+   map_public_ip_on_launch = true
+   owner_id              = (known after apply)
+   private_dns_hostname_type_on_launch = (known after apply)
+   tags                 = {
+     + "Name" = "PublicSubnet1"
+   }
+   tags_all     = {
+     + "Name" = "PublicSubnet1"
+   }
+   vpc_id       = (known after apply)
}

# aws_vpc.main_vpc will be created
+ resource "aws_vpc" "main_vpc" {
+   arn          = (known after apply)
+   cidr_block            = "10.0.0.0/16"
+   default_network_acl_id = (known after apply)
+   default_route_table_id = (known after apply)
+   default_security_group_id = (known after apply)
+   dhcp_options_id       = (known after apply)
+   enable_dns_hostnames = (known after apply)
+   enable_dns_support    = true
+   enable_network_address_usage_metrics = (known after apply)
+   id                   = (known after apply)
+   instance_tenancy     = "default"
+   ipv6_association_id  = (known after apply)
+   ipv6_cidr_block       = (known after apply)
}
```

```
Command Prompt x + v
+ instance_tenancy      = "default"
+ ipv6_association_id   = (known after apply)
+ ipv6_cidr_block        = (known after apply)
+ ipv6_cidr_block_network_border_group = (known after apply)
+ main_route_table_id   = (known after apply)
+ owner_id               = (known after apply)
+ tags                  = {
+   + "Name" = "MainVPC"
+ }
+ tags_all               = {
+   + "Name" = "MainVPC"
+ }

Plan: 7 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

aws_vpc.main_vpc: Creating...
aws_vpc.main_vpc: Creation complete after 2s [id=vpc-0ea415b9040f753dc]
aws_internet_gateway.igw: Creating...
aws_subnet.public_subnet_1: Creating...
aws_subnet.private_subnet_1: Creating...
aws_route_table.public_rt: Creating...
aws_internet_gateway.igw: Creation complete after 1s [id=igw-057bd0790bbd28a4a]
aws_route_table.public_rt: Creation complete after 1s [id=rtb-0c4fc0f97dfe205e5]
aws_route.internet_access: Creating...
aws_subnet.private_subnet_1: Creation complete after 1s [id=subnet-0c184948abfe6c733]
aws_route.internet_access: Creation complete after 1s [id=rth-0c4fc0f97dfe205e51080289494]
aws_subnet.public_subnet_1: Still creating... [10s elapsed]
aws_subnet.public_subnet_1: Creation complete after 19s [id=subnet-04b29a39893e7944d]
aws_route_table_association.public_subnet_assoc: Creating...
aws_route_table_association.public_subnet_assoc: Creation complete after 1s [id=rtbassoc-080529014332fac87]

apply complete! Resources: 7 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

VPC dashboard

Your VPCs (1/2) Info

Name	VPC ID	State	IPv4 CIDR	IPv6 CIDR	DHCP options
MainVPC	vpc-0ea415b9040f753dc	Available	10.0.0.0/16	-	dopt-028d2411c99231a5
-	vpc-01bb7cb53a9c88ffd	Available	172.31.0.0/16	-	dopt-00fce1099ded28b70

vpc-0ea415b9040f753dc / MainVPC

Details | Resource map | CIDs | Flow logs | Tags | Integrations

Details

VPC ID vpc-0ea415b9040f753dc	State Available	DNS hostnames Disabled	DNS resolution Enabled
Tenancy Default	DHCP option set dopt-028d2411c99231a5	Main route table rtb-00fce1099ded28b70	Main network ACL acl-003959774e89e0c5c
Default VPC No	IPv4 CIDR 10.0.0.0/16	IPv6 pool -	IPv6 CIDR (Network border group) -
Network Address Usage metrics Disabled	Route 53 Resolver DNS Firewall rule groups -	Owner ID 134575801911	-

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

VPC dashboard

Subnets (1/8) Info

Name	Subnet ID	State	VPC	IPv4 CIDR
-	subnet-0b29a39893e7944d	Available	vpc-0ea415b9040f753dc	10.0.1.0/24
-	subnet-038e4d458c75deeb7	Available	vpc-01bb7cb53a9c88ffd	172.31.16.0/20
-	subnet-087243af11c694cf9	Available	vpc-01bb7cb53a9c88ffd	172.31.32.0/20
PublicSubnet1	subnet-04b29a39893e7944d	Available	vpc-0ea415b9040f753dc	10.0.1.0/24

subnet-04b29a39893e7944d / PublicSubnet1

Details | Flow logs | Route table | Network ACL | CIDR reservations | Sharing | Tags

Details

Subnet ID subnet-04b29a39893e7944d	Subnet ARN arn:aws:ec2:us-east-1:134575801911:subnet/subnet-04b29a39893e7944d	State Available	IPv4 CIDR 10.0.1.0/24
Available IPv4 addresses 251	-	IPv6 CIDR association ID -	Availability zone us-east-1a
Availability Zone ID use1-az1	IPv6 CDR -	VPC vpc-0ea415b9040f753dc MainVPC	Route table rtb-04fcfd97fe205e5 PublicrouteTable
Network ACL acl-003959774e89e0c5c	Network border group us-east-1	Auto-assign public IPv4 address Yes	Auto-assign IPv6 address No
-	Default subnet	-	-

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

VPC dashboard

Subnets (1/8) Info

Name	Subnet ID	State	VPC	IPv4 CIDR
-	subnet-0941fb13260929842	Available	vpc-01bb7cb53a9c88ffd	172.31.80.0/20
-	subnet-053e112bb58c68cff	Available	vpc-01bb7cb53a9c88ffd	172.31.64.0/20
PrivateSubnet1	subnet-0c184948abfe6c733	Available	vpc-0ea415b9040f753dc	10.0.2.0/24

subnet-0c184948abfe6c733 / PrivateSubnet1

Details | Flow logs | Route table | Network ACL | CIDR reservations | Sharing | Tags

Details

Subnet ID subnet-0c184948abfe6c733	Subnet ARN arn:aws:ec2:us-east-1:134575801911:subnet/subnet-0c184948abfe6c733	State Available	IPv4 CIDR 10.0.2.0/24
Available IPv4 addresses 251	-	IPv6 CIDR association ID -	Availability zone us-east-1a
Availability Zone ID use1-az1	IPv6 CDR -	VPC vpc-0ea415b9040f753dc MainVPC	Route table rtb-00fce1099ded28b70
Network ACL acl-003959774e89e0c5c	Network border group us-east-1	Auto-assign public IPv4 address No	Auto-assign IPv6 address No
-	Default subnet No	-	-
-	Outpost ID	-	IPv4 CIDR reservations

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

I am viewing the AWS VPC Console. I have selected the 'Internet gateways' section under the 'Virtual private cloud' menu. There are two internet gateways listed:

Name	Internet gateway ID	State	VPC ID	Owner
igw-037d25060f2cd3ce6	Attached	vpc-01bb7cb53a9c88ff	134575801911	
MainIGW	igw-057bd0790bbd28a4a	Attached	vpc-0ea415b9040f753dc MainVPC	134575801911

The 'MainIGW' gateway is selected. Below it, the details for 'igw-057bd0790bbd28a4a / MainIGW' are shown:

Internet gateway ID	State	VPC ID	Owner
igw-057bd0790bbd28a4a	Attached	vpc-0ea415b9040f753dc MainVPC	134575801911

At the bottom right, there are links for 'CloudShell' and 'Feedback'.

I am viewing the AWS VPC Console. I have selected the 'Route tables' section under the 'Virtual private cloud' menu. There are two route tables listed:

Name	Route table ID	Explicit subnet assoc...	Edge associati...	Main	VPC
-	rtb-05d7cf17bd79188c	-	-	Yes	vpc-01bb7cb53a9c88ff
PublicRouteTable	rtb-0c4fc0f97dfe205e5	subnet-04b29a39893e7944d / PublicSubnet1	-	No	vpc-0ea415b9040f753dc MainVPC

The 'PublicRouteTable' is selected. Below it, the details for 'rtb-0c4fc0f97dfe205e5 / PublicRouteTable' are shown:

Route table ID	Main	Explicit subnet associations	Edge associations
rtb-0c4fc0f97dfe205e5	No	subnet-04b29a39893e7944d / PublicSubnet1	-
VPC	Owner ID	134575801911	

At the bottom right, there are links for 'CloudShell' and 'Feedback'.

I am viewing the AWS VPC Console. I have selected the 'Route tables' section under the 'Virtual private cloud' menu. The same two route tables are listed as before:

Name	Route table ID	Explicit subnet assoc...	Edge associati...	Main	VPC
-	rtb-05d7cf17bd79188c	-	-	Yes	vpc-01bb7cb53a9c88ff
PublicRouteTable	rtb-0c4fc0f97dfe205e5	subnet-04b29a39893e7944d / PublicSubnet1	-	No	vpc-0ea415b9040f753dc MainVPC

The 'PublicRouteTable' is selected. Below it, the details for 'rtb-0c4fc0f97dfe205e5 / PublicRouteTable' are shown, with the 'Routes' tab selected:

Destination	Target	Status	Propagated
0.0.0.0/0	igw-057bd0790bbd28a4a	Active	No
10.0.0.0/16	local	Active	No

At the bottom right, there are links for 'CloudShell' and 'Feedback'.

EC2 Instances and Auto Scaling

a) Overview

In this phase of the project, I focused on setting up EC2 instances within an Auto Scaling Group (ASG) to ensure scalability and high availability for the application. The setup includes creating a security group, a launch template with user data to install and configure NGINX and configuring scaling policies. All configurations were defined in a new Terraform file named `ec2_autoscaling.tf` in the Terraform project directory. I referenced both the Terraform Registry and AWS Documentation to ensure accuracy.

b) Security Group Configuration

To secure access to the instances, I created a security group allowing specific ingress and egress rules. This security group restricts SSH access to my IP address, allows HTTP traffic from anywhere, and permits all outbound traffic. Here's the code block for the security group:

```
# Security Group
resource "aws_security_group" "app_sg" {
  name          = "app-sg"
  description   = "Allow HTTP and SSH traffic"
  vpc_id        = aws_vpc.main_vpc.id

  ingress {
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
  }
}
```

```

        cidr_blocks = ["142.59.28.49/32"] # Restricted SSH to my IP only
    }

ingress {
    from_port   = 80
    to_port     = 80
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"] # Allows HTTP traffic from anywhere
}

egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"] # Allows all outbound traffic
}

tags = {
    Name = "AppSecurityGroup"
}
}

```

c) Launch Template with NGINX Installation

Next, I created a launch template that defines the configuration for instances launched by the ASG. The user data script installs and starts NGINX, ensuring that each instance is ready to serve web traffic immediately upon launch.

```

# Launch Template for EC2 Instances
resource "aws_launch_template" "app" {
    name_prefix      = "app-launch-template"
    image_id         = "ami-06b213ccaff8cd686" # Amazon Linux 2 AMI
    instance_type    = "t2.micro"                  # Cost-effective instance type

    network_interfaces {
        security_groups = [aws_security_group.app_sg.id]
    }

    user_data = base64encode(<<-EOF
        #!/bin/bash
        sudo yum update -y
        sudo amazon-linux-extras install nginx1 -y
        sudo systemctl start nginx
        sudo systemctl enable nginx
    EOF
    )

    lifecycle {
        create_before_destroy = true
    }

    tag_specifications {
        resource_type = "instance"
        tags = {
            Name = "AppInstance"
        }
    }
}

```

d) Auto Scaling Group (ASG) Setup

The ASG is configured to maintain a minimum of 1 instance, a desired capacity of 2 instances, and a maximum of 5 instances. The ASG utilizes the launch template defined above and deploys instances in the specified subnet.

```
# Auto Scaling Group (ASG) for us-east-1
resource "aws_autoscaling_group" "app_asg" {
  desired_capacity      = 2
  min_size              = 1
  max_size              = 5
  vpc_zone_identifier   = [aws_subnet.public_subnet_1.id]

  launch_template {
    id      = aws_launch_template.app.id
    version = "$Latest"
  }

  tag {
    key          = "Name"
    value         = "AppASGInstance"
    propagate_at_launch = true
  }

  health_check_type = "EC2"

  lifecycle {
    create_before_destroy = true
  }
}
```

e) Auto Scaling Policies

I configured scaling policies to dynamically adjust the number of instances based on demand. The scale-up policy increases the capacity by 1 instance, while the scale-down policy decreases it by 1 instance.

```
# Auto Scaling Policies

# Scaling Up Policy
resource "aws_autoscaling_policy" "scale_up" {
  name          = "scale-up"
  scaling_adjustment = 1 # Adds 1 instance when triggered
  adjustment_type = "ChangeInCapacity"
  autoscaling_group_name = aws_autoscaling_group.app_asg.name
}

# Scaling Down Policy
resource "aws_autoscaling_policy" "scale_down" {
  name          = "scale-down"
  scaling_adjustment = -1 # Removes 1 instance when triggered
  adjustment_type = "ChangeInCapacity"
  autoscaling_group_name = aws_autoscaling_group.app_asg.name
}
```

f) Deployment and Troubleshooting

After saving the `ec2_autoscaling.tf` file, I ran the following commands to initialize, plan, and apply the configuration:

```
terraform init  
terraform plan  
terraform apply -auto-approve
```

Upon verification, I encountered an issue while connecting to the EC2 instances. The EC2 Instance Connect IP range (`18.206.107.24/29`) was not authorized for SSH access. To fix this, I updated the security group in `ec2_autoscaling.tf`:

```
ingress {  
    from_port   = 22  
    to_port     = 22  
    protocol    = "tcp"  
    cidr_blocks = ["18.206.107.24/29"]  # EC2 Instance Connect IP range for  
    us-east-1  
}
```

After adding this rule, I re-ran `terraform plan` and `terraform apply -auto-approve` to update the security group.

g) Testing NGINX Installation

Once connected, I encountered another issue where NGINX failed to run due to an incorrect package command. I updated the user data block in the launch template:

```
user_data = base64encode(<<-EOF  
#!/bin/bash  
sudo dnf update -y  
sudo dnf install nginx -y  
sudo systemctl start nginx  
sudo systemctl enable nginx  
EOF  
)
```

To apply this change, I destroyed the existing launch template:

```
terraform destroy -target=aws_launch_template.app
```

Then, I re-ran `terraform apply -auto-approve` to create new instances with the updated NGINX configuration.

h) Verification

After successful deployment, I verified the NGINX installation by connecting to the EC2 instances and running:

```
sudo systemctl status nginx
```

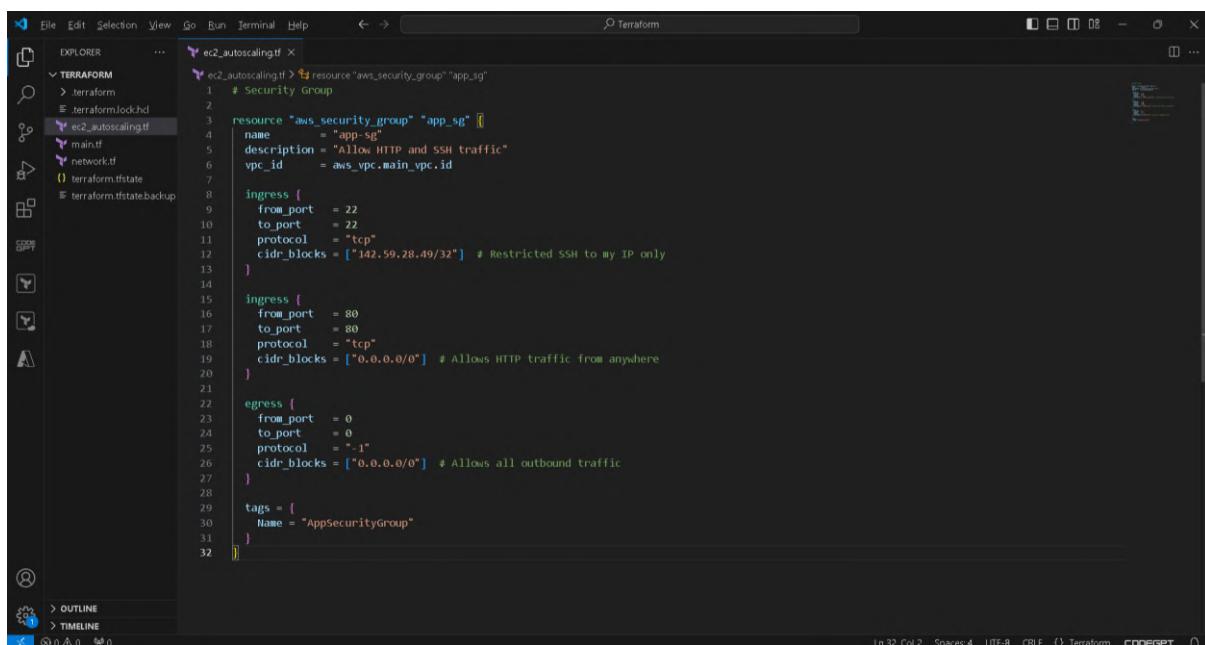
The status showed `active` (`running`), confirming that NGINX was successfully installed and running. Additionally, I accessed the NGINX web page by entering the instance's public IP in my browser.

ASG Verification

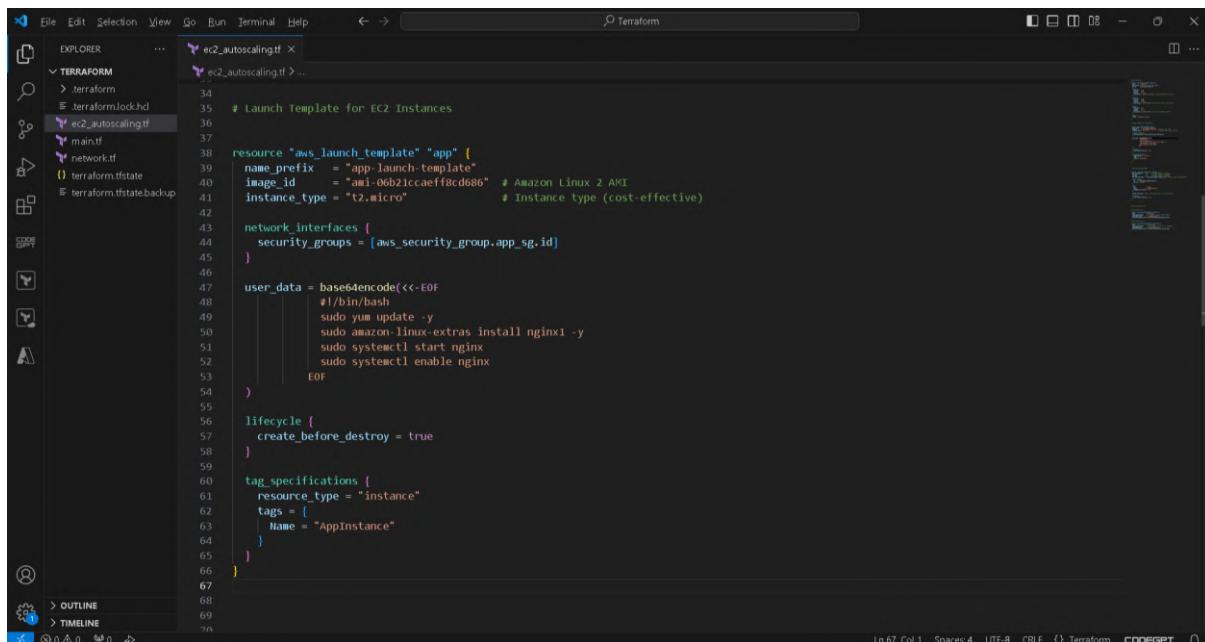
Finally, I checked the Auto Scaling Group in the AWS Console to ensure it was configured correctly with the expected settings.

This phase of the project was successfully completed, with the EC2 instances running NGINX and the ASG configured to scale resources as needed.

Screenshots



```
File Edit Selection View Go Run Terminal Help < - > ⚡ Terraform
EXPLORER ... ec2_autoscaling.tf ✘
TERRAFORM >_terraform
>_terramod.lock.hcl
main.tf
network.tf
terraform.tfstate
terraform.tfstate.backup
resource "aws_security_group" "app_sg" {
  name        = "app-sg"
  description = "Allow HTTP and SSH traffic"
  vpc_id      = aws_vpc.main_vpc.id
  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["142.59.28.49/32"] # Restricted SSH to my IP only
  }
  ingress {
    from_port  = 80
    to_port    = 80
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"] # Allows HTTP traffic from anywhere
  }
  egress {
    from_port  = 0
    to_port    = 0
    protocol   = "-1"
    cidr_blocks = ["0.0.0.0/0"] # Allows all outbound traffic
  }
  tags = [
    { Name = "AppSecurityGroup" }
  ]
}
```



```
File Edit Selection View Go Run Terminal Help < - > ⚡ Terraform
EXPLORER ... ec2_autoscaling.tf ...
TERRAFORM >_terraform
>_terramod.lock.hcl
main.tf
network.tf
terraform.tfstate
terraform.tfstate.backup
# Launch Template for EC2 Instances
resource "aws_launch_template" "app" {
  name_prefix  = "app-launch-template"
  image_id     = "ami-0eb21cae0ff8cd686" # Amazon Linux 2 AMI
  instance_type = "t2.micro" # Instance type (cost-effective)
  network_interfaces {
    security_groups = [aws_security_group.app_sg.id]
  }
  user_data = base64encode(<<<-EOF
#!/bin/bash
sudo yum update -y
sudo amazon-linux-extras install nginx1 -y
sudo systemctl start nginx
sudo systemctl enable nginx
EOF
)
  lifecycle {
    create_before_destroy = true
  }
  tag_specifications {
    resource_type = "instance"
    tags = [
      { Name = "AppInstance" }
    ]
  }
}
In 67, Col 1 Spaces:4 UTF-8 CRLF {} Terraform CODEGPT
```

```

file: ec2_autoscaling.tf
1 resource "aws_launch_template" "app" {
2   ...
3 }
4
5 resource "aws_autoscaling_group" "app_asg" {
6   desired_capacity = 2
7   min_size        = 1
8   max_size        = 5
9   vpc_zone_identifier = [aws_subnet.public_subnet_1.id]
10
11   launch_template {
12     id      = aws_launch_template.app.id
13     version = "$Latest"
14   }
15
16   tag {
17     key      = "Name"
18     value    = "AppASGInstance"
19     propagate_at_launch = true
20   }
21
22   health_check_type = "EC2"
23
24   lifecycle {
25     create_before_destroy = true
26   }
27 }
28
29 # Auto Scaling Policies
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117

```

This screenshot shows a code editor window with a dark theme. The main area displays a Terraform configuration file named `ec2_autoscaling.tf`. The code defines an `aws_launch_template` named `app` and an `aws_autoscaling_group` named `app_asg`. The `app_asg` has a `desired_capacity` of 2, a `min_size` of 1, and a `max_size` of 5. It uses the `app` launch template and includes a tag with the key `Name` and value `AppASGInstance`. The `health_check_type` is set to `EC2`, and the `lifecycle` block specifies that instances should be created before the group is destroyed.

```

file: ec2_autoscaling.tf
1 resource "aws_autoscaling_group" "app_asg" {
2   lifecycle {
3     create_before_destroy = true
4   }
5 }
6
7 # Auto Scaling Policies
8
9
10
11
12
13 resource "aws_autoscaling_policy" "scale_up" {
14   name          = "scale-up"
15   scaling_adjustment = 1 # Adds 1 instance when triggered
16   adjustment_type = "ChangeInCapacity"
17   autoscaling_group_name = aws_autoscaling_group.app_asg.name
18 }
19
20 # Scaling Down Policy
21
22 resource "aws_autoscaling_policy" "scale_down" {
23   name          = "scale-down"
24   scaling_adjustment = -1 # Removes 1 instance when triggered
25   adjustment_type = "ChangeInCapacity"
26   autoscaling_group_name = aws_autoscaling_group.app_asg.name
27 }
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118

```

This screenshot shows a code editor window with a dark theme. The main area displays a Terraform configuration file named `ec2_autoscaling.tf`. It defines two `aws_autoscaling_policy` resources: `scale_up` and `scale_down`. The `scale_up` policy adds one instance when triggered and is associated with the `app_asg` group. The `scale_down` policy removes one instance when triggered and is also associated with the `app_asg` group.

```

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
aws_vpc.main.vpc: Refreshing state... [id=vpc-0eaa15b904bf752dc]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-0c18948abfe6c733]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-0b29a39893e7944d]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bd0799bbd28a4a]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc0f97dfe205e]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_route_table_association.public_subnet_assoc: Refreshing state... [id=rbassoc-086529014332fac87]
aws_route.internet_access: Refreshing state... [id=rts-0c4fc0f97dfe205e51080289494]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be created
+ resource "aws_autoscaling_group" "app_asg" {
  + arn                                = (known after apply)
  + availability_zones                  = (known after apply)
  + defaultCooldown                    = (known after apply)
  + desired_capacity                   = 2
  + force_delete                       = false
  + force_delete_warm_pool             = false
}
```

This screenshot shows a terminal window with a dark theme. The output shows the results of running `terraform init` and `terraform plan`. The `terraform init` command initializes the backend and provider plugins. The `terraform plan` command generates an execution plan, indicating that a new `aws_autoscaling_group` resource named `app_asg` will be created. The plan details the configuration for the new group, including its desired capacity of 2, and lists various provider and resource refreshes.

```
Command Prompt + v
+ health_check_grace_period      = 300
+ health_check_type              = "EC2"
+ id                             = (known after apply)
+ ignore_failed_scaling_activities = false
+ load_balancers                 = (known after apply)
+ max_size                       = 5
+ metrics_granularity            = "1Minute"
+ min_size                       = 1
+ name                           = (known after apply)
+ name_prefix                     = (known after apply)
+ predicted_capacity              = (known after apply)
+ protect_from_scale_in          = false
+ service_linked_role_arn        = (known after apply)
+ target_group_arns               = (known after apply)
+ vpc_zone_identifier              = [
    + "subnet-04b29a39893e7944d",
]
+ wait_for_capacity_timeout       = "10m"
+ warm_pool_size                  = (known after apply)

+ launch_template {
    + id      = (known after apply)
    + name   = (known after apply)
    + version = "$Latest"
}

+ mixed_instances_policy (known after apply)

+ tag {
    + key      = "Name"
    + propagate_at_launch = true
    + value    = "AppASGInstance"
}

+ traffic_source (known after apply)

# aws_autoscaling_policy.scale_down will be created
resource "aws_autoscaling_policy" "scale_down" {
+ adjustment_type      = "ChangeInCapacity"
+ arn                   = (known after apply)
}
```

```
Command Prompt + v
# aws_autoscaling_policy.scale_down will be created
resource "aws_autoscaling_policy" "scale_down" {
+ adjustment_type      = "ChangeInCapacity"
+ arn                   = (known after apply)
+ autoscaling_group_name = (known after apply)
+ enabled               = true
+ id                   = (known after apply)
+ metric_aggregation_type = (known after apply)
+ name                 = "scale-down"
+ policy_type          = "SimpleScaling"
+ scaling_adjustment   = -1
}

# aws_autoscaling_policy.scale_up will be created
resource "aws_autoscaling_policy" "scale_up" {
+ adjustment_type      = "ChangeInCapacity"
+ arn                   = (known after apply)
+ autoscaling_group_name = (known after apply)
+ enabled               = true
+ id                   = (known after apply)
+ metric_aggregation_type = (known after apply)
+ name                 = "scale-up"
+ policy_type          = "SimpleScaling"
+ scaling_adjustment   = 1
}

# aws_launch_template.app will be created
resource "aws_launch_template" "app" {
+ arn                   = (known after apply)
+ default_version       = (known after apply)
+ id                   = (known after apply)
+ image_id              = "ami-06b21ccaeff8cd686"
+ instance_type         = "t2.micro"
+ latest_version        = (known after apply)
+ name                 = (known after apply)
+ name_prefix           = "app-launch-template"
+ tags.all              = (known after apply)
+ user_data             = "iYEvYmluL2Jhc2gNCnNzG8geXVtIHVwZGF0ZSAtQ0Kc3VkbhWF6b24tbGludXgtZXh0cmFzIGluc3RhbGwgmdpbngxIC15DQpzdWRvIHN5c3RlbWN0bCBzdGFydCBuZ2ueA0Kc3VkbYBeXN0ZW1jdGwgZWh5YmxlG5naW54DQo="
+ metadata_options (known after apply)
```

```
Command Prompt + v
# aws_security_group.app_sg will be created
resource "aws_security_group" "app_sg" {
+ arn                   = (known after apply)
+ description           = "Allow HTTP and SSH traffic"
+ egress                = [
    + {
        + cidr_blocks     = [
            + "0.0.0.0/0",
        ]
        + from_port       = 0
        + ipv6_cidr_blocks = []
        + prefix_list_ids = []
        + protocol        = "-1"
        + security_groups = []
        + self             = false
        + to_port          = 0
        # (1 unchanged attribute hidden)
    },
]
+ id                   = (known after apply)
+ ingress              = [
    + {
        + cidr_blocks     = [
            + "0.0.0.0/0",
        ]
        + from_port       = 80
        + ipv6_cidr_blocks = []
        + prefix_list_ids = []
        + protocol        = "tcp"
        + security_groups = []
        + self             = false
        + to_port          = 80
        # (1 unchanged attribute hidden)
    },
    + {
        + cidr_blocks     = [
            + "142.59.28.49/32",
        ]
        + from_port       = 22
        + ipv6_cidr_blocks = []
    }
]
```

```
Plan: 5 to add, 0 to change, 0 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup>terraform apply
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws_s3.bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv--bucket]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc0f97dfe205e5]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bbd0790bbd28a4]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-04a29a39893e7944d]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-0c18498abfe6c733]
aws_route_table.access: Refreshing state... [id=rtb-0c4fc0f97dfe205e510808289494]
aws_s3.bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv--bucket]
aws_route_table_association.public_subnet_assoc: Refreshing state... [id=rtbassoc-080529014332fac87]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be created
+ resource "aws_autoscaling_group" "app_asg" {
    + arn = (known after apply)
    + availability_zones = (known after apply)
    + default_cooldown = (known after apply)
    + desired_capacity = 2
    + force_delete = false
    + force_delete_warm_pool = false
    + health_check_grace_period = 300
    + health_check_type = "EC2"
    + id = (known after apply)
    + ignore_failed_scaling_activities = false
    + load_balancers = (known after apply)
    + max_size = 5
    + metrics_granularity = "1Minute"
    + min_size = 1
    + name = (known after apply)
    + name_prefix = (known after apply)
    + predicted_capacity = (known after apply)
```

```
# aws_autoscaling_policy.scale_down will be created
+ resource "aws_autoscaling_policy" "scale_down" {
+   adjustment_type      = "ChangeInCapacity"
+   arn                  = (known after apply)
+   autoscaling_group_name = (known after apply)
+   enabled               = true
+   id                   = (known after apply)
+   metric_aggregation_type = (known after apply)
+   name                 = "scale-down"
+   policy_type          = "SimpleScaling"
+   scaling_adjustment    = -1
}

# aws_autoscaling_policy.scale_up will be created
+ resource "aws_autoscaling_policy" "scale_up" {
+   adjustment_type      = "ChangeInCapacity"
+   arn                  = (known after apply)
+   autoscaling_group_name = (known after apply)
+   enabled               = true
+   id                   = (known after apply)
+   metric_aggregation_type = (known after apply)
+   name                 = "scale-up"
+   policy_type          = "SimpleScaling"
+   scaling_adjustment    = 1
}

# aws_launch_template.app will be created
+ resource "aws_launch_template" "app" {
+   arn                  = (known after apply)
+   default_version       = (known after apply)
+   id                   = (known after apply)
+   image_id              = "ami-06621caefff8cd686"
+   instance_type         = "t2.micro"
+   latest_version        = (known after apply)
+   name                 = (known after apply)
+   name_prefix           = "app-launch-template"
+   tags_all              = (known after apply)
+   user_data             = "i=EvYmlul2Jhc2gChN1ZG8geXVtIHvWzGF0ZSAteQ0Kc3VkbvBhbWF6b24tbGludXgtZXh0cmFzIGluc3RhbgwgbmdpbngxIC15DQpzdlRvIHN5c3RlbWN0CBzdGFydkC82ZlueA0Kc3VkbvBzeXN0ZWIjdGwgZw5hYmxLIG5naW54DQo="
```

```
# aws_security_group.app_sg will be created
+ resource "aws_security_group" "app_sg" {
+   arn                      = (known after apply)
+   description              = "Allow HTTP and SSH traffic"
+   egress                   = [
+     {
+       cidr_blocks          = [
+         "+0.0.0.0/0",
+       ]
+       from_port            = 0
+       ipv6_cidr_blocks     = []
+       prefix_list_ids      = []
+       protocol              = "-1"
+       security_groups       = []
+       self                  = false
+       to_port               = 0
+       # (1 unchanged attribute hidden)
+     },
+   ],
+   id                      = (known after apply)
+   ingress                 = [
+     {
+       cidr_blocks          = [
+         "+0.0.0.0/0",
+       ]
+       from_port            = 80
+       ipv6_cidr_blocks     = []
+       prefix_list_ids      = []
+       protocol              = "tcp"
+       security_groups       = []
+       self                  = false
+       to_port               = 80
+       # (1 unchanged attribute hidden)
+     },
+   ],
+   {
+     cidr_blocks          = [
+       "+192.59.28.49/32",
+     ]
+     from_port            = 22
+     ipv6_cidr_blocks     = []
+   }
]
```

```

Command Prompt x + -
+ self           = false
+ to_port        = 22
# (1 unchanged attribute hidden)
},
+ name          = "app-sg"
+ name_prefix   = (known after apply)
+ owner_id      = (known after apply)
+ revoke_rules_on_delete = false
+ tags          =
+ tags_all     =
+ "Name" = "AppSecurityGroup"
}
+ tags_all     =
+ "Name" = "AppSecurityGroup"
}
+ vpc_id        = "vpc-0ea415b9040f753dc"
}

Plan: 5 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

aws_security_group.app_sg: Creating...
aws_security_group.app_sg: Creation complete after 3s [id=sg-0bf07049429bd4a81]
aws_launch_template.app: Creating...
aws_launch_template.app: Creation complete after 1s [id=lt-04ebe7149bc419c5c]
aws_autoscaling_group.app_asg: Creating...
aws_autoscaling_group.app_asg: Still creating... [10s elapsed]
aws_autoscaling_group.app_asg: Creation complete after 16s [id=terraform-20241025062721849900000003]
aws_autoscaling_policy.scale_up: Creating...
aws_autoscaling_policy.scale_down: Creating...
aws_autoscaling_policy.scale_down: Creation complete after 0s [id=scale-down]
aws_autoscaling_policy.scale_up: Creating complete after 1s [id=scale-up]

Apply complete! Resources: 5 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

```

Name	Security group ID	Security group name	VPC ID	Description
-	sg-0a6e24127747c69dd	default	vpc-01bb7cb53a9c88ffd	default VPC sec
-	sg-026fe863d61fff21	default	vpc-0ea415b9040f753dc	default VPC sec
AppSecurityGroup	sg-0bf07049429bd4a81	app-sg	vpc-0ea415b9040f753dc	Allow HTTP and SSH traffic

Name	Security group ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-021312c63bcf3f6e7	IPv4	SSH	TCP	22	142.59.28.49/32	-
-	sgr-01684d236f9a14929	IPv4	HTTP	TCP	80	0.0.0.0/0	-

SecurityGroups | VPC Console

us-east-1.console.aws.amazon.com/vpcconsole/home?region=us-east-1#SecurityGroups

aws Services Search [Alt+S]

Security Groups (1/3) Info Actions Export security groups to CSV Create security group

Find resources by attribute or tag

Name	Security group ID	Security group name	VPC ID	Description
-	sg-0a6e24127747c69dd	default	vpc-01bb7cb53a9c88ffd	default VPC sec
-	sg-026fe863d61fff21	default	vpc-0ea415b9040f753dc	default VPC sec
<input checked="" type="checkbox"/> AppSecurityGroup	sg-0bf07049429bd4a81	app-sg	vpc-0ea415b9040f753dc	Allow HTTP and

Details Inbound rules Outbound rules Tags

Outbound rules (1/1)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Destination	Description
-	sgr-0fd4fe7ecf5c2d67f	IPv4	All traffic	All	All	0.0.0.0/0	-

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Launch templates | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#LaunchTemplates

aws Services Search [Alt+S]

Launch Templates (1/1) Info Actions Create launch template

Search

Launch Template ID	Launch Template Name	Default Version	Latest Version	Create Time	Created By
<input checked="" type="checkbox"/> lt-04ebe7149bc419c5c	app-launch-template202410...	1	1	2024-10-25T06:27:21.0...	arn:aws:iam::...

app-launch-template20241025062721490700000001 (lt-04ebe7149bc419c5c)

Launch template details Actions Delete template

Launch template ID lt-04ebe7149bc419c5c	Launch template name app-launch-template202410250627214907000000001	Default version 1	Owner arn:aws:iam:134575801911:user/vivekashish
--	--	----------------------	--

Details Versions Template tags

Launch template version details Actions Delete template version

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Launch templates | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#LaunchTemplates

aws Services Search [Alt+S]

Launch Templates (1/1) Info Actions Create launch template

Search

Launch Template ID	Launch Template Name	Default Version	Latest Version	Create Time	Created By
<input checked="" type="checkbox"/> lt-04ebe7149bc419c5c	app-launch-template202410...	1	1	2024-10-25T06:27:21.0...	arn:aws:iam::...

app-launch-template20241025062721490700000001 (lt-04ebe7149bc419c5c)

Version Description Date created Created by

1 (Default)	-	2024-10-25T06:27:21.000Z	arn:aws:iam:134575801911:user/vivekashish
-------------	---	--------------------------	---

Instance details Storage Resource tags Network interfaces Advanced details

AMI ID ami-06b21ccaeff8cd686	Instance type t2.micro	Availability Zone -	Key pair name -
Security groups -	Security group IDs sg-0bf07049429bd4a81		

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS EC2 Instances page showing two running t2.micro instances. The left sidebar shows navigation links for EC2 Dashboard, Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity, and Reservations. The main content area displays the instance details for 'i-0f3b3f351cef80bf4' (AppASGInstance).

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
AppASGInstance	i-0f3b3f351cef80bf4	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	-
AppASGInstance	i-0bb6c24776f8f7b24	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	-

i-0f3b3f351cef80bf4 (AppASGInstance)

Details | Status and alarms | Monitoring | Security | Networking | Storage | Tags

Instance summary

Instance ID	Public IPv4 address	Private IPv4 addresses
i-0f3b3f351cef80bf4	44.199.234.166 open address	10.0.1.61

IPv6 address	Instance state	Public IPv4 DNS
-	Running	-

Hostname type	Private IP DNS name (IPv4 only)	Elastic IP addresses
IP name: ip-10-0-1-61.ec2.internal	ip-10-0-1-61.ec2.internal	-

Answer private resource DNS name	Instance type	AWS Compute Optimizer finding
-	t2.micro	-

Auto-assigned IP address	VPC ID
-	-

Screenshot of the AWS EC2 Instances page showing the same two instances. The left sidebar is identical. The main content area displays the networking rules for 'i-0f3b3f351cef80bf4' (AppASGInstance). It shows inbound and outbound security group rules.

Inbound rules

Name	Security group rule ID	Port range	Protocol	Source	Security group
-	sgr-021312c63bfcf3f6e7	22	TCP	142.59.28.49/32	app-sg
-	sgr-01684d236f9a14929	80	TCP	0.0.0.0/0	app-sg

Outbound rules

Name	Security group rule ID	Port range	Protocol	Destination	Security group
-	sgr-0fd4fe7ecf5c2d67f	All	All	0.0.0.0/0	app-sg

Screenshot of the AWS EC2 Instances page showing the same two instances. The left sidebar is identical. The main content area displays the details for 'i-0bb6c24776f8f7b24' (AppASGInstance).

Details | Status and alarms | Monitoring | Security | Networking | Storage | Tags

Instance summary

Instance ID	Public IPv4 address	Private IPv4 addresses
i-0bb6c24776f8f7b24	44.201.59.142 open address	10.0.1.103

IPv6 address	Instance state	Public IPv4 DNS
-	Running	-

Hostname type	Private IP DNS name (IPv4 only)	Elastic IP addresses
IP name: ip-10-0-1-103.ec2.internal	ip-10-0-1-103.ec2.internal	-

Answer private resource DNS name	Instance type	AWS Compute Optimizer finding
-	t2.micro	-

Auto-assigned IP address	VPC ID
-	-

Screenshot of the AWS EC2 Instances page showing two running AppASG instances. The left sidebar shows various EC2-related services and regions.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
AppASGinstance	i-0f3b3f551cef80bf4	Running	t2.micro	2/2 checks pass	View alarms	us-east-1a	-
AppASGinstance	i-0bb6c24776f8f7b24	Running	t2.micro	2/2 checks pass	View alarms	us-east-1a	-

Inbound rules:

Name	Security group rule ID	Port range	Protocol	Source	Security group
-	sgr-021512c63bcf3f6e7	22	TCP	142.59.28.49/32	app-sg
-	sgr-01684d236f9a14929	80	TCP	0.0.0.0/0	app-sg

Outbound rules:

Name	Security group rule ID	Port range	Protocol	Destination	Security group
-	sqr-0fd4fe7ecf5c2d67f	All	All	0.0.0.0/0	app-sg

Screenshot of the AWS EC2 Connect to instance page for instance i-0f3b3f551cef80bf4.

Connect to instance info

Connect to your instance i-0f3b3f551cef80bf4 (AppASGinstance) using any of these options

EC2 Instance Connect Session Manager SSH client EC2 serial console

⚠ EC2 Instance Connect service IP addresses are not authorized
Port 22 (SSH) is authorized in [your security group](#). However, to use EC2 Instance Connect, it is recommended to also authorize port 22 for the EC2 Instance Connect service IP addresses in your Region: 18.206.107.24/29. [Learn more](#).

instance ID: i-0f3b3f551cef80bf4 (AppASGinstance)

Connection Type:

- Connect using EC2 Instance Connect: Connect using the EC2 Instance Connect browser-based client, with a public IPv4 or IPv6 address.
- Connect using EC2 Instance Connect Endpoint: Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint.

Public IPv4 address: 44.199.234.166

IPv6 address: -

Username: -

Screenshot of the Terraform code editor showing the ec2_autoscaling.tf file.

```

file: ec2_autoscaling.tf
resource "aws_security_group" "app_sg" {
  name     = "app-sg"
  description = "Allow HTTP and SSH traffic"
  vpc_id   = aws_vpc.main_vpc.id

  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["142.59.28.49/32"] # Restricted SSH to my IP only
  }

  ingress {
    from_port  = 80
    to_port    = 80
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"] # Allows HTTP traffic from anywhere
  }

  egress {
    from_port  = 0
    to_port    = 0
    protocol   = "-1"
    cidr_blocks = ["0.0.0.0/0"] # Allows all outbound traffic
  }
}

tags = [
  { Name = "AppSecurityGroup" }
]

```

Security Groups | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#SecurityGroups

aws Services Search [Alt+S] N. Virginia vivekashish@v3281

Images AMIs AMI Catalog

Elastic Block Store Volumes Snapshots Lifecycle Manager

Network & Security Security Groups

Elastic IPs Placement Groups Key Pairs Network Interfaces

Load Balancing Load Balancers Target Groups Trust Stores New

Auto Scaling Auto Scaling Groups

CloudShell Feedback

Security Groups (1/3) **Info**

Find resources by attribute or tag

Name Security group ID Security group name VPC ID Description

-	sg-0a6e24127747c69dd	default	vpc-01bb7ch53a9c88ff	default VPC sec
-	sg-026fe863d61fffc21	default	vpc-0ead15b9040f753dc	default VPC sec
<input checked="" type="checkbox"/> AppSecurityGroup	sg-0bf070d9429bd4a81	app-sg	vpc-0ea415b9040f753dc	Allow HTTP and

Inbound rules (1/3)

Search

Name Security group rule ID IP version Type Protocol Portrange Source Description

<input checked="" type="checkbox"/>	sgr-05141acb3cd090490	IPv4	SSH	TCP	22	18.206.107.24/29	-
<input type="checkbox"/>	sgr-021312c63bcf3f6e7	IPv4	SSH	TCP	22	142.59.28.49/32	-
<input type="checkbox"/>	sgr-01684d236f9a14929	IPv4	HTTP	TCP	80	0.0.0.0/0	-

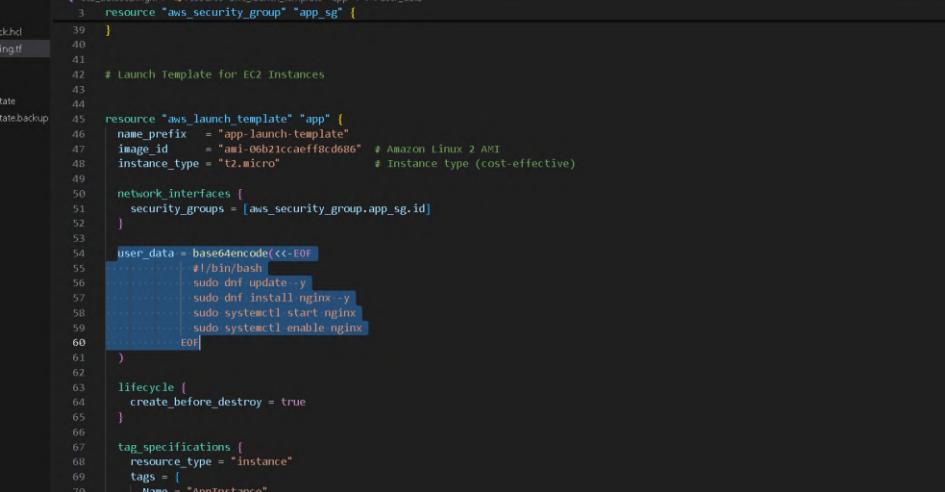
The screenshot shows a CloudShell terminal window titled "EC2 Instance Connect". The URL in the address bar is "us-east-1.console.aws.amazon.com/ec2-instance-connect/shell?region=us-east-1&connType=standard&instanceId=i-0f3b3f351cef80bf4". The terminal displays a root shell on an Amazon Linux 2023 instance. The user has run the command "curl https://aws.amazon.com/linux/amazon-linux-2023" and is viewing the resulting page. The page content includes the URL "https://aws.amazon.com/linux/amazon-linux-2023" and a large logo for Amazon Linux 2023. The terminal also shows the user's last login information: "Last login: Fri Oct 25 07:23:16 2024 from 18.206.107.28 [ec2-user@ip-10-0-1-61 ~]\$".

i-0f3b3f351cef80bf4 (AppASGInstance)
PublicIPs: 44.199.234.166 PrivateIPs: 10.0.1.61

```
# Amazon Linux 2023
# https://aws.amazon.com/linux/amazon-linux-2023

last login: Fri Oct 25 07:23:16 2024 from 18.206.107.28
[ec2-user@ip-10-0-1-61 ~]$ curl http://localhost
curl: (7) Failed to connect to localhost port 80 after 0 ms: Couldn't connect to server
[ec2-user@ip-10-0-1-61 ~]$
```

i-0f3b3f351cef80bf4 (AppASGInstance)
Public IPs: 44.199.234.166 Private IPs: 10.0.1.61



The screenshot shows the Visual Studio Code interface with the following details:

- File Explorer:** Shows the project structure with files like `main.tf`, `network.tf`, and `ec2_autoscaling.tf`.
- Terrafom Editor:** The main editor pane displays the `ec2_autoscaling.tf` file containing Terraform code.
- Code Snippets:** A sidebar on the right lists various AWS-related snippets for quick access.
- Bottom Status Bar:** Shows the current line (Ln 60), column (Col 16), character count (235 selected), and encoding (UTF-8). It also indicates the file type as `Terraform` and the provider as `CODEGPT`.

```
resource "aws_launch_template" "app" {
  name_prefix = "app-launch-template"
  image_id    = "ami-0db21ccaeff8cd686" # Amazon Linux 2 AMI
  instance_type = "t2.micro"           # Instance type (cost-effective)

  network_interfaces [
    security_groups = [aws_security_group.app_sg.id]
  ]

  user_data = base64encode(<<EOF
#!/bin/bash
sudo dnf update -y
sudo dnf install nginx -y
sudo systemctl start nginx
sudo systemctl enable nginx
EOF
)
}

lifecycle {
  create_before_destroy = true
}

tag_specifications [
  resource_type = "instance"
  tags = [
    Name = "AppInstance"
  ]
}
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -target=aws_launch_template.app

aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b99040f753dc]
aws_security_group.app_sg: Refreshing state... [id=sg-0bf07849429bd4a81]
aws_launch_template.app: Refreshing state... [id=lt-04ebe7149bc4195c]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
- destroy

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be destroyed
- resource "aws_autoscaling_group" "app_asg" {
    arn = "arn:aws:autoscaling:us-east-1:134575801911:autoScalingGroup:f03bdcbd-1b03-43dd-ae11-9e3aa43fec49:autoScalingGrou
pName"
  terraform="20241025062721849900000003" -> null
  - availability_zones = [
      - "us-east-1a",
    ] -> null
  - capacity_rebalance = false -> null
  - default_cooldown = 300 -> null
  - default_instance_warmup = 0 -> null
  - desired_capacity = 2 -> null
  - enabled_metrics = [] -> null
  - force_delete = false -> null
  - force_delete_warm_pool = false -> null
  - health_check_grace_period = 300 -> null
  - health_check_type = "EC2" -> null
  - id = "terraform-20241025062721849900000003" -> null
  - ignore_failed_scaling_activities = false -> null
  - load_balancers = [] -> null
  - max_instance_lifetime = 0 -> null
  - max_size = 5 -> null
  - metrics_granularity = "1Minute" -> null
  - min_size = 1 -> null
  - name = "terraform-20241025062721849900000003" -> null
  - name_prefix = "terraform-" -> null
  - predicted_capacity = 0 -> null
  - protect_from_scale_in = false -> null
  - service_linked_role_arn = "arn:aws:iam::134575801911:role/aws-service-role/autoscaling.amazonaws.com/AWSServiceRoleForAutoScaling" -> null
  - suspended_processes = [] -> null
  - target_group_arns = [] -> null
```

```
Configuration: configuration

The -target option is not for routine use, and is provided only for exceptional situations such as recovering from errors or mistakes, or when Terraform specifically suggests to use it as part of an error message.

Do you really want to destroy all resources?
Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

aws_autoscaling_policy.scale_up: Destroying... [id=scale-up]
aws_autoscaling_policy.scale_down: Destroying... [id=scale-down]
aws_autoscaling_policy.scale_up: Destruction complete after 1s
aws_autoscaling_policy.scale_down: Destruction complete after 1s
aws_autoscaling_group.app_asg: Destroying... [id=terraform-20241025062721849900000003]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 10s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 20s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 30s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 40s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 50s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 1m0s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 1m10s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 1m20s elapsed]
aws_autoscaling_group.app_asg: Still destroying... [id=terraform-20241025062721849900000003, 1m30s elapsed]
aws_launch_template.app: Destruction complete after 1m32s
aws_launch_template.app: Destruction complete after 1m32s

Warning: Applied changes may be incomplete

The plan was created with the -target option in effect, so some changes requested in the configuration may have been ignored and the output values may not be fully updated. Run the following command to verify that no other changes are pending:
  terraform plan

Note that the -target option is not suitable for routine use, and is provided only for exceptional situations such as recovering from errors or mistakes, or when Terraform specifically suggests to use it as part of an error message.

Destroy complete! Resources: 1 destroyed.
```

```

Command Prompt x + - o x
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform apply -auto-approve
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea15b9040f753dc]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bd0790bbd284a]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc0f97dfe205e5]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-0c184948abfe6c733]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-04b29a39893e7944d]
aws_security_group.app_sg: Refreshing state... [id=sg-0bf07049429bd4a81]
aws_route.internet_access: Refreshing state... [id=r-rtb-0c4fc0f97dfe205e51080289494]
aws_route_table_association.public_subnet_assoc: Refreshing state... [id=rbassoc-080529014332fac87]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be created
resource "aws_autoscaling_group" "app_asg" {
  # arm
  # availability_zones
  # default_coldown
  # desired_capacity
  # force_delete
  # force_delete_warm_pool
  # health_check_grace_period
  # health_check_type
  # id
  # ignore_failed_scaling_activities
  # load_balancers
  # max_size
  # metrics_granularity
  # min_size
  # name
  # name_prefix
  # predicted_capacity
  # protect_from_scale_in
  # service_linked_role_arn
  # target_group_arns
  # vpc_zone_identifier
}

```

```

Command Prompt x + - o x
+ default_version = (known after apply)
+ id = (known after apply)
+ image_id = "ami-06b21ccaeff8cd686"
+ instance_type = "t2.micro"
+ latest_version = (known after apply)
+ name = (known after apply)
+ name_prefix = "app-launch-template"
+ tags_all = (known after apply)
+ user_data = "IyEvYmluL2Jhc2gNCnN1ZG8gZG5mIHVwZGF0ZSAtQ0Kc3VkbmYgaW5zdGFsbCBuZ2lueCAteQ0Kc3VkbYBzeXN0ZW1jdGwg3RhcnQgbmdpbngNCnN1ZG8gc3lzdGVtY3RsIGVuYWJsZSBuZ2lueA0K"
+ metadata_options (known after apply)
+ network_interfaces {
    + security_groups = [
        + "sg-0bf07049429bd4a81",
    ]
}
+ tagSpecifications {
    + resource_type = "instance"
    + tags = {
        + "Name" = "AppInstance"
    }
}
}

Plan: 4 to add, 0 to change, 0 to destroy.
aws_launch_template.app: Creating...
aws_launch_template.app: Creation complete after 0s [id=lt-09d7e0b6e8deddeb]
aws_autoscaling_group.app_asg: Creating...
aws_autoscaling_group.app_asg: Still creating... [10s elapsed]
aws_autoscaling_group.app_asg: Creation complete after 16s [id=terraform-20241025074616245700000003]
aws_autoscaling_policy.scale_down: Creating...
aws_autoscaling_policy.scale_up: Creating...
aws_autoscaling_policy.scale_up: Creation complete after 0s [id=scale-up]
aws_autoscaling_policy.scale_down: Creation complete after 0s [id=scale-down]

Apply complete! Resources: 4 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

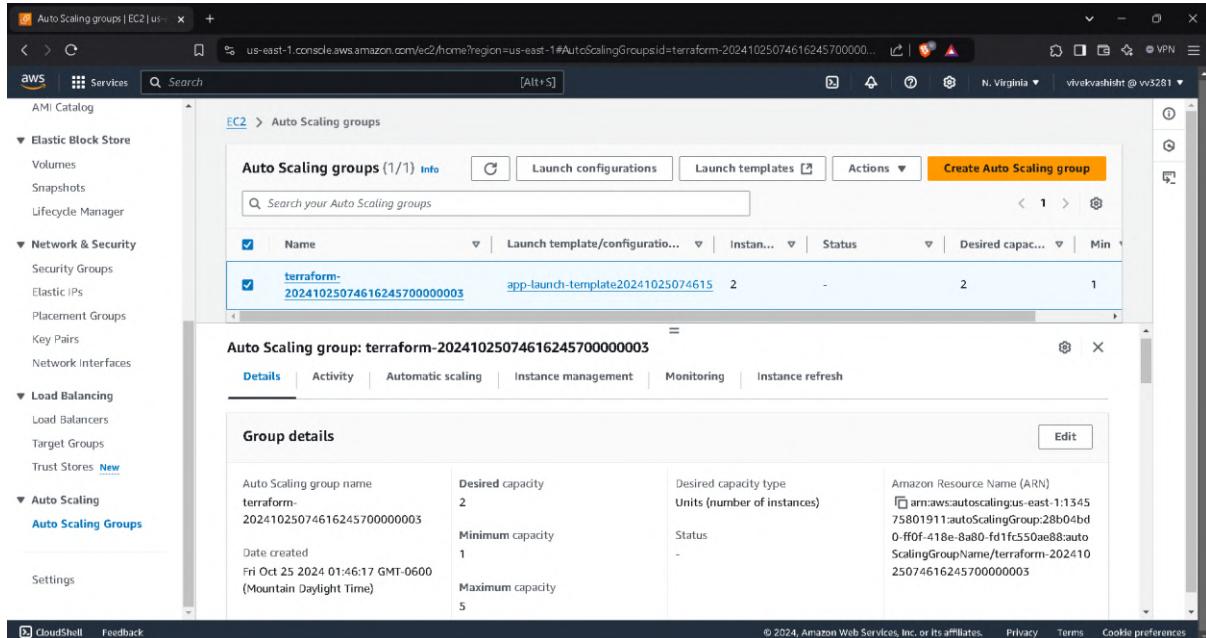
```

The screenshot shows the AWS EC2 Instances page. The left sidebar includes links for EC2 Dashboard, EC2 Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity, Reservations (New), Images (AMIs, AMI Catalog), and Elastic Block Store (Volumes, Snapshots, Lifecycle Manager). The main content area displays a table titled 'Instances (2/4) info' with columns: Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, and Public IPv4. There are four rows: two terminated instances (i-0fb3f351cef80bf4 and i-0bb6c24776f8fb24) and two running instances (i-0824eee43289a9... and i-059e693bc67d53ed5), both labeled 't2.micro'. The status check for the running instances shows '2/2 checks pass'. At the bottom, it says '2 instances selected'.

The screenshot shows the AWS Management Console with the EC2 Instances page open. The left sidebar shows navigation links for EC2 Dashboard, EC2 Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity, and Reservations. The main content area displays a table of instances with columns for Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, and Public IP. One instance, "AppASGinstance" with ID i-08242eee43289a9eb, is selected and shown in a detailed view below the table. The detailed view includes tabs for Details, Status and alarms, Monitoring, Security, Networking, Storage, and Tags. Under the Details tab, the Instance summary section shows fields for Instance ID (i-08242eee43289a9eb), IPv4 address (3.219.247.195), Private IPv4 addresses (10.0.1.136), IPv6 address (-), Instance state (Running), Public IP DNS name (ip-10-0-1-136.ec2.internal), Answer private resource DNS name (-), Instance type (t2.micro), and Elastic IP addresses (-).

```
instances | EC2 | us-east-1 EC2 Instance Connect Welcome to nginx! + us-east-1.console.aws.amazon.com/ec2-instance-connect/sh?region=us-east-1&connType=standard&instanceId=i-0824e... Search N. Virginia VPN [Alt+J] Services Search AWS # Amazon Linux 2023 https://aws.amazon.com/linux/amazon-linux-2023 [ec2-user@ip-10-0-1-136 ~]$ sudo systemctl status nginx nginx.service - The nginx HTTP and reverse proxy server Active: active (running) since Fri 2024-10-25 07:47:08 UTC; lmin 15s ago Main PID: 3594 (nginx) Tasks: 2 (limit: 1112) Memory: 1.2K CPU: 49ms CGroup: /system.slice/nginx.service ├─ 3594 nginx master process └─ 3600 nginx worker process [ec2-user@ip-10-0-1-136 ~]$ curl http://localhost <!DOCTYPE html> <html> <head> <title>Welcome to nginx!</title> <style> body { color: #000; background-color: #fff; margin: 0; padding: 0; } font-family: Tahoma, Verdana, Arial, sans-serif; </style> </head> <body> <h1>Welcome to nginx!</h1> <p>if you see this page, the nginx web server is successfully installed and working. Further configuration is required.</p> <p>For online documentation and support please refer to <a href="http://nginx.org">http://nginx.org</a><br> Commercial support is available at <a href="http://nginx.com/en/support.html">http://nginx.com/en/support.html</a>.</p> <p><em>Thank you for using nginx.</em></p> </body> </html> [ec2-user@ip-10-0-1-136 ~]$
```

A screenshot of a web browser window. The address bar shows the URL as "Not secure 3.219.247.193". The main content of the page is a large, bold, black "Welcome to nginx!" heading. Below it, smaller text reads: "If you see this page, the nginx web server is successfully installed and working. Further configuration is required." It also includes links to "nginx.org" and "nginx.com". At the bottom, there is a thank you message: "Thank you for using nginx.".



IAM Roles and S3 Integration

a) Overview

The objective of this phase was to set up IAM roles for EC2 instances to enable them to access S3 buckets securely, following AWS best practices. This setup involved defining IAM roles, updating instance configurations, and creating an additional S3 bucket for backups. Each step was carefully crafted to ensure the correct permissions and organization of resources in Terraform.

b) IAM Role Definition

To begin, I created a new file named `iam_roles.tf` in the Terraform project folder. This file defines an IAM role that allows EC2 instances to assume permissions required to access S3 resources. Here's the code for the IAM role, which includes a policy attachment to allow S3 read-only access:

```
# IAM Role for EC2 to assume
resource "aws_iam_role" "ec2_role" {
    name = "ec2-role"

    assume_role_policy = <<EOF
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Service": "ec2.amazonaws.com"
            },
            "Action": "sts:AssumeRole"
        }
    ]
}
```

```

}
EOF
}

# Attachment of S3 read-only policy to the role
resource "aws_iam_role_policy_attachment" "ec2_s3_readonly" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess"
}

```

This configuration defines the EC2 role and attaches the `AmazonS3ReadOnlyAccess` policy to ensure EC2 instances can only read from S3, following the principle of least privilege.

c) Updating `ec2_autoscaling.tf` for IAM Role Attachment

Next, I updated the `ec2_autoscaling.tf` file to include the IAM role for the EC2 instances by defining an instance profile and attaching it to the launch template:

```

# IAM instance profile for EC2
resource "aws_iam_instance_profile" "ec2_instance_profile" {
  name = "ec2-instance-profile"
  role = aws_iam_role.ec2_role.name
}

# Attach the IAM role to EC2 instances
aws_iam_instance_profile {
  name = aws_iam_instance_profile.ec2_instance_profile.name
}

```

This instance profile links the IAM role with EC2 instances, granting them the necessary permissions to access S3 resources.

d) S3 Bucket Configuration for Backups

For better resource organization and compliance with Terraform principles, I moved the configuration of the S3 bucket used for Terraform state storage from `main.tf` to a new file, `s3_bucket.tf`. This file also defines a new S3 bucket specifically for storing backups. The updated code for `s3_bucket.tf` is as follows:

```

# S3 bucket for storing Terraform state
resource "aws_s3_bucket" "terraform_state" {
  bucket = "aws-tf-backend-vv-bucket"

  tags = {
    Name      = "TerraformStateBucket"
    Environment = "Dev"
  }
}

# Public access block for the S3 bucket
resource "aws_s3_bucket_public_access_block" "terraform_state_block" {
  bucket = aws_s3_bucket.terraform_state.id

  block_public_acls      = true
  block_public_policy     = true
  restrict_public_buckets = true
}

```

```

        ignore_public_acls      = true
    }

# S3 bucket for backups
resource "aws_s3_bucket" "backup_bucket" {
    bucket = "terra-backup-vv-project-bucket"

    tags = {
        Name          = "BackupBucket"
        Environment   = "Dev"
    }
}

# Public access block for the backup S3 bucket
resource "aws_s3_bucket_public_access_block" "backup_bucket_block" {
    bucket = aws_s3_bucket.backup_bucket.id

    block_public_acls      = true
    block_public_policy     = true
    restrict_public_buckets = true
    ignore_public_acls      = true
}

```

The S3 bucket `terra-backup-vv-project-bucket` is configured for backup storage with public access blocked for security. Both S3 buckets have access controls to ensure they are secure and compliant.

e) Applying the Changes

After configuring `iam_roles.tf` and `s3_bucket.tf`, I followed these steps to apply the changes:

1. **Destroying the Existing Launch Template** – Since the updated instance profile would only take effect on new instances, I first destroyed the current launch template:

```
terraform destroy -target=aws_launch_template.app
```

2. **Re-initializing and Applying Terraform** – I then ran the following commands to initialize, plan, and apply the updated configurations:

```
terraform init
terraform plan
terraform apply -auto-approve
```

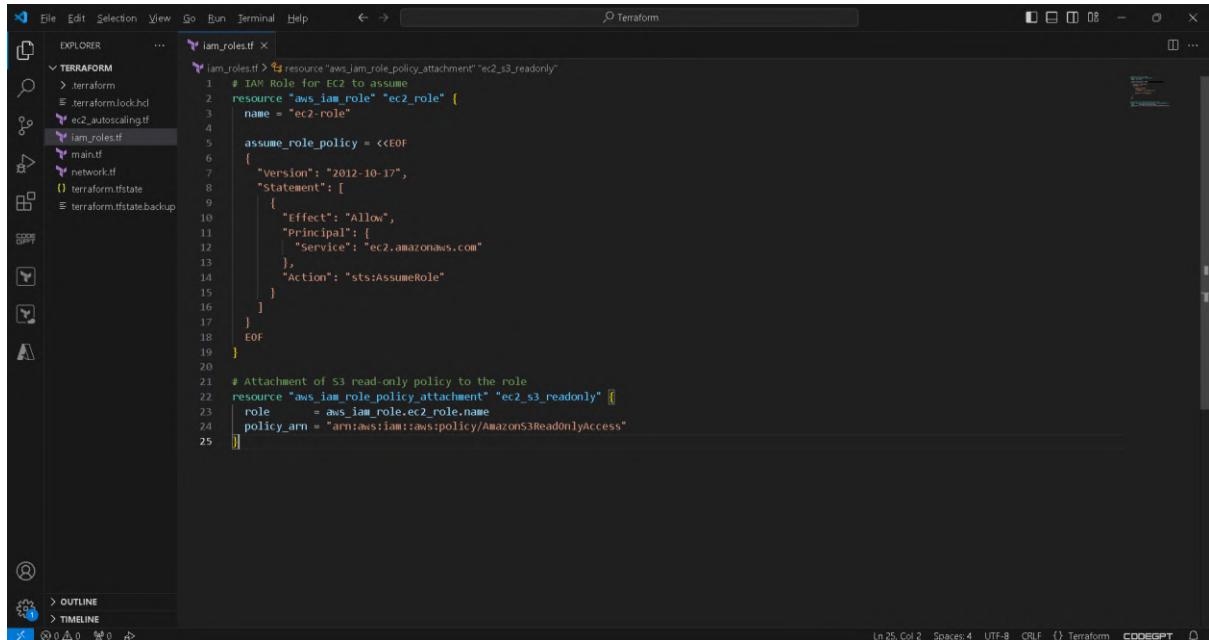
f) Verification

After a successful apply, I verified the resources in the AWS Console:

- **IAM Role and Policy** – The `ec2-role` with `AmazonS3ReadOnlyAccess` was successfully created and attached to the EC2 instances.
- **EC2 Instances** – Each instance had the `ec2-role` attached, allowing them to access the S3 bucket with read-only permissions.
- **S3 Buckets** – The `terra-backup-vv-project-bucket` was created and configured as expected, in addition to the Terraform state bucket.

This phase of the project was successfully completed with the EC2 instances securely configured to access S3 resources, following least-privilege access principles and terraform best practices.

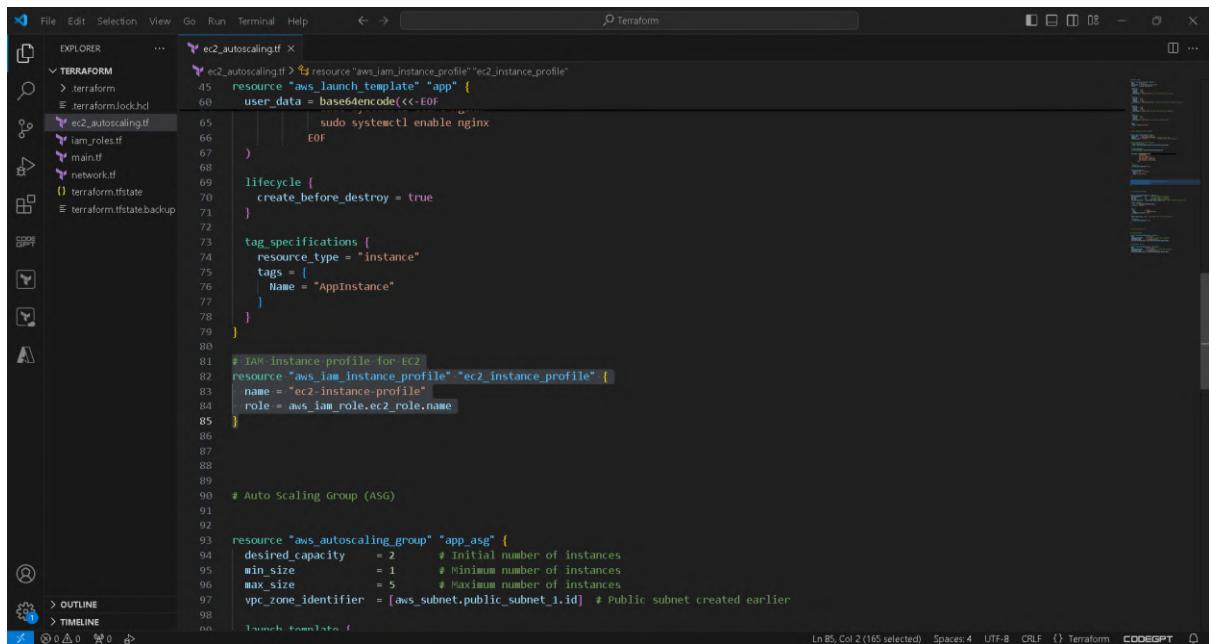
Screenshots



The screenshot shows the CodeGPT IDE interface with the Terraform extension active. The left sidebar displays the project structure under 'EXPLORER' with files like .terraform, .terraform.lock.hcl, ec2_autoscaling.tf, and iam_roles.tf. The main editor window shows the contents of the iam_roles.tf file:

```
File Edit Selection View Go Run Terminal Help < > Terraform
EXPLORER ... iam_roles.tf ...
TERRAFORM > .terraform > iam_roles.tf > resource "aws_iam_role_policy_attachment" "ec2_s3_readonly"
1 # IAM Role for EC2 to assume
2 resource "aws_iam_role" "ec2_role" {
3   name = "ec2-role"
4
5   assume_role_policy = <<EOF
6   {
7     "Version": "2012-10-17",
8     "Statement": [
9       {
10       "Effect": "Allow",
11       "Principal": "*",
12       "Service": "ec2.amazonaws.com"
13     },
14     "Action": "sts:AssumeRole"
15   ]
16 }
17 EOF
18
19
20 # Attachment of s3 read-only policy to the role
21 resource "aws_iam_role_policy_attachment" "ec2_s3_readonly" [
22   role      = aws_iam_role.ec2_role.name
23   policy_arn = "arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess"
24 ]
25 ]
```

The status bar at the bottom indicates 'Ln 25, Col 2' and 'Spaces:4'. The right side of the interface shows a preview of the AWS CloudFormation template generated from the Terraform code.



The screenshot shows the CodeGPT IDE interface with the Terraform extension active. The left sidebar displays the project structure under 'EXPLORER' with files like .terraform, .terraform.lock.hcl, iam_roles.tf, main.tf, network.tf, terraform.tfstate, and terraform.tfstate.backup. The main editor window shows the contents of the ec2_autoscaling.tf file:

```
File Edit Selection View Go Run Terminal Help < > Terraform
EXPLORER ... ec2_autoscaling.tf ...
TERRAFORM > .terraform > ec2_autoscaling.tf > resource "aws_launch_template" "app" {
1   user_data = base64encode(<<EOF
2
3   sudo systemctl enable nginx
4
5   EOF
6 )
7
8   lifecycle {
9     create_before_destroy = true
10   }
11
12   tagSpecifications {
13     resource_type = "instance"
14     tags = [
15       {
16         Name = "Appinstance"
17       }
18     ]
19   }
20
21 # IAM instance profile for EC2
22 resource "aws_iam_instance_profile" "ec2_instance_profile" [
23   name = "ec2-instance-profile"
24   role = aws_iam_role.ec2_role.name
25 ]
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45   resource "aws_autoscaling_group" "app_asg" {
46     desired_capacity    = 2      # Initial number of instances
47     min_size           = 1      # Minimum number of instances
48     max_size           = 5      # Maximum number of instances
49     vpc_zone_identifier = [aws_subnet.public_subnet_1.id] # Public subnet created earlier
50
51   }
52
53
54
55
56
57
58
59
60 }
```

The status bar at the bottom indicates 'Ln 85, Col 2 (165 selected)' and 'Spaces:4'. The right side of the interface shows a preview of the AWS CloudFormation template generated from the Terraform code.

The screenshot shows the CodeGPT extension integrated into the Visual Studio Code interface. The left sidebar features the 'EXPLORER' view with a tree structure containing files like 'main.tf', 's3_bucket.tf', 'iam_roles.tf', and 'network.tf'. The main editor area displays Terraform configuration for an S3 bucket named 'aws-tf-backend-vv-bucket'. The code includes a resource block for the bucket itself, which is identified as the 'terraform_state_block'. It also includes a block for public access, setting 'block_public_acls' to true and 'block_public_policy' to true. The status bar at the bottom indicates the code is at Line 21, Column 2, with 4 spaces, using UTF-8 encoding, and is associated with the 'Terraform' language.

```
s3_bucket.tf > main.tf
1 # S3 bucket for storing Terraform state
2
3 resource "aws_s3_bucket" "terraform_state" {
4   bucket = "aws-tf-backend-vv-bucket"
5
6   tags = [
7     Name      = "TerraformStateBucket"
8     Environment = "Dev"
9   ]
10 }
11
12 # Public access block for the s3 bucket
13 resource "aws_s3_bucket_public_access_block" "terraform_state_block" {
14   bucket = aws_s3_bucket.terraform_state.id
15
16   block_public_acls  = true
17   block_public_policy = true
18   restrict_public_buckets = true
19   ignore_public_acls = true
20 }
21
```

The screenshot shows the Codecept IDE interface with a Terraform project open. The left sidebar displays the file structure under 'TERRAFORM' with files like .terraform, .terraform.lock.hcl, ec2_autoscaling.tf, iam_roles.tf, main.tf, network.tf, s3_bucket.tf, terraform.state, and terraform.tfstate.backup. The right pane shows the contents of the main.tf file:

```
1 terraform {  
2   backend "s3" {  
3     bucket = "aws-tf-backend-vv-bucket"  
4     key    = "terraform/dev/terraform.tfstate"  
5     region = "us-east-1"  
6   }  
7 }  
8  
9 provider "aws" {  
10   region = "us-east-1"  
11 }
```

The screenshot shows the CodeGPT IDE interface with the Terraform configuration file 's3_bucket.tf' open in the main editor area. The code defines an AWS S3 bucket named 'terra-backup-vv-project-bucket' with specific access policies and tags.

```
resource "aws_s3_bucket" "backup_bucket" {
  bucket = "terra-backup-vv-project-bucket"
  tags = {
    Name      = "BackupBucket"
    Environment = "Dev"
  }
}

resource "aws_s3_bucket_public_access_block" "backup_bucket_block" {
  bucket = aws_s3_bucket.backup_bucket.id

  block_public_acls  = true
  block_public_policy = true
  restrict_public_buckets = true
  ignore_public_acls = true
}
```

The screenshot shows a Command Prompt window with the command `terraform destroy -target=aws_launch_template.app` being run. The output shows the command is still executing.

The screenshot shows a Command Prompt window with the command `terraform init` being run. The output indicates Terraform has been successfully initialized and provides instructions for working with the provider.

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!
```

The screenshot then shows the command `terraform plan` being run, displaying the execution plan for creating an AWS Auto Scaling group.

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup>terraform plan
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bd0796bbd28a4a]
aws_subnet_public.subnet_1: Refreshing state... [id=subnet-04b29a39893e7944c]
aws_subnet_private.subnet_1: Refreshing state... [id=subnet-0c184948abfe6c733]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc0f97dfe205e5]
aws_security_group.app_sg: Refreshing state... [id=sg-0bf07049429bd4a81]
aws_route_internet.access: Refreshing state... [id=rtrb-0c4fc0f97dfe205e51080289494]
aws_route_table_association.public_subnet_assoc: Refreshing state... [id=rtbassoc-080529014332fac87]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be created
+ resource "aws_autoscaling_group" "app_asg" {
  + arn                                = (known after apply)
  + availability_zones                  = (known after apply)
  + default_cooldown                   = (known after apply)
  + desired_capacity                   = 2
  + force_delete                       = false
  + force_delete_warm_pool             =
```

```
Command Prompt
# aws_autoscaling_policy.scale_down will be created
+ resource "aws_autoscaling_policy" "scale_down" {
+   adjustment_type      = "ChangeInCapacity"
+   arn                  = "(known after apply)"
+   autoscaling_group_name = "(known after apply)"
+   enabled               = true
+   id                   = "(known after apply)"
+   metric_aggregation_type = "(known after apply)"
+   name                 = "scale-down"
+   policy_type          = "SimpleScaling"
+   scaling_adjustment    = -1
}

# aws_autoscaling_policy.scale_up will be created
+ resource "aws_autoscaling_policy" "scale_up" {
+   adjustment_type      = "ChangeInCapacity"
+   arn                  = "(known after apply)"
+   autoscaling_group_name = "(known after apply)"
+   enabled               = true
+   id                   = "(known after apply)"
+   metric_aggregation_type = "(known after apply)"
+   name                 = "scale-up"
+   policy_type          = "SimpleScaling"
+   scaling_adjustment    = 1
}

# aws_iam_instance_profile.ec2_instance_profile will be created
+ resource "aws_iam_instance_profile" "ec2_instance_profile" {
+   arn                  = "(known after apply)"
+   create_date          = "(known after apply)"
+   id                   = "(known after apply)"
+   name                 = "ec2-instance-profile"
+   name_prefix          = "(known after apply)"
+   path                 = "/"
+   role                 = "ec2-role"
+   tags_all             = "(known after apply)"
+   unique_id            = "(known after apply)"
}

# aws_iam_role.ec2_role will be created
+ resource "aws_iam_role" "ec2_role" {
```

```
Command Prompt
# aws_iam_role_policy_attachment.ec2_s3_readonly will be created
+ resource "aws_iam_role_policy_attachment" "ec2_s3_readonly" {
+   id                  = "(known after apply)"
+   policy_arn          = "arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess"
+   role                = "ec2-role"
}

# aws_launch_template.app will be created
+ resource "aws_launch_template" "app" {
+   arn                  = "(known after apply)"
+   default_version      = "(known after apply)"
+   id                   = "(known after apply)"
+   image_id             = "ami-06b21ccaeff8cd686"
+   instance_type        = "t2.micro"
+   latest_version       = "(known after apply)"
+   name                 = "(known after apply)"
+   name_prefix          = "app-launch-template"
+   tags_all             = "(known after apply)"
+   user_data            = "TfyEvYmluL2Jhc2gNCnN1ZG8gZG5mIHVwZGF0ZSAteQ0Kc3VkbmYgaW5zdGFsbCBuZ2lueCAtQ0Kc3VkbByBzeXN0ZW1jdGwg3RhcnQgbmdpbngNCnN1ZG8gc3lzdGVtY3RsIGVuYmJzSBuZ2lueA0K"

+   iam_instance_profile {
+     name = "ec2-instance-profile"
+   }

+   metadata_options (known after apply)

+   network_interfaces {
+     security_groups = [
+       "+sg-0bf07049429bd4a81",
+     ]
+   }

+   tag_specifications {
+     resource_type = "instance"
+     tags          = {
+       "+Name" = "AppInstance"
+     }
+   }
}
```

```
Command Prompt
# aws_s3_bucket.backup_bucket will be created
+ resource "aws_s3_bucket" "backup_bucket" {
+   acceleration_status      = "(known after apply)"
+   acl                      = "(known after apply)"
+   arn                      = "(known after apply)"
+   bucket                   = "terra-backup-vv-project-bucket"
+   bucket_domain_name       = "(known after apply)"
+   bucket_prefix             = "(known after apply)"
+   bucketRegionalDomainName = "(known after apply)"
+   force_destroy             = false
+   hostedZoneId             = "(known after apply)"
+   id                       = "(known after apply)"
+   objectLockEnabled         = "(known after apply)"
+   policy                   = "(known after apply)"
+   region                   = "(known after apply)"
+   requestPayer              = "(known after apply)"
+   tags                     = {
+     "+Environment" = "Dev"
+     "+Name"         = "BackupBucket"
+   }
+   tags_all                 = {
+     "+Environment" = "Dev"
+     "+Name"         = "BackupBucket"
+   }
+   website_domain           = "(known after apply)"
+   website_endpoint          = "(known after apply)"
+   cors_rule (known after apply)
+   grant (known after apply)
+   lifecycle_rule (known after apply)
+   logging (known after apply)
+   objectLockConfiguration (known after apply)
+   replicationConfiguration (known after apply)
+   serverSideEncryptionConfiguration (known after apply)
```

```

Command Prompt x + - ×
# aws_s3_bucket_public_access_block.backup_bucket_block will be created
+ resource "aws_s3_bucket_public_access_block" "backup_bucket_block" {
  + block_public_acls      = true
  + block_public_policy     = true
  + bucket                  = (known after apply)
  + id                      = (known after apply)
  + ignore_public_acls      = true
  + restrict_public_buckets = true
}

Plan: 9 to add, 0 to change, 0 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

C:\Users\hp\OneDrive\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform apply
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-0c184948abfe6c733]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc8f97dfe205e5]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bd0799bbd28a4a]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-00b29a39893e7944d]
aws_security_group.app_sg: Refreshing state... [id=sg-0bf0709429b4d4a81]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_route.internet_access: Refreshing state... [id=r-rtb-0c4fc8f97dfe205e510880289494]
aws_route_table_association.public_subnet_assoc: Refreshing state... [id=rtbassoc-086529014332fac87]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be created
+ resource "aws_autoscaling_group" "app_asg" {
  + arn                                = (known after apply)
  + availability_zones                 = (known after apply)
  + default_cooldown                   = (known after apply)
  + desired_capacity                  = 2
  + force_delete                      = false
}

```

```

Command Prompt x + - ×
+ resource "aws_s3_bucket_public_access_block" "backup_bucket_block" {
  + block_public_acls      = true
  + block_public_policy     = true
  + bucket                  = (known after apply)
  + id                      = (known after apply)
  + ignore_public_acls      = true
  + restrict_public_buckets = true
}

Plan: 9 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

Enter a value: yes

aws_iam_role.ec2_role: Creating...
aws_s3_bucket.backup_bucket: Creating...
aws_iam_role.ec2_role: Creation complete after 0s [id=ec2-role]
aws_iam_role_policy_attachment.ec2_s3_readonly: Creating...
aws_iam_instance_profile.ec2_instance_profile: Creating...
aws_iam_role_policy_attachment.ec2_s3_readonly: Creation complete after 1s [id=ec2-role-20241026083002429100000001]
aws_iam_instance_profile.ec2_instance_profile: Creation complete after 1s [id=ec2-instance-profile]
aws_launch_template.app: Creating...
aws_launch_template.app: Creation complete after 1s [id=lt-0d1591b9b2e425b]
aws_s3_bucket.backup_bucket: Creation complete after 2s [id=terra-backup-vv-project-bucket]
aws_s3_bucket_public_access_block.backup_bucket_block: Creating...
aws_autoscaling_group.app_asg: Creating...
aws_s3_bucket_public_access_block.backup_bucket_block: Creation complete after 0s [id=terra-backup-vv-project-bucket]
aws_autoscaling_group.app_asg: Still creating... [10s elapsed]
aws_autoscaling_group.app_asg: Still creating... [20s elapsed]
aws_autoscaling_group.app_asg: Creation complete after 22s [id=terraform-20241026083003598900000004]
aws_autoscaling_policy.scale_down: Creating...
aws_autoscaling_policy.scale_up: Creating...
aws_autoscaling_policy.scale_up: Creation complete after 1s [id=scale-up]
aws_autoscaling_policy.scale_down: Creation complete after 1s [id=scale-down]

apply complete! Resources: 9 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

```

Identity and Access Management (IAM)

ec2-role Info

Summary

Creation date	ARN
October 26, 2024, 02:30 (UTC-06:00)	arn:aws:iam:134575801911:role/ec2-role
Last activity	Instance profile ARN
-	arn:aws:iam:134575801911:instance-profile/ec2-instance-profile
Maximum session duration	1 hour

Permissions Trust relationships Tags Last Accessed Revoke sessions

Permissions policies (1) Info

You can attach up to 10 managed policies.

Policy name	Type	Attached entities
AmazonSSReadonlyAccess	AWS managed	1

Screenshot of the AWS EC2 Instances page showing two running t2.micro instances. The left sidebar includes links for EC2 Dashboard, Global View, Events, Instances (with sub-links for Instances, Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity, and Reservations), Images (AMIs, Catalog), and Elastic Block Store (Volumes, Snapshots, Lifecycle Manager). The bottom of the page features CloudShell and Feedback buttons.

Instances (1/2) Info

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public
AppASGinstance	i-0540213890ad3b53a	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	-
AppASGinstance	i-0986608f4066935da	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	-

i-0540213890ad3b53a (AppASGinstance)

Security

Owner ID: 134575801911 | Launch time: Sat Oct 26 2024 02:30:16 GMT-0600 (Mountain Daylight Time)

Security groups: sg-0bf07049429bd4a81 (app-sg)

Inbound rules:

Screenshot of the AWS EC2 Instances page showing two running t2.micro instances. The left sidebar includes links for EC2 Dashboard, Global View, Events, Instances (with sub-links for Instances, Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity, and Reservations), Images (AMIs, Catalog), and Elastic Block Store (Volumes, Snapshots, Lifecycle Manager). The bottom of the page features CloudShell and Feedback buttons.

Instances (1/2) Info

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public
AppASGinstance	i-0540213890ad3b53a	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	-
AppASGinstance	i-0986608f4066935da	Running	t2.micro	2/2 checks pass	View alarms +	us-east-1a	-

i-0986608f4066935da (AppASGinstance)

Security

Owner ID: 134575801911 | Launch time: Sat Oct 26 2024 02:30:16 GMT-0600 (Mountain Daylight Time)

Security groups: sg-0bf07049429bd4a81 (app-sg)

Inbound rules:

Screenshot of the AWS S3 Buckets page showing two general purpose buckets: aws-tf-backend-vv-bucket and terra-backup-vv-project-bucket. The left sidebar includes links for Buckets, Access Grants, Access Points, Object Lambda Access Points, Multi-Region Access Points, Batch Operations, IAM Access Analyzer for S3, Block Public Access settings for this account, Storage Lens (Dashboards, Storage Lens groups), AWS Organizations settings, Feature spotlight, and AWS Marketplace for S3. The bottom of the page features CloudShell and Feedback buttons.

Amazon S3

General purpose buckets (2) Info

Name	AWS Region	IAM Access Analyzer	Creation date
aws-tf-backend-vv-bucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	October 23, 2024, 03:10:54 (UTC-06:00)
terra-backup-vv-project-bucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	October 26, 2024, 02:30:03 (UTC-06:00)

Multi-Region Deployment and Failover

a) Overview

In this phase, I set up a multi-region deployment with a failover mechanism. By establishing resources in two AWS regions, `us-east-1` (primary) and `us-west-2` (secondary), I ensured the application would remain available in case of failure in the primary region. To support this, I configured an Application Load Balancer (ALB) in each region and implemented route tables, VPCs, subnets, and Internet Gateways.

b) Configuring Providers for Multi-Region Deployment

I updated the `main.tf` file to define an additional provider alias for `us-west-2`. This allows us to deploy resources in two regions within the same Terraform configuration.

```
# Adding provider for us-west-2
provider "aws" {
  alias  = "west"
  region = "us-west-2"
}
```

The complete `main.tf` now includes both `us-east-1` and `us-west-2` configurations:

```
terraform {
  backend "s3" {
    bucket = "aws-tf-backend-vv-bucket"
    key    = "terraform/dev/terraform.tfstate"
    region = "us-east-1"
  }
}
```

```

}

provider "aws" {
  region = "us-east-1"
}

provider "aws" {
  alias  = "west"
  region = "us-west-2"
}

```

c) Network Configuration in us-east-1

In `network.tf`, I added a second public subnet for the primary region (`us-east-1`). Two public subnets are required to set up the Application Load Balancer for high availability.

```

# Additional Public Subnet in us-east-1b for ALB
resource "aws_subnet" "public_subnet_2" {
  vpc_id          = aws_vpc.main_vpc.id
  cidr_block      = "10.0.4.0/24"
  availability_zone = "us-east-1b"
  map_public_ip_on_launch = true

  tags = {
    Name = "PublicSubnet2"
  }
}

```

I also modified the CIDR block of the private subnet in `us-east-1a` to `10.0.3.0/24` to avoid any overlapping issues.

d) Configuring Network Resources in us-west-2

Following the setup for `us-east-1`, I configured similar resources for the `us-west-2` region. This includes creating a VPC, public and private subnets, Internet Gateway, and route tables. These resources allow the application to function independently within `us-west-2`, ensuring that traffic can be routed to this region if the primary region fails.

Here's the final `network.tf` configuration for multi-region deployment:

```

# VPC and Subnets for us-east-1
resource "aws_vpc" "main_vpc" {
  cidr_block = "10.0.0.0/16"

  tags = {
    Name = "MainVPC"
  }
}

resource "aws_subnet" "public_subnet_1" {
  vpc_id          = aws_vpc.main_vpc.id
  cidr_block      = "10.0.1.0/24"
  availability_zone = "us-east-1a"
  map_public_ip_on_launch = true

  tags = {

```

```

        Name = "PublicSubnet1"
    }
}

resource "aws_subnet" "public_subnet_2" {
    vpc_id           = aws_vpc.main_vpc.id
    cidr_block       = "10.0.4.0/24"
    availability_zone = "us-east-1b"
    map_public_ip_on_launch = true

    tags = {
        Name = "PublicSubnet2"
    }
}

resource "aws_subnet" "private_subnet_1" {
    vpc_id           = aws_vpc.main_vpc.id
    cidr_block       = "10.0.3.0/24"
    availability_zone = "us-east-1a"
    map_public_ip_on_launch = false

    tags = {
        Name = "PrivateSubnet1"
    }
}

# VPC and Subnets for us-west-2
resource "aws_vpc" "west_vpc" {
    provider = aws.west
    cidr_block = "10.1.0.0/16"

    tags = {
        Name = "WestVPC"
    }
}

resource "aws_subnet" "west_public_subnet_1" {
    provider = aws.west
    vpc_id = aws_vpc.west_vpc.id
    cidr_block = "10.1.1.0/24"
    availability_zone = "us-west-2a"
    map_public_ip_on_launch = true

    tags = {
        Name = "WestPublicSubnet1"
    }
}

resource "aws_subnet" "west_public_subnet_2" {
    provider = aws.west
    vpc_id = aws_vpc.west_vpc.id
    cidr_block = "10.1.2.0/24"
    availability_zone = "us-west-2b"
    map_public_ip_on_launch = true

    tags = {
        Name = "WestPublicSubnet2"
    }
}

resource "aws_subnet" "west_private_subnet_1" {

```

```

provider = aws.west
vpc_id = aws_vpc.west_vpc.id
cidr_block = "10.1.3.0/24"
availability_zone = "us-west-2a"

tags = {
  Name = "WestPrivateSubnet1"
}
}

```

e) Internet Gateways and Route Tables

I defined Internet Gateways for both regions and associated each with its corresponding VPC. For each public subnet in both regions, I configured route tables to allow outbound internet access through the Internet Gateway.

```

# Internet Gateway for us-east-1
resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main_vpc.id
  tags = {
    Name = "MainIGW"
  }
}

# Public Route Table and Route for us-east-1
resource "aws_route_table" "public_rt" {
  vpc_id = aws_vpc.main_vpc.id
  tags = {
    Name = "PublicRouteTable"
  }
}

resource "aws_route" "internet_access" {
  route_table_id = aws_route_table.public_rt.id
  destination_cidr_block = "0.0.0.0/0"
  gateway_id = aws_internet_gateway.igw.id
}

resource "aws_route_table_association" "public_subnet_assoc_1" {
  subnet_id = aws_subnet.public_subnet_1.id
  route_table_id = aws_route_table.public_rt.id
}

resource "aws_route_table_association" "public_subnet_assoc_2" {
  subnet_id = aws_subnet.public_subnet_2.id
  route_table_id = aws_route_table.public_rt.id
}

# Internet Gateway and Route Table for us-west-2
resource "aws_internet_gateway" "west_igw" {
  provider = aws.west
  vpc_id = aws_vpc.west_vpc.id
  tags = {
    Name = "WestIGW"
  }
}

resource "aws_route_table" "west_public_rt" {
  provider = aws.west

```

```

vpc_id = aws_vpc.west_vpc.id
tags = {
    Name = "WestPublicRouteTable"
}
}

resource "aws_route" "west_internet_access" {
    provider = aws.west
    route_table_id = aws_route_table.west_public_rt.id
    destination_cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.west_igw.id
}

resource "aws_route_table_association" "west_public_subnet_assoc_1" {
    provider = aws.west
    subnet_id = aws_subnet.west_public_subnet_1.id
    route_table_id = aws_route_table.west_public_rt.id
}

resource "aws_route_table_association" "west_public_subnet_assoc_2" {
    provider = aws.west
    subnet_id = aws_subnet.west_public_subnet_2.id
    route_table_id = aws_route_table.west_public_rt.id
}

```

f) Application Load Balancer (ALB) Setup

I created an ALB in each region to distribute incoming traffic across instances in the Auto Scaling Group. Each ALB is associated with two public subnets in its region, enabling it to handle high availability across different Availability Zones.

The `alb.tf` configuration for the primary ALB in `us-east-1` and the secondary ALB in `us-west-2` is as follows:

```

# Application Load Balancer (ALB) (us-east-1)
resource "aws_lb" "primary_alb" {
    name          = "primary-alb"
    internal      = false
    load_balancer_type = "application"
    security_groups = [aws_security_group.app_sg.id] # Attach the security
group
    subnets       = [aws_subnet.public_subnet_1.id,
aws_subnet.public_subnet_2.id]

    enable_deletion_protection = false
    idle_timeout               = 400

    tags = {
        Name = "PrimaryALB"
    }
}

# ALB Target Group for us-east-1
resource "aws_lb_target_group" "primary" {
    name      = "primary-target-group"
    port      = 80
    protocol = "HTTP"
    vpc_id   = aws_vpc.main_vpc.id
}

```

```

health_check {
  path          = "/"
  interval      = 30
  timeout       = 5
  healthy_threshold = 2
  unhealthy_threshold = 2
}

tags = {
  Name = "PrimaryTargetGroup"
}
}

# ALB Listener for HTTP traffic
resource "aws_lb_listener" "http" {
  load_balancer_arn = aws_lb.primary_alb.arn
  port            = 80
  protocol        = "HTTP"

  default_action {
    type      = "forward"
    target_group_arn = aws_lb_target_group.primary.arn
  }

  tags = {
    Name = "PrimaryALBLListener"
  }
}

# Application Load Balancer (ALB) for us-west-2
resource "aws_lb" "secondary_alb" {
  provider      = aws.west
  name          = "secondary-alb"
  internal      = false
  load_balancer_type = "application"
  security_groups = [aws_security_group.west_app_sg.id]
  subnets        = [aws_subnet.west_public_subnet_1.id,
aws_subnet.west_public_subnet_2.id]

  enable_deletion_protection = false
  idle_timeout               = 400

  tags = {
    Name = "SecondaryALB"
  }
}

# ALB Target Group for us-west-2
resource "aws_lb_target_group" "secondary" {
  provider = aws.west
  name     = "secondary-target-group"
  port     = 80
  protocol = "HTTP"
  vpc_id   = aws_vpc.west_vpc.id

  health_check {
    path          = "/"
    interval      = 30
    timeout       = 5
    healthy_threshold = 2
    unhealthy_threshold = 2
  }
}

```

```

    }

    tags = {
      Name = "SecondaryTargetGroup"
    }
}

# ALB Listener for HTTP traffic in us-west-2
resource "aws_lb_listener" "secondary_http" {
  provider          = aws.west
  load_balancer_arn = aws_lb.secondary_alb.arn
  port              = 80
  protocol          = "HTTP"

  default_action {
    type        = "forward"
    target_group_arn = aws_lb_target_group.secondary.arn
  }

  tags = {
    Name = "SecondaryALBLListener"
  }
}

```

g) Auto Scaling Group (ASG) Update for us-east-1 Region

After saving the `alb.tf` file, I proceeded to update the Auto Scaling Group (ASG) configuration for the `us-east-1` region in the `ec2_autoscaling.tf` file. Here, I added `public_subnet_2` to the `vpc_zone_identifier` to ensure that the ASG is configured to deploy instances across multiple subnets for high availability. Additionally, I included the `target_group_arns` configuration, attaching the ASG to the primary ALB Target Group, and set `health_check_grace_period` and `health_check_type` to use ELB-based health checks.

```

# Auto Scaling Group (ASG) for us-east-1
resource "aws_autoscaling_group" "app_asg" {
  desired_capacity      = 2
  min_size              = 1
  max_size              = 5
  vpc_zone_identifier   = [aws_subnet.public_subnet_1.id,
aws_subnet.public_subnet_2.id] # Attach to ALB Target Group
  target_group_arns     = [aws_lb_target_group.primary.arn]

  health_check_grace_period = 300 # Grace period
  health_check_type         = "ELB" # Change to ELB-based health check

  launch_template {
    id      = aws_launch_template.app.id
    version = "$Latest"
  }

  tag {
    key        = "Name"
    value      = "AppASGInstance"
    propagate_at_launch = true
  }
}

```

```

    lifecycle {
      create_before_destroy = true
    }
}

```

With these changes, the ASG in `us-east-1` is now configured to manage instances across both public subnets (`public_subnet_1` and `public_subnet_2`) and integrated with the Application Load Balancer (ALB) for efficient traffic distribution.

h) Multi-Region Resource Configuration in us-west-2

Following the `us-east-1` configuration, I replicated the necessary configurations for the `us-west-2` region, defining the Security Group, Launch Template, Auto Scaling Group (ASG), and Auto Scaling Policies. This ensures the project's high availability and resilience across both `us-east-1` and `us-west-2` regions.

Security Group for us-west-2

The security group for `us-west-2`, named `west-app-sg`, was created to manage access to instances in this region. This security group includes rules to allow SSH access from a specified IP, as well as HTTP access for web traffic.

```

# Security Group
resource "aws_security_group" "west_app_sg" {
  provider      = aws.west
  name          = "west-app-sg"
  description   = "Allow HTTP and SSH traffic in us-west-2"
  vpc_id        = aws_vpc.west_vpc.id

  ingress {
    from_port    = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["142.59.28.49/32"]  # Restrict SSH to your IP
  }

  ingress {
    from_port    = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["18.237.140.160/29"]  # EC2 Instance Connect IP range
  for us-west-2
  }

  ingress {
    from_port    = 80
    to_port     = 80
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]  # Allow HTTP traffic from anywhere
  }

  egress {
    from_port    = 0
    to_port     = 0
    protocol    = "-1"
  }
}

```

```

        cidr_blocks = [ "0.0.0.0/0" ] # Allow all outbound traffic
    }

    tags = {
        Name = "WestAppSecurityGroup"
    }
}

```

Launch Template for EC2 Instances in us-west-2

The launch template in us-west-2 enables the ASG to provision instances with a pre-defined configuration. This template includes user data to install and start NGINX upon instance initialization, ensuring each instance is ready to serve traffic.

```

# Launch Template for EC2 Instances in us-west-2
resource "aws_launch_template" "secondary_app" {
    provider      = aws.west
    name_prefix   = "secondary-app-launch-template"
    image_id      = "ami-07c5ecd84908c59d5" # Amazon Linux AMI
    instance_type = "t2.micro"

    # Attachment of the IAM role to EC2 instances in us-west-2
    iam_instance_profile {
        name = aws_iam_instance_profile.ec2_instance_profile.name
    }

    network_interfaces {
        security_groups = [aws_security_group.west_app_sg.id]
    }

    user_data = base64encode(<<-EOF
        #!/bin/bash
        sudo dnf update -y
        sudo dnf install nginx -y
        sudo systemctl start nginx
        sudo systemctl enable nginx
    EOF
    )

    lifecycle {
        create_before_destroy = true
    }

    tag_specifications {
        resource_type = "instance"
        tags = {
            Name = "SecondaryAppInstance"
        }
    }
}

```

Auto Scaling Group (ASG) for us-west-2

For high availability, I created an ASG in us-west-2 to manage instance deployment across public subnets and attach to the secondary ALB target group.

```
# Auto Scaling Group (ASG) for us-west-2
resource "aws_autoscaling_group" "secondary_asg" {
```

```

provider          = aws.west
desired_capacity = 2
min_size         = 1
max_size         = 5
vpc_zone_identifier = [aws_subnet.west_public_subnet_1.id,
aws_subnet.west_public_subnet_2.id] # Attach to ALB Target Group
target_group_arns = [aws_lb_target_group.secondary.arn]

health_check_grace_period = 300 # Grace period
health_check_type         = "ELB" # Use ELB-based health check for ALB

launch_template {
  id      = aws_launch_template.secondary_app.id
  version = "$latest"
}

tag {
  key        = "Name"
  value      = "SecondaryAppASGInstance"
  propagate_at_launch = true
}

lifecycle {
  create_before_destroy = true
}
}

```

Auto Scaling Policies for us-west-2

To enable dynamic scaling based on demand, I created scale-up and scale-down policies for us-west-2. These policies adjust the ASG capacity by one instance based on the load.

```

# Auto Scaling Policies

# Scaling Up Policy for us-west-2
resource "aws_autoscaling_policy" "west_scale_up" {
  provider          = aws.west
  name              = "west-scale-up"
  scaling_adjustment = 1 # Adds 1 instance when triggered
  adjustment_type   = "ChangeInCapacity"
  autoscaling_group_name = aws_autoscaling_group.secondary_asg.name
}

# Scaling Down Policy for us-west-2
resource "aws_autoscaling_policy" "west_scale_down" {
  provider          = aws.west
  name              = "west-scale-down"
  scaling_adjustment = -1 # Removes 1 instance when triggered
  adjustment_type   = "ChangeInCapacity"
  autoscaling_group_name = aws_autoscaling_group.secondary_asg.name
}

```

i) IAM Role Configuration and Permissions Update

After saving the `ec2_autoscaling.tf` file, I updated the `iam_roles.tf` file to ensure that EC2 instances had the necessary permissions to interact with AWS services required for this multi-region deployment.

```

# IAM Role for EC2 to assume
resource "aws_iam_role" "ec2_role" {
  name = "ec2-role"

  assume_role_policy = <<EOF
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "ec2.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
EOF
}

# IAM instance profile for EC2
resource "aws_iam_instance_profile" "ec2_instance_profile" {
  name = "ec2-instance-profile"
  role = aws_iam_role.ec2_role.name
}

# Attach required permissions
resource "aws_iam_role_policy_attachment" "ec2_full_access" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonEC2FullAccess"
}

resource "aws_iam_role_policy_attachment" "autoscaling_full_access" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AutoScalingFullAccess"
}

resource "aws_iam_role_policy_attachment" "elb_full_access" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/ElasticLoadBalancingFullAccess"
}

# Attachment of S3 read-only policy to the role
resource "aws_iam_role_policy_attachment" "ec2_s3_READONLY" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess"
}

```

j) Route 53 Configuration for Failover

To implement failover between the two regions, I created `route53.tf`, which defines the Route 53 hosted zone and configures failover routing between the primary (us-east-1) and secondary (us-west-2) regions. I set up health checks for each ALB to monitor availability and automatically switch traffic to the secondary region if needed.

```

# Route 53 Hosted Zone
resource "aws_route53_zone" "main" {
  name = "myinfra.local" # A fake domain for testing purposes
}

```

```

# Primary Route 53 Record (us-east-1)
resource "aws_route53_record" "primary" {
  zone_id = aws_route53_zone.main.zone_id
  name    = "app.myinfra.local"
  type    = "A"

  alias {
    name          = aws_lb.primary_alb.dns_name
    zone_id       = aws_lb.primary_alb.zone_id
    evaluate_target_health = true
  }

  set_identifier      = "Primary"
  failover_routing_policy {
    type = "PRIMARY"
  }
}

# Secondary Route 53 Record (us-west-2)
resource "aws_route53_record" "secondary" {
  zone_id = aws_route53_zone.main.zone_id
  name    = "app.myinfra.local"
  type    = "A"

  alias {
    name          = aws_lb.secondary_alb.dns_name
    zone_id       = aws_lb.secondary_alb.zone_id
    evaluate_target_health = true
  }

  set_identifier      = "Secondary"
  failover_routing_policy {
    type = "SECONDARY"
  }
}

# Route 53 Health Checks

# Primary Health Check (us-east-1)
resource "aws_route53_health_check" "primary_health_check" {
  fqdn      = aws_lb.primary_alb.dns_name # Use ALB DNS Name
  type      = "HTTP"
  port      = 80
  request_interval = 30
  failure_threshold = 3
}

# Secondary Health Check (us-west-2)
resource "aws_route53_health_check" "secondary_health_check" {
  fqdn      = aws_lb.secondary_alb.dns_name # Use ALB DNS Name
  type      = "HTTP"
  port      = 80
  request_interval = 30
  failure_threshold = 3
}

```

k) Verification and Testing of Resources

After saving `route53.tf`, I ran the following commands to verify the deployment:

```
terraform init  
terraform plan  
terraform apply -auto-approve
```

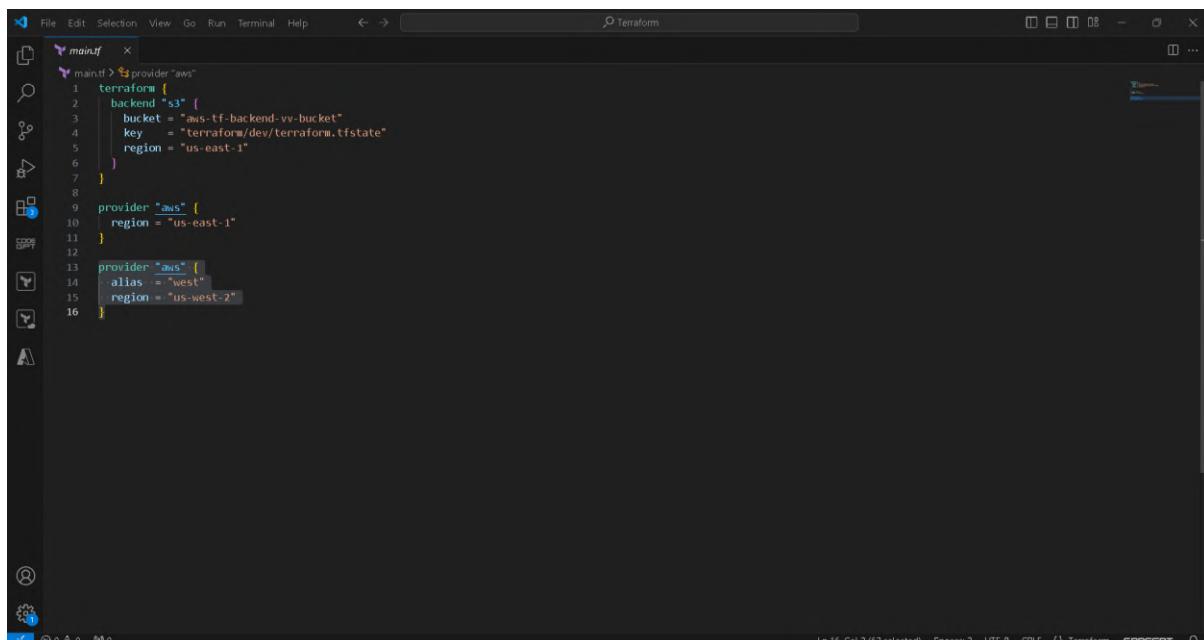
The plan showed 33 resources to add, confirming the configurations for this phase. After successful deployment, I checked the AWS Console to verify each of the resources as expected:

- **VPCs and Subnets:** Verified the `WestVPC`, `WestPublicSubnet1`, `WestPublicSubnet2`, `WestPrivateSubnet1`, and the route tables configured in both regions.
- **Internet Gateways and Route Tables:** Confirmed the presence and configuration of `WestIGW` and associated public route tables.
- **Security Groups:** Verified `WestAppSecurityGroup` and its inbound/outbound rules.
- **Instances:** Observed that 2 instances were running in `us-west-2` as expected.
- **Application Load Balancers:** Verified both `primary-alb` and `secondary-alb`, including listeners, target groups, and health checks.
- **Route 53:** Confirmed that the `myinfra.local` Route 53 hosted zone was created with failover routing configured between `us-east-1` and `us-west-2`, with healthy status for both health checks.

With all 33 resources created and verified as expected, this phase of the project was successfully completed.

This completes the multi-region deployment and failover setup, ensuring that the application can seamlessly handle failover between `us-east-1` and `us-west-2` in the event of an outage in the primary region.

Screenshots



```
main.tf  
provider "aws"  
  backend "s3" {  
    bucket = "aws-tf-backend-vv-bucket"  
    key    = "terraform/dev/terraform.tfstate"  
    region = "us-east-1"  
  }  
provider "aws" {  
  region = "us-east-1"  
}  
provider "aws" {  
  alias  = "west"  
  region = "us-west-2"  
}
```

```
network.tf
1 # VPC for us-east-1
2
3 resource "aws_vpc" "main_vpc" {
4   cidr_block = "10.0.0.0/16" # VPC with a /16 subnet range
5   tags = [
6     { Name = "MainVPC" }
7   ]
8 }
9
10 # Public Subnet in us-east-1a
11
12 resource "aws_subnet" "public_subnet_1" {
13   vpc_id          = aws_vpc.main_vpc.id
14   cidr_block      = "10.0.1.0/24" # Public subnet's CIDR block
15   availability_zone = "us-east-1a" # Availability zone
16   map_public_ip_on_launch = true # Ensures EC2 instances in this subnet get a public IP
17   tags = [
18     { Name = "PublicSubnet1" }
19   ]
20 }
21
22 # Public Subnet in us-east-1b
23
24 resource "aws_subnet" "public_subnet_2" {
25   vpc_id          = aws_vpc.main_vpc.id
26   cidr_block      = "10.0.4.0/24"
27   availability_zone = "us-east-1b"
28   map_public_ip_on_launch = true
29   tags = [
30     { Name = "PublicSubnet2" }
31   ]
32 }
33
34
35 # Private Subnet in us-east-1a
36
37 resource "aws_subnet" "private_subnet_1" {
```

```
network.tf
12 resource "aws_subnet" "public_subnet_1" {
13   availability_zone = "us-east-1a" # Availability zone
14   map_public_ip_on_launch = true # Ensures EC2 instances in this subnet get a public IP
15   tags = [
16     { Name = "PublicSubnet1" }
17   ]
18 }
19
20
21 # Public Subnet in us-east-1b
22
23 resource "aws_subnet" "public_subnet_2" {
24   vpc_id          = aws_vpc.main_vpc.id
25   cidr_block      = "10.0.4.0/24"
26   availability_zone = "us-east-1b"
27   map_public_ip_on_launch = true
28   tags = [
29     { Name = "PublicSubnet2" }
30   ]
31 }
32
33
34
35 # Private Subnet in us-east-1a
36
37 resource "aws_subnet" "private_subnet_1" {
38   vpc_id          = aws_vpc.main_vpc.id
39   cidr_block      = "10.0.3.0/24" # Private subnet's CIDR block
40   availability_zone = "us-east-1a" # Same AZ as the public subnet
41   map_public_ip_on_launch = false # Instances in this subnet will not have public IPs
42   tags = [
43     { Name = "PrivateSubnet1" }
44   ]
45 }
```

```
network.tf
51 # VPC for us-west-2
52
53 resource "aws_vpc" "west_vpc" {
54   provider    = aws.west
55   cidr_block  = "10.1.0.0/16"
56
57   tags = [
58     { Name = "WestVPC" }
59   ]
60 }
61
62 # Public Subnet in us-west-2a
63
64 resource "aws_subnet" "west_public_subnet_1" {
65   provider    = aws.west
66   vpc_id      = aws_vpc.west_vpc.id
67   cidr_block  = "10.1.1.0/24"
68   availability_zone = "us-west-2a"
69   map_public_ip_on_launch = true
70
71   tags = [
72     { Name = "WestPublicSubnet1" }
73   ]
74 }
75
76
77 # Public Subnet in us-west-2b
78 resource "aws_subnet" "west_public_subnet_2" {
79   provider    = aws.west
80   vpc_id      = aws_vpc.west_vpc.id
81   cidr_block  = "10.1.2.0/24"
82   availability_zone = "us-west-2b"
83   map_public_ip_on_launch = true
84   tags = [
85     { Name = "WestPublicSubnet2" }
86   ]
87 }
```

```
network.tf
64 resource "aws_subnet" "west_private_subnet_1" {
65   cidr_block          = "10.1.1.0/24"
66   availability_zone   = "us-west-2a"
67   map_public_ip_on_launch = true
68
69   tags = [
70     { Name = "WestPrivateSubnet1" }
71   ]
72 }
73
74 }
75
76
77 # Public Subnet in us-west-2b
78 resource "aws_subnet" "west_public_subnet_2" {
79   provider           = aws.west
80   vpc_id             = aws_vpc.west_vpc.id
81   cidr_block         = "10.1.2.0/24"
82   availability_zone = "us-west-2b"
83   map_public_ip_on_launch = true
84
85   tags = [
86     { Name = "WestPublicSubnet2" }
87   ]
88
89
90 # Private Subnet in us-west-2a
91 resource "aws_subnet" "west_private_subnet_1" {
92   provider           = aws.west
93   vpc_id             = aws_vpc.west_vpc.id
94   cidr_block         = "10.1.3.0/24"
95   availability_zone = "us-west-2a"
96
97   tags = [
98     { Name = "WestPrivateSubnet1" }
99   ]
100 }
101 }
```

```
network.tf
105 # Internet Gateway (IGW)
106
107 resource "aws_internet_gateway" "igw" {
108   vpc_id = aws_vpc.main_vpc.id
109
110   tags = [
111     { Name = "MainIGM" }
112   ]
113
114 # Public Route Table and Route
115
116 resource "aws_route_table" "public_rt" {
117   vpc_id = aws_vpc.main_vpc.id
118
119   tags = [
120     { Name = "PublicRouteTable" }
121   ]
122
123 # Associate the public subnets with the route table in us-east-1
124
125 resource "aws_route" "internet_access" {
126   route_table_id      = aws_route_table.public_rt.id
127   destination_cidr_block = "0.0.0.0/0" # sends traffic to the internet
128   gateway_id          = aws_internet_gateway.igw.id
129 }
130
131 resource "aws_route_table_association" "public_subnet_assoc_1" {
132   subnet_id           = aws_subnet.public_subnet_1.id
133   route_table_id      = aws_route_table.public_rt.id
134 }
135
136 resource "aws_route_table_association" "public_subnet_assoc_2" {
137   subnet_id           = aws_subnet.public_subnet_2.id
138   route_table_id      = aws_route_table.public_rt.id
139 }
```

```
network.tf
135
136 resource "aws_route_table_association" "public_subnet_assoc_2" {
137   subnet_id           = aws_subnet.public_subnet_2.id
138   route_table_id      = aws_route_table.public_rt.id
139 }
140
141
142
143
144
145
146 # Internet Gateway (IGW) for us-west-2
147
148 resource "aws_internet_gateway" "west_igw" {
149   provider           = aws.west
150   vpc_id             = aws_vpc.west_vpc.id
151
152   tags = [
153     { Name = "WestIGM" }
154   ]
155
156 # Public Route Table and Route for us-west-2
157
158 resource "aws_route_table" "west_public_rt" {
159   provider           = aws.west
160   vpc_id             = aws_vpc.west_vpc.id
161
162   tags = [
163     { Name = "WestPublicRouteTable" }
164   ]
165
166 resource "aws_route" "west_internet_access" {
167   provider           = aws.west
168   route_table_id     = aws_route_table.west_public_rt.id
169   destination_cidr_block = "0.0.0.0/0" # sends traffic to the internet
170   gateway_id          = aws_internet_gateway.west_igw.id
171 }
```

```
network.tf
158 resource "aws_route_table" "west_public_rt" {
159   route_table_id = aws_route_table.west_public_rt.id
160 }
161 }
162 }
163 }
164 }
165
166 resource "aws_route" "west_internet_access" {
167   provider      = aws.west
168   route_table_id = aws_route_table.west_public_rt.id
169   destination_cidr_block = "0.0.0.0/0" # sends traffic to the internet
170   gateway_id    = aws_internet_gateway.west_igw.id
171 }
172
173 # Associate the public subnets with the route table in us-west-2
174
175 resource "aws_route_table_association" "west_public_subnet_assoc_1" {
176   provider      = aws.west
177   subnet_id     = aws_subnet.west_public_subnet_1.id
178   route_table_id = aws_route_table.west_public_rt.id
179 }
180
181 resource "aws_route_table_association" "west_public_subnet_assoc_2" {
182   provider      = aws.west
183   subnet_id     = aws_subnet.west_public_subnet_2.id
184   route_table_id = aws_route_table.west_public_rt.id
185 }
```

In 184, Col 54 (495 selected) · Spaces: 2 · UTF-8 · CRLF · {} · Terraform · CODEGPT

```
alb.tf
1 # Application Load Balancer (ALB) (us-east-1)
2
3 resource "aws_lb" "primary_alb" {
4   name          = "primary-alb"
5   internal      = false
6   load_balancer_type = "application"
7   security_groups = [aws_security_group.app_sg.id] # Attach the security group
8   subnets        = [aws_subnet.public_subnet_1.id, aws_subnet.public_subnet_2.id]
9
10  enable_deletion_protection = false
11  idle_timeout               = 400
12
13  tags = [
14    { Name = "PrimaryALB" }
15  ]
16 }
17
18 # ALB Target Group for us-east-1
19
20 resource "aws_lb_target_group" "primary" {
21   name          = "primary-target-group"
22   port          = 80
23   protocol      = "HTTP"
24   vpc_id        = aws_vpc.main_vpc.id
25
26   health_check {
27     path          = "/"
28     interval     = 30
29     timeout      = 5
30     healthy_threshold = 2
31     unhealthy_threshold = 2
32   }
33
34   tags = [
35     { Name = "PrimaryTargetGroup" }
36   ]
37 }
```

In 69, Col 97 · Spaces: 2 · UTF-8 · CRLF · {} · Terraform · CODEGPT

```
alb.tf
38
39 # ALB Listener for HTTP traffic
40
41 resource "aws_lb_listener" "http" {
42   load_balancer_arn = aws_lb.primary_alb.arn
43   port              = 80
44   protocol          = "HTTP"
45
46   default_action {
47     type        = "forward"
48     target_group_arn = aws_lb_target_group.primary.arn
49   }
50
51   tags = [
52     { Name = "PrimaryALBListener" }
53   ]
54 }
55
56
57
58
59
60
61 # Application Load Balancer (ALB) for us-west-2
62
63 resource "aws_lb" "secondary_alb" {
64   provider      = aws.west
65   name          = "secondary-alb"
66   internal      = false
67   load_balancer_type = "application"
68   security_groups = [aws_security_group.west_app_sg.id]
69   subnets        = [aws_subnet.west_public_subnet_1.id, aws_subnet.west_public_subnet_2.id]
70
71   enable_deletion_protection = false
72   idle_timeout               = 400
73
74   tags = [
```

In 69, Col 97 · Spaces: 2 · UTF-8 · CRLF · {} · Terraform · CODEGPT

```
network.tf
alb.tf

61 # Application Load Balancer (ALB) for us-west-2
62
63 resource "aws_lb" "secondary_alb" {
64   provider = aws.west
65   name     = "secondary-alb"
66   internal = false
67   load_balancer_type = "application"
68   security_groups = [aws_security_group.west_app_sg.id]
69   subnets      = [aws_subnet.west_public_subnet_1.id, aws_subnet.west_public_subnet_2.id]
70
71   enable_deletion_protection = false
72   idle_timeout               = 400
73
74   tags = {
75     Name = "SecondaryALB"
76   }
77 }
78
79 # ALB Target Group for us-west-2
80
81 resource "aws_lb_target_group" "secondary" {
82   provider = aws.west
83   name     = "secondary-target-group"
84   port     = 80
85   protocol = "HTTP"
86   vpc_id   = aws_vpc.west_vpc.id
87
88   health_check {
89     path      = "/"
90     interval = 30
91     timeout   = 5
92     healthy_threshold = 2
93     unhealthy_threshold = 2
94   }
95
96   tags = {
97     Name = "SecondaryTargetGroup"
98   }
99 }
```

Ln 69, Col 97 Spaces: 2 UTF-8 CRLF {} Terraform CODEGPT

```
network.tf
alb.tf

81 resource "aws_lb_target_group" "secondary" {
82   tags = {
83     Name = "SecondaryTargetGroup"
84   }
85 }
86
87 # ALB Listener for HTTP traffic in us-west-2
88
89 resource "aws_lb_listener" "secondary_http" {
90   provider = aws.west
91   load_balancer_arn = aws_lb.secondary_alb.arn
92   port     = 80
93   protocol = "HTTP"
94
95   default_action {
96     type    = "forward"
97     target_group_arn = aws_lb_target_group.secondary.arn
98   }
99
100  tags = {
101    Name = "SecondaryALBLListener"
102  }
103 }
```

Ln 69, Col 97 Spaces: 2 UTF-8 CRLF {} Terraform CODEGPT

```
ec2_autoscaling.tf

85
86
87 # Auto Scaling Group (ASG) for us-east-1
88
89 resource "aws_autoscaling_group" "app_asg" {
90   desired_capacity    = 2
91   min_size            = 1
92   max_size            = 5
93   vpc_zone_identifier = [aws_subnet.public_subnet_1.id, aws_subnet.public_subnet_2.id]
94   target_group_arns... = [aws_lb_target_group.primary.arn] # Attach to ALB Target Group
95
96   health_check_grace_period = 300 # Grace period
97   health_check_type        = "ELB" # Change to ELB-based health check
98
99   launch_template {
100     id      = aws_launch_template.app.id
101     version = "$latest"
102   }
103
104   tag {
105     key      = "Name"
106     value    = "AppASGInstance"
107     propagate_at_launch = true
108   }
109
110   lifecycle {
111     create_before_destroy = true
112   }
113 }
114
115
116
117
118
119 # Auto Scaling Policies
120
121 }
```

Ln 97, Col 72 (302 selected) Spaces: 2 UTF-8 CRLF {} Terraform CODEGPT

```
File Edit Selection View Go Run Terminal Help ⏎ ⏎ Terraform

ec2_autoscaling.tf ×
└── e2e_ec2_autoscalingtf > ...

144 # For us-west-2:
145
146 # Security Group
147
148 resource "aws_security_group" "west_app_sg" {
149   provider = aws.west
150   name     = "west-app-sg"
151   description = "Allow HTTP and SSH traffic in us-west-2"
152   vpc_id   = aws_vpc.west_vpc.id
153
154   ingress {
155     from_port  = 22
156     to_port    = 22
157     protocol   = "tcp"
158     cidr_blocks = ["142.59.28.49/32"] # Restrict SSH to your IP
159   }
160
161   ingress {
162     from_port  = 22
163     to_port    = 22
164     protocol   = "tcp"
165     cidr_blocks = ["18.237.140.160/29"] # EC2 Instance Connect IP range for us-west-2
166   }
167
168   ingress {
169     from_port  = 80
170     to_port    = 80
171     protocol   = "tcp"
172     cidr_blocks = ["0.0.0.0/0"] # Allow HTTP traffic from anywhere
173   }
174
175   egress {
176     from_port  = 0
177     to_port    = 0
178     protocol   = "-1"
179     cidr_blocks = ["0.0.0.0/0"] # Allow all outbound traffic
180   }
}

Ln 144, Col 17 (16 selected) Spaces: 2 UTF-8 CRLF {} Terraform CODEGPT
```

```
File Edit Selection View Go Run Terminal Help ⏎ ⏎ Terraform

ec2_autoscalingtf ×
└── e2e_ec2_autoscalingtf > ...

148 resource "aws_security_group" "west_app_sg" {
149   tags = {
150     Name = "WestAppSecurityGroup"
151   }
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189 # Launch Template for EC2 Instances in us-west-2
190
191
192 resource "aws_launch_template" "secondary_app" {
193   provider      = aws.west
194   name_prefix   = "secondary-app-launch-template"
195   image_id      = "ami-07c5ecd8498c59db5" # Amazon Linux AMI
196   instance_type = "t2.micro"
197
198 # Attachment of the IAM role to EC2 instances in us-west-2
199
200   iam_instance_profile {
201     name = aws_iam_instance_profile.ec2_instance_profile.name
202   }
203
204   network_interfaces {
205     security_groups = [aws_security_group.west_app_sg.id]
206   }
207
208   user_data = base64encode(<<-EOF
209     #!/bin/bash
210     sudo dnf update -y
211     sudo dnf install nginx -y
212     sudo systemctl start nginx
213     sudo systemctl enable nginx
214   EOF
215 )
216
217   lifecycle {
218     create_before_destroy = true
219   }
220
221   tag_specifications {
222     resource_type = "instance"
223     tags = [
224       {name = "SecondaryAppInstance"}
225     ]
226   }
}

Ln 144, Col 17 (16 selected) Spaces: 2 UTF-8 CRLF {} Terraform CODEGPT
```

```
File Edit Selection View Go Run Terminal Help ⏎ ⏎ Terraform

ec2_autoscalingtf ×
└── e2e_ec2_autoscalingtf > ...

189 # Launch Template for EC2 Instances in us-west-2
190
191
192 resource "aws_launch_template" "secondary_app" {
193   provider      = aws.west
194   name_prefix   = "secondary-app-launch-template"
195   image_id      = "ami-07c5ecd8498c59db5" # Amazon Linux AMI
196   instance_type = "t2.micro"
197
198 # Attachment of the IAM role to EC2 instances in us-west-2
199
200   iam_instance_profile {
201     name = aws_iam_instance_profile.ec2_instance_profile.name
202   }
203
204   network_interfaces {
205     security_groups = [aws_security_group.west_app_sg.id]
206   }
207
208   user_data = base64encode(<<-EOF
209     #!/bin/bash
210     sudo dnf update -y
211     sudo dnf install nginx -y
212     sudo systemctl start nginx
213     sudo systemctl enable nginx
214   EOF
215 )
216
217   lifecycle {
218     create_before_destroy = true
219   }
220
221   tag_specifications {
222     resource_type = "instance"
223     tags = [
224       {name = "SecondaryAppInstance"}
225     ]
226   }
}

Ln 144, Col 17 (16 selected) Spaces: 2 UTF-8 CRLF {} Terraform CODEGPT
```

The screenshot shows a code editor with a dark theme, displaying a Terraform configuration file. The file defines two scaling policies for an Auto Scaling group named 'secondary_asg'. The first policy, 'west_scale_up', adds one instance when triggered by a change in capacity. The second policy, 'west_scale_down', removes one instance when triggered by a change in capacity. Both policies are associated with the 'aws.west' provider.

```
File Edit Selection View Go Run Terminal Help ⏴ ⏵ ⏴ Terraform
```

```
ec2_autoscalingtf x
ec2_autoscalingtf ... 258
259
260 # Auto Scaling Policies
261
262
263 # Scaling Up Policy for us-west-2
264
265 resource "aws_autoscaling_policy" "west_scale_up" {
266   provider      = aws.west
267   name          = "west-scale-up"
268   scaling_adjustment = 1 # Adds 1 instance when triggered
269   adjustment_type = "ChangeInCapacity"
270   autoscaling_group_name = aws_autoscaling_group.secondary_asg.name
271 }
272
273
274 # Scaling Down Policy for us-west-2
275
276 resource "aws_autoscaling_policy" "west_scale_down" {
277   provider      = aws.west
278   name          = "west-scale-down"
279   scaling_adjustment = -1 # Removes 1 instance when triggered
280   adjustment_type = "ChangeInCapacity"
281   autoscaling_group_name = aws_autoscaling_group.secondary_asg.name
282 }
```

```
File Edit Selection View Go Run Terminal Help ⏴ ⏵ Terraform
```

```
iam_roles.tf
iam_roles > ...
1 # IAM Role for EC2 to assume
2 resource "aws_iam_role" "ec2_role" {
3   name = "ec2-role"
4
5   assume_role_policy = <<EOF
6   [
7     "Version": "2012-10-17",
8     "Statement": [
9       {
10         "Effect": "Allow",
11         "Principal": "*",
12         "Service": "ec2.amazonaws.com"
13       },
14       "Action": "sts:AssumeRole"
15     ]
16   ]
17 }
18 EOF
19 }
20
21 # IAM instance profile for EC2
22 resource "aws_iam_instance_profile" "ec2_instance_profile" {
23   name = "ec2-instance-profile"
24   role = aws_iam_role.ec2_role.name
25 }
26
27 # Attach required permissions
28 resource "aws_iam_role_policy_attachment" "ec2_full_access" {
29   role      = aws_iam_role.ec2_role.name
30   policyArn = "arn:aws:iam::aws:policy/AmazonEC2FullAccess"
31 }
32
33 resource "aws_iam_role_policy_attachment" "autoscaling_full_access" {
34   role      = aws_iam_role.ec2_role.name
35   policyArn = "arn:aws:iam::aws:policy/AutoScalingFullAccess"
36 }
37
```

A screenshot of a code editor displaying a Terraform configuration file named `iam_roles.tf`. The code defines an IAM role named `ec2_role` with various permissions attached, including `elb_full_access`, `ec2_full_access`, `autoscaling_full_access`, and `elb_full_access`. It also creates an instance profile named `ec2_instance_profile` and attaches it to the role.

```
resource "aws_iam_role" "ec2_role" {
  # ...
}

resource "aws_iam_instance_profile" "ec2_instance_profile" {
  name = "ec2-instance-profile"
  role = aws_iam_role.ec2_role.name
}

resource "aws_iam_role_policy_attachment" "elb_full_access" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonEC2FullAccess"
}

resource "aws_iam_role_policy_attachment" "autoscaling_full_access" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AutoScalingFullAccess"
}

resource "aws_iam_role_policy_attachment" "elb_full_access" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/ElasticLoadBalancingFullAccess"
}

resource "aws_iam_role_policy_attachment" "ec2_s3_readonly" {
  role      = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess"
}
```

A screenshot of a code editor displaying a Terraform configuration file named `route53.tf`. The code creates a main Route 53 zone named `myinfra.local` and two Route 53 records: a primary record `app.myinfra.local` and a secondary record `app.myinfra.local`. Both records point to an ALB named `aws_lb`.

```
resource "aws_route53_zone" "main" {
  name = "myinfra.local" # A fake domain for testing purposes
}

resource "aws_route53_record" "primary" {
  zone_id = aws_route53_zone.main.zone_id
  name   = "app.myinfra.local"
  type   = "A"

  alias {
    name           = aws_lb.primary.alb.dns_name
    zone_id        = aws_lb.primary.alb.zone_id
    evaluate_target_health = true
  }

  set_identifier = "Primary"
  failover_routing_policy {
    type = "PRIMARY"
  }
}

resource "aws_route53_record" "secondary" {
  zone_id = aws_route53_zone.main.zone_id
  name   = "app.myinfra.local"
  type   = "A"

  alias {
    name           = aws_lb.secondary.alb.dns_name
    zone_id        = aws_lb.secondary.alb.zone_id
    evaluate_target_health = true
  }

  set_identifier = "Secondary"
  failover_routing_policy {
    type = "SECONDARY"
  }
}
```

A screenshot of a code editor displaying a Terraform configuration file named `route53.tf`. This version of the code includes detailed annotations for the `alias` block, specifically for the `name` and `zone_id` fields, and the `evaluate_target_health` field. The code structure is identical to the previous screenshot but with additional annotations.

```
resource "aws_route53_record" "primary" {
  type   = "A"

  alias {
    name           = aws_lb.primary.alb.dns_name
    zone_id        = aws_lb.primary.alb.zone_id
    evaluate_target_health = true
  }

  set_identifier = "Primary"
  failover_routing_policy {
    type = "PRIMARY"
  }
}

resource "aws_route53_record" "secondary" {
  zone_id = aws_route53_zone.main.zone_id
  name   = "app.myinfra.local"
  type   = "A"

  alias {
    name           = aws_lb.secondary.alb.dns_name
    zone_id        = aws_lb.secondary.alb.zone_id
    evaluate_target_health = true
  }

  set_identifier = "Secondary"
  failover_routing_policy {
    type = "SECONDARY"
  }
}
```

The screenshot shows a code editor with a dark theme, displaying a Terraform configuration file named `route53.tf`. The file contains resources for AWS Route 53 and AWS Lambda.

```
resource "aws_route53_record" "secondary" {
  31  failover_routing_policy {
  32    43
  33    44
  34    45
  35    46
  36    47
  37    48
  38    49
  39    50
  40    51
  41    52  # Route 53 Health Checks
  42    53
  43    54
  44    55
  45    56  # Primary Health Check (us-east-1)
  46    57
  47    58  resource "aws_route53_health_check" "primary_health_check" {
  48      59    fqdn      = aws_lb.primary_alb.dns_name # Use ALB DNS Name
  49      60    type      = "HTTP"
  50      61    port      = 80
  51      62    request_interval = 30
  52      63    failure_threshold = 3
  53    64  }
  54    65
  55    66  # Secondary Health Check (us-west-2)
  56    67
  57    68  resource "aws_route53_health_check" "secondary_health_check" {
  58      69    fqdn      = aws_lb.secondary_alb.dns_name # Use ALB DNS Name
  59      70    type      = "HTTP"
  60      71    port      = 80
  61      72    request_interval = 30
  62      73    failure_threshold = 3
  63    74  }
  64
  65
  66
  67
  68
  69
  70
  71
  72
  73
  74
  75}
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
aws_route_table_association.public_subnet_assoc: Refreshing state... [id=rtbassoc-080529014332fac87]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b90f0f753dc]
aws_iam_role.ec2_role: Refreshing state... [id=ec2-role]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_iam_role_policy_attachment.ec2_s3readonly: Refreshing state... [id=ec2-role-20241026083002429100000001]
aws_iam_role_policy_attachment.ec2_full_access: Refreshing state... [id=ec2-role-2024103009417037793000000003]
aws_iam_role_policy_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-20241030041703774400000002]
aws_iam_role_policy_attachment.elb_full_access: Refreshing state... [id=ec2-role-202410300941703813700000004]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_s3_bucket_public_access_block.backup_bucket_block: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bd799bb28a4a]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-074ccceeed1aaaf778f]
aws_route_table_public_rt: Refreshing state... [id=rtb-0c4fc0f97fe205e5]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-04b29a39893e79544d]
aws_route.internet_access: Refreshing state... [id=r-rtb-0c4fc0f97fe205e51080289494]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create
- destroy

Terraform will perform the following actions:
```

```
Command Prompt + - x

+ owner_id = (known after apply)
+ private_dns_hostname_type_on_launch = (known after apply)
+ tags = {
    + "Name" = "WestPublicSubnet2"
}
+ tags_all = {
    + "Name" = "WestPublicSubnet2"
}
+ vpc_id = (known after apply)
}

# aws_vpc.west_vpc will be created
+ resource "aws_vpc" "west_vpc" {
    + arn = (known after apply)
    + cidr_block = "10.1.0.0/16"
    + default_network_acl_id = (known after apply)
    + default_route_table_id = (known after apply)
    + default_security_group_id = (known after apply)
    + dhcp_options_id = (known after apply)
    + enable_dns_hostnames = (known after apply)
    + enable_dns_support = true
    + enable_network_address_usage_metrics = (known after apply)
    + id = (known after apply)
    + instance_tenancy = "default"
    + ipv6_association_id = (known after apply)
    + ipv6_cidr_block = (known after apply)
    + ipv6_cidr_block_network_border_group = (known after apply)
    + main_route_table_id = (known after apply)
    + owner_id = (known after apply)
    + tags = {
        + "Name" = "WestVPC"
    }
    + tags_all = {
        + "Name" = "WestVPC"
    }
}

Plan: 33 to add, 0 to change, 1 to destroy.
```

```
Command Prompt + - o x
aws_vpc.west_vpc: Creating...
aws_route_table_association.public_subnet_assoc: Destroying... [id=rtbassoc-080529014332fac87]
aws_route_table_association.public_subnet_assoc_1: Creating...
aws_route53_zone.main: Creating...
aws_subnet.public_subnet_2: Creating...
aws_lb_target_group.primary: Creating...
aws_security_group.app_sg: Creating...
aws_route_table_association.public_subnet_assoc: Destruction complete after 1s
aws_vpc.west_vpc: Creation complete after 1s [id=rtbassoc-062c3dc339ec6eb03]
aws_internet_gateway.west_igw: Creating...
aws_subnet.west_public_subnet_2: Creating...
aws_route_table.west_public_rt: Creating...
aws_lb_target_group.secondary: Creating...
aws_subnet.west_public_subnet_1: Creating...
aws_security_group.west_app_sg: Creating...
aws_lb_target_group.primary: Creation complete after 2s [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:targetgroup/primary-target-group/a9447d338de6a3e5]
aws_subnet.west_private_subnet_1: Creating...
aws_internet_gateway.west_igw: Creation complete after 1s [id=igw-093c80565bfdb413c]
aws_route_table.west_public_rt: Creation complete after 1s [id=rtb-0bf4cfa31f877e69]
aws_route.west.internet.access: Creating...
aws_subnet.west_private_subnet_1: Creation complete after 0s [id=subnet-0de0b39d268f143ac]
aws_lb_target_group.secondary: Creation complete after 2s [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba401c8967]
aws_route.west.internet.access: Creation complete after 1s [id=rtb-0bf4cfa31f877e6910880289494]
aws_security_group.app_sg: Creation complete after 3s [id=sg-002a28cc779899cb]
aws_launch_template.app: Creating...
aws_launch_template.app: Creation complete after 1s [id=lt-074a269bf44eba96]
aws_security_group.west_app_sg: Creation complete after 3s [id=sg-03d04109977fd80d0]
aws_launch_template.secondary.app: Creating...
aws_launch_template.secondary.app: Creation complete after 0s [id=lt-0c625394b30f2746d]
aws_subnet.public_subnet_2: Still creating... [10s elapsed]
aws_route53_zone.main: Still creating... [10s elapsed]
aws_subnet.west_public_subnet_2: Still creating... [10s elapsed]
aws_subnet.west_public_subnet_1: Still creating... [10s elapsed]
aws_subnet.public_subnet_2: Creation complete after 1s [id=subnet-0d89362f71d4a5adf]
aws_lb.primary.alb: Creating...
aws_autoscaling_group.app_asg: Creating...
aws_route_table_association.public_subnet_assoc_2: Creation complete after 1s [id=rtbassoc-0c1e016c8860f95a6]
```

```
Command Prompt + - o x
aws_autoscaling_group.app_asg: Creating...
aws_route_table_association.public_subnet_assoc_2: Creation complete after 1s [id=rtbassoc-0c1e016c8860f95a6]
aws_subnet.west_public_subnet_1: Creation complete after 12s [id=subnet-0a1940534f9cb413c]
aws_route_table_association.west_public_subnet_assoc_1: Creating...
aws_subnet.west_public_subnet_2: Creation complete after 12s [id=subnet-05cf33fc2aa1e38da]
aws_route_table_association.west_public_subnet_assoc_2: Creating...
aws_lb.secondary.alb: Creating...
aws_autoscaling_group.secondary_asg: Creating...
aws_route_table_association.west_public_subnet_assoc_1: Creation complete after 0s [id=rtbassoc-05dc480a71aa84b77]
aws_route_table_association.west_public_subnet_assoc_2: Creation complete after 0s [id=rtbassoc-08e51fdd4ebe4bee8]
aws_route53_zone.main: Still creating... [20s elapsed]
aws_lb.primary.alb: Still creating... [10s elapsed]
aws_autoscaling_group.app_asg: Still creating... [10s elapsed]
aws_autoscaling_group.secondary_asg: Still creating... [10s elapsed]
aws_lb.secondary.alb: Still creating... [10s elapsed]
aws_autoscaling_group.app_asg: Creation complete after 17s [id=terraform-20241030094125457300000004]
aws_autoscaling_group.secondary_asg: Creation complete after 16s [id=terraform-20241030094126584900000004]
aws_autoscaling_policy.scale_up: Creating...
aws_autoscaling_policy.scale_down: Creating...
aws_autoscaling_policy.west_scale_up: Creating...
aws_autoscaling_policy.west_scale_down: Creating...
aws_route53_zone.main: Still creating... [30s elapsed]
aws_lb.primary.alb: Still creating... [20s elapsed]
aws_lb.secondary.alb: Still creating... [20s elapsed]
aws_route53_zone.main: Still creating... [40s elapsed]
aws_lb.primary.alb: Still creating... [40s elapsed]
aws_lb.secondary.alb: Still creating... [40s elapsed]
aws_lb.primary.alb: Still creating... [50s elapsed]
aws_lb.secondary.alb: Still creating... [50s elapsed]
aws_lb.primary.alb: Still creating... [1m1s elapsed]
aws_lb.secondary.alb: Still creating... [1m0s elapsed]
aws_lb.primary.alb: Still creating... [1m1s elapsed]
aws_lb.secondary.alb: Still creating... [1m0s elapsed]
aws_lb.primary.alb: Still creating... [1m21s elapsed]
```

```
Command Prompt + - o x
aws_lb.secondary.alb: Still creating... [2m10s elapsed]
aws_lb.primary.alb: Still creating... [2m21s elapsed]
aws_lb.secondary.alb: Still creating... [2m20s elapsed]
aws_lb.primary.alb: Still creating... [2m31s elapsed]
aws_lb.secondary.alb: Still creating... [2m30s elapsed]
aws_lb.primary.alb: Still creating... [2m41s elapsed]
aws_lb.secondary.alb: Still creating... [2m40s elapsed]
aws_lb.primary.alb: Still creating... [2m51s elapsed]
aws_lb.secondary.alb: Still creating... [2m50s elapsed]
aws_lb.primary.alb: Still creating... [3m1s elapsed]
aws_lb.secondary.alb: Still creating... [3m0s elapsed]
aws_lb.primary.alb: Still creating... [3m1s elapsed]
aws_lb.secondary.alb: Still creating... [3m10s elapsed]
aws_lb.secondary.alb: Creation complete after 9m12s [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:loadbalancer/app/secondary-alb/4b7a0a637b386685]
aws_lb.primary.alb: Creation complete after 3m15s [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:loadbalancer/app/primary-alb/55f8876b41516ba9]
aws_route53_record.secondary: Creating...
aws_route53_health_check.primary_health_check: Creating...
aws_route53_record.primary: Creating...
aws_route53_health_check.secondary_health_check: Creating...
aws_lb_listener.secondary_http: Creating...
aws_lb_listener.http: Creating...
aws_lb_listener.secondary_http: Creation complete after 0s [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:listener/app/secondary-alb/4b7a0a637b386685/938584833b8318cd]
aws_lb_listener.http: Creation complete after 0s [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:listener/app/primary-alb/55f8876b41516ba9/4f2d5bb1d8ba2df1]
aws_route53_health_check.primary_health_check: Creation complete after 0s [id=227ecc39-8273-4708-bca6-befef7ba7e1a]
aws_route53_health_check.secondary_health_check: Creation complete after 0s [id=5ecefa87-32b1-4d77-9ef1-e8cd7faad582]
aws_route53_record.primary: Still creating... [10s elapsed]
aws_route53_record.secondary: Still creating... [10s elapsed]
aws_route53_record.primary: Still creating... [20s elapsed]
aws_route53_record.secondary: Still creating... [20s elapsed]
aws_route53_record.primary: Still creating... [30s elapsed]
aws_route53_record.secondary: Still creating... [30s elapsed]
aws_route53_record.primary: Still creating... [40s elapsed]
aws_route53_record.secondary: Still creating... [40s elapsed]
aws_route53_record.primary: Creation complete after 43s [id=2074361515QXJV7STDIT_app.myinfra.local_A_Primary]
aws_route53_record.secondary: Creation complete after 44s [id=2074361515QXJV7STDIT_app.myinfra.local_A_Secondary]

apply complete! Resources: 33 added, 0 changed, 1 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

vcps | VPC Console

us-west-2.console.aws.amazon.com/vpcconsole/home?region=us-west-2#vpc:

aws Services Search [Alt + S] Actions Create VPC

Last updated less than a minute ago

VPC dashboard EC2 Global View Filter by VPC Virtual private cloud Your VPCs Subnets Route tables Internet gateways Egress-only internet gateways Carrier gateways DHCP option sets Elastic IPs Managed prefix lists Endpoints Endpoint services NAT gateways Peering connections Security Network ACLs CloudShell Feedback

Your VPCs (1/2) Info

Name	VPC ID	State	IPv4 CIDR	IPv6 CIDR	DHCP opt
-	ypc-06c75fed0e2a80ad0	Available	172.31.0.0/16	-	dopt-0e5f
WestVPC	ypc-0bc8cffa4878b083b	Available	10.1.0.0/16	-	dopt-0e5f

vpc-0bc8cffa4878b083b / WestVPC

Details Resource map CIDRs Flow logs Tags Integrations

Details

VPC ID ypc-0bc8cffa4878b083b	State Available	DNS hostnames Disabled	DNS resolution Enabled
Tenancy Default	DHCP option set dopt-0e5fb6b90c00db1f6	Main route table rtb-0566796a03d18a5e0	Main network ACL acl-03989dd5cb6405b41
Default VPC No	IPv4 CIDR 10.1.0.0/16	IPv6 pool -	IPv6 CIDR (Network border group) -
Network Address Usage metrics Disabled	Route 53 Resolver DNS Firewall rule groups -	Owner ID 134575801911	

The screenshot shows the AWS VPC Subnets page. The left sidebar includes 'EC2 Global View' and a 'Filter by VPC' dropdown. Under 'Virtual private cloud', 'Subnets' is selected. The main content area displays a table of subnets:

Name	Subnet ID	State	VPC	IPv4 CIDR
WestPublicSubnet1	subnet-0a1940534f9cb413c	Available	vpc-0bc8cffa4878b083b We...	10.1.0.0/24
-	subnet-0946ffe19b0abb1e5	Available	vpc-06c75fed0e2a80ad0	172.31.0.0/20
-	subnet-082908115138dc0f	Available	vpc-06c75fed0e2a80ad0	172.31.16.0/20
-	subnet-089e4e025bcd9abf7	Available	vpc-06c75fed0e2a80ad0	172.31.48.0/20
WestPublicSubnet2	subnet-05cf33fc2aa1e38da	Available	vpc-0bc8cffa4878b083b We...	10.1.2.0/24
WestPrivateSubnet1	subnet-0de0b39d26bf143ac	Available	vpc-0bc8cffa4878b083b We...	10.1.3.0/24
-	subnet-0623bb1597795b8aa	Available	vpc-06c75fed0e2a80ad0	172.31.32.0/20

Below the table, a message states: "Subnets: subnet-0a1940534f9cb413c, subnet-05cf33fc2aa1e38da, subnet-0de0b39d26bf143ac".

The screenshot shows the AWS VPC Route Tables page. On the left, there's a sidebar with 'VPC dashboard' and 'Virtual private cloud' sections. The main area displays a table of route tables:

Name	Route table ID	Explicit subnet assoc...	Main	VPC
-	rtb-087f1558472945f88	-	Yes	vpc-06c75fed0e2a80ad0
-	rtb-0566796a03d18a5e0	-	Yes	vpc-0bc8cffa4878b083b W
<input checked="" type="checkbox"/> WestPublicRouteTable	rtb-0bf4cfa31ff877e69	2 subnets	-	vpc-0bc8cffa4878b083b W

Below the table, the details for the selected route table ('rtb-0bf4cfa31ff877e69 / WestPublicRouteTable') are shown. The 'Details' tab is selected, displaying the following information:

Route table ID rtb-0bf4cfa31ff877e69	Main No	Explicit subnet associations 2 subnets	Edge associations -
VPC vpc-0bc8cffa4878b083b WestVPC	Owner ID 134575801911		

RouteTables | VPC Console

us-west-2.console.aws.amazon.com/vpcconsole/home?region=us-west-2#RouteTables

VPC dashboard

EC2 Global View

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

Endpoints

Endpoint services

NAT gateways

Peering connections

Security

Network ACLs

CloudShell Feedback

Last updated 4 minutes ago

Actions Create route table

Route tables (1/3) Info

Find resources by attribute or tag

Name	Route table ID	Explicit subnet assoc...	Edge associati...	Main	VPC
-	rtb-087f1558472945f88	-	-	Yes	vpc-06c75fed0e2a80ad0
-	rtb-0566796a03d18a5e0	-	-	Yes	vpc-0bc8cffa487bb083b V
<input checked="" type="checkbox"/> WestPublicRouteTable	rtb-0bf4cfa31ff877e69	2 subnets	-	No	vpc-0bc8cffa487bb083b V

rtb-0bf4cfa31ff877e69 / WestPublicRouteTable

Details Routes Subnet associations Edge associations Route propagation Tags

Routes (2)

Filter routes

Destination	Target	Status	Propagated
0.0.0.0/0	igw-093c80565bfdb459c	Active	No
10.1.0.0/16	local	Active	No

Both Edit routes

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

igws | VPC Console

us-west-2.console.aws.amazon.com/vpcconsole/home?region=us-west-2#gws

VPC dashboard

EC2 Global View

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

Endpoints

Endpoint services

NAT gateways

Peering connections

Security

Network ACLs

CloudShell Feedback

Create internet gateway

Internet gateways (1/2) Info

Find resources by attribute or tag

Name	Internet gateway ID	State	VPC ID	Owner
-	igw-0570f58551e6a75ac	Attached	vpc-06c75fed0e2a80ad0	134575801911
<input checked="" type="checkbox"/> WestIGW	igw-093c80565bfdb459c	Attached	vpc-0bc8cffa487bb083b WestVPC	134575801911

igw-093c80565bfdb459c / WestIGW

Details Tags

Details

Internet gateway ID igw-093c80565bfdb459c	State Attached	VPC ID vpc-0bc8cffa487bb083b WestVPC	Owner 134575801911
--	-------------------	---	-----------------------

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

SecurityGroups | VPC Console

us-west-2.console.aws.amazon.com/vpcconsole/home?region=us-west-2#SecurityGroups

VPC dashboard

EC2 Global View

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

Endpoints

Endpoint services

NAT gateways

Peering connections

Security

Network ACLs

Security groups

DNS firewall

Rule groups

Domain lists

Network Firewall

Firewalls

Firewall policies

Network Firewall rule groups

CloudShell Feedback

Create security group

Security Groups (1/4) Info

Find resources by attribute or tag

Name	Security group ID	Security group name	VPC ID
-	sg-058a5007c6ef0fd8e	default	vpc-06c75fed0e2a80ad0
-	sg-0998e8bd09229cd46	default	vpc-0bc8cffa487bb083b
-	sg-0ffebbb805f877ac14	default_elb_8cb7cedb-2b76-3e94-8...	vpc-06c75fed0e2a80ad0
<input checked="" type="checkbox"/> WestAppSecurityGroup	sg-03d0410997fd80d0	west-app-sg	vpc-0bc8cffa487bb083b

Details Inbound rules Outbound rules Tags

Details

Security group name west-app-sg	Security group ID sg-03d0410997fd80d0	Description Allow HTTP and SSH traffic in us-west-2	VPC ID vpc-0bc8cffa487bb083b
Owner 134575801911	Inbound rules count 3 Permission entries	Outbound rules count 1 Permission entry	

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS VPC Console showing the Security Groups page. The left sidebar shows navigation for Egress-only internet gateways, Carrier gateways, DHCP option sets, Managed prefix lists, Endpoints, Endpoint services, NAT gateways, Peering connections, Security groups, DNS firewall, Network Firewall, and CloudShell.

Security Groups (1/4) Info

Name	Security group ID	Security group name	VPC ID
-	sg-058a5007c6ef0fd4e	default	vpc-06c75fed0e2a80ad0
-	sg-0998e8bd09229cd46	default	vpc-0bc8cffa4878b083b
-	sg-0f6bb805fb77ac14	default_elb_8cb7cedb-2b76-3e94-8...	vpc-06c75fed0e2a80ad0
WestAppSecurityGroup	sg-03d04109977fd80d0	west-app-sg	vpc-0bc8cffa4878b083b

Inbound rules (3)

Name	Security group rule ID	IP version	Type	Protocol	Port range	Source	Description
-	sgr-0f03d581f9e487913	IPv4	SSH	TCP	22	18.237.140.160/29	-
-	sgr-05176265acd69b3...	IPv4	SSH	TCP	22	142.59.28.49/32	-
-	sgr-09b6a3878d57a88f4	IPv4	HTTP	TCP	80	0.0.0.0/0	-

Screenshot of the AWS EC2 Instances page. The left sidebar shows navigation for Dashboard, EC2 Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity, Reservations (New), Images (AMIs, AMI Catalog), and Elastic Block Store (Volumes, Snapshots, Lifecycle Manager).

Instances (1/2) Info

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
SecondaryAp...	i-03ebc0ea0e83ea805	Running	t2.micro	2/2 checks passed	View alarms	us-west-2a	-
SecondaryAp...	i-0eac73c68375c452e	Running	t2.micro	2/2 checks passed	View alarms	us-west-2b	-

i-03ebc0ea0e83ea805 (SecondaryAppASGInstance)

Details | Status and alarms | Monitoring | Security | Networking | Storage | Tags

Instance summary

instance ID	Public IPv4 address	Private IPv4 addresses
i-03ebc0ea0e83ea805	55.88.183.255 open address	10.1.1.15
IPv6 address	Instance state	Public IPv4 DNS
-	Running	-
Hostname type	Private IP DNS name (IPv4 only)	
IP name: ip-10-1-1-15.us-west-2.compute.internal	ip-10-1-1-15.us-west-2.compute.internal	

Screenshot of the AWS EC2 Instances page. The left sidebar shows navigation for Dashboard, EC2 Global View, Events, Instances (selected), Instance Types, Launch Templates, Spot Requests, and EBS volumes.

Instances (1/2) Info

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
SecondaryAp...	i-03ebc0ea0e83ea805	Running	t2.micro	2/2 checks passed	View alarms	us-west-2a	-
SecondaryAp...	i-0eac73c68375c452e	Running	t2.micro	2/2 checks passed	View alarms	us-west-2b	-

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org. Commercial support is available at nginx.com.

Thank you for using nginx.

EC2 Instance Connect

```

last login: Wed Oct 30 09:47:28 2024 from 18.237.140.164
[ec2-user@ip-10-1-1-15 ~]$ curl http://localhost
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto; font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>

```

Auto Scaling groups | EC2 | us-west-2

Auto Scaling groups (1/1) Info Launch configurations Launch templates Actions Create Auto Scaling group

Search your Auto Scaling groups

Name	Launch template/configuration	Instances	Status	Desired capacity	Min
terraform-20241030094126584900000004	secondary-app-launch-template20241030094126584900000004	2	-	2	1

Auto Scaling group: terraform-20241030094126584900000004

Details Activity Automatic scaling Instance management Monitoring Instance refresh

Group details

Auto Scaling group name	Desired capacity	Desired capacity type	Amazon Resource Name (ARN)
terraform-20241030094126584900000004	2	Units (number of instances)	arn:aws:autoscaling:us-west-2:134575801911:autoScalingGroup:774d14fa-ec71-45e0-99ed-fd086ede3d7:autoScalingGroupName/terraform-20241030094126584900000004
	Minimum capacity	Status	
	1	-	
	Maximum capacity		
	5		

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

CloudShell Feedback

Load balancers | EC2 | us-west-2

Load balancers (1/1) Actions Create load balancer

Filter load balancers

Name	DNS name	State	VPC ID	Availability Zones	Type	Data
secondary-alb	secondary-alb-18570568...	Active	vpc-0bc8cffa4878b0...	2 Availability Zones	application	Oct 30, 2024

Load balancer: secondary-alb

Details

Load balancer type	Status	Load balancer IP address type
Application	Active	IPv4
Scheme	Hosted zone	Date created
Internet-facing	Z1H1FLSHABSFS	October 30, 2024, 03:41 (UTC-06:00)
Load balancer ARN	DNS name	
	Info	

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

CloudShell Feedback

Load balancers | EC2 | us-west-2

Load balancers (1/1) Actions Create load balancer

Filter load balancers

Name	DNS name	State	VPC ID	Availability Zones	Type	Data
secondary-alb	secondary-alb-18570568...	Active	vpc-0bc8cffa4878b0...	2 Availability Zones	application	Oct 30, 2024

Load balancer: secondary-alb

Listeners and rules (1/1) Info

A listener checks for connection requests on its configured protocol and port. Traffic received by the listener is routed according to the default action and any additional rules.

Manage rules Manage listener Add listener

Filter listeners

Protocol:Port	Default action	Rules	ARN	Security policy	Defa...
HTTP:80	Forward to target group	1 rule	ARN	Not applicable	Not...
	• secondary-target-group [2]	1 (100%)			
	• Target group stickiness: Off				

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

CloudShell Feedback

Target groups | EC2 us-west-2

us-west-2.console.aws.amazon.com/ec2/home?region=us-west-2#TargetGroupsv=3;\$case=tagsfalse%5Cclientfalse\$regex=tags... | 🔍 | 🛡️ | VPN | Oregon | vivekvashist @ vv3281 | ☰

AMI Catalog

Elastic Block Store

- Volumes
- Snapshots
- Lifecycle Manager

Network & Security

- Security Groups
- Elastic IPs
- Placement Groups
- Key Pairs
- Network Interfaces

Load Balancing

- Load Balancers
- Target Groups**
- Trust Stores [New](#)

Auto Scaling

- Auto Scaling Groups

Settings

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

EC2 > Target groups

Target groups (1/1) Info

Name	ARN	Port	Protocol	Target type	Load balancer
secondary-target-group	arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46eba4401c8967	80	HTTP	Instance	secondary-alb

Target group: secondary-target-group

Details

Target type: Instance	Protocol: Port: HTTP: 80	Protocol version: HTTP1	VPC: vpc-0bc8cffa4878b083b
IP address type: IPv4	Load balancer: secondary-alb		

2 Total targets	2 Healthy	0 Unhealthy	0 Unused	0 initial	0 Draining
-----------------	-----------	-------------	----------	-----------	------------

Target group: secondary-target-group

Registered targets (2/2) Info

Anomaly mitigation: Not applicable

Instance ID	Name	Port	Zone	Health status	Health status details
i-05ebc0ea0e83ea005	SecondaryAppASGInstance	80	us-west-2a (usw2-az2)	Healthy	-
i-0eac73c68375c452e	SecondaryAppASGInstance	80	us-west-2b (usw2-az1)	Healthy	-

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Target groups | EC2 us-west-2

us-west-2.console.aws.amazon.com/ec2/home?region=us-west-2#TargetGroupsv=3;\$case=tagsfalse%5Cclientfalse\$regex=tags... | 🔍 | 🛡️ | VPN | Oregon | vivekvashist @ vv3281 | ☰

AMI Catalog

Elastic Block Store

- Volumes
- Snapshots
- Lifecycle Manager

Network & Security

- Security Groups
- Elastic IPs
- Placement Groups
- Key Pairs
- Network Interfaces

Load Balancing

- Load Balancers
- Target Groups**
- Trust Stores [New](#)

Auto Scaling

- Auto Scaling Groups

Settings

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

EC2 > Target groups

Target groups (1/1) Info

Name	ARN	Port	Protocol	Target type	Load balancer
secondary-target-group	arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46eba4401c8967	80	HTTP	Instance	secondary-alb

Target group: secondary-target-group

Registered targets (2/2) Info

Anomaly mitigation: Not applicable

Instance ID	Name	Port	Zone	Health status	Health status details
i-05ebc0ea0e83ea005	SecondaryAppASGInstance	80	us-west-2a (usw2-az2)	Healthy	-
i-0eac73c68375c452e	SecondaryAppASGInstance	80	us-west-2b (usw2-az1)	Healthy	-

Target group: secondary-target-group

Health check settings

Protocol: HTTP	Path: /	Port: Traffic port	Healthy threshold: 2 consecutive health check successes
Unhealthy threshold: 2 consecutive health check failures	Timeout: 5 seconds	Interval: 30 seconds	Success codes: 200

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Target groups | EC2 us-west-2

us-west-2.console.aws.amazon.com/ec2/home?region=us-west-2#TargetGroupsv=3;\$case=tagsfalse%5Cclientfalse\$regex=tags... | 🔍 | 🛡️ | VPN | Oregon | vivekvashist @ vv3281 | ☰

AMI Catalog

Elastic Block Store

- Volumes
- Snapshots
- Lifecycle Manager

Network & Security

- Security Groups
- Elastic IPs
- Placement Groups
- Key Pairs
- Network Interfaces

Load Balancing

- Load Balancers
- Target Groups**
- Trust Stores [New](#)

Auto Scaling

- Auto Scaling Groups

Settings

CloudShell Feedback

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

EC2 > Target groups

Target groups (1/1) Info

Name	ARN	Port	Protocol	Target type	Load balancer
secondary-target-group	arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46eba4401c8967	80	HTTP	Instance	secondary-alb

Target group: secondary-target-group

Health checks

Protocol: HTTP	Path: /	Port: Traffic port	Healthy threshold: 2 consecutive health check successes
Unhealthy threshold: 2 consecutive health check failures	Timeout: 5 seconds	Interval: 30 seconds	Success codes: 200

Target group: secondary-target-group

Health check settings

Protocol: HTTP	Path: /	Port: Traffic port	Healthy threshold: 2 consecutive health check successes
Unhealthy threshold: 2 consecutive health check failures	Timeout: 5 seconds	Interval: 30 seconds	Success codes: 200

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS EC2 Load Balancers console showing the 'primary-alb' load balancer details.

Load balancers (1/1)

Name	DNS name	Status	VPC ID	Availability Zones	Type
primary-alb	primary-alb-136165955...	Active	vpc-0ea415b9040f75...	2 Availability Zones	application

Load balancer: primary-alb

Details

Load balancer type	Status	VPC	Load balancer IP address type
Application	Active	vpc-0ea415b9040f753dc	IPv4
Scheme	Hosted zone	Availability Zones	Date created
Internet-facing	Z35SXDOTRQ7X7K	subnet-0d89362f71d4a5ad, us-east-1b (use1-az2) subnet-04b29a39893e7944d, us-east-1a (use1-az1)	October 30, 2024, 03:41 (UTC-06:00)

Screenshot of the AWS EC2 Load Balancers console showing the 'primary-alb' load balancer listeners and rules.

Listeners and rules (1/1) Info

Protocol:Port	Default action	Rules	ARN	Security policy	Defa
HTTP:80	Forward to target group • primary-target-group (1, 100%) • Target group stickiness: Off	1 rule	ARN	Not applicable	Not applicable

Screenshot of the AWS EC2 Target Groups console showing the 'primary-target-group' target group details.

Target groups (1/1) Info

Name	ARN	Port	Protocol	Target type	Load balancer
primary-target-group	arn:aws:elasticloadbalanc...	80	HTTP	Instance	primary-alb

Target group: primary-target-group

Details

Target type	Protocol : Port	Protocol version	VPC
Instance	HTTP: 80	HTTP1	vpc-0ea415b9040f753dc
IP address type	Load balancer		
IPv4	primary-alb		

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
2	2	0	0	0	0
Total targets	Healthy	Unhealthy	Unused	Initial	Draining
0 Anomalous					

Target groups | EC2 | us-east-1

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#TargetGroups:

EC2 > Target groups

Target groups (1/1) Info

Name	ARN	Port	Protocol	Target type	Load balancer
primary-target-group	arn:aws:elasticloadbalanc...	80	HTTP	instance	primary-alb

Target group: primary-target-group

Registered targets (2/2) Info

Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTP2 target groups with at least 5 healthy targets.

Instance ID	Name	Port	Zone	Health status	Health status details	Launch...
i-00347e579d3947ab7	AppSGinstance	80	us-east-1b (us...)	Healthy	-	October ..
i-0028fb710309986f1	AppSGinstance	80	us-east-1a (us...)	Healthy	-	October ..

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Route 53 - dashboard

us-east-1.console.aws.amazon.com/route53/v2/home?region=us-east-1#Dashboard

Route 53

Route 53 Dashboard Info

DNS management

1 Hosted zone

Traffic management

A visual tool that lets you easily create policies for multiple endpoints in complex configurations.

Create policy

Availability monitoring

2 Health checks

Domain registration

A domain is the name, such as example.com, that your users use to access your application.

Register domain

Register domain

Find and register an available domain, or transfer your existing domains to Route 53.

Enter a domain name

Check

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Route 53 Console hosted zones

us-east-1.console.aws.amazon.com/route53/v2/hostedzones?region=us-east-1#

Route 53

Hosted zones

Hosted zones (1/1)

Hosted zone name	Type	Created by	Record count
myinfra.local	Public	Route 53	4

Hosted zone details

Hosted zone name: myinfra.local

Hosted zone ID: Z074361515QXJV7STD1T

Description: Managed by Terraform

Query log

Type: Public hosted zone

Record count: 4

Name servers:

- ns-893.awsdns-47.net
- ns-162.awsdns-20.com
- ns-1553.awsdns-02.co.uk
- ns-1053.awsdns-03.org

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Route 53 Management Console

us-east-1.console.aws.amazon.com/route53/healthchecks/home?region=us-east-1#/

Create health check Delete health check Edit health check

Name	Status	Description	Alarms	ID
<input checked="" type="checkbox"/> primary-alb-1361659550.us-east-1.elb.amazonaws.com	Healthy an hour ago now	http://primary-alb-1361659550.us-east-1.elb.amazonaws.com:80/	No alarms configured.	227ecc39-8273-4708-bca6-beef7ba7e1a
<input type="checkbox"/> secondary-alb-1857056899.us-west-2.elb.amazonaws.com	Healthy an hour ago now	http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com:80/	No alarms configured.	5ece7a87-32b1-4d77-9ef1-e8cd7faad582

Filter by keyword

Info Monitoring Alarms Tags Health checkers Latency

ID: 227ecc39-8273-4708-bca6-beef7ba7e1a

URL: http://primary-alb-1361659550.us-east-1.elb.amazonaws.com:80/

Specify endpoint by Domain name

Protocol: HTTP

Domain name: primary-alb-1361659550.us-east-1.elb.amazonaws.com

Advanced configuration

Request interval: 30 seconds

Failure threshold: 3

Search string: -

Latency graphs: No

Invert health check: No

CloudShell Feedback

Route 53 Management Console

us-east-1.console.aws.amazon.com/route53/healthchecks/home?region=us-east-1#/

Create health check Delete health check Edit health check

Name	Status	Description	Alarms	ID
<input type="checkbox"/> primary-alb-1361659550.us-east-1.elb.amazonaws.com	Healthy an hour ago a minute ago	http://primary-alb-1361659550.us-east-1.elb.amazonaws.com:80/	No alarms configured.	227ecc39-8273-4708-bca6-beef7ba7e1a
<input checked="" type="checkbox"/> secondary-alb-1857056899.us-west-2.elb.amazonaws.com	Healthy an hour ago a minute ago	http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com:80/	No alarms configured.	5ece7a87-32b1-4d77-9ef1-e8cd7faad582

Filter by keyword

Info Monitoring Alarms Tags Health checkers Latency

ID: 5ece7a87-32b1-4d77-9ef1-e8cd7faad582

URL: http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com:80/

Specify endpoint by Domain name

Protocol: HTTP

Domain name: secondary-alb-1857056899.us-west-2.elb.amazonaws.com

Advanced configuration

Request interval: 30 seconds

Failure threshold: 3

Search string: -

Latency graphs: No

Invert health check: No

CloudShell Feedback

Monitoring and Alerts with CloudWatch

a) Overview

In this phase, I set up monitoring and alerting for EC2 instances in both us-east-1 and us-west-2 to ensure visibility into resource usage and automated response to any high-utilization events. To enable this, I first updated the IAM role in `iam_roles.tf` to grant EC2 instances permission to run the CloudWatch Agent. Then, I configured the EC2 launch templates in us-east-1 and us-west-2 with user data scripts to pre-install both NGINX and the CloudWatch Agent, along with a configuration for tracking CPU, memory, and disk usage. For alerting, I created `cloudwatch_alarms.tf`, defining alarms for high CPU, memory, and disk utilization, each set to trigger notifications via SNS when metrics exceed thresholds. I also added an IAM policy to allow EC2 instances to publish to SNS topics in both regions. After deployment, I verified that CloudWatch alarms were correctly set up and linked to SNS, and I confirmed alert functionality by subscribing to SNS topics via email. This monitoring and alerting setup provides automated insights into resource utilization and prompt alerts for effective performance management across both regions.

b) Attaching CloudWatch Agent Permissions

To enable CloudWatch monitoring, I needed to attach the CloudWatch agent policy to the IAM role for the EC2 instances.

Code Block: IAM Role Policy Attachment

```
# Attachment of CloudWatch Agent Policy to the IAM role for the EC2 instances
resource "aws_iam_role_policy_attachment" "ec2_cloudwatch_agent" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/CloudWatchAgentServerPolicy"
}
```

After saving the updated `iam_roles.tf`, I ran `terraform plan` and `terraform apply --auto-approve` to implement the changes. After successful deployment, I verified in the AWS Console under IAM roles that the `CloudWatchAgentServerPolicy` was indeed attached to the `ec2-role`, confirming proper configuration.

c) Updating EC2 User Data for CloudWatch Agent Installation

To automate the installation of the CloudWatch agent on EC2 instances, I updated the user data in the Launch Template configuration. This ensures that each instance has both `nginx` and `amazon-cloudwatch-agent` pre-installed.

Code Block: Launch Template for us-east-1 with CloudWatch Agent Installation

```
# Launch Template for EC2 Instances for us-east-1
resource "aws_launch_template" "app" {
  name_prefix = "app-launch-template"
```

```

image_id      = "ami-06b21ccaeff8cd686" # Amazon Linux AMI
instance_type = var.instance_type

# Attach IAM role to EC2 instances
iam_instance_profile {
    name = aws_iam_instance_profile.ec2_instance_profile.name
}

network_interfaces {
    security_groups = [aws_security_group.app_sg.id]
}

user_data = base64encode(<<-EOF
#!/bin/bash
sudo dnf update -y
sudo dnf install -y nginx amazon-cloudwatch-agent
sudo systemctl start nginx
sudo systemctl enable nginx

# Configure the CloudWatch Agent
sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json << 'EOF2'
{
    "metrics": {
        "append_dimensions": {
            "AutoScalingGroupName": "AppASGInstance"
        },
        "metrics_collected": {
            "mem": {
                "measurement": [
                    "mem_used_percent"
                ],
                "metrics_collection_interval": 60
            },
            "disk": {
                "measurement": [
                    "used_percent"
                ],
                "resources": [
                    "/"
                ],
                "metrics_collection_interval": 60
            }
        }
    }
}
EOF2

# Start the CloudWatch Agent
sudo /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-
agent-ctl -a start -c file:/opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json -m ec2
EOF
)

lifecycle {
    create_before_destroy = true
}

tag_specifications {
    resource_type = "instance"
}

```

```

        tags = {
            Name = "AppInstance"
        }
    }
}

```

d) Adding Instance Refresh for Auto Scaling Group (ASG) in us-east-1

To apply this updated configuration to all existing instances, I updated the Auto Scaling Group (ASG) in `us-east-1` to perform a rolling instance refresh. This ensures minimal downtime while applying the new settings.

Code Block: Instance Refresh Configuration

```

instance_refresh {
    strategy = "Rolling"
    preferences {
        min_healthy_percentage = 0
        instance_warmup         = 120
    }
}

```

e) Updating the Launch Template and ASG for us-west-2

Similarly, I applied the same configuration updates to the Launch Template and ASG in `us-west-2` to ensure CloudWatch monitoring for instances in both regions.

Code Block: Launch Template for us-west-2 with CloudWatch Agent Installation

```

# Launch Template for EC2 Instances in us-west-2
resource "aws_launch_template" "secondary_app" {
    provider      = aws.west
    name_prefix   = "secondary-app-launch-template"
    image_id      = "ami-07c5ecd8498c59db5" # Amazon Linux AMI
    instance_type = var.instance_type

    # Attach IAM role to EC2 instances
    iam_instance_profile {
        name = aws_iam_instance_profile.ec2_instance_profile.name
    }

    network_interfaces {
        security_groups = [aws_security_group.west_app_sg.id]
    }

    user_data = base64encode(<<-EOF
        #!/bin/bash
        sudo dnf update -y
        sudo dnf install -y nginx amazon-cloudwatch-agent
        sudo systemctl start nginx
        sudo systemctl enable nginx

        # Configure the CloudWatch Agent
        sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-
        cloudwatch-agent.json << 'EOF2'
        {
            "metrics": {

```

```

    "append_dimensions": {
        "AutoScalingGroupName": "SecondaryAppASGInstance"
    },
    "metrics_collected": {
        "mem": {
            "measurement": [
                "mem_used_percent"
            ],
            "metrics_collection_interval": 60
        },
        "disk": {
            "measurement": [
                "used_percent"
            ],
            "resources": [
                "/"
            ],
            "metrics_collection_interval": 60
        }
    }
}
EOF2

# Start the CloudWatch Agent
sudo /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-
agent-ctl -a start -c file:/opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json -m ec2
EOF
)

lifecycle {
    create_before_destroy = true
}

tag_specifications {
    resource_type = "instance"
    tags = {
        Name = "SecondaryAppInstance"
    }
}
}

```

Updating ASG for us-west-2 with Instance Refresh

```

instance_refresh {
    strategy = "Rolling"
    preferences {
        min_healthy_percentage = 0
        instance_warmup        = 120
    }
}

```

f) Redeploying with Updated Configurations

To apply these configurations, I first destroyed the existing ASGs and launch templates to ensure fresh provisioning with the updated settings. I ran:

```
terraform destroy -target=aws_autoscaling_group.app_asg -  
target=aws_autoscaling_group.secondary_asg -target=aws_launch_template.app  
-target=aws_launch_template.secondary_app
```

After successfully destroying the resources, I re-ran:

```
terraform init  
terraform plan  
terraform apply -auto-approve
```

g) Verification

Once the resources were recreated, I verified the setup by connecting to instances in both us-east-1 and us-west-2 and running the following commands:

```
curl http://localhost  
rpm -qa | grep amazon-cloudwatch-agent
```

These commands confirmed that both nginx and the cloudwatch-agent were correctly installed and active on all instances in both regions.

h) Creating CloudWatch Alarms and SNS Topics

To enable real-time alerts for CPU, memory, and disk utilization across both us-east-1 and us-west-2 regions, I created a new Terraform file named `cloudwatch_alarms.tf`. This file defines the CloudWatch alarms for each metric, as well as SNS topics for alert notifications.

Code Block: CloudWatch Alarms and SNS Topics Configuration

```
# us-east-1

# CloudWatch Alarm for CPU Utilization
resource "aws_cloudwatch_metric_alarm" "high_cpu_utilization" {
  alarm_name          = "HighCPUUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUUtilization"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 80 # Threshold at 80% CPU Utilization
  alarm_description    = "Triggers if CPU utilization exceeds 80%."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  }
  actions_enabled = true
  alarm_actions    = [aws sns topic.alert_topic.arn]
}

# CloudWatch Alarm for Memory Utilization
resource "aws_cloudwatch_metric_alarm" "high_memory_utilization" {
  alarm_name          = "HighMemoryUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods  = 2
  metric_name         = "mem_used_percent" # Custom metric from CloudWatch
  Agent
```

```

namespace          = "CWAgent"
period            = 300
statistic         = "Average"
threshold         = 80 # Threshold at 80% Memory Utilization
alarm_description = "Triggers if memory utilization exceeds 80%."
dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
}
actions_enabled = true
alarm_actions   = [aws_sns_topic.alert_topic.arn]
}

# CloudWatch Alarm for Disk Utilization
resource "aws_cloudwatch_metric_alarm" "high_disk_utilization" {
    alarm_name          = "HighDiskUtilizationAlarm"
    comparison_operator = "GreaterThanThreshold"
    evaluation_periods  = 2
    metric_name         = "disk_used_percent" # Custom metric from CloudWatch
Agent
    namespace          = "CWAgent"
    period            = 300
    statistic         = "Average"
    threshold         = 80 # Threshold at 80% Disk Utilization
    alarm_description = "Triggers if disk utilization exceeds 80%."
    dimensions = {
        AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
    actions_enabled = true
    alarm_actions   = [aws_sns_topic.alert_topic.arn]
}

# SNS Topic for Alarm Notifications in us-east-1
resource "aws_sns_topic" "alert_topic" {
    name = "EC2AlertTopic"
}

resource "aws_sns_topic_subscription" "email_subscription" {
    topic_arn = aws_sns_topic.alert_topic.arn
    protocol  = "email"
    endpoint   = "vivekvash1507@gmail.com"
}

```

Similarly, I configured alarms and SNS topics for the us-west-2 region, following the same structure.

Code Block: CloudWatch Alarms and SNS Topics Configuration for us-west-2

```

# us-west-2

# CloudWatch Alarm for CPU Utilization
resource "aws_cloudwatch_metric_alarm" "west_high_cpu_utilization" {
    provider           = aws.west
    alarm_name         = "WestHighCPUUtilizationAlarm"
    comparison_operator = "GreaterThanThreshold"
    evaluation_periods = 2
    metric_name        = "CPUUtilization"
    namespace          = "AWS/EC2"
    period             = 300
    statistic          = "Average"
    threshold          = 80
}

```

```

alarm_description      = "Triggers if CPU utilization exceeds 80%."
dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
}
actions_enabled = true
alarm_actions     = [aws_sns_topic.west_alert_topic.arn]
}

# CloudWatch Alarm for Memory Utilization in us-west-2
resource "aws_cloudwatch_metric_alarm" "west_high_memory_utilization" {
    provider          = aws.west
    alarm_name        = "WestHighMemoryUtilizationAlarm"
    comparison_operator = "GreaterThanOrEqualToThreshold"
    evaluation_periods = 2
    metric_name       = "mem_used_percent"
    namespace         = "CWAgent"
    period            = 300
    statistic         = "Average"
    threshold          = 80 # Set threshold based on your requirements
    alarm_description = "This alarm triggers if memory utilization
exceeds 80%."
    dimensions = {
        AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
    }
    actions_enabled = true
    alarm_actions     = [aws_sns_topic.west_alert_topic.arn]
}

# CloudWatch Alarm for Disk Utilization in us-west-2
resource "aws_cloudwatch_metric_alarm" "west_high_disk_utilization" {
    provider          = aws.west
    alarm_name        = "WestHighDiskUtilizationAlarm"
    comparison_operator = "GreaterThanOrEqualToThreshold"
    evaluation_periods = 2
    metric_name       = "disk_used_percent"
    namespace         = "CWAgent"
    period            = 300
    statistic         = "Average"
    threshold          = 80 # Set threshold based on your requirements
    alarm_description = "This alarm triggers if disk utilization exceeds
80%."
    dimensions = {
        AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
    }
    actions_enabled = true
    alarm_actions     = [aws_sns_topic.west_alert_topic.arn]
}

# SNS Topic for Alarm Notifications in us-west-2
resource "aws sns topic" "west_alert_topic" {
    provider = aws.west
    name      = "WestEC2AlertTopic"
}

resource "aws sns topic subscription" "west_email_subscription" {
    provider   = aws.west
    topic_arn  = aws_sns_topic.west_alert_topic.arn
    protocol   = "email"
    endpoint   = "vivekvash1507@gmail.com"
}

```

i) Updating IAM Role Permissions for SNS Publishing

To allow EC2 instances to publish alerts to the SNS topics, I added an additional policy to the IAM role in `iam_roles.tf`.

Code Block: SNS Publish Policy Attachment

```
# Allow EC2 instances to publish to both SNS topics for alerts
resource "aws_iam_role_policy" "ec2_sns_publish_policy" {
  name = "ec2-sns-publish-policy"
  role = aws_iam_role.ec2_role.name

  policy = jsonencode({
    "Version": "2012-10-17",
    "Statement": [
      {
        "Effect": "Allow",
        "Action": "sns:Publish",
        "Resource": [
          "${aws sns topic.alert_topic.arn}",           # SNS topic in us-east-
1
          "${aws sns topic.west_alert_topic.arn}"         # SNS topic in us-west-
2
        ]
      }
    ]
  })
}
```

j) Deploying and Verifying CloudWatch Alarms and SNS Notifications

After saving the updated `iam_roles.tf`, I executed the following commands to apply the changes:

```
terraform init
terraform plan
terraform apply -auto-approve
```

Following successful deployment, I verified each configuration:

1. **IAM Role Verification:** In the AWS Console, I confirmed that the `ec2-sns-publish-policy` was attached to the `ec2-role`, granting SNS publish permissions.
2. **CloudWatch Alarm Verification:** In both `us-east-1` and `us-west-2`, I accessed CloudWatch and confirmed the presence of the alarms for CPU, memory, and disk utilization.
3. **Email Subscription Confirmation:** I received email confirmations from SNS for both `EC2AlertTopic` and `WestEC2AlertTopic`. I clicked the provided links to confirm each subscription.

With these configurations, I successfully enabled CloudWatch monitoring and alerting for key utilization metrics across both regions. The EC2 instances are now capable of sending alerts to the configured SNS topics, ensuring timely notifications for any critical utilization thresholds. This marks the successful completion of the Monitoring and Alerts phase of the project.

Screenshots

The screenshot shows the Codeberg IDE interface with the 'TERRAFORM' workspace selected. The 'EXPLORER' panel on the left lists several Terraform files: .terraform, alb.tf, ec2_autoscaling.tf, iam_roles.tf (which is currently open), main.tf, network.tf, route53.tf, and s3_bucket.tf. The 'CODE' tab is active, displaying the contents of the 'iam_roles.tf' file:

```
resource "aws_iam_role_policy_attachment" "ec2_cloudwatch_agent" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonCloudWatchLogsFullAccess"
}

resource "aws_iam_role_policy_attachment" "elb_full_access" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/ElasticLoadBalancingFullAccess"
}

# Attachment of S3 read-only policy to the role
resource "aws_iam_role_policy_attachment" "ec2_s3_readonly" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/AmazonS3ReadOnlyAccess"
}

# Attachment of CloudWatch Agent Policy to the IAM role for the EC2 instances
resource "aws_iam_role_policy_attachment" "ec2_cloudwatch_agent" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:iam::aws:policy/CloudWatchAgentServerPolicy"
}
```

The screenshot shows a Command Prompt window with the title 'Command Prompt'. It displays the output of a Terraform plan command, showing the execution plan for the resources defined in the configuration files. The output includes resource names, IDs, and status messages such as 'Refreshing state...' and 'will be created'. The plan summary at the bottom indicates 1 resource to add, 0 to change, and 0 to destroy.

```
aws_autoscaling_group.secondary_asg: Refreshing state... [id=terrafrom-20241030094126584900000004]
aws_route53_record.secondary: Refreshing state... [id=Z0743615150XJV7STD1_app.myinfra.local_A_Secondary]
aws_lb_listener.secondary_http: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:listener/app/secondary-alb/4b7a0a637b386685/93858403368318cd]
aws_route53_health_check.secondary_health_check: Refreshing state... [id=Sece7a87-32b1-4d77-9ef1-e8cd7faad582]
aws_s3_bucket_public_access_block.backup_bucket_block: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_route_table_association.public_subnet_assoc_1: Refreshing state... [id=rbaassoc-062c3dc39ec6eb03]
aws_route_table_association.public_subnet_assoc_2: Refreshing state... [id=rbaassoc-0c1e016c8860f95a6]
aws_route.internet_access: Refreshing state... [id=-rtb-0c4fc097dfc205e1080289494]
aws_lb.primary_alb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:loadbalancer/app/primary-alb/55f8876b41516ba9]
aws_launch_template.app: Refreshing state... [id=t-074d209bf44eba96]
aws_autoscaling_policy.west_scale_down: Refreshing state... [id=west-scale-down]
aws_autoscaling_policy.west_scale_up: Refreshing state... [id=west-scale-up]
aws_autoscaling_group.app_asg: Refreshing state... [id=terrafrom-20241030094125457300000004]
aws_route53_health_check.primary_health_check: Refreshing state... [id=27ecc39-8273-4708-bca6-befef7ba7e1a]
aws_route53_record.primary: Refreshing state... [id=Z0743615150XJV7STD1_app.myinfra.local_A_Primary]
aws_lb_listener.http: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:listener/app/primary-alb/55f8876b41516ba9/4f2d5bb1d8ba2df1]
aws_autoscaling_policy.scale_down: Refreshing state... [id=scale-down]
aws_autoscaling_policy.scale_up: Refreshing state... [id=scale-up]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# aws_iam_role_policy_attachment.ec2_cloudwatch_agent will be created
+ resource "aws_iam_role_policy_attachment" "ec2_cloudwatch_agent" {
  + id          = (known after apply)
  + policy_arn = "arn:aws:iam::aws:policy/CloudWatchAgentServerPolicy"
  + role        = "ec2-role"
}

Plan: 1 to add, 0 to change, 0 to destroy.
aws_iam_role_policy_attachment.ec2_cloudwatch_agent: Creating...
aws_iam_role_policy_attachment.ec2_cloudwatch_agent: Creation complete after 1s [id=ec2-role-20241031075603479000000001]

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

Screenshot of the AWS IAM console showing the 'ec2-role' details page.

Summary

- Creation date: October 26, 2024, 02:30 (UTC-06:00)
- ARN: arn:aws:iam::134575801911:role/ec2-role
- Last activity: 39 minutes ago
- Maximum session duration: 1 hour

Permissions

Permissions policies (1/5) Info

You can attach up to 10 managed policies.

Policy name	Type	Attached entities
AmazonEC2FullAccess	AWS managed	1
AmazonS3ReadOnlyAccess	AWS managed	1
AutoScalingFullAccess	AWS managed	1
CloudWatchAgentServerPolicy	AWS managed	1
ElasticLoadBalancingFullAccess	AWS managed	1

Screenshot of the AWS IAM console showing the 'ec2-role' details page.

Permissions

Permissions policies (1/5) Info

You can attach up to 10 managed policies.

Policy name	Type	Attached entities
AmazonEC2FullAccess	AWS managed	1
AmazonS3ReadOnlyAccess	AWS managed	1
AutoScalingFullAccess	AWS managed	1
CloudWatchAgentServerPolicy	AWS managed	1
ElasticLoadBalancingFullAccess	AWS managed	1

Permissions boundary (not set)

Screenshot of the Terraform code editor showing the 'ec2_autoscaling.tf' file.

```

resource "aws_launch_template" "app" {
  name_prefix = "app-launch-template"
  image_id    = "ami-0db21ccaeffcd686" # Amazon Linux AMI
  instance_type = "t2.micro" # Instance type (cost-effective)

  # Attach the IAM role to EC2 instances
  iam_instance_profile {
    name = aws_iam_instance_profile.ec2_instance_profile.name
  }

  network_interfaces {
    security_groups = [aws_security_group.app_sg.id]
  }

  user_data = base64encode(<<EOF
#!/bin/bash
sudo dnf update -y
sudo dnf install -y nginx amazon-cloudwatch-agent
sudo systemctl start nginx
sudo systemctl enable nginx

# Configure the CloudWatch Agent
sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.json << 'EOF2'
{
  "metrics": {
    "append_dimensions": [
      {"AutoScalingGroupName": "AppASGInstance"}
    ],
    "metrics_collected": [
      {
        "measurement": [
          "mem_used_percent"
        ],
        "metrics_collection_interval": 60
      }
    ]
  }
}
EOF
)
}

```

This screenshot shows the CodeGPT IDE interface with the Terraform extension. The left sidebar displays a file tree with files like `ec2_autoscaling.tf`, `iam_roles.tf`, and `main.tf`. The main editor area contains Terraform code for configuring CloudWatch Metrics collection:

```
resource "aws_launch_template" "app" {
    user_data = base64encode(<<-EOF
        # Configure the CloudWatch Agent
        sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.json << 'EOF2'
        {
            "metrics": [
                "append_dimensions": [
                    "AutoScalingGroupName": "AppASGInstance"
                ],
                "metrics_collected": [
                    "mem": [
                        "measurement": [
                            "mem_used_percent"
                        ],
                        "metrics_collection_interval": 60
                    ],
                    "disk": [
                        "measurement": [
                            "used_percent"
                        ],
                        "resources": [
                            "/"
                        ],
                        "metrics_collection_interval": 60
                    ]
                ]
            ]
        }
        EOF2

        # Start the CloudWatch Agent
        sudo /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -a start -c file:/opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.json
    EOF
}
```

The status bar at the bottom indicates the code is in `Terraform` mode.

This screenshot shows the CodeGPT IDE interface with the Terraform extension. The left sidebar displays a file tree with files like `ec2_autoscaling.tf`, `iam_roles.tf`, and `main.tf`. The main editor area contains Terraform code for an Auto Scaling Group (ASG) configuration:

```
resource "aws_autoscaling_group" "app_asg" {
    desired_capacity      = 2
    min_size              = 1
    max_size              = 5
    vpc_zone_identifier   = [aws_subnet.public_subnet_1.id, aws_subnet.public_subnet_2.id]
    target_group_arns     = [aws_lb_target_group.primary.arn] # Attach to ALB Target Group

    # Replacement of instances
    instance_refresh {
        strategy = "Rolling"
        preferences {
            min_healthy_percentage = 0
            instance_warmup        = 120
        }
    }

    health_check_grace_period = 300 # Grace period
    health_check_type         = "ELB" # Change to ELB-based health check

    launch_template {
        id           = aws_launch_template.app.id
        version      = "$Latest"
    }

    tag {
        key          = "Name"
        value         = "AppASGInstance"
        propagate_at_launch = true
    }

    lifecycle {
        create_before_destroy = true
    }
}
```

The status bar at the bottom indicates the code is in `Terraform` mode.

This screenshot shows the CodeGPT IDE interface with the Terraform extension. The left sidebar displays a file tree with files like `ec2_autoscaling.tf`, `iam_roles.tf`, and `main.tf`. The main editor area contains Terraform code for a secondary AWS Launch Template:

```
resource "aws_launch_template" "secondary_app" {
    provider      = aws.west
    name_prefix   = "secondary-app-launch-template"
    image_id      = "ami-07c5ed849e59d65" # Amazon Linux AMI
    instance_type = "t2.micro"

    # Attach the IAM role to EC2 instances
    iam_instance_profile {
        name = aws_iam_instance_profile.ec2_instance_profile.name
    }

    network_interfaces {
        security_groups = [aws_security_group.west_app_sg.id]
    }

    user_data = base64encode(<<-EOF
        #!/bin/bash
        sudo dnf update -y
        sudo dnf install -y nginx amazon-cloudwatch-agent
        sudo systemctl start nginx
        sudo systemctl enable nginx

        # configure the CloudWatch Agent
        sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.json << 'EOF2'
        {
            "metrics": [
                "append_dimensions": [
                    "AutoScalingGroupName": "SecondaryAppASGInstance"
                ],
                "metrics_collected": [
                    "mem": [
                        "measurement": [
                            "mem_used_percent"
                        ],
                        "metrics_collection_interval": 60
                    ]
                ]
            ]
        }
        EOF2
    }
}
```

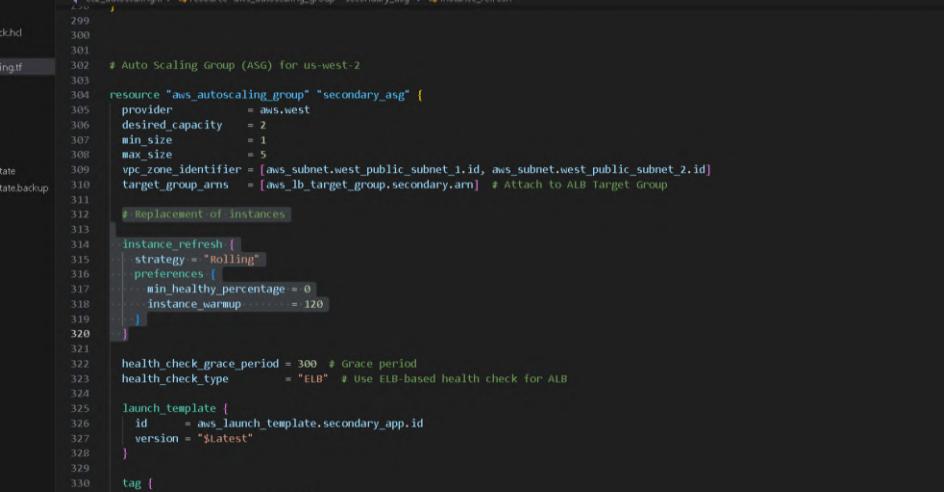
The status bar at the bottom indicates the code is in `Terraform` mode.

The screenshot shows the CodeGPT IDE interface with a Terraform configuration file open. The left sidebar displays the project structure with files like `EXPLORER`, `TERRAFORM`, `alb.tf`, `ec2_autoscaling.tf`, `iam_roles.tf`, `main.tf`, `network.tf`, `route53.tf`, `s3_bucket.tf`, `terraform.tfstate`, and `terraform.tfstate.backup`. The main editor area shows the following Terraform code:

```
resource "aws_launch_template" "secondary_app" {
  user_data = base64encode(<><>)
  sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-cloudwatch-agent.json << 'EOF2'
  {
    "metrics": [
      "append_dimensions": {
        "AutoScalingGroupName": "SecondaryAppASGInstance"
      }
    ],
    "metrics_collected": {
      "mem": {
        "measurement": [
          "mem_used_percent"
        ],
        "metrics_collection_interval": 60
      },
      "disk": {
        "measurement": [
          "used_percent"
        ],
        "resources": [
          "/"
        ],
        "metrics_collection_interval": 60
      }
    }
  }
  EOF2

  # Start the CloudWatch Agent
  sudo /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-agent-ctl -a start -c file:/opt/aws/amazon-cloudwatch-agent/etc/amaz
EOF

lifecycle {
  create_before_destroy = true
}
```



The screenshot shows the Visual Studio Code interface with the following details:

- File Explorer:** Shows the project structure with files like `EXPLORER`, `TERRAFORM`, and `ec2_autoscaling.tf`.
- Code Editor:** The main editor area displays the `ec2_autoscaling.tf` Terraform configuration file.
- Terminal:** A terminal window is visible at the bottom left.
- Status Bar:** Shows the current line (Ln 320), column (Col 4), and character count (177 selected). It also indicates spaces (Spaces: 2), UTF-8 encoding (UTF-8), and CRLF line endings (CRLF).
- Bottom Bar:** Includes icons for outline, timeline, and other development tools.

```
resource "aws_autoscaling_group" "secondary_asg" {
  provider           = aws.west
  desired_capacity   = 2
  min_size           = 1
  max_size           = 5
  vpc_zone_identifier = [aws_subnet.west_public_subnet_1.id, aws_subnet.west_public_subnet_2.id]
  target_group_arns  = [aws_lb_target_group.secondary.arn] # Attach to ALB Target Group

  # Replacement of instances
  instance_refresh {
    strategy = "Rolling"
    preferences {
      min_healthy_percentage = 0
      instance_warmup        = 120
    }
  }

  health_check_grace_period = 300 # Grace period
  health_check_type         = "ELB" # Use ELB-based health check for ALB

  launch_template {
    id      = aws_launch_template.secondary_app.id
    version = "$Latest"
  }

  tag {
    key     = "Name"
    value   = "SecondaryAppASGInstance"
    propagate_at_launch = true
  }
}
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -target=aws_autoscaling_group.app_asg -target=aws_autoscaling_group.secondary_asg -target=aws_launch_template.app -target=aws_launch_template.secondary_app
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b683b]
aws_iam_role.ec2_role: Refreshing state... [id=ec2-role]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_subnet.west_public_subnet_2: Refreshing state... [id=subnet-05cf3fc2aa1e38da]
aws_subnet.west_public_subnet_1: Refreshing state... [id=subnet-0a1940534af9cb413c]
aws_security_group.west_app_sg: Refreshing state... [id=sg-03d04109977fd80d0]
aws_lb_target_group.secondary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba4401c8967]
aws_launch_template.secondary_app: Refreshing state... [id=lt-0c625394b30f2746d]
aws_autoscaling_group.secondary_asg: Refreshing state... [id=terraform-20241030094126584900000004]
aws_subnet.public_subnet_2: Refreshing state... [id=subnet-0d89362f71d4a5df]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-04b29a39893e7944d]
aws_lb_target_group.primary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:targetgroup/primary-target-group/a9447d338de6a3e5]
aws_security_group.app_sg: Refreshing state... [id=sg-0d072a28cc779890cb]
aws_launch_template.app: Refreshing state... [id=lt-074d209bf44eaba96]
aws_autoscaling_group.app_asg: Refreshing state... [id=terraform-20241030094125457300000004]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
- = destroy

Terraform will perform the following actions:

# aws_autoscaling_group.app_asg will be destroyed
- resource "aws_autoscaling_group" "app_asg" {
  - arn = "arn:aws:autoscaling:us-east-1:134575801911:autoScalingGroup:71dfa62f-4579-4e12-8e71-ff9cca91e452:autoScalingGroup"
  pName=terraform-20241030094125457300000004" >> null
    - availability_zones = [
      - "us-east-1a",
      - "us-east-1b",
    ] >> null
    - capacity_rebalance = false >> null
    - default_cooldown = 300 >> null
    - default_instance_warmup = 0 >> null
    - desired_capacity = 2 >> null
    - enabled_metrics = [] >> null
    - force_delete = false >> null
    - force_delete_warm_pool = false >> null
}
```

```
Command Prompt x + ×
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 5m0s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 5m1s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 5m10s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 5m11s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 5m20s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 5m21s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 5m30s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 5m31s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 5m40s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 5m41s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 5m50s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 5m51s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m0s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 6m1s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m10s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094125457300000004, 6m11s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m20s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m21s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m30s elapsed]
aws autoscaling_group.app_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m31s elapsed]
aws_launch_template.app: Destruction complete after 6m36s
aws_launch_template.app: Destruction complete after 0s
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m40s elapsed]
aws autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241030094126584900000004, 6m50s elapsed]
aws autoscaling_group.secondary_asg: Destruction complete after 6m51s
aws_launch_template.secondary_app: Destroying... [id=lt-0c625394b30f2746d]
aws_launch_template.secondary_app: Destruction complete after 0s

Warning: Applied changes may be incomplete

The plan was created with the -target option in effect, so some changes requested in the configuration may have been ignored and the output values may not be fully updated. Run the following command to verify that no other changes are pending:
  terraform plan

Note that the -target option is not suitable for routine use, and is provided only for exceptional situations such as recovering from errors or mistakes, or when Terraform specifically suggests to use it as part of an error message.

Destroy complete! resources: 8 destroyed.
```

```
Command Prompt x + ×
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see any changes that are required for your infrastructure. All Terraform commands should now work.

If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory. If you forget, other commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws_route53_zone.main: Refreshing state... [id=z0743615150XJV7STD1T]
aws iam_role.ec2_role: Refreshing state... [id=ec2-role]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b904ff753dc]
aws_s3.bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_s3.bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_iam_role_policy_attachment.ec2_cloudwatch_agent: Refreshing state... [id=ec2-role-20241031075603479000000001]
aws_iam_role_policy_attachment.ec2_full_access: Refreshing state... [id=ec2-role-20241030041783779300000003]
aws_iam_role_policy_attachment.ec2_s3_readonly: Refreshing state... [id=ec2-role-20241026083002429100000001]
aws_iam_role_policy_attachment.elb_full_access: Refreshing state... [id=ec2-role-20241030041783813700000004]
aws_iam_role_policy_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-20241030041703774400000002]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_subnet.west_public_subnet.1: Refreshing state... [id=subnet-0a1940534f9cb413c]
aws_subnet.west_private_subnet.1: Refreshing state... [id=subnet-0de0b30d268f143ac]
aws_route_table.west_public_rt: Refreshing state... [id=rt-0b4fcfa31ff87e691088028949]
aws_lb_target_group.secondary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba4401c8967]
aws_security_group.west_app_sg: Refreshing state... [id=sg-03d04109977fd80d0]
aws_internet_gateway.west_igw: Refreshing state... [id=igw-093c80565bfdb459c]
aws_subnet.west_public_subnet.2: Refreshing state... [id=subnet-05cf33fc2aae38da]
aws_route_table_association.west_public_subnet_assoc.1: Refreshing state... [id=rtbassoc-05dc480a71aa04b77]
aws_route.west_internet_access: Refreshing state... [id=rtb-0b4fcfa31ff87e691088028949]
aws_route_table_association.west_public_subnet_assoc.2: Refreshing state... [id=rtbassoc-08e5ffd4ebe4bee8]
aws_lb.secondary_alb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:loadbalancer/app/secondary-alb/4b7a0a637b386685]
```

```
Command Prompt x + ×
}
}

Plan: 8 to add, 0 to change, 0 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform apply -auto-approve
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws_route53_zone.main: Refreshing state... [id=z0743615150XJV7STD1T]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b904ff753dc]
aws iam_role.ec2_role: Refreshing state... [id=ec2-role]
aws_s3.bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3.bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_internet_gateway.west_igw: Refreshing state... [id=igw-093c80565bfdb459c]
aws_subnet.west_public_subnet.2: Refreshing state... [id=subnet-05cf33fc2aae38da]
aws_subnet.west_private_subnet.1: Refreshing state... [id=subnet-0de0b30d268f143ac]
aws_route_table.west_public_rt: Refreshing state... [id=rt-0b4fcfa31ff87e691088028949]
aws_security_group.west_app_sg: Refreshing state... [id=sg-03d04109977fd80d0]
aws_lb_target_group.secondary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba4401c8967]
aws_iam_role_policy_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-20241030041783774400000002]
aws_iam_role_policy_attachment.elb_full_access: Refreshing state... [id=ec2-role-20241030041703812700000004]
aws_iam_role_policy_attachment.ec2_s3_readonly: Refreshing state... [id=ec2-role-20241026083002429100000001]
aws_iam_role_policy_attachment.ec2_cloudwatch_agent: Refreshing state... [id=ec2-role-20241031075603479000000001]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_iam_role_policy_attachment.ec2_full_access: Refreshing state... [id=ec2-role-20241030041703779300000003]
aws_iam_role_policy_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-20241030041703774400000002]
aws_iam_role_policy_attachment.elb_full_access: Refreshing state... [id=ec2-role-20241030041703812700000004]
aws_iam_role_policy_attachment.ec2_s3_readonly: Refreshing state... [id=ec2-role-20241026083002429100000001]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_iam_role_policy_attachment.ec2_full_access: Refreshing state... [id=ec2-role-20241030041703779300000003]
aws_route_table_association.west_public_subnet_assoc.1: Refreshing state... [id=rtbassoc-05dc480a71aa04b77]
aws_route_table_association.west_public_subnet_assoc.2: Refreshing state... [id=rtbassoc-08e5ffd4ebe4bee8]
aws_lb.secondary_alb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:loadbalancer/app/secondary-alb/4b7a0a637b386685]
aws_route53_health_check.secondary_health_check: Refreshing state... [id=secce-87-32b1-4d77-9ef1-e8cd7faad582]
aws_route53_record.secondary: Refreshing state... [id=z0743615150XJV7STD1T_app.myinfra.local_A_Secondary]
aws_lb_listener.secondary_http: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:listener/app/secondary-alb/4b7a0a637b386685/938584833b8318cd]
aws_subnet.private_subnet.1: Refreshing state... [id=subnet-074ccceeedaaaf778f]
aws_subnet.public_subnet.1: Refreshing state... [id=subnet-04b29a39893e794d]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc0f97df205e5]
```

```

        }
    }

    + metadata_options (known after apply)

    + network_interfaces {
        + security_groups = [
            + "sg-03d04109977fd80d0",
        ]
    }

    + tag_specifications {
        + resource_type = "instance"
        + tags = [
            + "Name" = "SecondaryAppInstance"
        ]
    }
}

Plan: 8 to add, 0 to change, 0 to destroy.
aws_launch_template.app: Creating...
aws_launch_template.secondary_app: Creating...
aws_launch_template.secondary_app: Creation complete after 0s [id=lt-060730da2d5f4f76a]
aws_autoscaling_group.secondary_asg: Creating...
aws_launch_template.app: Creation complete after 0s [id=lt-08a5c2c9fe206473d]
aws_autoscaling_group.app_asg: Creating...
aws_autoscaling_group.secondary_asg: Still creating... [10s elapsed]
aws_autoscaling_group.app_asg: Still creating... [10s elapsed]
aws_autoscaling_group.secondary_asg: Creation complete after 16s [id=terraform-20241031095009964500000003]
aws_autoscaling_policy.west_scale_down: Creating...
aws_autoscaling_policy.west_scale_up: Creating...
aws_autoscaling_policy.west_scale_down: Creation complete after 0s [id=west-scale-down]
aws_autoscaling_policy.west_scale_up: Creation complete after 0s [id=west-scale-up]
aws_autoscaling_group.app_asg: Creation complete after 17s [id=terraform-20241031095010218800000003]
aws_autoscaling_policy.scale_down: Creating...
aws_autoscaling_policy.scale_up: Creating...
aws_autoscaling_policy.scale_down: Creation complete after 1s [id=scale-down]
aws_autoscaling_policy.scale_up: Creation complete after 1s [id=scale-up]

Apply complete! Resources: 8 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

```

The screenshot shows the AWS CloudWatch Metrics Insights interface. A query is being run against the 'CloudWatch Metrics' dataset. The query is:

```

    @m1:1 SELECT * FROM CloudWatchMetrics WHERE metric.name = 'CPUUtilization'
    @m1:2 AND metric.namespace = 'AWS/EC2'
    @m1:3 AND metric.stats = 'Average'
    @m1:4 AND metric.dimensions.instanceId = 'i-0d7733c4d268fc705'
    @m1:5 AND metric.stats = 'Average'
    @m1:6 AND metric.dimensions.instanceId = 'i-00347e579d3947ab7'
    @m1:7 AND metric.stats = 'Average'
    @m1:8 AND metric.dimensions.instanceId = 'i-0028fb710309986f1'
    @m1:9 AND metric.stats = 'Average'
    @m1:10 AND metric.dimensions.instanceId = 'i-007062986c229675c'
    @m1:11 AND metric.stats = 'Average'
    @m1:12 AND metric.dimensions.instanceId = 'i-0e84fc1a84d2e0060'
    @m1:13 AND metric.stats = 'Average'
    @m1:14 AND metric.dimensions.instanceId = 'i-058f21b32ea3690fb'
  
```

The results show CPU utilization data for multiple EC2 instances over time.

The screenshot shows the AWS CloudWatch Metrics Insights interface. A query is being run against the 'CloudWatch Metrics' dataset. The query is:

```

    @m1:1 SELECT * FROM CloudWatchMetrics WHERE metric.name = 'CPUUtilization'
    @m1:2 AND metric.namespace = 'AWS/EC2'
    @m1:3 AND metric.stats = 'Average'
    @m1:4 AND metric.dimensions.instanceId = 'i-0d7733c4d268fc705'
    @m1:5 AND metric.stats = 'Average'
    @m1:6 AND metric.dimensions.instanceId = 'i-00347e579d3947ab7'
    @m1:7 AND metric.stats = 'Average'
    @m1:8 AND metric.dimensions.instanceId = 'i-0028fb710309986f1'
    @m1:9 AND metric.stats = 'Average'
    @m1:10 AND metric.dimensions.instanceId = 'i-007062986c229675c'
    @m1:11 AND metric.stats = 'Average'
    @m1:12 AND metric.dimensions.instanceId = 'i-0e84fc1a84d2e0060'
    @m1:13 AND metric.stats = 'Average'
    @m1:14 AND metric.dimensions.instanceId = 'i-058f21b32ea3690fb'
  
```

The results show CPU utilization data for multiple EC2 instances over time.

The screenshot shows the VS Code interface with the Terraform extension active. The left sidebar displays the project structure under 'EXPLORER' with files like '.terraform', 'cloudwatch_alarms.tf', 'iam_roles.tf', etc. The main editor area contains the following Terraform code:

```
cloudwatch_alarms.tf > ...
resource "aws_cloudwatch_metric_alarm" "high_cpu_utilization" {
  alarm_name        = "HighCPUUtilizationAlarm"
  comparison_operator = "greaterthanthreshold"
  evaluation_periods = 2
  metric_name       = "CPUUtilization"
  namespace          = "AWS/EC2"
  period             = 300
  statistic          = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description  = "This alarm triggers if CPU utilization exceeds 80%."
  dimensions = [
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  ]
  actions_enabled     = true
  alarm_actions       = [aws_sns_topic.alert_topic.arn]
}

# CloudWatch Alarm for Memory Utilization

resource "aws_cloudwatch_metric_alarm" "high_memory_utilization" {
  alarm_name        = "HighMemoryUtilizationAlarm"
  comparison_operator = "Greaterthanthreshold"
  evaluation_periods = 2
  metric_name       = "mem_used_percent"
  namespace          = "CMQAgent"
  period             = 300
  statistic          = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description  = "This alarm triggers if memory utilization exceeds 80%."
  dimensions = [
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  ]
}
```

The screenshot shows the VS Code interface with the Terraform extension active. The left sidebar displays the project structure under 'EXPLORER' with files like '.terraform', 'cloudwatch_alarms.tf', 'iam_roles.tf', etc. The main editor area continues the Terraform code from the previous screenshot:

```
resource "aws_cloudwatch_metric_alarm" "high_memory_utilization" {
  namespace          = "CMQAgent"
  period             = 300
  statistic          = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description  = "This alarm triggers if memory utilization exceeds 80%."
  dimensions = [
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  ]
  actions_enabled     = true
  alarm_actions       = [aws_sns_topic.alert_topic.arn]
}

# CloudWatch Alarm for Disk Utilization

resource "aws_cloudwatch_metric_alarm" "high_disk_utilization" {
  alarm_name        = "HighDiskUtilizationAlarm"
  comparison_operator = "Greaterthanthreshold"
  evaluation_periods = 2
  metric_name       = "disk_used_percent"
  namespace          = "CMQAgent"
  period             = 300
  statistic          = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description  = "This alarm triggers if disk utilization exceeds 80%."
  dimensions = [
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  ]
  actions_enabled     = true
  alarm_actions       = [aws_sns_topic.alert_topic.arn]
}
```

The screenshot shows the VS Code interface with the Terraform extension active. The left sidebar displays the project structure under 'EXPLORER' with files like '.terraform', 'cloudwatch_alarms.tf', 'iam_roles.tf', etc. The main editor area continues the Terraform code from the previous screenshots and adds an SNS topic:

```
# CloudWatch Alarm for Disk Utilization

resource "aws_cloudwatch_metric_alarm" "high_disk_utilization" {
  alarm_name        = "HighDiskUtilizationAlarm"
  comparison_operator = "Greaterthanthreshold"
  evaluation_periods = 2
  metric_name       = "disk_used_percent"
  namespace          = "CMQAgent"
  period             = 300
  statistic          = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description  = "This alarm triggers if disk utilization exceeds 80%."
  dimensions = [
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  ]
  actions_enabled     = true
  alarm_actions       = [aws_sns_topic.alert_topic.arn]
}

# SNS Topic for Alarm Notifications in us-east-1

resource "aws sns topic" "alert_topic" {
  name = "EC2AlertTopic"
}

resource "aws sns topic subscription" "email_subscription" {
  topic_arn = aws sns topic.alert_topic.arn
  protocol = "email"
  endpoint = "vivekvash1507@gmail.com"
}
```

```
File Edit Selection View Go Run Terminal Help ⏪ ⏩ 🔍 Terraform

EXPLORER ... cloudwatch_alarms.tf x iam_roles.tf
TERRAFORM > .terraform & cloudwatch_alarms > ...
  > .terraform.lock.hcl
  & alb.tf
  & cloudwatch_alarms.tf
  & ec2_autoscaling.tf
  & iam_roles.tf
  & main.tf
  & network.tf
  & route53.tf
  & s3_bucket.tf
  & terraform.state
  & terraform.state.backup

CODEGPT

cloudwatch_alarms.tf
83 # us-west-2
84
85
86 # CloudWatch Alarm for CPU Utilization
87
88 resource "aws_cloudwatch_metric_alarm" "west_high_cpu_utilization" {
89   provider      = aws.west
90   alarm_name    = "WestHighCPUUtilizationAlarm"
91   comparison_operator = "GreaterThanOrEqualToThreshold"
92   evaluation_periods = 2
93   metric_name   = "CPUUtilization"
94   namespace     = "AWS/EC2"
95   period        = 300
96   statistic     = "Average"
97   threshold     = 80 # Set threshold based on your requirements
98   alarm_description = "This alarm triggers if CPU utilization exceeds 80%."
99   dimensions = [
100     {AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name}
101   ]
102   actions_enabled = true
103   alarm_actions   = [aws sns topic west_alert_topic.arn]
104 }

105
106
107 # CloudWatch Alarm for Memory Utilization
108
109 resource "aws_cloudwatch_metric_alarm" "west_high_memory_utilization" {
110   provider      = aws.west
111   alarm_name    = "WestHighMemoryUtilizationAlarm"
112   comparison_operator = "GreaterThanOrEqualToThreshold"
113   evaluation_periods = 2
114   metric_name   = "MemUsedPercent"
115   namespace     = "CMagent"
116   period        = 300
117   statistic     = "Average"
118   threshold     = 80 # Set threshold based on your requirements
119   alarm_description = "This alarm triggers if memory utilization exceeds 80%."
120   dimensions = [
121     {AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name}
122   ]
123   actions_enabled = true
124   alarm_actions   = [aws sns topic west_alert_topic.arn]
125 }

126
127
128
129
130 # CloudWatch Alarm for Disk Utilization
131
132 resource "aws_cloudwatch_metric_alarm" "west_high_disk_utilization" {
133   provider      = aws.west
134   alarm_name    = "WestHighDiskUtilizationAlarm"
135   comparison_operator = "GreaterThanOrEqualToThreshold"
136   evaluation_periods = 2
137   .
138   .
139   .
140   .
141   .
142   .
143   .
144   .
145   .
146   .
147   .
148   .

Ln 83, Col 12 (11 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT
```

```
File Edit Selection View Go Run Terminal Help ⏪ ⏩ 🔍 Terraform

EXPLORER ... cloudwatch_alarms.tf x iam_roles.tf
TERRAFORM > .terraform & cloudwatch_alarms > ...
  > .terraform.lock.hcl
  & alb.tf
  & cloudwatch_alarms.tf
  & ec2_autoscaling.tf
  & iam_roles.tf
  & main.tf
  & network.tf
  & route53.tf
  & s3_bucket.tf
  & terraform.state
  & terraform.state.backup

CODEGPT

cloudwatch_alarms.tf
88 resource "aws_cloudwatch_metric_alarm" "west_high_cpu_utilization" {
89   provider      = aws.west
90   alarm_name    = "WestHighCPUUtilizationAlarm"
91   comparison_operator = "GreaterThanOrEqualToThreshold"
92   evaluation_periods = 2
93   metric_name   = "CPUUtilization"
94   namespace     = "AWS/EC2"
95   period        = 300
96   statistic     = "Average"
97   threshold     = 80 # Set threshold based on your requirements
98   alarm_description = "This alarm triggers if CPU utilization exceeds 80%."
99   dimensions = [
100     {AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name}
101   ]
102   actions_enabled = true
103   alarm_actions   = [aws sns topic west_alert_topic.arn]
104 }

105
106
107 # CloudWatch Alarm for Memory Utilization
108
109 resource "aws_cloudwatch_metric_alarm" "west_high_memory_utilization" {
110   provider      = aws.west
111   alarm_name    = "WestHighMemoryUtilizationAlarm"
112   comparison_operator = "GreaterThanOrEqualToThreshold"
113   evaluation_periods = 2
114   metric_name   = "MemUsedPercent"
115   namespace     = "CMagent"
116   period        = 300
117   statistic     = "Average"
118   threshold     = 80 # Set threshold based on your requirements
119   alarm_description = "This alarm triggers if memory utilization exceeds 80%."
120   dimensions = [
121     {AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name}
122   ]
123   actions_enabled = true
124   alarm_actions   = [aws sns topic west_alert_topic.arn]
125 }

126
127
128
129
130 # CloudWatch Alarm for Disk Utilization
131
132 resource "aws_cloudwatch_metric_alarm" "west_high_disk_utilization" {
133   provider      = aws.west
134   alarm_name    = "WestHighDiskUtilizationAlarm"
135   comparison_operator = "GreaterThanOrEqualToThreshold"
136   evaluation_periods = 2
137   .
138   .
139   .
140   .
141   .
142   .
143   .
144   .
145   .
146   .
147   .
148   .

Ln 83, Col 12 (11 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT
```

```
File Edit Selection View Go Run Terminal Help ⏪ ⏩ 🔍 Terraform

EXPLORER ... cloudwatch_alarms.tf x iam_roles.tf
TERRAFORM > .terraform & cloudwatch_alarms > ...
  > .terraform.lock.hcl
  & alb.tf
  & cloudwatch_alarms.tf
  & ec2_autoscaling.tf
  & iam_roles.tf
  & main.tf
  & network.tf
  & route53.tf
  & s3_bucket.tf
  & terraform.state
  & terraform.state.backup

CODEGPT

cloudwatch_alarms.tf
129
130 # CloudWatch Alarm for Disk Utilization
131
132 resource "aws_cloudwatch_metric_alarm" "west_high_disk_utilization" {
133   provider      = aws.west
134   alarm_name    = "WestHighDiskUtilizationAlarm"
135   comparison_operator = "GreaterThanOrEqualToThreshold"
136   evaluation_periods = 2
137   metric_name   = "diskUsedPercent"
138   namespace     = "CMagent"
139   period        = 300
140   statistic     = "Average"
141   threshold     = 80 # Set threshold based on your requirements
142   alarm_description = "This alarm triggers if disk utilization exceeds 80%."
143   dimensions = [
144     {AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name}
145   ]
146   actions_enabled = true
147   alarm_actions   = [aws sns topic west_alert_topic.arn]
148 }

149
150
151
152 # SNS Topic for Alarm Notifications in us-west-2
153
154 resource "aws sns topic" "west_alert_topic" {
155   provider = aws.west
156   name     = "WestEC2AlertTopic"
157 }
158
159
160 resource "aws sns topic subscription" "west_email_subscription" {
161   provider = aws.west
162   topic_arn = aws sns topic west_alert_topic.arn
163   protocol = "email"
164   endpoint = "vivekvashis158@gmail.com"
165 }

Ln 83, Col 12 (11 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT
```

The screenshot shows the Terraform Cloud interface with a code editor displaying a Terraform configuration file. The file, named `iam_roles.tf`, defines two AWS IAM roles: `ec2_cloudwatch_agent` and `ec2_sns_publish_policy`. The `ec2_cloudwatch_agent` role is attached to the `aws_iam_policy_attachment` resource, which also specifies the `ec2_sns_publish_policy` as its policy. The `ec2_sns_publish_policy` grants the `sns:Publish` action on specific SNS topics.

```
resource "aws_iam_role_policy_attachment" "ec2_cloudwatch_agent" {
  role        = aws_iam_role.ec2_role.name
  policy_arn = "arn:aws:policy/CloudWatchAgentServerPolicy"
}

# Allow EC2 instances to publish to both SNS topics for alerts
resource "aws_iam_role_policy" "ec2_sns_publish_policy" {
  name      = "ec2-sns-publish-policy"
  role      = aws_iam_role.ec2_role.name

  policy = jsonencode({
    "Version": "2012-10-17",
    "Statement": [
      {
        "Effect": "Allow",
        "Action": "sns:Publish",
        "Resource": [
          "${aws sns topic.alert_topic.arn}", # SNS topic in us-east-1
          "${aws sns topic.west_alert_topic.arn}" # SNS topic in us-west-2
        ]
      }
    ]
  })
}
```

The screenshot shows the AWS Identity and Access Management (IAM) console. The left sidebar is titled "Identity and Access Management (IAM)" and includes sections for Dashboard, Access management (User groups, Users, Roles, Policies, Identity providers, Account settings), and Access reports (Access Analyzer, External access, Unused access, Analyzer settings). The main content area is titled "Permissions policies (1/6)" and displays a table of policies. The table has columns for Policy name, Type, and Attached entities. It lists six policies: "AmazonEC2FullAccess" (AWS managed, 1 entity), "AmazonS3ReadOnlyAccess" (AWS managed, 1 entity), "AutoScalingFullAccess" (AWS managed, 1 entity), "CloudWatchAgentServerPolicy" (AWS managed, 1 entity), "ec2-sns-publish-policy" (Customer inline, 0 entities, selected), and "ElasticLoadBalancingFullAccess" (AWS managed, 1 entity). Below the table, there are sections for "Permissions boundary (not set)" and "Generate policy based on CloudTrail events".

Policy name	Type	Attached entities
AmazonEC2FullAccess	AWS managed	1
AmazonS3ReadOnlyAccess	AWS managed	1
AutoScalingFullAccess	AWS managed	1
CloudWatchAgentServerPolicy	AWS managed	1
ec2-sns-publish-policy	Customer inline	0
ElasticLoadBalancingFullAccess	AWS managed	1

The screenshot shows the AWS CloudWatch Alarms interface. The left sidebar includes 'CloudWatch' and 'Metrics' sections. The main area displays four alarms:

Name	State	Last state update (UTC)	Conditions	Actions
HighCPUUtilizationAlarm	OK	2024-10-31 10:41:03	CPUUtilization > 80 for 2 datapoints within 10 minutes	Actions enabled
HighDiskUtilizationAlarm	Insufficient data	2024-10-31 10:40:26	disk_used_percent > 80 for 2 datapoints within 10 minutes	Actions enabled
HighMemoryUtilizationAlarm	Insufficient data	2024-10-31 10:40:26	mem_used_percent > 80 for 2 datapoints within 10 minutes	Actions enabled
MySimpleBillingAlarm	Insufficient data	2023-09-07 05:08:59	EstimatedCharges > 50 for 1 datapoints within 6 hours	Actions enabled

The screenshot shows the AWS CloudWatch Alarms console. On the left, there's a navigation sidebar with options like Dashboards, Alarms, Logs, Metrics, X-Ray traces, Events, Application Signals, Network monitoring, and Insights. The main area displays three alarms:

Name	State	Last state update (UTC)	Conditions
WestHighCPUUtilizationAlarm	OK	2024-10-31 10:41:26	CPUUtilization > 80 for 2 datapoints within 10 minutes
WestHighDiskUtilizationAlarm	Insufficient data	2024-10-31 10:40:25	disk_used_percent > 80 for 2 datapoints within 10 minutes
WestHighMemoryUtilizationAlarm	Insufficient data	2024-10-31 10:40:25	mem_used_percent > 80 for 2 datapoints within 10 minutes

The screenshot shows an email from AWS Notifications (no-reply@sns.amazonaws.com) to me. The subject is "AWS Notifications". The email body contains the following text:

You have chosen to subscribe to the topic:
arn:aws:sns:us-east-1:134575801911:EC2AlertTopic

To confirm this subscription, click or visit the link below (if this was in error no action is necessary):
[Confirm subscription](#)

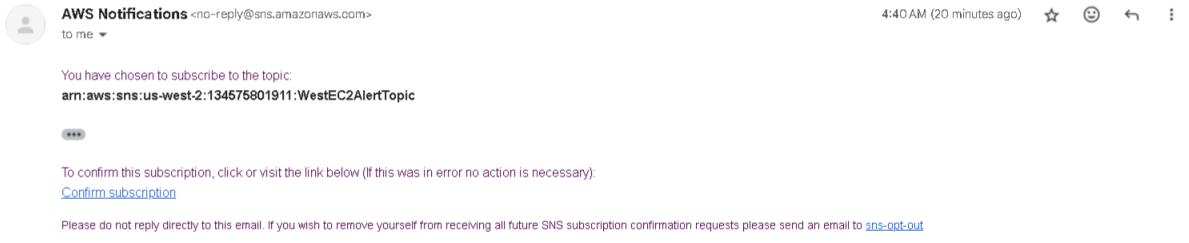
Please do not reply directly to this email. If you wish to remove yourself from receiving all future SNS subscription confirmation requests please send an email to [sns-opt-out](#)

The screenshot shows the AWS SNS Topics console. On the left, there's a navigation sidebar with options like Dashboard, Topics, Subscriptions, and Mobile (Push notifications, Text messaging (SMS)). The main area displays the details for the EC2AlertTopic:

Name	Display name
EC2AlertTopic	-

Below the details, there's a "Subscriptions" tab showing one confirmed subscription:

ID	Endpoint	Status	Protocol
45bab4c9-67d8-4001-97e7-52...	vivekvash1507@gmail.com	Confirmed	EMAIL



ID	Endpoint	Status	Protocol
52f0ce9a-1e0a-4ce2-964d-20...	vivekvash1507@gmail.com	Confirmed	EMAIL

Parameterization and Validation

a) Overview

In this phase, I enhanced the flexibility and maintainability of the Terraform configurations by introducing variables and adding validation constraints. This adjustment allows easy modifications to configurations, making the infrastructure adaptable while ensuring inputs meet specific criteria. This centralized approach also enables consistent updates across configurations by managing values from a single source.

b) Defining Variables in variables.tf

To facilitate customization, I created a `variables.tf` file to define key variables that could be easily modified to adapt the infrastructure. This file includes primary and secondary AWS

regions, EC2 instance types, S3 bucket names for Terraform state and backups, Application Load Balancer (ALB) names, a Route 53 domain, and an email address for SNS notifications.

Additionally, I introduced validation rules for specific variables to enforce constraints:

- **Instance Type Validation:** Limits `instance_type` to "t2.micro", "t2.small", or "t3.micro", ensuring only compatible options are used.
- **Email Format Validation:** Enforces a valid format for the `notification_email` variable.

The following is the complete structure of `variables.tf`:

```
# AWS Regions
variable "aws_region" {
  description = "Primary AWS region"
  type        = string
  default     = "us-east-1"
}

variable "aws_west_region" {
  description = "Secondary AWS region"
  type        = string
  default     = "us-west-2"
}

# Instance Type with Validation
variable "instance_type" {
  description = "EC2 instance type"
  type        = string
  default     = "t2.micro"
  validation {
    condition      = contains(["t2.micro", "t2.small", "t3.micro"], var.instance_type)
    error_message = "The instance type must be t2.micro, t2.small, or t3.micro."
  }
}

# S3 Bucket Names
variable "terraform_state_bucket" {
  description = "S3 bucket name for storing Terraform state"
  type        = string
  default     = "aws-tf-backend-vv-bucket"
}

variable "backup_bucket" {
  description = "S3 bucket name for backups"
  type        = string
  default     = "terra-backup-vv-project-bucket"
}

# ALB Names
variable "primary_alb_name" {
  description = "Name of the primary ALB"
  type        = string
  default     = "primary-alb"
}
```

```

variable "secondary_alb_name" {
  description = "Name of the secondary ALB"
  type        = string
  default     = "secondary-alb"
}

# Route 53 Zone Domain
variable "route53_zone_domain" {
  description = "Domain name for Route 53 hosted zone"
  type        = string
  default     = "myinfra.local"
}

# Notification Email for SNS with Validation
variable "notification_email" {
  description = "Email address for SNS notifications"
  type        = string
  default     = "vivekvash1507@gmail.com"
  validation {
    condition      = length(regexall("^[a-zA-Z0-9._%+-]+@[a-zA-Z0-9.-]+\.\[a-zA-Z\]{2,}$", var.notification_email)) > 0
    error_message  = "The email format is invalid."
  }
}

```

Each variable's `description` attribute provides clarity on its purpose, making it easier for anyone to understand its role in the configuration.

c) Modifying alb.tf to Use Variables

With `variables.tf` in place, I proceeded to modify the `alb.tf` file to utilize these variables. This step involved replacing hardcoded values with variable references for increased flexibility and ease of maintenance.

1. Primary ALB Configuration (`us-east-1`)

The primary Application Load Balancer configuration now references

`primary_alb_name` for the ALB name, `aws_region` for the primary AWS region, and integrates the security group and subnets using previously defined resources. This setup enhances clarity and adaptability.

```

# Application Load Balancer (ALB) (us-east-1)
resource "aws_lb" "primary_alb" {
  name          = var.primary_alb_name
  internal      = false
  load_balancer_type = "application"
  security_groups = [aws_security_group.app_sg.id]
  subnets       = [aws_subnet.public_subnet_1.id,
                 aws_subnet.public_subnet_2.id]
  enable_deletion_protection = false
  idle_timeout           = 400

  tags = {
    Name = "PrimaryALB"
  }
}

```

```

# ALB Target Group for us-east-1
resource "aws_lb_target_group" "primary" {
    name      = "primary-target-group"
    port      = 80
    protocol = "HTTP"
    vpc_id   = aws_vpc.main_vpc.id

    health_check {
        path          = "/"
        interval     = 30
        timeout      = 5
        healthy_threshold = 2
        unhealthy_threshold = 2
    }

    tags = {
        Name = "PrimaryTargetGroup"
    }
}

# ALB Listener for HTTP traffic
resource "aws_lb_listener" "http" {
    load_balancer_arn = aws_lb.primary_alb.arn
    port            = 80
    protocol        = "HTTP"

    default_action {
        type          = "forward"
        target_group_arn = aws_lb_target_group.primary.arn
    }

    tags = {
        Name = "PrimaryALBLListener"
    }
}

```

2. Secondary ALB Configuration (us-west-2)

The secondary ALB configuration, located in us-west-2, references `secondary_alb_name` for the ALB name and `aws_west_region` for the AWS region, thus centralizing region and name parameters.

```

# Application Load Balancer (ALB) for us-west-2
resource "aws_lb" "secondary_alb" {
    provider      = aws.west
    name          = var.secondary_alb_name
    internal      = false
    load_balancer_type = "application"
    security_groups = [aws_security_group.west_app_sg.id]
    subnets       = [aws_subnet.west_public_subnet_1.id,
aws_subnet.west_public_subnet_2.id]
    enable_deletion_protection = false
    idle_timeout      = 400

    tags = {
        Name = "SecondaryALB"
    }
}

# ALB Target Group for us-west-2
resource "aws_lb_target_group" "secondary" {

```

```

provider = aws.west
name     = "secondary-target-group"
port      = 80
protocol = "HTTP"
vpc_id   = aws_vpc.west_vpc.id

health_check {
  path          = "/"
  interval      = 30
  timeout       = 5
  healthy_threshold = 2
  unhealthy_threshold = 2
}

tags = {
  Name = "SecondaryTargetGroup"
}
}

# ALB Listener for HTTP traffic in us-west-2
resource "aws_lb_listener" "secondary_http" {
  provider        = aws.west
  load_balancer_arn = aws_lb.secondary_alb.arn
  port           = 80
  protocol       = "HTTP"

  default_action {
    type            = "forward"
    target_group_arn = aws_lb_target_group.secondary.arn
  }

  tags = {
    Name = "SecondaryALBLListener"
  }
}

```

d) Updated cloudwatch_alarms.tf to Use Variables

In this step, I updated the `cloudwatch_alarms.tf` file to use the variables defined in `variables.tf`, enhancing flexibility and consistency across the Terraform configuration. Below is the updated configuration file.

```

# us-east-1

# CloudWatch Alarm for CPU Utilization
resource "aws_cloudwatch_metric_alarm" "high_cpu_utilization" {
  alarm_name          = "HighCPUUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUUtilization"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 80 # Set threshold to 80% CPU Utilization
  alarm_description    = "This alarm triggers if CPU utilization exceeds
80%."

```

```

dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
}
actions_enabled = true
alarm_actions   = [aws_sns_topic.alert_topic.arn]
}

# CloudWatch Alarm for Memory Utilization
resource "aws_cloudwatch_metric_alarm" "high_memory_utilization" {
    alarm_name          = "HighMemoryUtilizationAlarm"
    comparison_operator = "GreaterThanOrEqualToThreshold"
    evaluation_periods  = 2
    metric_name         = "mem_used_percent" # Assuming this is a custom
metric from CloudWatch Agent
    namespace           = "CWAgent"
    period              = 300
    statistic           = "Average"
    threshold           = 80 # Set threshold to 80% Memory Utilization
    alarm_description   = "This alarm triggers if memory utilization exceeds
80%."
    dimensions = {
        AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
    actions_enabled = true
    alarm_actions   = [aws_sns_topic.alert_topic.arn]
}

# CloudWatch Alarm for Disk Utilization
resource "aws_cloudwatch_metric_alarm" "high_disk_utilization" {
    alarm_name          = "HighDiskUtilizationAlarm"
    comparison_operator = "GreaterThanOrEqualToThreshold"
    evaluation_periods  = 2
    metric_name         = "disk_used_percent" # Custom metric from CloudWatch
Agent
    namespace           = "CWAgent"
    period              = 300
    statistic           = "Average"
    threshold           = 80 # Set threshold based on your requirements
    alarm_description   = "This alarm triggers if disk utilization exceeds
80%."
    dimensions = {
        AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
    actions_enabled = true
    alarm_actions   = [aws_sns_topic.alert_topic.arn]
}

# SNS Topic for Alarm Notifications in us-east-1
resource "aws sns topic" "alert_topic" {
    name = "EC2AlertTopic"
}

resource "aws sns topic subscription" "email_subscription" {
    topic_arn = aws_sns_topic.alert_topic.arn
    protocol = "email"
    endpoint = var.notification_email
}

# us-west-2

# CloudWatch Alarm for CPU Utilization

```

```

resource "aws_cloudwatch_metric_alarm" "west_high_cpu_utilization" {
  provider          = aws.west
  alarm_name        = "WestHighCPUUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods = 2
  metric_name       = "CPUUtilization"
  namespace         = "AWS/EC2"
  period            = 300
  statistic         = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description = "This alarm triggers if CPU utilization exceeds
80%."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
  }
  actions_enabled = true
  alarm_actions    = [aws_sns_topic.west_alert_topic.arn]
}

# CloudWatch Alarm for Memory Utilization
resource "aws_cloudwatch_metric_alarm" "west_high_memory_utilization" {
  provider          = aws.west
  alarm_name        = "WestHighMemoryUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods = 2
  metric_name       = "mem_used_percent"
  namespace         = "CWAgent"
  period            = 300
  statistic         = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description = "This alarm triggers if memory utilization
exceeds 80%."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
  }
  actions_enabled = true
  alarm_actions    = [aws_sns_topic.west_alert_topic.arn]
}

# CloudWatch Alarm for Disk Utilization
resource "aws_cloudwatch_metric_alarm" "west_high_disk_utilization" {
  provider          = aws.west
  alarm_name        = "WestHighDiskUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods = 2
  metric_name       = "disk_used_percent"
  namespace         = "CWAgent"
  period            = 300
  statistic         = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description = "This alarm triggers if disk utilization exceeds
80%."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
  }
  actions_enabled = true
  alarm_actions    = [aws_sns_topic.west_alert_topic.arn]
}

# SNS Topic for Alarm Notifications in us-west-2
resource "aws_sns_topic" "west_alert_topic" {

```

```

provider = aws.west
name      = "WestEC2AlertTopic"
}

resource "aws sns topic_subscription" "west_email_subscription" {
  provider  = aws.west
  topic_arn = aws_sns_topic.west_alert_topic.arn
  protocol  = "email"
  endpoint   = var.notification_email
}

```

e) Updated ec2_autoscaling.tf to Use Variables

In this step, I updated the `ec2_autoscaling.tf` file to use the variables defined in `variables.tf`, making the configuration more adaptable and easier to manage. Below is the updated Terraform configuration file.

```

# For us-east-1:

# Security Group
resource "aws_security_group" "app_sg" {
  name      = "app-sg"
  description = "Allow HTTP and SSH traffic"
  vpc_id    = aws_vpc.main_vpc.id

  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["142.59.28.49/32"] # Restricted SSH to my IP only
  }

  ingress {
    from_port  = 22
    to_port    = 22
    protocol   = "tcp"
    cidr_blocks = ["18.206.107.24/29"] # EC2 Instance Connect IP range for
  us-east-1
  }

  ingress {
    from_port  = 80
    to_port    = 80
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"] # Allows HTTP traffic from anywhere
  }

  egress {
    from_port  = 0
    to_port    = 0
    protocol   = "-1"
    cidr_blocks = ["0.0.0.0/0"] # Allows all outbound traffic
  }

  tags = {
    Name = "AppSecurityGroup"
  }
}

```

```

        }

}

# Launch Template for EC2 Instances for us-east-1
resource "aws_launch_template" "app" {
  name_prefix      = "app-launch-template"
  image_id         = "ami-06b21ccaeff8cd686"  # Amazon Linux AMI
  instance_type    = var.instance_type
  iam_instance_profile {
    name = aws_iam_instance_profile.ec2_instance_profile.name
  }

  network_interfaces {
    security_groups = [aws_security_group.app_sg.id]
  }

  user_data = base64encode(<<-EOF
  #!/bin/bash
  sudo dnf update -y
  sudo dnf install -y nginx amazon-cloudwatch-agent
  sudo systemctl start nginx
  sudo systemctl enable nginx

  # Configure the CloudWatch Agent
  sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json << 'EOF2'
  {
    "metrics": {
      "append_dimensions": {
        "AutoScalingGroupName": "AppASGInstance"
      },
      "metrics_collected": {
        "mem": {
          "measurement": [
            "mem_used_percent"
          ],
          "metrics_collection_interval": 60
        },
        "disk": {
          "measurement": [
            "used_percent"
          ],
          "resources": [
            "/"
          ],
          "metrics_collection_interval": 60
        }
      }
    }
  }
EOF2

  # Start the CloudWatch Agent
  sudo /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-
agent-ctl -a start -c file:/opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json -m ec2
  EOF
}

lifecycle {
  create_before_destroy = true
}

```

```

    }

    tag_specifications {
        resource_type = "instance"
        tags = {
            Name = "AppInstance"
        }
    }
}

# Auto Scaling Group (ASG) for us-east-1
resource "aws_autoscaling_group" "app_asg" {
    desired_capacity      = 2
    min_size              = 1
    max_size              = 5
    vpc_zone_identifier   = [aws_subnet.public_subnet_1.id,
aws_subnet.public_subnet_2.id]
    target_group_arns     = [aws_lb_target_group.primary.arn]

    instance_refresh {
        strategy = "Rolling"
        preferences {
            min_healthy_percentage = 0
            instance_warmup       = 120
        }
    }

    health_check_grace_period = 300
    health_check_type         = "ELB"

    launch_template {
        id      = aws_launch_template.app.id
        version = "$Latest"
    }

    tag {
        key          = "Name"
        value         = "AppASGInstance"
        propagate_at_launch = true
    }

    lifecycle {
        create_before_destroy = true
    }
}

# Auto Scaling Policies
resource "aws_autoscaling_policy" "scale_up" {
    name           = "scale-up"
    scaling_adjustment = 1
    adjustment_type = "ChangeInCapacity"
    autoscaling_group_name = aws_autoscaling_group.app_asg.name
}

resource "aws_autoscaling_policy" "scale_down" {
    name           = "scale-down"
    scaling_adjustment = -1
    adjustment_type = "ChangeInCapacity"
    autoscaling_group_name = aws_autoscaling_group.app_asg.name
}

```

```

# For us-west-2:

# Security Group
resource "aws_security_group" "west_app_sg" {
  provider      = aws.west
  name          = "west-app-sg"
  description   = "Allow HTTP and SSH traffic in us-west-2"
  vpc_id        = aws_vpc.west_vpc.id

  ingress {
    from_port    = 22
    to_port      = 22
    protocol     = "tcp"
    cidr_blocks = ["142.59.28.49/32"] # Restrict SSH to your IP
  }

  ingress {
    from_port    = 22
    to_port      = 22
    protocol     = "tcp"
    cidr_blocks = ["18.237.140.160/29"] # EC2 Instance Connect IP range
  }
  for us-west-2
  }

  ingress {
    from_port    = 80
    to_port      = 80
    protocol     = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = {
    Name = "WestAppSecurityGroup"
  }
}

# Launch Template for EC2 Instances in us-west-2
resource "aws_launch_template" "secondary_app" {
  provider      = aws.west
  name_prefix   = "secondary-app-launch-template"
  image_id      = "ami-07c5ecd8498c59db5"
  instance_type = var.instance_type

  iam_instance_profile {
    name = aws_iam_instance_profile.ec2_instance_profile.name
  }

  network_interfaces {
    security_groups = [aws_security_group.west_app_sg.id]
  }

  user_data = base64encode(<<-EOF
      #!/bin/bash
      sudo dnf update -y

```

```

        sudo dnf install -y nginx amazon-cloudwatch-agent
        sudo systemctl start nginx
        sudo systemctl enable nginx

        # Configure the CloudWatch Agent
        sudo tee /opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json << 'EOF2'
{
    "metrics": {
        "append_dimensions": {
            "AutoScalingGroupName": "SecondaryAppASGInstance"
        },
        "metrics_collected": {
            "mem": {
                "measurement": [
                    "mem_used_percent"
                ],
                "metrics_collection_interval": 60
            },
            "disk": {
                "measurement": [
                    "used_percent"
                ],
                "resources": [
                    "/"
                ],
                "metrics_collection_interval": 60
            }
        }
    }
}
EOF2

# Start the CloudWatch Agent
sudo /opt/aws/amazon-cloudwatch-agent/bin/amazon-cloudwatch-
agent-ctl -a start -c file:/opt/aws/amazon-cloudwatch-agent/etc/amazon-
cloudwatch-agent.json -m ec2
EOF
)

lifecycle {
    create_before_destroy = true
}

tag_specifications {
    resource_type = "instance"
    tags = {
        Name = "SecondaryAppInstance"
    }
}

# Auto Scaling Group (ASG) for us-west-2
resource "aws_autoscaling_group" "secondary_asg" {
    provider          = aws.west
    desired_capacity = 2
    min_size         = 1
    max_size         = 5
    vpc_zone_identifier = [aws_subnet.west_public_subnet_1.id,
aws_subnet.west_public_subnet_2.id]
    target_group_arns = [aws_lb_target_group.secondary.arn]
}
```

```

instance_refresh {
  strategy = "Rolling"
  preferences {
    min_healthy_percentage = 0
    instance_warmup        = 120
  }
}

health_check_grace_period = 300
health_check_type         = "ELB"

launch_template {
  id      = aws_launch_template.secondary_app.id
  version = "$Latest"
}

tag {
  key          = "Name"
  value         = "SecondaryAppASGInstance"
  propagate_at_launch = true
}

lifecycle {
  create_before_destroy = true
}
}

# Auto Scaling Policies
resource "aws_autoscaling_policy" "west_scale_up" {
  provider           = aws.west
  name               = "west-scale-up"
  scaling_adjustment = 1
  adjustment_type   = "ChangeInCapacity"
  autoscaling_group_name = aws_autoscaling_group.secondary_asg.name
}

resource "aws_autoscaling_policy" "west_scale_down" {
  provider           = aws.west
  name               = "west-scale-down"
  scaling_adjustment = -1
  adjustment_type   = "ChangeInCapacity"
  autoscaling_group_name = aws_autoscaling_group.secondary_asg.name
}

```

f) Updated network.tf to Use Variables

I updated the `network.tf` file to utilize variables defined in `variables.tf`, enhancing flexibility and allowing consistent management across configurations. Below is the updated Terraform configuration file:

```

# VPC for us-east-1
resource "aws_vpc" "main_vpc" {
  cidr_block = "10.0.0.0/16"  # VPC with a /16 subnet range
  tags = {
    Name = "MainVPC"
  }

```

```

}

# Public Subnet in us-east-1a
resource "aws_subnet" "public_subnet_1" {
    vpc_id          = aws_vpc.main_vpc.id
    cidr_block      = "10.0.1.0/24" # Public subnet's CIDR block
    availability_zone = "${var.aws_region}a"
    map_public_ip_on_launch = true
    tags = {
        Name = "PublicSubnet1"
    }
}

# Public Subnet in us-east-1b
resource "aws_subnet" "public_subnet_2" {
    vpc_id          = aws_vpc.main_vpc.id
    cidr_block      = "10.0.4.0/24"
    availability_zone = "${var.aws_region}b"
    map_public_ip_on_launch = true
    tags = {
        Name = "PublicSubnet2"
    }
}

# Private Subnet in us-east-1a
resource "aws_subnet" "private_subnet_1" {
    vpc_id          = aws_vpc.main_vpc.id
    cidr_block      = "10.0.3.0/24" # Private subnet's CIDR block
    availability_zone = "${var.aws_region}a"
    map_public_ip_on_launch = false
    tags = {
        Name = "PrivateSubnet1"
    }
}

# VPC for us-west-2
resource "aws_vpc" "west_vpc" {
    provider     = aws.west
    cidr_block   = "10.1.0.0/16"
    tags = {
        Name = "WestVPC"
    }
}

# Public Subnet in us-west-2a
resource "aws_subnet" "west_public_subnet_1" {
    provider       = aws.west
    vpc_id         = aws_vpc.west_vpc.id
    cidr_block     = "10.1.1.0/24"
    availability_zone = "${var.aws_west_region}a"
    map_public_ip_on_launch = true
    tags = {
        Name = "WestPublicSubnet1"
    }
}

# Public Subnet in us-west-2b
resource "aws_subnet" "west_public_subnet_2" {
    provider       = aws.west
    vpc_id         = aws_vpc.west_vpc.id
    cidr_block     = "10.1.2.0/24"
}

```

```

availability_zone      = "${var.aws_west_region}b"
map_public_ip_on_launch = true
tags = {
  Name = "WestPublicSubnet2"
}
}

# Private Subnet in us-west-2a
resource "aws_subnet" "west_private_subnet_1" {
  provider          = aws.west
  vpc_id            = aws_vpc.west_vpc.id
  cidr_block        = "10.1.3.0/24"
  availability_zone = "${var.aws_west_region}a"
  tags = {
    Name = "WestPrivateSubnet1"
  }
}

# Internet Gateway (IGW) for us-east-1
resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main_vpc.id
  tags = {
    Name = "MainIGW"
  }
}

# Public Route Table and Route for us-east-1
resource "aws_route_table" "public_rt" {
  vpc_id = aws_vpc.main_vpc.id
  tags = {
    Name = "PublicRouteTable"
  }
}

resource "aws_route" "internet_access" {
  route_table_id      = aws_route_table.public_rt.id
  destination_cidr_block = "0.0.0.0/0"
  gateway_id          = aws_internet_gateway.igw.id
}

resource "aws_route_table_association" "public_subnet_assoc_1" {
  subnet_id           = aws_subnet.public_subnet_1.id
  route_table_id      = aws_route_table.public_rt.id
}

resource "aws_route_table_association" "public_subnet_assoc_2" {
  subnet_id           = aws_subnet.public_subnet_2.id
  route_table_id      = aws_route_table.public_rt.id
}

# Internet Gateway (IGW) for us-west-2
resource "aws_internet_gateway" "west_igw" {
  provider = aws.west
  vpc_id   = aws_vpc.west_vpc.id
  tags = {
    Name = "WestIGW"
  }
}

# Public Route Table and Route for us-west-2
resource "aws_route_table" "west_public_rt" {

```

```

provider = aws.west
vpc_id   = aws_vpc.west_vpc.id
tags = {
  Name = "WestPublicRouteTable"
}
}

resource "aws_route" "west_internet_access" {
  provider          = aws.west
  route_table_id    = aws_route_table.west_public_rt.id
  destination_cidr_block = "0.0.0.0/0"
  gateway_id        = aws_internet_gateway.west_igw.id
}

resource "aws_route_table_association" "west_public_subnet_assoc_1" {
  provider          = aws.west
  subnet_id         = aws_subnet.west_public_subnet_1.id
  route_table_id    = aws_route_table.west_public_rt.id
}

resource "aws_route_table_association" "west_public_subnet_assoc_2" {
  provider          = aws.west
  subnet_id         = aws_subnet.west_public_subnet_2.id
  route_table_id    = aws_route_table.west_public_rt.id
}

```

g) Updated route53.tf to Use Variables

I updated the `route53.tf` file to utilize variables from `variables.tf`, making the configuration adaptable for different environments. Below is the updated Terraform configuration file:

```

# Route 53 Hosted Zone
resource "aws_route53_zone" "main" {
  name = var.route53_zone_domain
}

# Primary Route 53 Record (us-east-1)
resource "aws_route53_record" "primary" {
  zone_id = aws_route53_zone.main.zone_id
  name    = "app.${var.route53_zone_domain}"
  type    = "A"

  alias {
    name                  = aws_lb.primary_alb.dns_name
    zone_id               = aws_lb.primary_alb.zone_id
    evaluate_target_health = true
  }

  set_identifier          = "Primary"
  failover_routing_policy {
    type = "PRIMARY"
  }
}

```

```

# Secondary Route 53 Record (us-west-2)
resource "aws_route53_record" "secondary" {
  zone_id = aws_route53_zone.main.zone_id
  name    = "app.${var.route53_zone_domain}"
  type    = "A"

  alias {
    name          = aws_lb.secondary_alb.dns_name
    zone_id       = aws_lb.secondary_alb.zone_id
    evaluate_target_health = true
  }

  set_identifier      = "Secondary"
  failover_routing_policy {
    type = "SECONDARY"
  }
}

# Route 53 Health Checks

# Primary Health Check (us-east-1)
resource "aws_route53_health_check" "primary_health_check" {
  fqdn        = aws_lb.primary_alb.dns_name # Use ALB DNS Name
  type        = "HTTP"
  port        = 80
  request_interval = 30
  failure_threshold = 3
}

# Secondary Health Check (us-west-2)
resource "aws_route53_health_check" "secondary_health_check" {
  fqdn        = aws_lb.secondary_alb.dns_name # Use ALB DNS Name
  type        = "HTTP"
  port        = 80
  request_interval = 30
  failure_threshold = 3
}

```

h) Updated s3_bucket.tf to Use Variables

I updated the `s3_bucket.tf` file to reference variables defined in `variables.tf`, allowing for adaptable S3 bucket names and configurations. Below is the updated configuration:

```

# S3 bucket for storing Terraform state
resource "aws_s3_bucket" "terraform_state" {
  bucket = var.terraform_state_bucket # Use variable for Terraform state
  bucket name

  tags = {
    Name        = "TerraformStateBucket"
    Environment = "Dev"
  }
}

# Public access block for the S3 bucket
resource "aws_s3_bucket_public_access_block" "terraform_state_block" {

```

```

bucket = aws_s3_bucket.terraform_state.id

block_public_acls      = true
block_public_policy    = true
restrict_public_buckets = true
ignore_public_acls     = true
}

# S3 bucket for backups
resource "aws_s3_bucket" "backup_bucket" {
  bucket = var.backup_bucket # Use variable for backup bucket name

  tags = {
    Name      = "BackupBucket"
    Environment = "Dev"
  }
}

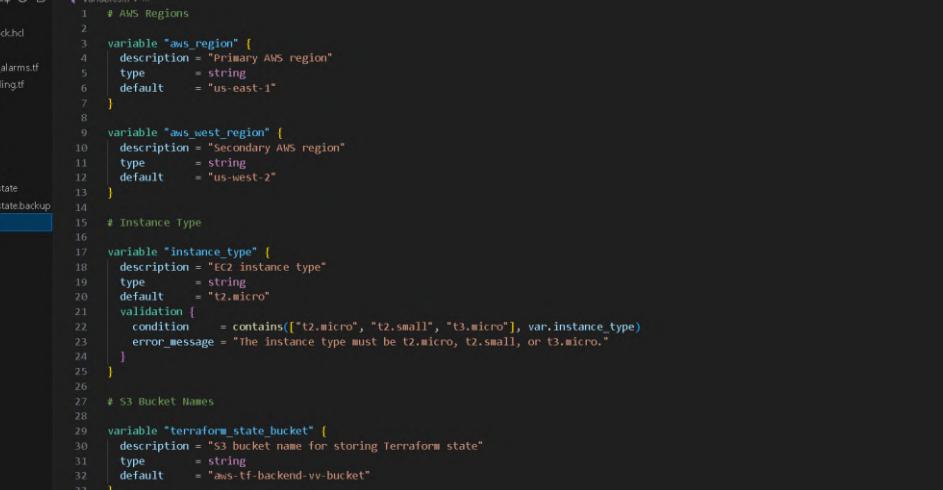
# Public access block for the backup S3 bucket
resource "aws_s3_bucket_public_access_block" "backup_bucket_block" {
  bucket = aws_s3_bucket.backup_bucket.id

  block_public_acls      = true
  block_public_policy    = true
  restrict_public_buckets = true
  ignore_public_acls     = true
}

```

I executed terraform init to initialize the working directory and ensure that all required Terraform modules and provider plugins were properly set up. Following that, I ran terraform plan to generate and review the execution plan, which outlined any anticipated infrastructure changes based on the updated configurations across all Terraform configuration files. The plan output confirmed that no modifications were needed, indicating that the configurations using variables matched the current state. Finally, I executed terraform apply, which applied the plan without making any changes to the infrastructure, further verifying that all configurations were in sync with the existing resources. This process confirmed that the updates were successful, marking the completion of this phase.

Screenshots



The screenshot shows the Visual Studio Code interface with a Terraform configuration file open in the center editor pane. The file is titled 'variables.tf' and contains the following code:

```
variable "aws_region" {
  description = "Primary AWS region"
  type        = string
  default     = "us-east-1"
}

variable "aws_west_region" {
  description = "Secondary AWS region"
  type        = string
  default     = "us-west-2"
}

variable "instance_type" {
  description = "EC2 instance type"
  type        = string
  default     = "t2.micro"
  validation {
    condition = contains(["t2.micro", "t2.small", "t3.micro"], var.instance_type)
    error_message = "The instance type must be t2.micro, t2.small, or t3.micro."
  }
}

variable "terraform_state_bucket" {
  description = "S3 bucket name for storing Terraform state"
  type        = string
  default     = "aws-tf-backend-vv-bucket"
}

variable "backup_bucket" {
  description = "S3 bucket name for backups"
  type        = string
}
```

```
variable "backup_bucket" {
  type        = string
  default    = "terra-backup-vv-project-bucket"
}

variable "primary_alb_name" {
  description = "Name of the primary ALB"
  type        = string
  default    = "primary-alb"
}

variable "secondary_alb_name" {
  description = "Name of the secondary ALB"
  type        = string
  default    = "secondary-alb"
}

variable "route53_zone_domain" {
  description = "Domain name for Route 53 hosted zone"
  type        = string
  default    = "myinfra.local"
}

variable "notification_email" {
  description = "Email address for SMS notifications"
  type        = string
  default    = "vivekashish1507@gmail.com"
  validation {
    condition = length(regexall("^(a-zA-Z0-9_\\.\\-]+@[a-zA-Z0-9_\\-\\.]+\\.[a-zA-Z]{2,}$", var.notification_email)) > 0
    error_message = "The email format is invalid."
  }
}
```

The screenshot shows a code editor interface with multiple tabs open. The active tab is titled 'alb.tf' and contains Terraform configuration code for creating an Application Load Balancer (ALB) and its target group. The code defines resources like 'aws_lb' for the ALB and 'aws_lb_target_group' for the primary target group, specifying details such as port, protocol, and health check parameters. Other tabs visible include 'aws_vpc.tf', 'iam_roles.tf', 'main.tf', 'network.tf', 'route53.tf', 's3_bucket.tf', and 'variables.tf'. The left sidebar displays a tree view of the project structure under 'TERRAFORM'.

```
resource "aws_lb" "primary_lb" {
  name      = var.primary_lb_name
  internal   = false
  load_balancer_type = "application"
  security_groups = [aws_security_group.app_sg.id] # attach the security group
  subnets      = [aws_subnet.public_subnet_1.id, aws_subnet.public_subnet_2.id]

  enable_deletion_protection = false
  idle_timeout                = 400

  tags = [
    { Name = "PrimaryALB" }
  ]
}

# ALB Target Group for us-east-1

resource "aws_lb_target_group" "primary" {
  name      = "primary-target-group"
  port      = 80
  protocol = "HTTP"
  vpc_id    = aws_vpc.main_vpc.id

  health_check {
    path      = "/"
    interval = 30
    timeout   = 5
    healthy_threshold = 2
    unhealthy_threshold = 2
  }

  tags = [
    { Name = "PrimaryTargetGroup" }
  ]
}
```

```
File Edit Selection View Go Run Terminal Help <- > O Terraform
EXPLORER ... alb.tf ...
TERRAFORM > _terraform > alb.tf > resource "aws_lb" "secondary_lb" > name
> _terraform.lock.hcl > port = 80
> alb.tf > protocol = "HTTP"
> cloudwatch_alarms.tf > default_action {
> ec2_autoscaling.tf > type = "forward"
> iam_roles.tf > target_group_arn = aws_lb_target_group.primary.arn
> main.tf > }
> network.tf > tags = [
> route53.tf > Name = "PrimaryALBLListener"
> s3_bucket.tf > ]
> terraform.tfstate > }
> terraform.state.backup > variables.tf > }

41 resource "aws_lb_listener" "http" {
42   port          = 80
43   protocol      = "HTTP"
44
45   default_action {
46     type        = "forward"
47     target_group_arn = aws_lb_target_group.primary.arn
48   }
49
50   tags = [
51     Name = "PrimaryALBLListener"
52   ]
53 }
54
55 # Application Load Balancer (ALB) for us-west-2
56
57 resource "aws_lb" "secondary_lb" {
58   provider      = aws.west
59   name          = var.secondary_lb_name
60   internal      = false
61   load_balancer_type = "application"
62   security_groups = [aws_security_group.west_app_sg.id]
63   subnets       = [aws_subnet.west_public_subnet_1.id, aws_subnet.west_public_subnet_2.id]
64
65   enable_deletion_protection = false
66   idle_timeout           = 400
67
68   tags = [
69     Name = "SecondaryALB"
70   ]
71 }
72
73
74 # ALB Target Group for us-west-2
75
76 resource "aws_lb_target_group" "secondary" {
77   provider = aws.west
78   name    = "secondary_target_group"
79 }
```

Ln 60, Col 46 (43 selected) Spaces:2 UTF-8 CRLF {} Terraform CODEGPT

```
File Edit Selection View Go Run Terminal Help <- > O Terraform
EXPLORER ... cloudwatch_alarms.tf ...
TERRAFORM > _terraform > cloudwatch_alarms.tf > resource "aws_sns_topic_subscription" "email_subscription" > endpoint
> _terraform.lock.hcl > provider = aws.west
> alb.tf > name = var.notification_email
> cloudwatch_alarms.tf > dimensions = [
> ec2_autoscaling.tf > AutoScalingGroupName = aws_autoscaling_group.app_asg.name
> iam_roles.tf > ]
> main.tf > actions_enabled = true
> network.tf > alarm_actions = [aws sns topic.alert_topic.arn]
> route53.tf > ]
> s3_bucket.tf > }
> terraform.tfstate > }
> terraform.state.backup > variables.tf > }

43 resource "aws_cloudwatch_metric_alarm" "high_disk_utilization" {
44   alarm_description = "This alarm triggers if disk utilization exceeds 80%."
45   dimensions = [
46     AutoScalingGroupName = aws_autoscaling_group.app_asg.name
47   ]
48   actions_enabled = true
49   alarm_actions = [aws sns topic.alert_topic.arn]
50 }
51
52 # SNS Topic for Alarm Notifications in us-east-1
53
54 resource "aws sns topic" "alert_topic" {
55   name = "EC2AlertTopic"
56 }
57
58 resource "aws sns topic_subscription" "email_subscription" {
59   topic_arn = aws sns topic.alert_topic.arn
60   protocol = "email"
61   endpoint = var.notification_email
62 }
63
64
65 # us-west-2
66
67 # CloudWatch Alarm for CPU Utilization in us-west-2
68
69 resource "aws_cloudwatch_metric_alarm" "west_high_cpu_utilization" {
70   provider      = aws.west
71   alarm_name    = "WestHighCPUUtilizationAlarm"
72   comparison_operator = "GreaterThanOrEqualToThreshold"
73   evaluation_periods = 2
74   metric_name   = "CPUUtilization"
75   namespace     = "AWS/EC2"
76   period        = 300
77   statistic     = "Average"
78   threshold     = 80 # Set threshold based on your requirements
79   alarm_description = "This alarm triggers if CPU utilization exceeds 80%."
80   dimensions = [
81     AutoScalingGroupName = aws_autoscaling_group.app_asg.name
82   ]
83   actions_enabled = true
84   alarm_actions = [aws sns topic.west_alert_topic.arn]
85 }
```

Ln 60, Col 37 (34 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT

```
File Edit Selection View Go Run Terminal Help <- > O Terraform
EXPLORER ... cloudwatch_alarms.tf ...
TERRAFORM > _terraform > cloudwatch_alarms.tf > resource "aws_sns_topic_subscription" "west_email_subscription" > endpoint
> _terraform.lock.hcl > provider = aws.west
> alb.tf > name = var.notification_email
> cloudwatch_alarms.tf > dimensions = [
> ec2_autoscaling.tf > AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
> iam_roles.tf > ]
> main.tf > actions_enabled = true
> network.tf > alarm_actions = [aws sns topic.west_alert_topic.arn]
> route53.tf > ]
> s3_bucket.tf > }
> terraform.tfstate > }
> terraform.state.backup > variables.tf > }

116 resource "aws_cloudwatch_metric_alarm" "west_high_disk_utilization" {
117   metric_name      = "disk_used_percent"
118   namespace        = "CMAgent"
119   period           = 300
120   statistic        = "Average"
121   threshold        = 80 # Set threshold based on your requirements
122   alarm_description = "This alarm triggers if disk utilization exceeds 80%."
123   dimensions = [
124     AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
125   ]
126   actions_enabled = true
127   alarm_actions = [aws sns topic.west_alert_topic.arn]
128 }
129
130
131 # SNS Topic for Alarm Notifications in us-west-2
132
133 resource "aws sns topic" "west_alert_topic" {
134   provider = aws.west
135   name     = "WestEC2AlertTopic"
136 }
137
138
139 resource "aws sns topic_subscription" "west_email_subscription" {
140   provider = aws.west
141   topic_arn = aws sns topic.west_alert_topic.arn
142   protocol = "email"
143   endpoint = var.notification_email
144 }
```

Ln 145, Col 37 (64 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT

This screenshot shows the Visual Studio Code interface with the Terraform extension open. The left sidebar displays a file tree under 'TERRAFORM' containing files like .terraform, .terraform.lock.hcl, alb.tf, cloudwatch_alarms.tf, ec2_autoscaling.tf, iam_roles.tf, main.tf, network.tf, route53.tf, s3_bucket.tf, terraform.tfstate, and variables.tf. The main editor pane shows Terraform code for an EC2 Auto Scaling configuration. It includes resources for an AWS Security Group ('app_sg') and an AWS Launch Template ('app'). The security group has an egress rule allowing all outbound traffic. The launch template specifies an Amazon Linux AMI, an IAM instance profile, network interfaces (with security groups), and user data for installing nginx and the CloudWatch agent.

```
resource "aws_security_group" "app_sg" {
  egress {
    from_port = 0
    to_port = 0
    protocol = "-1"
    cidr_blocks = ["0.0.0.0/0"] # Allows all outbound traffic
  }

  tags = [
    { Name = "AppSecurityGroup" }
  ]
}

# Launch Template for EC2 Instances for us-east-1

resource "aws_launch_template" "app" {
  name_prefix = "app-launch-template"
  image_id = "ami-0eb21caeff8d686" # Amazon Linux AMI
  instance_type = var.instance_type
  # Attach the IAM role to EC2 instances
  iam_instance_profile {
    name = aws_iam_instance_profile.ec2_instance_profile.name
  }

  network_interfaces {
    security_groups = [aws_security_group.app_sg.id]
  }

  user_data = base64encode(<<<-EOF
#!/bin/bash
sudo dnf update -y
sudo dnf install -y nginx amazon-cloudwatch-agent
sudo systemctl start nginx
sudo systemctl enable nginx
  <<<EOF)
}

# Configure the CloudWatch Agent
```

This screenshot shows the Visual Studio Code interface with the Terraform extension open. The left sidebar displays a file tree under 'TERRAFORM' containing files like .terraform, .terraform.lock.hcl, alb.tf, cloudwatch_alarms.tf, ec2_autoscaling.tf, iam_roles.tf, main.tf, network.tf, route53.tf, s3_bucket.tf, terraform.tfstate, and variables.tf. The main editor pane shows Terraform code for a secondary EC2 Auto Scaling configuration. It includes resources for an AWS Security Group ('west_app_sg') and an AWS Launch Template ('secondary_app'). The security group has an egress rule allowing all outbound traffic. The launch template specifies an Amazon Linux AMI, an IAM instance profile, network interfaces (with security groups), and user data for installing nginx and the CloudWatch agent.

```
resource "aws_security_group" "west_app_sg" {
  egress {
    to_port = 0
    protocol = "-1"
    cidr_blocks = ["0.0.0.0/0"] # Allow all outbound traffic
  }

  tags = [
    { Name = "WestAppSecurityGroup" }
  ]
}

# Launch Template for EC2 Instances in us-west-2

resource "aws_launch_template" "secondary_app" {
  provider = aws.west
  name_prefix = "secondary-app-launch-template"
  image_id = "ami-07c5ecd8498c59db5" # Amazon Linux AMI
  instance_type = var.instance_type

  # Attach the IAM role to EC2 instances
  iam_instance_profile {
    name = aws_iam_instance_profile.ec2_instance_profile.name
  }

  network_interfaces {
    security_groups = [aws_security_group.west_app_sg.id]
  }

  user_data = base64encode(<<<-EOF
#!/bin/bash
sudo dnf update -y
sudo dnf install -y nginx amazon-cloudwatch-agent
sudo systemctl start nginx
sudo systemctl enable nginx
  <<<EOF)
}
```

This screenshot shows the Visual Studio Code interface with the Terraform extension open. The left sidebar displays a file tree under 'TERRAFORM' containing files like .terraform, .terraform.lock.hcl, alb.tf, cloudwatch_alarms.tf, ec2_autoscaling.tf, iam_roles.tf, main.tf, network.tf, route53.tf, s3_bucket.tf, terraform.tfstate, and variables.tf. The main editor pane shows Terraform code for VPC subnet configuration. It includes resources for public subnets ('public_subnet_1', 'public_subnet_1', 'public_subnet_2') and a private subnet ('private_subnet_1'). The public subnets are in the 'us-east-1a' availability zone and have public IP allocation enabled. The private subnet is in the 'us-east-1a' availability zone and does not have public IP allocation.

```
# Public subnet in us-east-1a

resource "aws_subnet" "public_subnet_1" {
  vpc_id = aws_vpc.main_vpc.id
  cidr_block = "10.0.1.0/24" # Public subnet's CIDR block
  availability_zone = "${var.aws_region}a" # Availability zone
  map_public_ip_on_launch = true # Ensures EC2 instances in this subnet get a public IP
  tags = [
    { Name = "PublicSubnet1" }
  ]
}

# Public subnet in us-east-1b

resource "aws_subnet" "public_subnet_2" {
  vpc_id = aws_vpc.main_vpc.id
  cidr_block = "10.0.4.0/24"
  availability_zone = "${var.aws_region}b"
  map_public_ip_on_launch = true
  tags = [
    { Name = "PublicSubnet2" }
  ]
}

# Private subnet in us-east-1a

resource "aws_subnet" "private_subnet_1" {
  vpc_id = aws_vpc.main_vpc.id
  cidr_block = "10.0.3.0/24" # Private subnet's CIDR block
  availability_zone = "${var.aws_region}a" # Same AZ as the public subnet
  map_public_ip_on_launch = false # Instances in this subnet will not have public IPs
  tags = [
    { Name = "PrivateSubnet1" }
  ]
}
```

```
58 # Public Subnet in us-west-2a
59
60 resource "aws_subnet" "west_private_subnet_1" {
61   provider      = aws.west
62   vpc_id        = aws_vpc.west_vpc.id
63   cidr_block    = "10.1.1.0/24"
64   availability_zone = "${var.aws_west_region}a"
65   map_public_ip_on_launch = true
66
67   tags = [
68     { Name = "WestPrivateSubnet1" }
69   ]
70 }
71
72 # Public Subnet in us-west-2b
73
74 resource "aws_subnet" "west_public_subnet_1" {
75   provider      = aws.west
76   vpc_id        = aws_vpc.west_vpc.id
77   cidr_block    = "10.1.2.0/24"
78   availability_zone = "${var.aws_west_region}b"
79   map_public_ip_on_launch = true
80
81   tags = [
82     { Name = "WestPublicSubnet1" }
83   ]
84 }
85
86 # Private Subnet in us-west-2a
87
88 resource "aws_subnet" "west_private_subnet_2" {
89   provider      = aws.west
90   vpc_id        = aws_vpc.west_vpc.id
91   cidr_block    = "10.1.3.0/24"
92   availability_zone = "${var.aws_west_region}a"
93
94   tags = [
```

```
1 # Route 53 Hosted Zone
2
3 resource "aws_route53_zone" "main" {
4   name = var.route53_zone_domain # Domain name variable for Route 53
5 }
6
7 # Primary Route 53 Record (us-east-1)
8
9 resource "aws_route53_record" "primary" {
10   zone_id = aws_route53_zone.main.zone_id
11   name   = "app.${var.route53_zone_domain}"
12   type   = "A"
13
14   alias {
15     name      = aws_lb.primary_alb.dns_name
16     zone_id   = aws_lb.primary_alb.zone_id
17     evaluate_target_health = true
18   }
19
20   set_identifier = "Primary"
21   failover_routing_policy {
22     type = "PRIMARY"
23   }
24 }
25
26 # Secondary Route 53 Record (us-west-2)
27
28 resource "aws_route53_record" "secondary" {
29   zone_id = aws_route53_zone.main.zone_id
30   name   = "app.${var.route53_zone_domain}"
31   type   = "A"
32
33   alias {
34     name      = aws_lb.secondary_alb.dns_name
35     zone_id   = aws_lb.secondary_alb.zone_id
36     evaluate_target_health = true
37 }
```

```
1 # S3 bucket for storing terraform state
2
3 resource "aws_s3_bucket" "terraform_state" {
4   bucket = var.terraform_state_bucket # Use variable for terraform state bucket name
5
6   tags = [
7     { Name = "TerraformStateBucket" }
8     { Environment = "Dev" }
9   ]
10
11 # Public access block for the S3 bucket
12 resource "aws_s3_bucket_public_access_block" "terraform_state_block" {
13   bucket = aws_s3_bucket.terraform_state.id
14
15   block_public_acls      = true
16   block_public_policy     = true
17   restrict_public_buckets = true
18   ignore_public_acls     = true
19 }
20
21 # S3 bucket for backups
22
23 resource "aws_s3_bucket" "backup_bucket" {
24   bucket = var.backup_bucket # Use variable for backup bucket name
25
26   tags = [
27     { Name = "BackupBucket" }
28     { Environment = "Dev" }
29   ]
30 }
31
32 # Public access block for the backup S3 bucket
33 resource "aws_s3_bucket_public_access_block" "backup_bucket_block" {
34   bucket = aws_s3_bucket.backup_bucket.id
35
36   block_public_acls      = true
37 }
```

```
Command Prompt + ×

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws sns_topic.west_alert_topic: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea15b9040f753dc]
aws_s3_bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws sns_topic.alert_topic: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic]
aws_route53_zone.main: Refreshing state... [id=2074361515QJV7STD1T]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws sns_topic.subscription.west_email_subscription: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic:52f0ce9a-1e0a-4ce2-964d-20c8633c413a]
aws sns_topic.subscription.email_subscription: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic:45bab4c9-67d8-4001-97e7-524b1a0e81c0]
aws_route_table.west_public_rt: Refreshing state... [id=rtb-0bf4cfca31f877e69]
aws_subnet.west_public_subnet_1: Refreshing state... [id=subnet-0a1940534f9cb413c]
aws_subnet.west_public_subnet_2: Refreshing state... [id=subnet-05cf33fc2aa1e38da]
aws_subnet.west_private_subnet_1: Refreshing state... [id=subnet-0de0b39d268f143ac]
aws_security_group.west_app_sg: Refreshing state... [id=sg-03d0418997f788d0]
aws_lb_target_group.secondary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba44801c8967]
aws_internet_gateway.west_igw: Refreshing state... [id=igw-093c80565bfdb459c]
aws_iam_role_policy_attachment.ec2_s3_readonly: Refreshing state... [id=ec2-role-20241026083002429100000001]
aws_iam_role_policy_attachment.ec2_cloudwatch_agent: Refreshing state... [id=ec2-role-20241031075603479000000001]
aws_iam_role_policy_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-20241030041703774400000002]


```

```
Command Prompt + ×

aws_autoscaling_group.secondary_asg: Refreshing state... [id=terraform-20241031095009964500000003]
aws_route.internet_access: Refreshing state... [id=r-tb-0c4fc0f97dfe205e51088289494]
aws_launch_template.app: Refreshing state... [id=lt-08a5c2c9fe206473d]
aws_route_table_association.public_subnet_assoc_1: Refreshing state... [id=rbassoc-062c3dc339ec6eb03]
aws_lb.primary_lb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:loadbalancer/app/primary-alb/55f8876b41516ba9]
aws_route_table_association.public_subnet_assoc_1: Refreshing state... [id=rbassoc-0c1e016c8860f95a6]
aws_autoscaling_policy.west_scale_up: Refreshing state... [id=west-scale-up]
aws_cloudwatch_metric_alarm.west_high_disk_utilization: Refreshing state... [id=WestHighDiskUtilizationAlarm]
aws_cloudwatch_metric_alarm.west_high_memory_utilization: Refreshing state... [id=WestHighMemoryUtilizationAlarm]
aws_cloudwatch_metric_alarm.west_high_cpu_utilization: Refreshing state... [id=WestHighCPUUtilizationAlarm]
aws_autoscaling_policy.west_scale_down: Refreshing state... [id=west-scale-down]
aws_autoscaling_group.app_asg: Refreshing state... [id=terraform-20241031095010218800000003]
aws_route53_record.primary: Refreshing state... [id=2074361515QJV7STD1T_app.myinfra.local_A_Primary]
aws_route53_health_check.primary_health_check: Refreshing state... [id=27ecc39-8273-4708-bca6-beef7ba7e1a]
aws_lb_listener.http: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:listener/app/primary-alb/55f8876b41516ba9/4f2d5bb1d8ba2df1]
aws_autoscaling_policy.scale_down: Refreshing state... [id=scale-down]
aws_autoscaling_policy.scale_up: Refreshing state... [id=scale-up]
aws_cloudwatch_metric_alarm.high_cpu_utilization: Refreshing state... [id=HighCPUUtilizationAlarm]
aws_cloudwatch_metric_alarm.high_disk_utilization: Refreshing state... [id=HighDiskUtilizationAlarm]
aws_cloudwatch_metric_alarm.high_memory_utilization: Refreshing state... [id=HighMemoryUtilizationAlarm]

No changes. Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes are needed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform apply -auto-approve
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws sns_topic.west_alert_topic: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic]
aws_route53_zone.main: Refreshing state... [id=2074361515QJV7STD1T]
aws sns_topic.alert_topic: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic]
aws_iam_role.ec2_role: Refreshing state... [id=ec2-role-05cf33fc2aa1e38da]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea15b9040f753dc]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws sns_topic.subscription.west_email_subscription: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic:52f0ce9a-1e0a-4ce2-964d-20c8633c413a]
aws sns_topic.subscription.email_subscription: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic:45bab4c9-67d8-4001-97e7-524b1a0e81c0]
aws_route_table.west_public_rt: Refreshing state... [id=igw-093c80565bfdb459c]
aws_subnet.west_public_subnet_1: Refreshing state... [id=subnet-0a1940534f9cb413c]


```

```
Command Prompt + ×

aws_subnet.private_subnet_1: Refreshing state... [id=subnet-074ccceedaaf778f]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-04b29a39893e794d]
aws_lb.target_group.primary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:targetgroup/primary-target-group/a9447d338de6a3e5]
aws_security_group.app_sg: Refreshing state... [id=sg-0d02a28c779890c]
aws_subnet.public_subnet_2: Refreshing state... [id=subnet-0d89362f71d4a5adf]
aws_launch_template.secondary_app: Refreshing state... [id=lt-060730d25f476a]
aws_route53_record.secondary: Refreshing state... [id=2074361515QJV7STD1T_app.myinfra.local_A_Secondary]
aws_route53_health_check.secondary_health_check: Refreshing state... [id=5ece7a87-32b1-4d77-9fe1-8ecd7fa582]
aws_lb.listener.secondary_http: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:listener/app/secondary-alb/4b7a0a637b386685/938584833b8318cd]
aws_s3_bucket.public_access_block.backup_bucket.block: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_s3_bucket_public_access_block.terraform_state.block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_autoscaling_group.secondary_asg: Refreshing state... [id=terraform-20241031095009964500000003]
aws_route.internet_access: Refreshing state... [id=r-tb-0c4fc0f97dfe205e51088289494]
aws_route_table_association.public_subnet_assoc_1: Refreshing state... [id=rbassoc-062c3dc339ec6eb03]
aws_launch_template.app: Refreshing state... [id=lt-08a5c2c9fe206473d]
aws_route_table_association.public_subnet_assoc_2: Refreshing state... [id=rbassoc-0c1e016c8860f95a6]
aws_lb.primary_lb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:loadbalancer/app/primary-alb/55f8876b41516ba9]
aws_cloudwatch_metric_alarm.west_high_disk_utilization: Refreshing state... [id=WestHighDiskUtilizationAlarm]
aws_cloudwatch_metric_alarm.west_high_memory_utilization: Refreshing state... [id=WestHighMemoryUtilizationAlarm]
aws_autoscaling_policy.west_scale_up: Refreshing state... [id=west-scale-up]
aws_cloudwatch_metric_alarm.high_cpu_utilization: Refreshing state... [id=HighCPUUtilizationAlarm]
aws_autoscaling_group.app_asg: Refreshing state... [id=terraform-20241031095010218800000003]
aws_route53_health_check.primary_health_check: Refreshing state... [id=27ecc39-8273-4708-bca6-beef7ba7e1a]
aws_route53_record.primary: Refreshing state... [id=2074361515QJV7STD1T_app.myinfra.local_A_Primary]
aws_lb_listener.http: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:listener/app/primary-alb/55f8876b41516ba9/4f2d5bb1d8ba2df1]
aws_cloudwatch_metric_alarm.high_memory_utilization: Refreshing state... [id=HighMemoryUtilizationAlarm]
aws_cloudwatch_metric_alarm.high_cpu_utilization: Refreshing state... [id=HighCPUUtilizationAlarm]
aws_autoscaling_policy.scale_up: Refreshing state... [id=scale-up]
aws_cloudwatch_metric_alarm.high_disk_utilization: Refreshing state... [id=HighDiskUtilizationAlarm]
aws_autoscaling_policy.scale_down: Refreshing state... [id=scale-down]

No changes. Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes are needed.

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

Testing

1) Application Load Balancer (ALB) Testing

a) Overview

In this phase, I focused on testing the functionality of the Application Load Balancers (ALBs) deployed in both the primary (us-east-1) and secondary (us-west-2) regions. The primary goal was to confirm that the ALBs distribute incoming traffic evenly across all EC2 instances within each region, ensuring load balancing is functioning as expected.

b) Setting Up Test Pages on EC2 Instances

To identify which instance was serving the traffic during testing, I customized the default NGINX web page on each instance in both regions with unique identifiers.

- **us-east-1 Instances:**

- **Instance 1:**

```
echo "Welcome to nginx! Served by instance ID: i-  
0aad6808354626dab (us-east-1) INSTANCE 1" | sudo tee  
/usr/share/nginx/html/index.html  
curl http://localhost  
curl http://44.201.4.236
```

- **Instance 2:**

```
echo "Welcome to nginx! Served by instance ID: i-  
0003f08c9125643ef (us-east-1) INSTANCE 2" | sudo tee  
/usr/share/nginx/html/index.html  
curl http://localhost  
curl http://18.207.173.25
```

- **us-west-2 Instances:**

- **Instance 1:**

```
echo "Welcome to nginx! Served by instance ID: i-  
0848e26fb476e924b (us-west-2) INSTANCE 1" | sudo tee  
/usr/share/nginx/html/index.html  
curl http://localhost  
curl http://18.236.133.124
```

- **Instance 2:**

```
echo "Welcome to nginx! Served by instance ID: i-  
070b3c508dbbb30d5 (us-west-2) INSTANCE 2" | sudo tee  
/usr/share/nginx/html/index.html  
curl http://localhost
```

```
curl http://34.213.180.233
```

These commands modified the index page on each instance to reflect its unique instance ID, allowing for clear tracking of traffic distribution during testing.

c) Retrieving ALB DNS Names

After setting up the EC2 instances, I obtained the DNS names of both ALBs from the EC2 Console:

- **Primary ALB (us-east-1):** primary-alb-1361659550.us-east-1.elb.amazonaws.com
- **Secondary ALB (us-west-2):** secondary-alb-1857056899.us-west-2.elb.amazonaws.com

d) Testing ALB Load Balancing via Command Line

Using the ALB DNS names, I tested traffic distribution by repeatedly sending requests from the command line:

- **Primary ALB Test (us-east-1):**

```
curl http://primary-alb-1361659550.us-east-1.elb.amazonaws.com
```

- **Secondary ALB Test (us-west-2):**

```
curl http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com
```

By running these commands multiple times, I confirmed that traffic was equally balanced across the two instances in each region, as each instance took turns responding to the requests.

e) Testing ALB Load Balancing via Browser

To further verify load balancing, I accessed each ALB's DNS URL in my local browser:

- **Primary ALB (us-east-1):** <http://primary-alb-1361659550.us-east-1.elb.amazonaws.com>
- **Secondary ALB (us-west-2):** <http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com>

By refreshing the pages multiple times, I observed the instances taking turns to serve the web page, further confirming that the load balancer was distributing the traffic evenly. The ALB testing was successfully completed, with both the primary and secondary ALBs functioning as expected, distributing traffic equally across their respective EC2 instances in each region. This validated the load balancer setup and ensured high availability across the deployment.

Screenshots

A screenshot of the AWS CloudShell interface. The title bar shows "Instances | EC2 | us-east-1". The main terminal window displays the following text:

```
A newer release of "Amazon Linux" is available.  
Version 2023.6.20241028;  
Version 2023.6.20241031;  
Run '/usr/bin/dnf check-release-update' for full release and version update info  
[ec2-user@ip-10-0-1-114 ~]$ echo "Welcome to nginx! Served by instance ID: i-0aad60808354626dab (us-east-1) INSTANCE 1" | sudo tee /usr/share/nginx/html/index.html  
Welcome to nginx! Served by instance ID: i-0aad60808354626dab (us-east-1) INSTANCE 1  
[ec2-user@ip-10-0-1-114 ~]$ curl http://localhost  
Welcome to nginx! Served by instance ID: i-0aad60808354626dab (us-east-1) INSTANCE 1  
[ec2-user@ip-10-0-1-114 ~]$ curl http://44.201.4.236  
Welcome to nginx! Served by instance ID: i-0aad60808354626dab (us-east-1) INSTANCE 1  
[ec2-user@ip-10-0-1-114 ~]$
```

The bottom of the screen shows the CloudShell logo and "Feedback" link.

A screenshot of the AWS CloudShell interface. The title bar shows "Instances | EC2 | us-east-1". The main terminal window displays the following text:

```
A newer release of "Amazon Linux" is available.  
Version 2023.6.20241028;  
Version 2023.6.20241031;  
Run '/usr/bin/dnf check-release-update' for full release and version update info  
[ec2-user@ip-10-0-4-53 ~]$ echo "Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2" | sudo tee /usr/share/nginx/html/index.html  
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2  
[ec2-user@ip-10-0-4-53 ~]$ curl http://localhost  
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2  
[ec2-user@ip-10-0-4-53 ~]$ curl http://18.207.173.25  
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2  
[ec2-user@ip-10-0-4-53 ~]$
```

The bottom of the screen shows the CloudShell logo and "Feedback" link.

```
A newer release of "Amazon Linux" is available.  
Version 2023.6.20241028;  
Version 2023.6.20241031;  
Run '/usr/bin/dnf check-release-update' for full release and version update info  
#  
Amazon Linux 2023  
https://aws.amazon.com/linux/amazon-linux-2023  
Last login: Sat Nov 2 13:49:50 2024 from 18.237.140.163  
[ec2-user@ip-10-1-2-70 ~]$ echo "Welcome to nginx! Served by instance ID: i-0840e26fb476e924b (us-west-2) INSTANCE 1" | sudo tee /usr/share/nginx/html/index.htm  
1  
Welcome to nginx! Served by instance ID: i-0840e26fb476e924b (us-west-2) INSTANCE 1  
[ec2-user@ip-10-1-2-70 ~]$ curl http://localhost  
Welcome to nginx! Served by instance ID: i-0840e26fb476e924b (us-west-2) INSTANCE 1  
[ec2-user@ip-10-1-2-70 ~]$ curl http://18.236.133.124  
Welcome to nginx! Served by instance ID: i-0840e26fb476e924b (us-west-2) INSTANCE 1  
[ec2-user@ip-10-1-2-70 ~]$ [ ]
```

```
A newer release of "Amazon Linux" is available.  
Version 2023.6.20241028;  
Version 2023.6.20241031;  
Run '/usr/bin/dnf check-release-update' for full release and version update info  
#  
Amazon Linux 2023  
https://aws.amazon.com/linux/amazon-linux-2023  
Last login: Sat Nov 2 13:51:03 2024 from 18.237.140.164  
[ec2-user@ip-10-1-1-46 ~]$ echo "Welcome to nginx! Served by instance ID: i-070b3c508dbbb30d5 (us-west-2) INSTANCE 2" | sudo tee /usr/share/nginx/html/index.htm  
1  
Welcome to nginx! Served by instance ID: i-070b3c508dbbb30d5 (us-west-2) INSTANCE 2  
[ec2-user@ip-10-1-1-46 ~]$ curl http://localhost  
Welcome to nginx! Served by instance ID: i-070b3c508dbbb30d5 (us-west-2) INSTANCE 2  
[ec2-user@ip-10-1-1-46 ~]$ curl http://34.213.180.233  
Welcome to nginx! Served by instance ID: i-070b3c508dbbb30d5 (us-west-2) INSTANCE 2  
[ec2-user@ip-10-1-1-46 ~]$ [ ]
```

The screenshot shows the AWS Management Console interface for managing load balancers. The left sidebar navigation includes 'Instances | EC2 | us-west-2', 'EC2 Instance Connect', 'EC2 Instance Connect', 'CloudShell', and 'Feedback'. The main content area is titled 'Load balancers (1/1)' under the 'EC2 > Load balancers' section. It displays a table with one row for 'primary-alb'. The table columns include 'Name' (primary-alb), 'DNS name' (primary-alb-136165955...), 'State' (Active), 'VPC ID' (vpc-0ea415b9040f75...), 'Availability Zones' (2 Availability Zones), and 'Type' (application). Below the table, a detailed view for 'Load balancer: primary-alb' is shown, listing 'Scheme' (Internet-facing), 'Hosted zone' (Z35SXDDOTRQ7X7K), 'Availability Zones' (subnet-0db9562f71d4a5adf in us-east-1b and subnet-04b29a39895e7944d in us-east-1a), and 'Date created' (October 30, 2024, 03:41 (UTC-06:00)). The 'Load balancer ARN' is listed as arn:aws:elasticloadbalancing:us-east-1:134575801911:loadbalancer/app/pri... and the 'DNS name' is primary-alb-1361659550.us-east-1.elb.amazonaws.com (A Record).

```
Command Prompt

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>curl http://primary-alb-1361659550.us-east-1.elb.amazonaws.com
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2

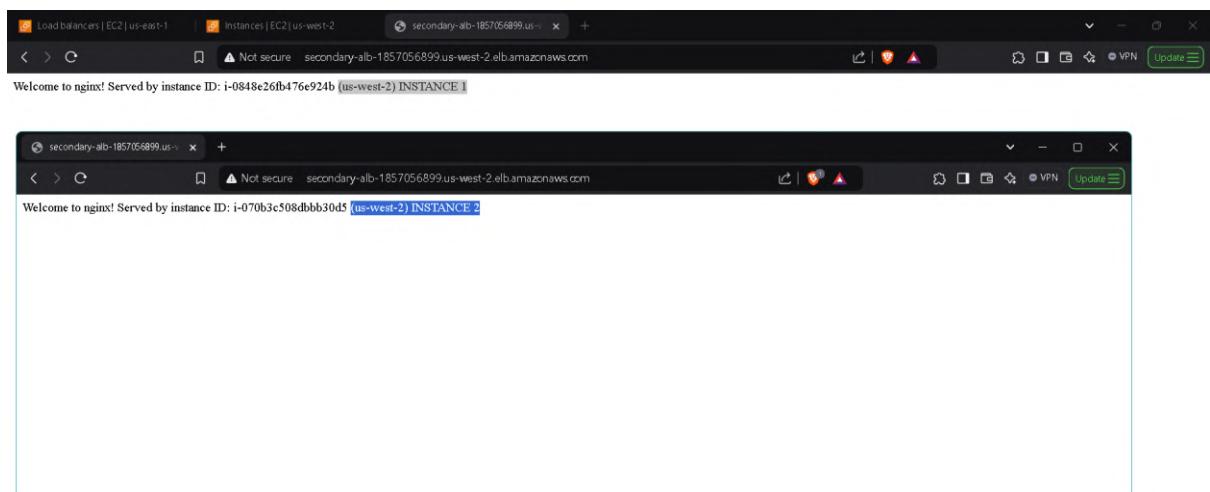
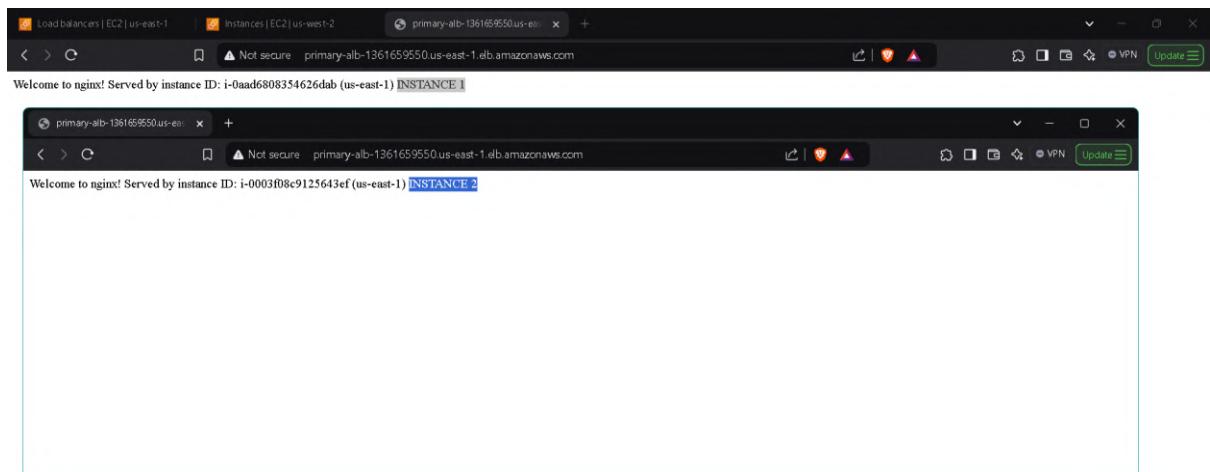
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>curl http://primary-alb-1361659550.us-east-1.elb.amazonaws.com
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>curl http://primary-alb-1361659550.us-east-1.elb.amazonaws.com
Welcome to nginx! Served by instance ID: i-0aad6808354626dab (us-east-1) INSTANCE 1

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>curl http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com
Welcome to nginx! Served by instance ID: i-0848e26fb476e924b (us-west-2) INSTANCE 1

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>curl http://secondary-alb-1857056899.us-west-2.elb.amazonaws.com
Welcome to nginx! Served by instance ID: i-070b3c508dbbb30d5 (us-west-2) INSTANCE 2

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```



2) Route 53 Failover Testing

a) Overview

This phase focused on testing the Route 53 failover configuration to ensure that traffic seamlessly switches to the secondary region (us-west-2) if instances in the primary region (us-east-1) become unavailable. Since this was for testing purposes, I updated `route53.tf` to create a private Route 53 hosted zone. This allowed testing failover without requiring a public domain.

b) Updating `route53.tf` for Private Hosted Zone

To work with a private zone in both VPCs across regions, I updated `route53.tf` to associate the `myinfra.local` private hosted zone with `MainVPC` in us-east-1 and `WestVPC` in us-west-2. This configuration enabled VPC connectivity in each region to handle the failover.

```
# Route 53 Hosted Zone
resource "aws_route53_zone" "main" {
    name = var.route53_zone_domain

    # Associate with the VPC in us-east-1
    vpc {
        vpc_id = aws_vpc.main_vpc.id  # us-east-1 VPC
    }

    # Associate with the VPC in us-west-2
    vpc {
        vpc_id      = aws_vpc.west_vpc.id  # us-west-2 VPC
        vpc_region = "us-west-2"
    }
}

# Primary Route 53 Record (us-east-1)
resource "aws_route53_record" "primary" {
    zone_id = aws_route53_zone.main.zone_id
    name    = "app.${var.route53_zone_domain}"
    type    = "A"

    alias {
        name          = aws_lb.primary_alb.dns_name
        zone_id       = aws_lb.primary_alb.zone_id
        evaluate_target_health = true
    }

    set_identifier = "Primary"
    failover_routing_policy {
        type = "PRIMARY"
    }
}

# Secondary Route 53 Record (us-west-2)
resource "aws_route53_record" "secondary" {
    zone_id = aws_route53_zone.main.zone_id
    name    = "app.${var.route53_zone_domain}"
```

```

type      = "A"

alias {
  name          = aws_lb.secondary_alb.dns_name
  zone_id       = aws_lb.secondary_alb.zone_id
  evaluate_target_health = true
}

set_identifier = "Secondary"
failover_routing_policy {
  type = "SECONDARY"
}
}

# Route 53 Health Checks

# Primary Health Check (us-east-1)
resource "aws_route53_health_check" "primary_health_check" {
  fqdn      = aws_lb.primary_alb.dns_name  # Use ALB DNS Name
  type      = "HTTP"
  port      = 80
  request_interval = 30
  failure_threshold = 3
}

# Secondary Health Check (us-west-2)
resource "aws_route53_health_check" "secondary_health_check" {
  fqdn      = aws_lb.secondary_alb.dns_name  # Use ALB DNS Name
  type      = "HTTP"
  port      = 80
  request_interval = 30
  failure_threshold = 3
}

```

c) Applying Updated route53.tf Configuration

After saving the updated `route53.tf`, I ran the following commands in the terminal to destroy the existing resources and apply the new configuration:

```

terraform destroy -target=aws_route53_zone.main -
target=aws_route53_health_check.primary_health_check -
target=aws_route53_health_check.secondary_health_check
terraform apply -auto-approve

```

After the successful application, the private hosted zone and health checks were redeployed.

d) Verifying Route 53 Setup

In the AWS Console:

1. Under Route 53, I confirmed that the `myinfra.local` private hosted zone was created with A records for Primary and Secondary ALBs.
2. I verified that `MainVPC` in us-east-1 and `WestVPC` in us-west-2 were associated with the hosted zone.
3. Both health checks for the primary and secondary ALBs were in a healthy state.

e) Enabling DNS Hostnames for VPCs

To enable instance-level DNS resolution, I edited the DNS settings for both `MainVPC` and `WestVPC`, enabling DNS hostnames.

f) Initial Testing with Failover Disabled

To validate the initial setup, I connected to an EC2 instance in us-west-2 and issued the following command multiple times:

```
curl http://app.myinfra.local
```

The output showed responses from the primary region instances in us-east-1:

```
Welcome to nginx! Served by instance ID: i-0aad6808354626dab (us-east-1)
INSTANCE 1
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1)
INSTANCE 2
```

g) Simulating Primary Region Outage

To test the failover:

1. I disabled both scale-up and scale-down policies in `AppASGInstance` in us-east-1.
2. Detached both instances in `AppASGInstance`, set the desired capacity to 0, and stopped the instances from the AWS Console.

h) Observing Failover Behavior

With primary instances offline, I checked the health checks in Route 53 and saw `primary-alb` transitioning to an unhealthy state. I then issued the `curl` command again from the us-west-2 instance:

```
curl http://app.myinfra.local
```

This time, the output showed responses from the secondary region instances in us-west-2:

```
Welcome to nginx! Served by instance ID: i-0848e26fb476e924b (us-west-2)
INSTANCE 1
Welcome to nginx! Served by instance ID: i-070b3c508dbbb30d5 (us-west-2)
INSTANCE 2
```

i) Testing Failback to Primary Region

To confirm the failback process:

1. I re-enabled the Auto Scaling policies in `AppASGInstance` and restarted the stopped instances in us-east-1.
2. Once the `primary-alb` health check returned to a healthy state, I ran the `curl` command from the us-west-2 instance:

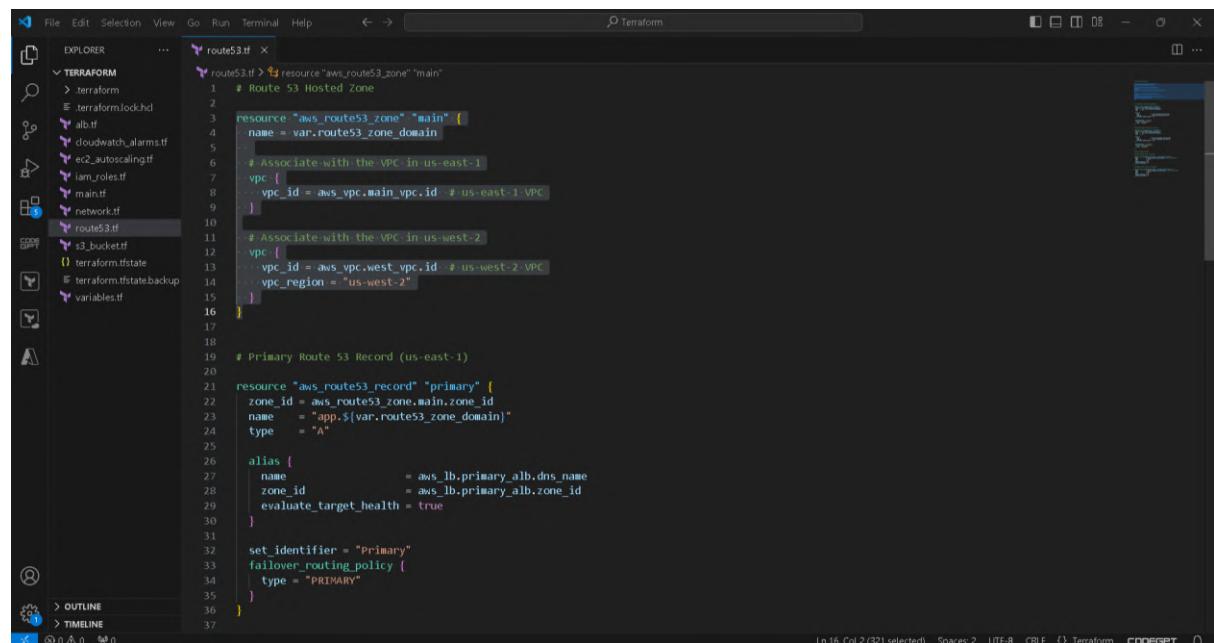
```
curl http://app.myinfra.local
```

The response returned to instances in us-east-1, confirming failback:

```
Welcome to nginx! Served by instance ID: i-0aad6808354626dab (us-east-1)
INSTANCE 1
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1)
INSTANCE 2
```

The Route 53 failover testing was successful, with traffic automatically redirecting to us-west-2 when us-east-1 was down and reverting once us-east-1 was available again. This testing confirmed the setup's robustness and reliability in supporting failover for high availability.

Screenshots



The screenshot shows the VS Code interface with the Terraform extension open. The left sidebar displays the file structure with files like `.terraform`, `alb.tf`, `cloudwatch_alarms.tf`, `ec2_autoscaling.tf`, `iam_roles.tf`, `main.tf`, `network.tf`, `route53.tf`, `s3_bucket.tf`, `variables.tf`, and `.terraform.lock.hcl`. The main editor area contains the Terraform code for setting up a Route 53 Hosted Zone and a Route 53 Record. The code defines two VPCs (us-east-1 and us-west-2), two Route 53 zones (main and backup), and a primary Route 53 record pointing to the primary ALB in us-east-1. A tooltip is visible over the `aws_vpc.main_vpc.id` variable, indicating it is a reference to the VPC ID.

```
resource "aws_route53_zone" "main" {
  name = var.route53_zone_domain
}

# Associate with the VPC in us-east-1
vpc {
  vpc_id = aws_vpc.main_vpc.id # us-east-1 VPC
}

# Associate with the VPC in us-west-2
vpc {
  vpc_id = aws_vpc.west_vpc.id # us-west-2 VPC
  vpc_region = "us-west-2"
}

# Primary Route 53 Record (us-east-1)
resource "aws_route53_record" "primary" {
  zone_id = aws_route53_zone.main.zone_id
  name   = "app.${var.route53_zone_domain}"
  type   = "A"

  alias {
    name           = aws_lb.primary_alb.dns_name
    zone_id       = aws_lb.primary_alb.zone_id
    evaluate_target_health = true
  }

  set_identifier = "Primary"
  fallover_routing_policy {
    type = "PRIMARY"
  }
}
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -target=aws_route53_zone.main
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws_route53_zone.main: Refreshing state... [id=Z074361515QXJV7STDIT]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
- destroy

Terraform will perform the following actions:

# aws_route53_record.primary will be destroyed
- resource "aws_route53_record" "primary" {
  - fqdn           = "app.myinfra.local" -> null
  - id             = "Z074361515QXJV7STDIT.app.myinfra.local_A_Primary" -> null
  - multivalue_answer_routing_policy = false -> null
  - name           = "app.myinfra.local" -> null
  - records        = [] -> null
  - set_identifier = "Primary" -> null
  - ttl            = 0 -> null
  - type           = "A" -> null
  - zone_id        = "Z074361515QXJV7STDIT" -> null
  # (1 unchanged attribute hidden)

  - alias {
    - evaluate_target_health = true -> null
    - name                  = "primary-alb-1361659550.us-east-1.elb.amazonaws.com" -> null
    - zone_id               = "Z35SXDOTRQ7X7K" -> null
  }

  - failover_routing_policy {
    - type = "PRIMARY" -> null
  }
}

# aws_route53_record.secondary will be destroyed
- resource "aws_route53_record" "secondary" {
  - fqdn           = "app.myinfra.local" -> null
  - id             = "Z074361515QXJV7STDIT.app.myinfra.local_A_Secondary" -> null
  - multivalue_answer_routing_policy = false -> null
  - name           = "app.myinfra.local" -> null
}
```

```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -target=aws_route53_health_check
aws_route_table_association.public_subnet_assoc_2
aws_route_table_association.west_public_subnet_assoc_1
aws_route_table_association.west_public_subnet_assoc_2
aws_s3_bucket.backup_bucket
aws_s3_bucket.terraform_state
aws_s3_bucket_public_access_block.backup_bucket_block
aws_s3_bucket_public_access_block.terraform_state_block
aws_security_group.app_sg
aws_security_group.west_app_sg
aws sns topic.alert_topic
aws sns topic.west_alert_topic
aws sns topic_subscription.email_subscription
aws sns topic_subscription.west_email_subscription
aws_subnet.private_subnet_1
aws_subnet.public_subnet_1
aws_subnet.public_subnet_2
aws_subnet.west_private_subnet_1
aws_subnet.west_public_subnet_1
aws_subnet.west_public_subnet_2
aws_vpc.main_vpc
aws_vpc.west_vpc

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -target=aws_route53_health_check
primary_health_check -target=aws_route53_health_check.secondary_health_check
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws_subnet.west_public_subnet_2: Refreshing state... [id=subnet-05cf33fc2aae38da]
aws_subnet.west_public_subnet_1: Refreshing state... [id=subnet-0a19045349cb413c]
aws_security_group.west_app_sg: Refreshing state... [id=sg-02d04109977fd88d0]
aws_lb.secondary_lb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:loadbalancer/app/secondary-alb/4b7a0a637b386685]
aws_route53_health_check.secondary_health_check: Refreshing state... [id=5ece7a87-32b1-4d77-9ef1-e8cd7faad582]
aws_subnet.public_subnet_2: Refreshing state... [id=subnet-0d89362f71d4a5adf]
aws_subnet.public_subnet_1: Refreshing state... [id=subnet-00b2939893e79404d]
aws_security_group.app_sg: Refreshing state... [id=sg-0d02a28cc779890cb]
aws_lb.primary_lb: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:loadbalancer/app/primary-alb/55f8876b41516ba9]
aws_route53_health_check.primary_health_check: Refreshing state... [id=227ecc39-8273-4708-bca6-befe7ba7e1a]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
- destroy

Terraform will perform the following actions:
```

```
+ vpc {
  + vpc_id      = "vpc-0bc8cffa4878b083b"
  + vpc_region  = "us-west-2"
}
+ vpc {
  + vpc_id      = "vpc-0ea415b9040f753dc"
  + vpc_region  = (known after apply)
}

Plan: 5 to add, 0 to change, 0 to destroy.
aws_route53_health_check.secondary_health_check: Creating...
aws_route53_health_check.primary_health_check: Creating...
aws_route53_zone.main: Creating...
aws_route53_health_check.secondary_health_check: Creation complete after 0s [id=f0cbd99-c3f2-4380-81f7-fd99bdf117d2]
aws_route53_health_check.primary_health_check: Creation complete after 0s [id=f34c8718-ef57-42af-bee6-fe807f3f04a1]
aws_route53_zone.main: Still creating... [10s elapsed]
aws_route53_zone.main: Still creating... [20s elapsed]
aws_route53_zone.main: Still creating... [30s elapsed]
aws_route53_zone.main: Still creating... [40s elapsed]
aws_route53_zone.main: Still creating... [50s elapsed]
aws_route53_zone.main: Still creating... [1m0s elapsed]
aws_route53_zone.main: Still creating... [1m10s elapsed]
aws_route53_zone.main: Still creating... [1m20s elapsed]
aws_route53_zone.main: Still creating... [1m30s elapsed]
aws_route53_zone.main: Still creating... [1m40s elapsed]
aws_route53_zone.main: Still creating... [1m50s elapsed]
aws_route53_zone.main: Still creating... [2m0s elapsed]
aws_route53_zone.main: Creation complete after 2m7s [id=Z04964961BUHHZ3TYTHQG]
aws_route53_record.secondary: Creating...
aws_route53_record.primary: Creating...
aws_route53_record.secondary: Still creating... [10s elapsed]
aws_route53_record.primary: Still creating... [10s elapsed]
aws_route53_record.secondary: Still creating... [20s elapsed]
aws_route53_record.primary: Still creating... [20s elapsed]
aws_route53_record.primary: Creation complete after 22s [id=Z04964961BUHHZ3TYTHQG.app.myinfra.local_A_Primary]
aws_route53_record.secondary: Creation complete after 25s [id=Z04964961BUHHZ3TYTHQG.app.myinfra.local_A_Secondary]

Apply complete! Resources: 5 added, 0 changed, 0 destroyed.
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

Route 53 Management Console | us-east-1.console.aws.amazon.com/route53/v2/hostedzones#ListRecordSets/Z04964961BUHHZTYHQG

myinfra.local

Hosted zone details

Records (4)

Record...	Type	Routine...	Differ...	Alias	Value/Route traffic to
myinfra...	NS	Simple	-	No	ns-1536.awsdns-00.co.uk. ns-0.awsdns-00.com. ns-1024.awsdns-00.org. ns-512.awsdns-00.net.
myinfra...	SOA	Simple	-	No	ns-1536.awsdns-00.co.uk. awsdns-hostmaster.amazon.com....
app.myin...	A	Failover	Primary	Yes	primary-alb-1361659550.us-east-1.elb.amazonaws.com.
app.myin...	A	Failover	Secondary	Yes	secondary-alb-1857056899.us-west-2.elb.amazonaws.com.

Route 53 Resolver | us-east-1 | us-east-1.console.aws.amazon.com/route53resolver/home?region=us-east-1#/vpcs

VPCs

You are signed in to the following Region: us-east-1 (N. Virginia)

To change your Region, use the Region selector in the upper-right corner.

VPCs (2)

ID	Name	Rules	Outbound endpoints	Inbound endpoints	Query logging status	Query logging configs	DNS Firewall rule groups	DNS Firewall fail open
vpc-01bb7cb53a9c88ffd	-	1	0	0	Not set up	0	0	Disabled
vpc-0ea415b9040f753d	MainVPC	1	0	0	Not set up	0	0	Disabled

Route 53 Resolver | us-west-2 | us-west-2.console.aws.amazon.com/route53resolver/home?region=us-west-2#/vpcs

VPCs

You are signed in to the following Region: us-west-2 (Oregon)

To change your Region, use the Region selector in the upper-right corner.

VPCs (2)

ID	Name	Rules	Outbound endpoints	Inbound endpoints	Query logging status	Query logging configs	DNS Firewall rule groups	DNS Firewall fail open
vpc-06c75fed0e2a80ad0	-	1	0	0	Not set up	0	0	Disabled
vpc-0bc8cffa4878b083b	WestVPC	1	0	0	Not set up	0	0	Disabled

Route 53 Management Console

us-east-1.console.aws.amazon.com/route53/healthchecks/home?region=us-east-1#

Services Search [Alt+S]

Global vivekvashisth @ vv3281 Update

Dashboard Hosted zones Health checks Profiles IP-based routing CIDR collections Traffic flow Traffic policies Policy records Domains Registered domains Pending requests Resolver VPCs Inbound endpoints Outbound endpoints Rules Query logging Outposts

Create health check Delete health check Edit health check

Filter by keyword

Name	Status	Description	Alarms	ID
http://secondary-alb-1857056899.us-west-2.amazonaws.com/	Healthy	http://secondary-alb-1857056899.us-west-2.amazonaws.com/	No alarms configured.	f0cbda99-c3f2-4380-81f7-fd91
http://primary-alb-1361659550.us-east-1.amazonaws.com/	Healthy	http://primary-alb-1361659550.us-east-1.amazonaws.com/	No alarms configured.	f34c8718-ef57-42af-bee6-fe81

Info Monitoring Alarms Tags Health checkers Latency

No health check selected.

No health check selected.

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

EditVpcSettings | VPC Console

us-east-1.console.aws.amazon.com/vpcconsole/home?region=us-east-1#EditVpcSettingsVpdd=vpc-0ea415b9040f753dc

Services Search [Alt+S]

N. Virginia vivekvashisth @ vv3281 Update

VPC > Your VPCs > vpc-0ea415b9040f753dc > Edit VPC settings

Edit VPC settings Info

VPC details

VPC ID: vpc-0ea415b9040f753dc Name: MainVPC

DHCP settings

DHCP option set: dopt-028d2411c992321a5

DNS settings

Enable DNS resolution: Info Enable DNS hostnames: Info

Network Address Usage metrics settings

Enable Network Address Usage metrics: Info

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Console Home | Instances | EC2 us-west-2 | EC2 Instance Connect

us-west-2.console.aws.amazon.com/ec2-instance-connect/ssh?region=us-west-2&connType=standard&instanceId=i-084...

Services Search [Alt+S]

Oregon vivekvashisth @ vv3281 Update

```
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx!  Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx!  Served by instance ID: i-0aad6808354626dab (us-east-1) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$
```

CloudShell Feedback © 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS Cloud Console showing the Auto Scaling groups page. A context menu is open over the 'scale-up' scaling policy, with the 'Disable' option selected.

Auto Scaling group: terraform-20241031095010218800000003

Dynamic scaling policies (2/2)

scale-down	scale-up
Policy type: Simple scaling	Policy type: Simple scaling
Enabled or disabled: Enabled	Enabled or disabled: Enabled

Actions

- Enable
- Disable**
- Execute
- Edit
- Delete

Screenshot of the AWS Cloud Console showing the Auto Scaling groups page. A context menu is open over the first instance in the 'Instances' table, with the 'Detach' option selected.

Auto Scaling group: terraform-20241031095010218800000003

Instances (2/2)

Instance ID	Lifecycle	Instance...	Weight...	Launch...	Available...
i-0003f08c9125643ef	InService	t2.micro	-	app-launch-ten	us-east-1b
i-0aad6808354626dab	InService	t2.micro	-	app-launch-ten	us-east-1a

Actions

- Detach
- Set to Standby
- Set to InService
- Instance scale-in protection
- Set scale-in protection
- Healthy**

Screenshot of the AWS Cloud Console showing the 'Create Auto Scaling group' wizard. The 'Group size' step is displayed.

Group size

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum scaling limits.

Desired capacity type

Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed Instances groups configured with a set of instance attributes.

Units (number of instances)

Desired capacity

Specify your group size.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity	Max desired capacity
0	5

Equal or less than desired capacity

Equal or greater than desired capacity

Maximum capacity

Create Auto Scaling group

Screenshot of the AWS EC2 Instances page showing two instances selected for stopping.

Instances (2/4) Info

Stop instances

Stopping your instance allows you to reduce costs, modify settings, and troubleshoot problems.

Instance ID	Stop protection
i-0aad6808354626dab (AppASGinstance)	Off (Can stop instance)
i-0003f08c9125643ef (AppASGinstance)	Off (Can stop instance)

Associated resources

You will continue to incur charges for these resources while the instance is stopped.

CPU utilization (%)

Network in (bytes)

Network out (bytes)

Network packets in (cou...

Cancel **Stop**

Screenshot of the AWS EC2 Instances page showing the same two instances now in a stopped state.

Instances (2) Info

Last updated less than a minute ago

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv
AppASGinstance	i-0aad6808354626dab	Stopped	t2.micro	-	View alarms +	us-east-1a	-
AppASGinstance	i-0003f08c9125643ef	Stopped	t2.micro	-	View alarms +	us-east-1b	-

Select an instance

Screenshot of the AWS Route 53 Management Console Health checks page.

Create health check **Delete health check** **Edit health check**

Filter by keyword

Name	Status	Description	Alarms	ID
an hour ago	Healthy	http://secondary-alb-1857056899.us-w...	No alarms configured.	f0cbda99-c3f2-4380-91f7-fd9!
an hour ago	Unhealthy	http://primary-alb-1361659550.us-east-...	No alarms configured.	f34c8718-ef57-42af-bee6-fe8!

Info **Monitoring** **Alarms** **Tags** **Health checkers** **Latency**

No health check selected.

No health check selected.

```

[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0aad6808354626dab (us-east-1) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-070b2c508dbbb30d5 (us-west-2) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-070b2c508dbbb30d5 (us-west-2) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0848e26fb476e924b (us-west-2) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$ 

```

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
AppASGinstance	i-0aad6808354626dab	Running	t2.micro	2/2 checks pass	View alarms	us-east-1a	ec2-3-237
AppASGinstance	i-0003f08c9125643ef	Running	t2.micro	2/2 checks pass	View alarms	us-east-1b	ec2-44-20

2 instances selected

Monitoring

Alarm recommendations

3h 1d 1w 1h UTC timezone

CPU utilization (%) Network in (bytes) Network out (bytes) Network packets in (cou...)

```

[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0aad6808354626dab (us-east-1) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-070b2c508dbbb30d5 (us-west-2) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-070b2c508dbbb30d5 (us-west-2) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0848e26fb476e924b (us-west-2) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0848e26fb476e924b (us-west-2) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-070b2c508dbbb30d5 (us-west-2) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0848e26fb476e924b (us-west-2) INSTANCE 1
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-070b2c508dbbb30d5 (us-west-2) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ curl http://app.myinfra.local
Welcome to nginx! Served by instance ID: i-0003f08c9125643ef (us-east-1) INSTANCE 2
[ec2-user@ip-10-1-2-70 ~]$ 

```

3) Auto Scaling Group (ASG) Validation

a) Overview

In this phase, the focus was on testing and validating the Auto Scaling Group (ASG) configurations by simulating conditions to trigger scaling actions. To achieve this, I first modified specific CloudWatch alarms in the `cloudwatch_alarms.tf` file to define scale-up and scale-down conditions for CPU utilization. This testing phase ensured that the ASG could automatically increase or decrease the number of instances based on the load, simulating real-world scenarios.

b) CloudWatch Alarm Adjustments

To set up alarms that would trigger scaling actions, I first removed the existing general CPU utilization alarms from `cloudwatch_alarms.tf`, which included:

```
# Removed - CPU Utilization Alarm in us-east-1
resource "aws_cloudwatch_metric_alarm" "high_cpu_utilization" {
  alarm_name          = "HighCPUUtilizationAlarm"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUUtilization"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 80  # Threshold for alarm
  alarm_description    = "Triggers if CPU utilization exceeds 80%."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  }
  actions_enabled = true
  alarm_actions    = [aws_sns_topic.alert_topic.arn]
}

# Removed - CPU Utilization Alarm in us-west-2
resource "aws_cloudwatch_metric_alarm" "west_high_cpu_utilization" {
  provider          = aws.west
  alarm_name        = "WestHighCPUUtilizationAlarm"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods = 2
  metric_name       = "CPUUtilization"
  namespace         = "AWS/EC2"
  period            = 300
  statistic          = "Average"
  threshold          = 80
  alarm_description   = "Triggers if CPU utilization exceeds 80%."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
  }
  actions_enabled = true
  alarm_actions    = [aws_sns_topic.west_alert_topic.arn]
}
```

c) New Alarms for Scale-Up and Scale-Down Triggers

After removing the original alarms, I added new CPU utilization alarms to enable automatic scaling. These alarms trigger based on predefined CPU thresholds.

For us-east-1 Region (Scale Up and Scale Down):

```
# Scale-Up Alarm in us-east-1
resource "aws_cloudwatch_metric_alarm" "scale_up_cpu_utilization" {
  alarm_name          = "ScaleUpCPUUtilizationAlarm"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUUtilization"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 80
  alarm_description    = "Triggers if CPU utilization exceeds 80% to scale up."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  }
  actions_enabled = true
  alarm_actions   = [aws_autoscaling_policy.scale_up.arn,
aws_sns_topic.alert_topic.arn]
}

# Scale-Down Alarm in us-east-1
resource "aws_cloudwatch_metric_alarm" "scale_down_cpu_utilization" {
  alarm_name          = "ScaleDownCPUUtilizationAlarm"
  comparison_operator = "LessThanThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUUtilization"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 30
  alarm_description    = "Triggers if CPU utilization drops below 30% to scale down."
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.app_asg.name
  }
  actions_enabled = true
  alarm_actions   = [aws_autoscaling_policy.scale_down.arn,
aws_sns_topic.alert_topic.arn]
}
```

For us-west-2 Region (Scale Up and Scale Down):

```
# Scale-Up Alarm in us-west-2
resource "aws_cloudwatch_metric_alarm" "west_scale_up_cpu_utilization" {
  provider          = aws.west
  alarm_name        = "WestScaleUpCPUUtilizationAlarm"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods = 2
  metric_name       = "CPUUtilization"
  namespace          = "AWS/EC2"
  period             = 300
  statistic          = "Average"
```

```

threshold          = 80
alarm_description = "Triggers if CPU utilization exceeds 80% to
scale up."
dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
}
actions_enabled = true
alarm_actions   = [aws_autoscaling_policy.west_scale_up.arn,
aws sns topic.west_alert_topic.arn]
}

# Scale-Down Alarm in us-west-2
resource "aws_cloudwatch_metric_alarm" "west_scale_down_cpu_utilization" {
provider           = aws.west
alarm_name         = "WestScaleDownCPUUtilizationAlarm"
comparison_operator = "LessThanThreshold"
evaluation_periods = 2
metric_name        = "CPUUtilization"
namespace          = "AWS/EC2"
period             = 300
statistic          = "Average"
threshold          = 30
alarm_description  = "Triggers if CPU utilization drops below 30% to
scale down."
dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
}
actions_enabled = true
alarm_actions   = [aws_autoscaling_policy.west_scale_down.arn,
aws sns topic.west_alert_topic.arn]
}

```

d) Terraform Commands and Resource Verification

After modifying `cloudwatch_alarms.tf`, I ran the following commands:

```

terraform init
terraform plan
terraform apply -auto-approve

```

The apply command successfully executed, creating 4 new resources and removing 2 resources (the initial CPU utilization alarms).

e) Auto Scaling Group Update

In `ec2_autoscaling.tf`, I updated the `desired_capacity` and `min_size` of the ASG to 1 to align with testing requirements.

f) Verification in AWS Console

I verified the creation of the new CloudWatch alarms under the “Alarms” section of CloudWatch in both the us-east-1 and us-west-2 regions. The scale-up and scale-down alarms were visible as expected.

g) Testing the Scale-Up Alarm

To simulate high CPU usage and test the scale-up alarm, I connected to an EC2 instance and installed the Apache HTTP server benchmarking tool:

```
sudo yum install -y httpd-tools
```

Then, I initiated a load test:

```
ab -n 5000000 -c 500 http://primary-alb-1361659550.us-east-1.elb.amazonaws.com/
```

As the CPU utilization rose above 80%, the ScaleUpCPUUtilizationAlarm triggered, and the ASG increased its capacity from 1 to 2. This was confirmed by reviewing the ASG event logs:

```
Status: Success
Description: Launching a new EC2 instance: i-058967e362272ea8. Triggered by ScaleUpCPUUtilizationAlarm in state ALARM.
```

h) Testing the Scale-Down Alarm

To test the scale-down alarm, I stopped the load test by pressing `Ctrl + c`. As CPU utilization dropped below 30%, the ScaleDownCPUUtilizationAlarm triggered, reducing the ASG capacity from 2 back to 1:

```
Status: Success
Description: Terminating EC2 instance: i-058967362272ea8 due to ScaleDownCPUUtilizationAlarm.
```

This phase confirmed that the ASG dynamically adjusts based on load conditions, successfully scaling up and down according to the set CloudWatch alarms.

Screenshots

```
File Edit Selection View Go Run Terminal Help < >
TERRAFORM ... route53.tf cloudwatch_alarms.tf
aws_cloudwatch_metric_alarm "high_cpu_utilization" {
  alarm_name        = "HighCPUUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods = 2
  metric_name       = "CPUUtilization"
  namespace         = "AWS/EC2"
  period            = 300
  statistic         = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description = "This alarm triggers if CPU utilization exceeds 80%."
  dimensions = [
    {
      AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
  ]
  actions_enabled    = true
  alarm_actions      = [aws_sns_topic.alert_topicarn]
}

# cloudWatch Alarm for Memory Utilization

aws_cloudwatch_metric_alarm "high_memory_utilization" {
  alarm_name        = "HighMemoryUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods = 2
  metric_name       = "mem_used_percent"
  namespace         = "CMAgent"
  period            = 300
  statistic         = "Average"
  threshold          = 80 # Set threshold based on your requirements
  alarm_description = "This alarm triggers if memory utilization exceeds 80%."
  dimensions = [
    {
      AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
  ]
  actions_enabled    = true
  alarm_actions      = [aws_sns_topic.tonic_alert_tonicarn]
}
```

The screenshot shows the VS Code interface with the Terraform extension active. The left sidebar displays the file structure under 'TERRAFORM'. The main editor area contains the following Terraform code:

```

resource "aws_cloudwatch_metric_alarm" "scale_up_cpu_utilization" {
  alarm_name          = "ScaleUpCPUUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUMotion"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 80
  alarm_description   = "This alarm triggers if CPU utilization exceeds 80% to scale up."
  dimensions           = [
    {
      AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
  ]
  actions_enabled      = true
  alarm_actions        = [aws_autoscaling_policy.scale_up.arm, aws_sns_topic.alert_topic.arm]
}

# CloudWatch Alarm for CPU Utilization to Trigger Scaling Down in us-east-1

resource "aws_cloudwatch_metric_alarm" "scale_down_cpu_utilization" {
  alarm_name          = "ScaleDownCPUUtilizationAlarm"
  comparison_operator = "LessThanThreshold"
  evaluation_periods  = 2
  metric_name         = "CPUMotion"
  namespace           = "AWS/EC2"
  period              = 300
  statistic            = "Average"
  threshold            = 30
  alarm_description   = "This alarm triggers if CPU utilization drops below 30% to scale down."
  dimensions           = [
    {
      AutoScalingGroupName = aws_autoscaling_group.app_asg.name
    }
  ]
  actions_enabled      = true
  alarm_actions        = [aws_autoscaling_policy.scale_down.arm, aws_sns_topic.alert_topic.arm]
}

```

Bottom status bar: Ln 40, Col 2 (1495 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT

The screenshot shows the VS Code interface with the Terraform extension active. The left sidebar displays the file structure under 'TERRAFORM'. The main editor area contains the following Terraform code:

```

resource "aws_cloudwatch_metric_alarm" "west_scale_up_cpu_utilization" {
  provider          = aws.west
  alarm_name        = "WestScaleUpCPUUtilizationAlarm"
  comparison_operator = "GreaterThanOrEqualToThreshold"
  evaluation_periods = 2
  metric_name       = "CPUMotion"
  namespace         = "AWS/EC2"
  period            = 300
  statistic          = "Average"
  threshold          = 80
  alarm_description = "This alarm triggers if CPU utilization exceeds 80% to scale up."
  dimensions         = [
    {
      AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
    }
  ]
  actions_enabled    = true
  alarm_actions      = [aws_autoscaling_policy.west_scale_up.arm, aws_sns_topic.west_alert_topic.arm]
}

# CloudWatch Alarm for CPU Utilization to Trigger Scaling Down in us-west-2

resource "aws_cloudwatch_metric_alarm" "west_scale_down_cpu_utilization" {
  provider          = aws.west
  alarm_name        = "WestScaleDownCPUUtilizationAlarm"
  comparison_operator = "LessThanThreshold"
  evaluation_periods = 2
  metric_name       = "CPUMotion"
  namespace         = "AWS/EC2"
  period            = 300
  statistic          = "Average"
  threshold          = 30
  alarm_description = "This alarm triggers if CPU utilization drops below 30% to scale down."
  dimensions         = [
    {
      AutoScalingGroupName = aws_autoscaling_group.secondary_asg.name
    }
  ]
  actions_enabled    = true
  alarm_actions      = [aws_autoscaling_policy.west_scale_down.arm, aws_sns_topic.west_alert_topic.arm]
}

```

Bottom status bar: Ln 134, Col 2 (1609 selected) Spaces:4 UTF-8 CRLF {} Terraform CODEGPT

The screenshot shows a Command Prompt window with the following output:

```

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Using previously-installed hashicorp/aws v5.72.1

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup>terraform plan
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cff4a878b003b]
aws sns.topic.west_alert_topic: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic]
aws cloudwatch_metric_alarm.west_high_cpu_utilization: Refreshing state... [id=WestHighCPUUtilizationAlarm]
aws sns.topic.alert_topic: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic]
aws iam.role.ec2_role: Refreshing state... [id=ec2-role]
aws vpc.main.vpc: Refreshing state... [id=vpc-0ea415b9040f753dc]
aws s3.bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws cloudwatch_metric_alarm.high_cpu_utilization: Refreshing state... [id=HighCPUUtilizationAlarm]
aws s3.bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws sns.topic.subscription.west_email_subscription: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic:52f0ce9a-1e8a-4ce2-964d-20c8633c413a]
aws sns.topic.subscription.email_subscription: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic:45bab4c9-67d8-4001-97e7-524b1a0e81c0]
aws iam.role.attachment.ec2_full_access: Refreshing state... [id=ec2-role-202410300417038178000000084]
aws iam.instance.profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws iam.role.attachment.ec2_s3_readonly: Refreshing state... [id=ec2-role-20241026083002429100000001]
aws iam.role.policy.ec2_sns_publish_policy: Refreshing state... [id=ec2-role:ec2-sns-publish-policy]
aws iam.role.policy.attachment.ec2_cloudwatch_agent: Refreshing state... [id=ec2-role-20241031075603479000000001]
aws iam.role.policy.attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-2024103004170377400000002]
aws iam.role.policy.attachment.ec2_full_access: Refreshing state... [id=ec2-role-2024103004170377400000003]
aws route_table.west_public_rt: Refreshing state... [id=rtb-0bf4cf3a1f8977e69]
aws subnet.west_private_subnet_1: Refreshing state... [id=subnet-0de0b39d268f143ac]
aws internet_gateway.west_igw: Refreshing state... [id=igw-093c80565bfdb459c]

```

```
# aws_cloudwatch_metric_alarm.high_cpu_utilization will be destroyed
# (because aws_cloudwatch_metric_alarm.high_cpu_utilization is not in configuration)
resource "aws_cloudwatch_metric_alarm" "high_cpu_utilization" {
  - actions_enabled = true -> null
  - alarm_actions = [
    - "arn:aws:sns:us-east-1:134575801911:EC2AlertTopic",
  ] -> null
  - alarm_description = "This alarm triggers if CPU utilization exceeds 80%." -> null
  - alarm_name = "HighCPUUtilizationAlarm" -> null
  - arn = "arn:aws:cloudwatch:us-east-1:134575801911:alarm:HighCPUUtilizationAlarm" -> null
  - comparison_operator = "GreaterThanThreshold" -> null
  - datapoints_to_alarm = 0 -> null
  - dimensions = {
    - "AutoScalingGroupName" = "terraform-20241031095010218800000003"
  } -> null
  - evaluation_periods = 2 -> null
  - id = "HighCPUUtilizationAlarm" -> null
  - insufficient_data_actions = [] -> null
  - metric_name = "CPUUtilization" -> null
  - namespace = "AWS/EC2" -> null
  - ok_actions = []
  - period = 300 -> null
  - statistic = "Average" -> null
  - tags = {}
  - tags_all = {}
  - threshold = 80 -> null
  - treat_missing_data = "missing" -> null
}
# (4 unchanged attributes hidden)

# aws_cloudwatch_metric_alarm.scale_down_cpu_utilization will be created
+ resource "aws_cloudwatch_metric_alarm" "scale_down_cpu_utilization" {
  + actions_enabled = true
  + alarm_actions = [
    + "arn:aws:autoscaling:us-east-1:134575801911:scalingPolicy:a3919ef3-ae04-4f07-b99d-239c9c1cc0b5:autoScalingGroupName/terraform-20241031095010218800003:policyName/scale-down",
    + "arn:aws:sns:us-east-1:134575801911:EC2AlertTopic",
  ]
  + alarm_description = "This alarm triggers if CPU utilization drops below 30% to scale down."
  + alarm_name = "ScaleDownCPUUtilizationAlarm"
```

```
# aws_cloudwatch_metric_alarm.scale_up_cpu_utilization will be created
+ resource "aws_cloudwatch_metric_alarm" "scale_up_cpu_utilization" {
  + actions_enabled = true
  + alarm_actions = [
    + "arn:aws:autoscaling:us-east-1:134575801911:scalingPolicy:609861e2-ec7f-4ba0-9511-74455f76e686:autoScalingGroupName/terraform-20241031095010218800003:policyName/scale-up",
    + "arn:aws:sns:us-east-1:134575801911:EC2AlertTopic",
  ]
  + alarm_description = "This alarm triggers if CPU utilization exceeds 80% to scale up."
  + alarm_name = "ScaleUpCPUUtilizationAlarm"
  + arn = "(Known after apply)"
  + comparison_operator = "GreaterThanThreshold"
  + dimensions = {
    + "AutoScalingGroupName" = "terraform-20241031095010218800000003"
  }
  + evaluate_low_sample_count_percentiles = (known after apply)
  + evaluation_periods = 2
  + id = "(Known after apply)"
  + metric_name = "CPUUtilization"
  + namespace = "AWS/EC2"
  + period = 300
  + statistic = "Average"
  + tags_all = (known after apply)
  + threshold = 80
  + treat_missing_data = "missing"
}

# aws_cloudwatch_metric_alarm.west_high_cpu.utilization will be destroyed
# (because aws_cloudwatch_metric_alarm.west_high_cpu.utilization is not in configuration)
resource "aws_cloudwatch_metric_alarm" "west_high_cpu.utilization" {
  - actions_enabled = true -> null
  - alarm_actions = [
    - "arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic",
  ] -> null
  - alarm_description = "This alarm triggers if CPU utilization exceeds 80%." -> null
  - alarm_name = "WestHighCPUUtilizationAlarm" -> null
  - arn = "arn:aws:cloudwatch:us-west-2:134575801911:alarm:WestHighCPUUtilizationAlarm" -> null
  - comparison_operator = "GreaterThanThreshold" -> null
  - datapoints_to_alarm = 0 -> null
  - dimensions = {}
```

```
# aws_cloudwatch_metric_alarm.west_scale_down_cpu.utilization will be created
+ resource "aws_cloudwatch_metric_alarm" "west_scale_down_cpu.utilization" {
  + actions_enabled = true
  + alarm_actions = [
    + "arn:aws:autoscaling:us-west-2:134575801911:scalingPolicy:de7f7206-dd73-4077-9914-f442cc4ec6c1:autoScalingGroupName/terraform-202410310950099645:policyName/west-scale-down",
    + "arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic",
  ]
  + alarm_description = "This alarm triggers if CPU utilization drops below 30% to scale down."
  + alarm_name = "WestScaleDownCPUUtilizationAlarm"
  + arn = "(Known after apply)"
  + comparison_operator = "LessThanThreshold"
  + dimensions = {
    + "AutoScalingGroupName" = "terraform-202410310950099645000000003"
  }
  + evaluate_low_sample_count_percentiles = (known after apply)
  + evaluation_periods = 2
  + id = "(Known after apply)"
  + metric_name = "CPUUtilization"
  + namespace = "AWS/EC2"
  + period = 300
  + statistic = "Average"
  + tags_all = (known after apply)
  + threshold = 30
  + treat_missing_data = "missing"
}

# aws_cloudwatch_metric_alarm.west_scale_up_cpu.utilization will be created
+ resource "aws_cloudwatch_metric_alarm" "west_scale_up_cpu.utilization" {
  + actions_enabled = true
  + alarm_actions = [
    + "arn:aws:autoscaling:us-west-2:134575801911:scalingPolicy:919d7700-17aa-4352-8693-caa2c0f62d18:autoScalingGroupName/terraform-202410310950099645:policyName/west-scale-up",
    + "arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic",
  ]
  + alarm_description = "This alarm triggers if CPU utilization exceeds 80% to scale up."
  + alarm_name = "WestScaleUpCPUUtilizationAlarm"
  + arn = "(Known after apply)"
  + comparison_operator = "GreaterThanThreshold"
  + dimensions = {}
```

```
Command Prompt + - x

+ alarm_actions = [
+   "arn:aws:autoscaling:us-west-2:134575801911:scalingPolicy:919d7700-17aa-4352-8693-caa2c0f62d18:autoScalingGroupName/terraform-20241031095009964500000003:policyName/west-scale-up",
+   "arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic",
]
+ alarm_description = "This alarm triggers if CPU utilization exceeds 80% to scale up."
+ alarm_name = "WestScaledUpCPUUtilizationAlarm"
+ arn = "(known after apply)"
+ comparison_operator = "GreaterThanThreshold"
+ dimensions = {
+   "AutoScalingGroupName" = "terraform-20241031095009964500000003"
}
+ evaluate_low_sample_count_percentiles = (known after apply)
+ evaluation_periods = 2
+ id = (known after apply)
+ metric_name = "CPUUtilization"
+ namespace = "AWS/EC2"
+ period = 300
+ statistic = "Average"
+ tags_all = (known after apply)
+ threshold = 80
+ treat_missing_data = "missing"
}

Plan: 4 to add, 0 to change, 2 to destroy.
aws_cloudwatch_metric_alarm.west_high_cpu_utilization: Destroying... [id=WestHighCPUUtilizationAlarm]
aws_cloudwatch_metric_alarm.west_scale_down_cpu_utilization: Creating...
aws_cloudwatch_metric_alarm.west_scale_up_cpu_utilization: Creating...
aws_cloudwatch_metric_alarm.high_cpu_utilization: Destroying... [id=HighCPUUtilizationAlarm]
aws_cloudwatch_metric_alarm.scale_down_cpu_utilization: Creating...
aws_cloudwatch_metric_alarm.scale_up_cpu_utilization: Creating...
aws_cloudwatch_metric_alarm.west_high_cpu_utilization: Destruction complete after 1s
aws_cloudwatch_metric_alarm.west_scale_down_cpu_utilization: Creation complete after 1s [id=WestScaleDownCPUUtilizationAlarm]
aws_cloudwatch_metric_alarm.west_scale_up_cpu_utilization: Creation complete after 1s [id=WestScaleUpCPUUtilizationAlarm]
aws_cloudwatch_metric_alarm.high_cpu_utilization: Destruction complete after 1s
aws_cloudwatch_metric_alarm.scale_down_cpu_utilization: Creation complete after 1s [id=ScaleDownCPUUtilizationAlarm]
aws_cloudwatch_metric_alarm.scale_up_cpu_utilization: Creation complete after 1s [id=ScaleUpCPUUtilizationAlarm]

Apply complete! Resources: 4 added, 0 changed, 2 destroyed.

C:\Users\hp\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

The screenshot shows the CodeGPT IDE interface with a central code editor displaying Terraform configuration. The code defines an Auto Scaling Group (ASG) named 'app_asg' with a desired capacity of 1, min size of 1, and max size of 5. It uses a launch template with the latest version and attaches it to an ALB target group. The ASG uses a rolling strategy with a 120-second instance warmup period. A health check grace period of 300 seconds is specified, and the health check type is set to 'ELB'. A tag is defined for the instances with the key 'Name' and value 'AppASGInstance', and propagation at launch is enabled.

```
resource "aws_launch_template" "app" {
  tag_specifications {
    }

    # Auto Scaling Group (ASG) for us-east-1
    resource "aws_autoscaling_group" "app_asg" {
      desired_capacity     = 1
      min_size             = 1
      max_size             = 5
      vpc_zone_identifier = [aws_subnet.public_subnet_1.id, aws_subnet.public_subnet_2.id]
      target_group_arns   = [aws_lb_target_group.primary.arn] # Attach to ALB target group

      instance_refresh {
        strategy = "rolling"
        preferences {
          min_healthy_percentage = 0
          instance_warmup        = 120
        }
      }

      health_check_grace_period = 300 # Grace period
      health_check_type         = "ELB" # Change to ELB-based health check

      launch_template {
        id      = aws_launch_template.app.id
        version = "$Latest"
      }

      tag {
        key      = "Name"
        value    = "AppASGInstance"
        propagate_at_launch = true
      }
    }
}
```

The screenshot displays two separate AWS CloudWatch Alarms interfaces side-by-side, illustrating the state of alarms across different regions.

Top Region: us-east-1

- Alarms Overview:** Shows 2/5 alarms. One alarm is in an OK state (ScaleUpCPUUtilizationAlarm) and one is in an alarm state (ScaleDownCPUUtilizationAlarm).
- Table:** Details the alarms with columns: Name, State, Last state update (UTC), Conditions, and Actions.

Name	State	Last state update (UTC)	Conditions	Actions
ScaleUpCPUUtilizationAlarm	OK	2024-11-03 10:59:01	CPUUtilization > 80 for 2 datapoints within 10 minutes	Actions enabled
ScaleDownCPUUtilizationAlarm	In alarm	2024-11-03 10:58:26	CPUUtilization < 30 for 2 datapoints within 10 minutes	Actions enabled

Bottom Region: us-west-2

- Alarms Overview:** Shows 2/4 alarms. One alarm is in an alarm state (WestScaleDownCPUUtilizationAlarm) and one is in an OK state (WestScaleUpCPUUtilizationAlarm).
- Table:** Details the alarms with columns: Name, State, Last state update (UTC), Conditions, and Actions.

Name	State	Last state update (UTC)	Conditions	Actions
WestScaleDownCPUUtilizationAlarm	In alarm	2024-11-03 10:58:46	CPUUtilization < 30 for 2 datapoints within 10 minutes	Actions enabled
WestScaleUpCPUUtilizationAlarm	OK	2024-11-03 10:58:36	CPUUtilization > 80 for 2 datapoints within 10 minutes	Actions enabled

```
A newer release of "Amazon Linux" is available.
 Version 2023.6.20241028;
 Version 2023.6.20241031;
Run '/usr/bin/dnf check-release-update' for full release and version update info
  _\_ ##_
  ~~ \###\
  ~~ \|#
  ~~ \#_
  ~~ \|_-->
  ~~ \|_/
  ~~ \|_/
  /my\

Last login: Sat Nov  2 15:21:47 2024 from 18.206.107.27
[ec2-user@ip-10-0-1-160 ~]$ sudo yum install -y httpd-tools
Last metadata expiration check: 19:40:09 ago on Sat Nov  2 15:20:48 2024.
Dependencies resolved.

=====
 Package                               Architecture   Version           Repository      Size
 ==
Installing:
 httpd-tools                           x86_64        2.4.62-1.amzn2023      amazonlinux    81
Installing dependencies:
 apr                                    x86_64        1.7.2-2.amzn2023.0.2      amazonlinux    129
 apr-util                             x86_64        1.6.3-1.amzn2023.0.1      amazonlinux    98
Installing weak dependencies:
 apr-util-openssl                      x86_64        1.6.3-1.amzn2023.0.1      amazonlinux    17
Transaction Summary
 0 CloudShell  Feedback
© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences
```

```
Verifying : apr-util-1.6.3-1.amzn2023.0.1.x86_64          2/
4 Verifying : apr-util-openssl-1.6.3-1.amzn2023.0.1.x86_64  3/
4 Verifying : httpd-tools-2.4.62-1.amzn2023.x86_64         4/
=====
WARNING:
 A newer release of "Amazon Linux" is available.

Available Versions:

Version 2023.6.20241028;
 Run the following command to upgrade to 2023.6.20241028:
 dnf upgrade --releasever=2023.6.20241028

Release notes:
 https://docs.aws.amazon.com/linux/al2023/release-notes/relnotes-2023.6.20241028.html

Version 2023.6.20241031;
 Run the following command to upgrade to 2023.6.20241031:
 dnf upgrade --releasever=2023.6.20241031

Release notes:
 https://docs.aws.amazon.com/linux/al2023/release-notes/relnotes-2023.6.20241031.html

=====
Installed:
 apr-1.7.2-2.amzn2023.0.2.x86_64  apr-util-1.6.3-1.amzn2023.0.1.x86_64  apr-util-openssl-1.6.3-1.amzn2023.0.1.x86_64  httpd-tools-2.4.62-1.amzn2023.x86_64
[ec2-user@ip-10-0-1-160 ~]$ 
 0 CloudShell  Feedback
© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences
```

```
ScaleUpCPUUtilizationAlarm | Alarms | CloudWatch | us-east-1 | Auto Scaling groups | EC2 | us-east-1 | EC2 Instance Connect
< > C  us-east-1.console.aws.amazon.com/ec2-instance-connect/ssh?addressFamily=ipv4&connType=standard&instanceId=i-05...
aws Services Search [Alt+S] N. Virginia vivekvarshith @ vv3281
[ec2-user@ip-10-0-1-160 ~]$ ab -n 500000 -c 500 http://primary-alb-1361659550.us-east-1.elb.amazonaws.com/
This is ApacheBench, Version 2.3 <$Revision: 1913912 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
Benchmarking primary-alb-1361659550.us-east-1.elb.amazonaws.com (be patient)
Completed 500000 requests
Completed 1000000 requests
Completed 1500000 requests
Completed 2000000 requests
Completed 2500000 requests
Completed 3000000 requests
[ec2-user@ip-10-0-1-160 ~]$ 
 0 CloudShell  Feedback
© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences
```

Screenshot of the AWS CloudWatch Alarms console showing the ScaleUpCPUUtilizationAlarm.

CloudWatch Alarms

ScaleUpCPUUtilizationAlarm

Graph

CPUUtilization

Graph showing CPUUtilization over time. A red horizontal line at 53.3 represents the threshold. The graph shows data points from 11:15 to 12:10 UTC on November 3, 2024. A tooltip indicates "CPUUtilization > 53.3 for 2 datapoints within 10 minutes".

Actions

OK

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS Auto Scaling groups console showing the terraform-20241031095010218800000003 group.

Auto Scaling groups

terrafarm-20241031095010218800000003

Auto Scaling group: terraform-20241031095010218800000003

Table showing instance events:

Status	Description	Cause	Start time
Success	Launching a new EC2 instance: i-058b967e362272ea8	At 2024-11-03T12:11:01Z a monitor alarm ScaleUpCPUUtilizationAlarm in state ALARM triggered policy scale-up changing the desired capacity from 1 to 2. At 2024-11-03T12:11:03Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 2.	2024 November 03, 05:11:05 AM -07:00
Success	Terminating EC2 instance: i-09c157a92c351dcdb	At 2024-11-03T10:58:26Z a monitor alarm ScaleDownCPUUtilizationAlarm in state ALARM triggered policy scale-down changing the desired capacity from 2 to 1. At 2024-11-03T10:58:31Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 2 to 1. At 2024-11-03T10:58:31Z instance i-09c157a92c351dcdb was selected for termination.	2024 November 03, 05:58:31 AM -07:00

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Screenshot of the AWS Instances console showing the AppASGInstance.

Instances

i-058b967e362272ea8 (AppASGInstance)

Details

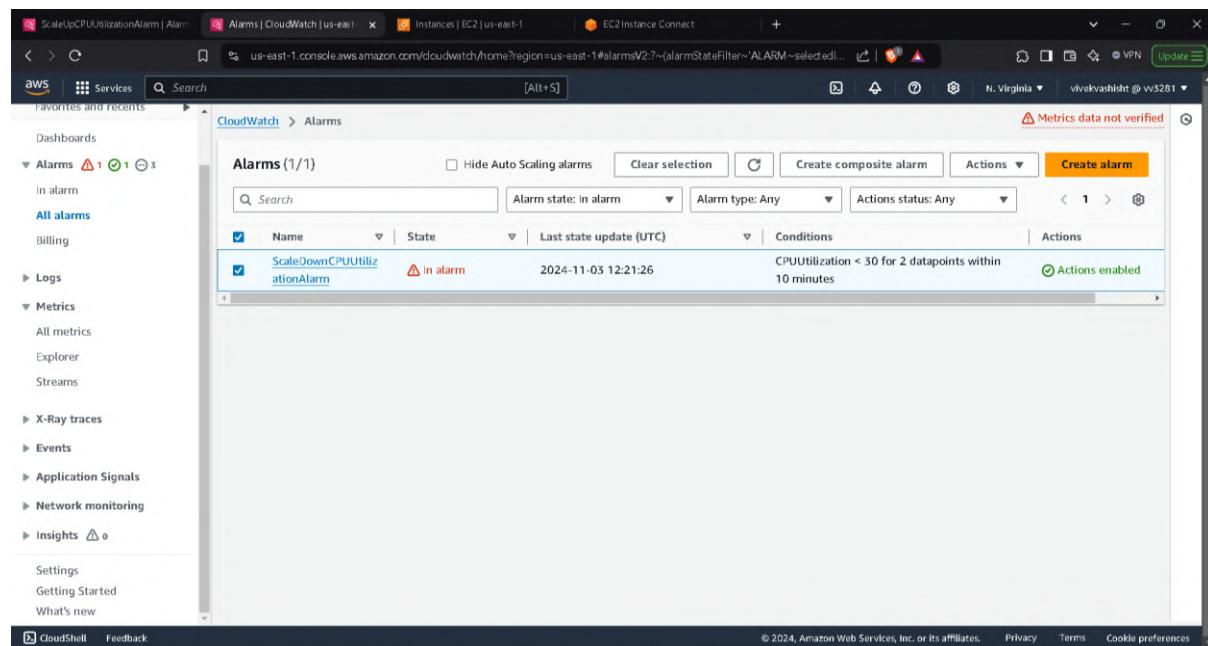
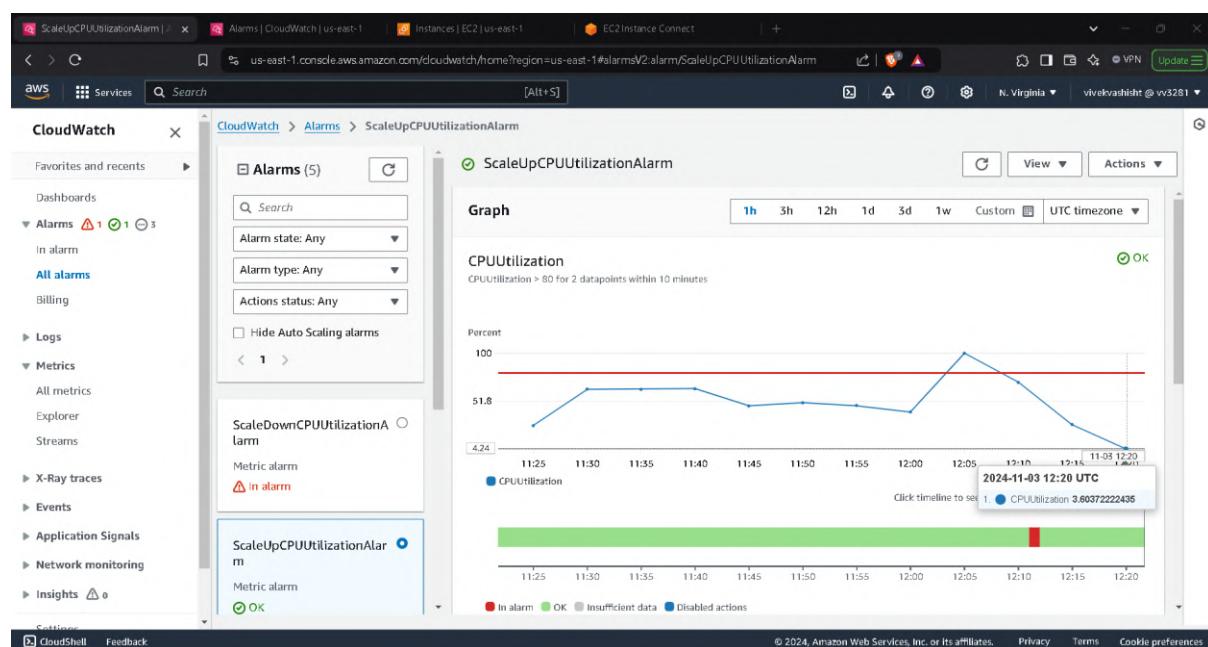
Public IPv4 address: 3.88.193.141 | open address

Private IPv4 addresses: 10.0.4.202

Public IPv4 DNS: ec2-3-88-193-141.compute-1.amazonaws.com

Instance state: Running

© 2024, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences



Screenshot of the AWS CloudWatch Metrics console showing the ScaleUpCPUUtilizationAlarm metric for the Auto Scaling group.

The chart shows the CPU utilization over time, with a sharp peak reaching approximately 100% utilization around November 3, 2024, at 11:01 AM. This triggers the ScaleUpCPUUtilizationAlarm, which then triggers the Auto Scaling group to increase its capacity from 1 to 2 instances.

Auto Scaling group details:

- Name:** terraform-
- Desired capacity:** 2
- Min capacity:** 1
- Max capacity:** 2
- Health check type:** ELB
- Termination policy:** Oldest instances first

Event history:

Status	Description	Cause	Start time
Terminating EC2 instance	Waiting For ELB Connection Draining	At 2024-11-03T12:21:26Z a monitor alarm ScaleDownCPUUtilizationAlarm in state ALARM triggered policy scale-down changing the desired capacity from 2 to 1. At 2024-11-03T12:21:51Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 2 to 1. At 2024-11-03T12:21:51Z instance i-058b967e362272ea8 was selected for termination.	2024 November 03, 05:21:51 AM -0:00
Launching a new EC2 instance		At 2024-11-03T12:11:01Z a monitor alarm ScaleUpCPUUtilizationAlarm in state ALARM triggered policy scale-up changing the desired capacity from 1 to 2. At 2024-11-03T12:11:03Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 2.	2024 November 03, 05:11:05 AM -0:00

Screenshot of the AWS Instances page showing the two EC2 instances created by the Auto Scaling group.

The table lists two instances:

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
AppASGinstance	i-05a5b8a278d2f615c	Running	t2.micro	2/2 checks pass	View alarms	us-east-1a	ec2-3-236
AppASGinstance	i-058b967e362272ea8	Terminated	t2.micro	-	View alarms	us-east-1b	-

Details for terminated instance:

Details	Status and alarms	Monitoring	Security	Networking	Storage	Tags
Instance summary	Info					
Instance ID	i-058b967e362272ea8	Public IPv4 address		Private IPv4 addresses		
IPv6 address	-	Instance state		Public IPv4 DNS		
	-	Terminated		-		

4) CloudWatch Alarms and Notifications Testing

a) Overview

This phase of the project focused on verifying the functionality of the CloudWatch alarms and ensuring that notifications were properly sent via email when specific thresholds were crossed. For this test, I chose to trigger the `ScaleUpCPUUtilizationAlarm` in the us-east-1 (N. Virginia) region, as it provided a quick way to confirm that the configured Simple Notification Service (SNS) would send an alert email when the alarm was triggered.

b) Preparation: Installing Stress Tool for Load Testing

To simulate high CPU usage on the EC2 instance and trigger the alarm, I connected to one of the EC2 instances in us-east-1 and installed the `stress` tool, which is commonly used for generating CPU load.

```
sudo yum install stress -y
```

This tool installation completed successfully, enabling me to proceed with the testing.

c) Triggering the Scale-Up Alarm

After installing the tool, I ran the following command to place CPU load on the instance:

```
stress --cpu 2
```

With this command running, the instance CPU utilization increased significantly. I accessed the `ScaleUpCPUUtilizationAlarm` in the CloudWatch Console in us-east-1 to monitor the alarm's status. The CPU utilization showed a spike to 99.1%, which met the alarm's threshold conditions (greater than 80% for two consecutive data points within 10 minutes).

d) Verifying the Notification

As expected, the `ScaleUpCPUUtilizationAlarm` entered the "ALARM" state due to the high CPU load. I received an email notification from SNS to the configured email address. The notification email included key details, confirming that the alarm was working as configured. Here's an excerpt from the email notification:

Email Notification:

AWS Notifications <no-reply@sns.amazonaws.com>
9:41 AM (4 minutes ago)
You are receiving this email because your Amazon CloudWatch Alarm "ScaleUpCPUUtilizationAlarm" in the US East (N. Virginia) region has entered the ALARM state, because "Threshold Crossed: 2 datapoints [99.20501324977917 (03/11/24 16:36:00), 98.81166107758926 (03/11/24 16:31:00)] were greater than the threshold (80.0)." at "Sunday 03 November, 2024 16:41:01 UTC".

View this alarm in the AWS Management Console:

<https://us-east-1.console.aws.amazon.com/cloudwatch/deeplink.js?region=us-east-1/alarm/ScaleUpCPUUtilizationAlarm>

Alarm Details:

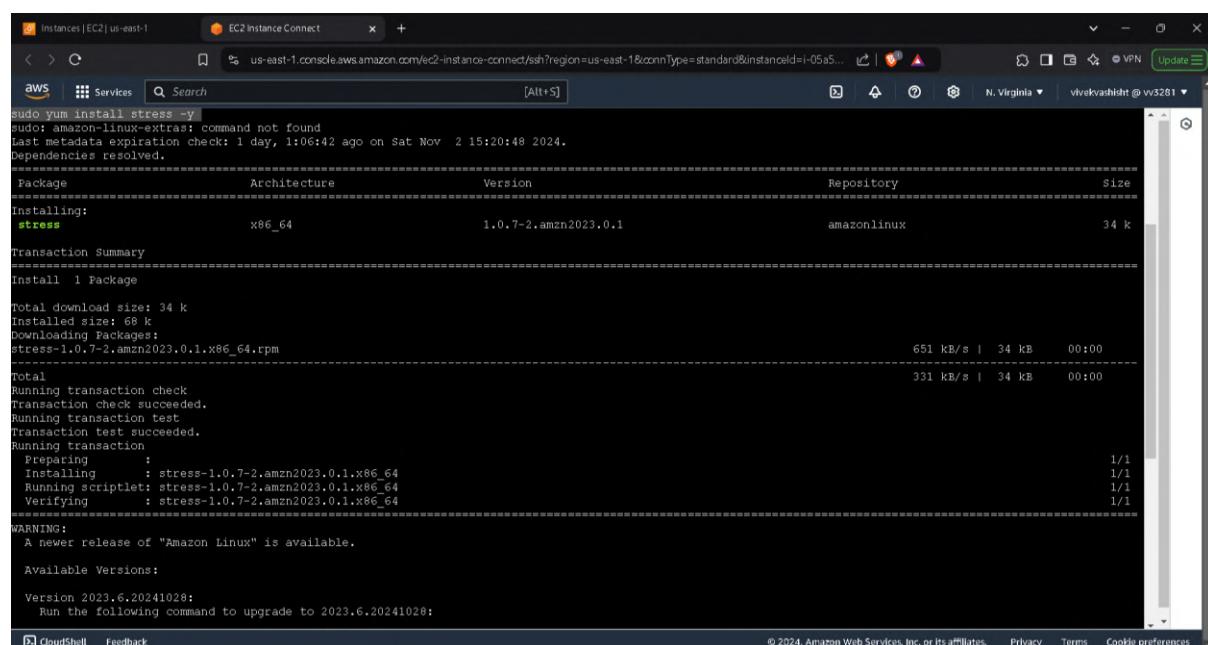
- **Name:** ScaleUpCPUUtilizationAlarm
- **Description:** This alarm triggers if CPU utilization exceeds 80% to scale up.
- **State Change:** OK -> ALARM
- **Timestamp:** Sunday 03 November, 2024 16:41:01 UTC
- **Reason for State Change:** Threshold Crossed: 2 datapoints [99.20501324977917 (03/11/24 16:36:00), 98.81166107758926 (03/11/24 16:31:00)] were greater than the threshold (80.0).
- **AWS Account:** 134575801911
- **Alarm Arn:** arn:aws:cloudwatch:us-east-1:134575801911:alarm:ScaleUpCPUUtilizationAlarm

Threshold:

The alarm is in the ALARM state when the metric is Greater Than Threshold for at least 2 of the last 2 period(s) of 300 seconds.

The testing was successful, confirming that the CloudWatch alarm triggers as expected under high CPU load and sends an email notification through SNS to the designated email. This test validated both the alarm configuration and the email notification setup.

Screenshots



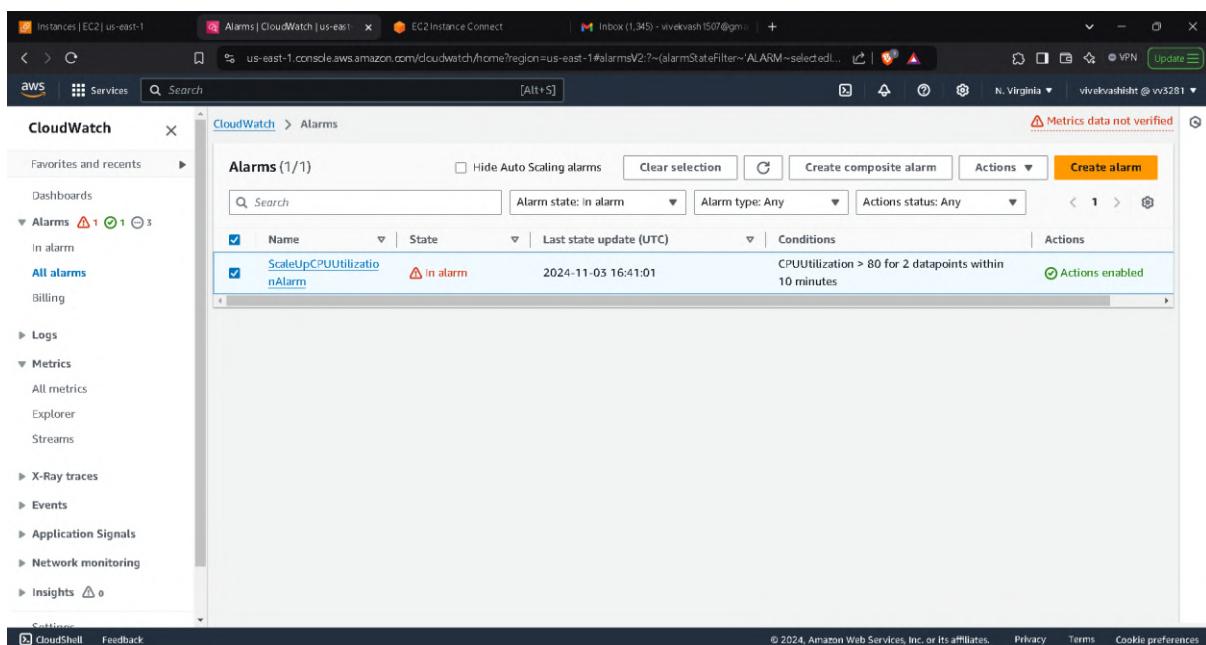
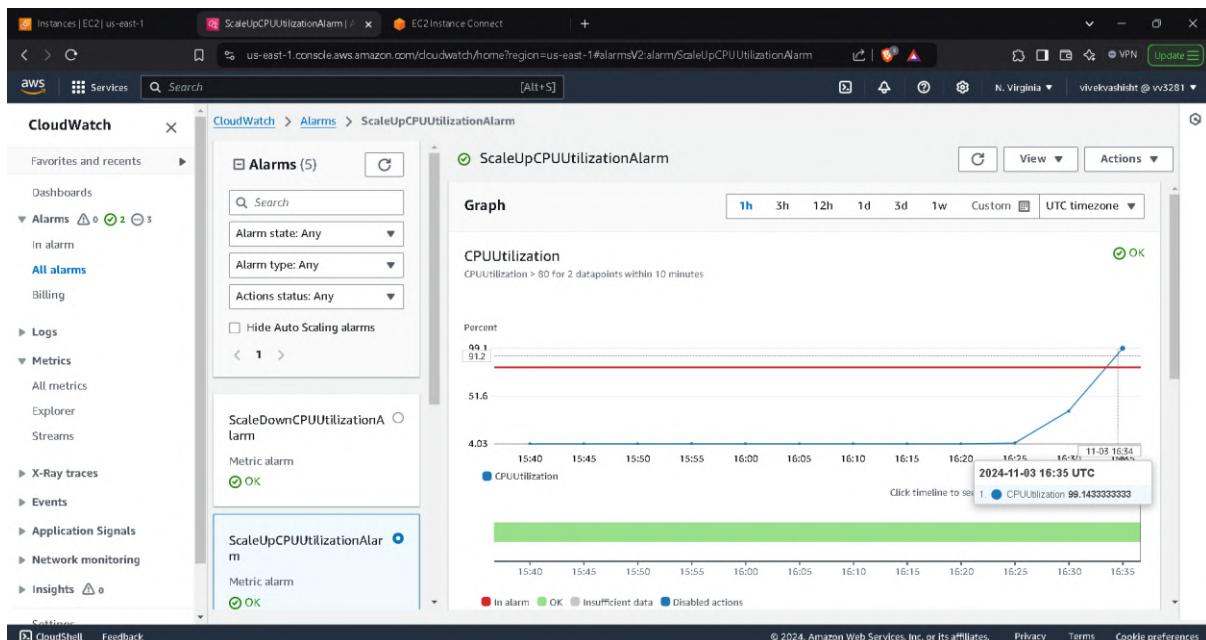
```
sudo yum install stress -y
sudo: amazon-linux-extras: command not found
Last metadata expiration check: 1 day, 1:06:42 ago on Sat Nov  2 15:20:48 2024.
Dependencies resolved.
=====
Package           Architecture Version       Repository      Size
stress            x86_64      1.0.7-2.amzn2023.0.1   amazonlinux    34 k
=====
Installing:
stress            x86_64      1.0.7-2.amzn2023.0.1   amazonlinux    34 k
=====
Transaction Summary
=====
Install 1 Package

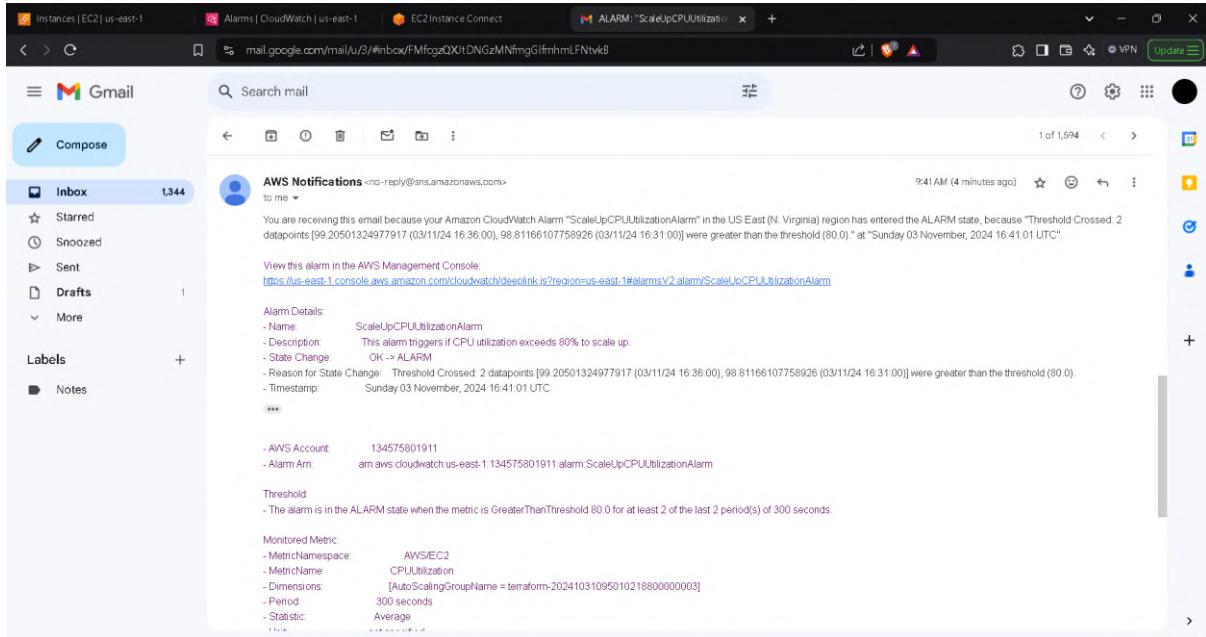
total download size: 34 k
installed size: 68 k
downloading Packages:
stress-1.0.7-2.amzn2023.0.1.x86_64.rpm
                                                               651 kB/s | 34 kB  00:00
                                                               331 kB/s | 34 kB  00:00

Total
Running transaction check
transaction check succeeded.
Running transaction test
transaction test succeeded.
Running transaction
Preparing : 1/1
Installing : stress-1.0.7-2.amzn2023.0.1.x86_64 1/1
Running scriptlet: stress-1.0.7-2.amzn2023.0.1.x86_64 1/1
Verifying  : stress-1.0.7-2.amzn2023.0.1.x86_64 1/1
=====
WARNING:
A newer release of "Amazon Linux" is available.

Available Versions:
Version 2023.6.20241028:
Run the following command to upgrade to 2023.6.20241028:
```

```
[ec2-user@ip-10-0-1-160 ~]$ stress --cpu 2
stress: info: [73711] dispatching hogs: 2 cpu, 0 io, 0 vmy, 0 hdd
```





Cleanup

a) Overview

In this final phase, the objective was to clean up and delete all AWS resources created during the project. Running `terraform destroy` ensures that all infrastructure is decommissioned, preventing unnecessary costs by terminating any active resources across all regions. This included resources like EC2 instances, S3 buckets, Route 53 records, and more.

b) Initial Destroy Plan

To start, I opened the command line interface (CLI) and ran the following command to view a detailed plan of the resources scheduled for destruction:

```
terraform plan -destroy
```

The output displayed a total of **63 resources** that Terraform would destroy. Reviewing this plan allowed me to verify that all resources were identified and included in the cleanup process.

c) Executing the Destroy Command

Once I confirmed the destroy plan, I executed the following command to initiate the cleanup:

```
terraform destroy -auto-approve
```

This command started the process of tearing down the infrastructure. However, an issue arose during the destruction process when Terraform encountered an error while attempting to

delete the S3 bucket named `aws-tf-backend-vv-bucket`. This error occurred because the bucket was not empty; it contained the `terraform.tfstate` file, which needed to be removed before the bucket could be deleted.

d) Resolving the S3 Bucket Issue

To resolve the issue, I navigated to the AWS Management Console, accessed the `aws-tf-backend-vv-bucket` in the S3 service, and manually deleted its contents. Once the bucket was empty, it was ready for deletion.

e) Re-running the Destroy Command

After clearing the S3 bucket, I ran the `terraform destroy` command once more:

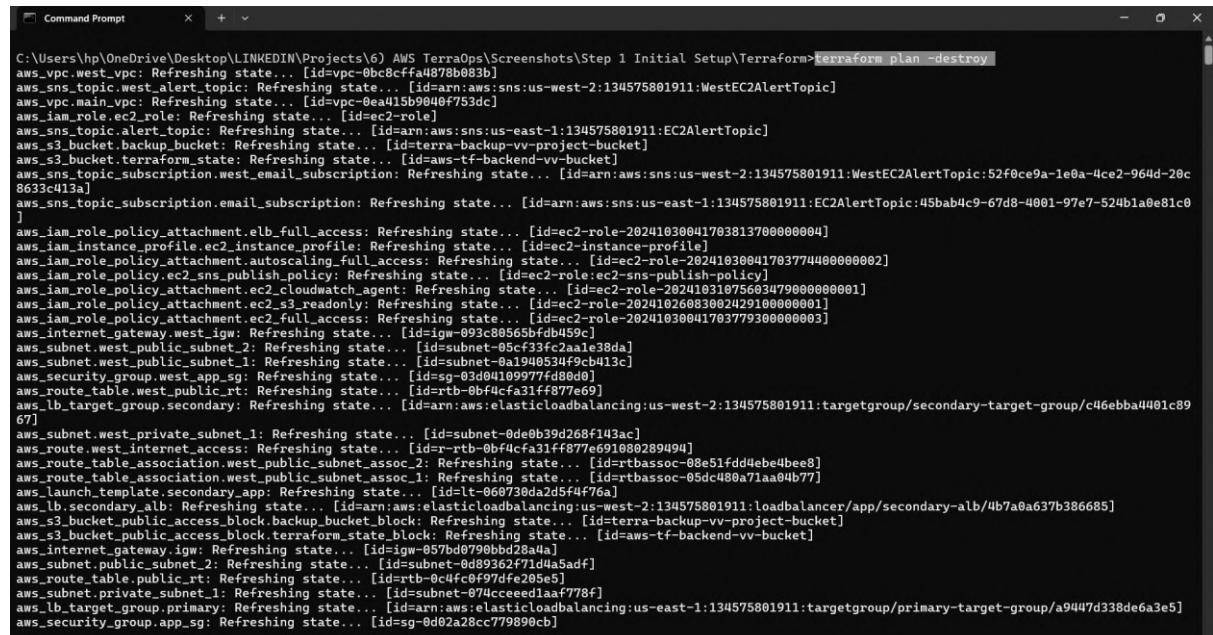
```
terraform destroy -auto-approve
```

This time, the destruction process completed successfully. The command output confirmed that all resources were deleted, displaying the message:

```
No changes. No objects need to be destroyed.
```

With the successful execution of `terraform destroy`, all resources associated with the project were fully decommissioned. This ensured a complete cleanup of the project infrastructure, marking the final and successful completion of the project.

Screenshots



```
C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform plan -destroy
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4078b083b]
aws_sns_topic.west_alert_topic: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic]
aws_vpc.main_vpc: Refreshing state... [id=vpc-0ea15b9040f753dc]
aws_iam_role.ec2_role: Refreshing state... [id=ec2-role]
aws_sns_topic.alert_topic: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic]
aws_s3_bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_sns_topic_subscription.west_email_subscription: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic:52f0ce9a-1e0a-4ce2-964d-28c8633c413a]
aws_sns_topic_subscription.email_subscription: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic:45bab4c9-67d8-4001-97e7-524b1a0e81c0]
aws_iam_role_policy_attachment.elb_full_access: Refreshing state... [id=ec2-role-202410300417038137000000004]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_iam_role_policy_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-202410300417037744000000002]
aws_iam_role_policy.ec2_sns_publish_policy: Refreshing state... [id=ec2-role:ec2-sns-publish-policy]
aws_iam_role_policy_attachment.ec2_cloudwatch_agent: Refreshing state... [id=ec2-role-202410310756034790000000001]
aws_iam_role_policy_attachment.ec2_s3_ready: Refreshing state... [id=ec2-role-202410260830024291000000001]
aws_iam_role_policy_attachment.ec2_full_access: Refreshing state... [id=ec2-role-202410300417037793000000003]
aws_internet_gateway.west_igw: Refreshing state... [id=igw-093c808565bfdb49c]
aws_subnet.west_public_subnet_2: Refreshing state... [id=subnet-05cf33fc2aae38da]
aws_subnet.west_public_subnet_1: Refreshing state... [id=subnet-0a1940534f9cb413c]
aws_security_group.west_app_sg: Refreshing state... [id=sg-03d041099977fd88d0]
aws_route_table.west_public_rt: Refreshing state... [id=rtb-0bf4cfa31f877e69]
aws_lb_target_group.secondary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba4401c8967]
aws_subnet.west_private_subnet_1: Refreshing state... [id=subnet-0de0b39d268f143ac]
aws_route.west_internet_access: Refreshing state... [id=r-rtb-0bf4cfa31f877e691080289494]
aws_route_table_association.west_public_subnet_assoc_2: Refreshing state... [id=rtbassoc-08e51fdd4ebe4bee8]
aws_route_table_association.west_public_subnet_assoc_1: Refreshing state... [id=rtbassoc-05d4c80a71aa04b77]
aws_launch_template.secondary_app: Refreshing state... [id=lt-060738da2d5f4f76a]
aws_load_balancer.app_secondary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:loadbalancer/app/secondary-alb/4b7a0a637b386685]
aws_s3_bucket_public_access_block.backup_bucket_block: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_s3_bucket_public_access_block.terraform_state_block: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_internet_gateway.igw: Refreshing state... [id=igw-057bd0796bbd28a4a]
aws_subnet.public_subnet_2: Refreshing state... [id=subnet-0d89362f71d4a5adf]
aws_route_table.public_rt: Refreshing state... [id=rtb-0c4fc8f97f6e205e5]
aws_subnet.private_subnet_1: Refreshing state... [id=subnet-074ccceeed1aaaf778f]
aws_lb_target_group.primary: Refreshing state... [id=arn:aws:elasticloadbalancing:us-east-1:134575801911:targetgroup/primary-target-group/a9447d338de6a3e5]
aws_security_group.app_sg: Refreshing state... [id=sg-0d02a28cc779890cb]
```

```

Command Prompt x + v
  - "Name" = "MainVPC"
  } -> null
  - tags_all
    - "Name" = "MainVPC"
    } -> null
    # (4 unchanged attributes hidden)
}

# aws_vpc.west_vpc will be destroyed
resource "aws_vpc" "west_vpc" {
  - arn
  - assign_generated_ipv6_cidr_block
  - cidr_block
  - default_network_acl_id
  - default_route_table_id
  - default_security_group_id
  - dhcp_options_id
  - enable_dns_hostnames
  - enable_dns_support
  - enable_network_address_usage_metrics
  - id
  - instance_tenancy
  - ipv6_netmask_length
  - main_route_table_id
  - owner_id
  - tags
    - "Name" = "WestVPC"
  } -> null
  - tags_all
    - "Name" = "WestVPC"
  } -> null
  # (4 unchanged attributes hidden)
}

Plan: 0 to add, 0 to change, 63 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

```

```

Command Prompt x + v
  - enable_network_address_usage_metrics = false -> null
  - id
  - instance_tenancy
  - ipv6_netmask_length
  - main_route_table_id
  - owner_id
  - tags
    - "Name" = "WestVPC"
  } -> null
  - tags_all
    - "Name" = "WestVPC"
  } -> null
  # (4 unchanged attributes hidden)
}

Plan: 0 to add, 0 to change, 63 to destroy.

Note: You didn't use the -out option to save this plan, so Terraform can't guarantee to take exactly these actions if you run "terraform apply" now.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -auto-approve
aws_vpc.west_vpc: Refreshing state... [id=vpc-0bc8cffa4878b083b]
aws_sns_topic.west_alert_topic: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic]
aws_sns_main_vpc: Refreshing state... [id=vpc-0bea415b9040f753dc]
aws_sns_topic.alert_topic: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic]
aws_iam_role.ec2_role: Refreshing state... [id=ec2-role]
aws_s3_bucket.terraform_state: Refreshing state... [id=aws-tf-backend-vv-bucket]
aws_s3_bucket.backup_bucket: Refreshing state... [id=terra-backup-vv-project-bucket]
aws_sns_topic_subscription.west_email_subscription: Refreshing state... [id=arn:aws:sns:us-west-2:134575801911:WestEC2AlertTopic:52f0ce9a-1e0a-4ce2-964d-20c863c413a]
aws_sns_topic_subscription.email_subscription: Refreshing state... [id=arn:aws:sns:us-east-1:134575801911:EC2AlertTopic:45bab4c9-67d8-4001-97e7-524b1a0e81c0]
aws_iam_role_attachment.autoscaling_full_access: Refreshing state... [id=ec2-role-20241030041703774400000002]
aws_iam_role_attachment.ec2_full_access: Refreshing state... [id=ec2-role-20241030041703779300000003]
aws_iam_role_attachment.ec2_cloudwatch_agent: Refreshing state... [id=ec2-role-20241031075603479000000001]
aws_iam_role_attachment.elb_full_access: Refreshing state... [id=ec2-role-202410300417038127000000004]
aws_iam_role_policy.ec2_sns_publish_policy: Refreshing state... [id=ec2-role:ec2-sns-publish-policy]
aws_iam_instance_profile.ec2_instance_profile: Refreshing state... [id=ec2-instance-profile]
aws_iam_role_attachment.ec2_s3_readonly: Refreshing state... [id=sec2-role-20241026083002429100000001]
aws_internet_gateway.west_igw: Refreshing state... [id=igw-093c80565bfdb459c]


```

```

Command Prompt x + v
aws_lb_target_group.primary: Destruction complete after 0s
aws_launch_template.app: Destruction complete after 0s
aws_security_group.app_sg: Destroying... [id=sg-0d02a28cc779890cb]
aws_subnet_public_subnet_1: Destruction complete after 0s
aws_subnet_public_subnet_2: Destruction complete after 0s
aws_security_group.app_sg: Destruction complete after 1s
aws_route53_zone.main: Still destroying... [id=2049649961BUHZ3TYTHQG, 1m0s elapsed]
aws_autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241031095009964500000003, 1m30s elapsed]
aws_route53_zone.main: Still destroying... [id=2049649961BUHZ3TYTHQG, 1m10s elapsed]
aws_autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241031095009964500000003, 1m40s elapsed]
aws_route53_zone.main: Still destroying... [id=2049649961BUHZ3TYTHQG, 1m20s elapsed]
aws_autoscaling_group.secondary_asg: Still destroying... [id=terraform-20241031095009964500000003, 1m50s elapsed]
aws_autoscaling_group.secondary_asg: Destruction complete after 1m52s
aws_subnet.west_public_subnet_2: Destroying... [id=subnet-05cf3fc2aa1e38da]
aws_subnet.west_public_subnet_1: Destroying... [id=subnet-0a1940534f9cb413c]
aws_lb_target_group.secondary: Destroying... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba4401c8967]
aws_launch_template.secondary_app: Destroying... [id=lt-060738da2d5f4f76a]
aws_launch_template.secondary_app: Destruction complete after 0s
aws_iam_instance_profile.ec2_instance_profile: Destroying... [id=ec2-instance-profile]
aws_security_group.west_app_sg: Destroying... [id=sg-03d00109977fd80d0]
aws_lb_target_group.secondary: Destruction complete after 0s
aws_subnet.west_public_subnet_2: Destruction complete after 0s
aws_subnet.west_public_subnet_1: Destruction complete after 0s
aws_security_group.west_app_sg: Destruction complete after 1s
aws_iam_instance_profile.ec2_instance_profile: Destruction complete after 1s
aws_iam_role.ec2_role: Destroying... [id=ec2-role]
aws_iam_role.ec2_role: Destruction complete after 0s
aws_route53_zone.main: Destruction complete after 1m25s
aws_vpc.west_vpc: Destroying... [id=vpc-0bc8cffa4878b083b]
aws_vpc.main_vpc: Destroying... [id=vpc-0bea415b9040f753dc]
aws_vpc.west_vpc: Destruction complete after 1s
aws_vpc.main_vpc: Destruction complete after 1s

Error: deleting S3 Bucket (aws-tf-backend-vv-bucket): operation error S3: DeleteBucket, https response error StatusCode: 409, RequestID: 889322GX633Q0324, HostID: bxuY06QLK6+fMqUPzox8GXCO2WsyLde6WetPuff+V1wf0F2oItI15QsujS4E/FggHvySPGquN+AzeRe9NXKAqPlc, api error BucketNotEmpty: The bucket you tried to delete is not empty

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6 AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>

```

Amazon S3

Buckets, Access Grants, Access Points, Object Lambda Access Points, Multi-Region Access Points, Batch Operations, IAM Access Analyzer for S3, Block Public Access settings for this account, Storage Lens, Dashboards, Storage Lens groups, AWS Organizations settings, Feature spotlight, AWS Marketplace for S3.

General purpose buckets (1) All AWS Regions

aws-tf-backend-vv-bucket

Name	AWS Region	IAM Access Analyzer	Creation date
aws-tf-backend-vv-bucket	US East (N. Virginia) us-east-1	View analyzer for us-east-1	October 23, 2024, 03:10:54 (UTC-06:00)

Amazon S3 > Buckets > aws-tf-backend-vv-bucket

aws-tf-backend-vv-bucket Info

Objects (1) All AWS Regions

Name	Type	Last modified	Size	Storage class
terraform/	Folder	-	-	-

```
aws_route53_zone_main: Still destroying... [id=204964961BUUH23TYTHQG, 1m10s elapsed]
aws_autoscaling_group_secondary_asg: Still destroying... [id=terraform-20241031095009964500000003, 1m40s elapsed]
aws_route53_zone_main: Still destroying... [id=204964961BUUH23TYTHQG, 1m20s elapsed]
aws_autoscaling_group_secondary_asg: Still destroying... [id=terraform-20241031095009964500000003, 1m50s elapsed]
aws_autoscaling_group_secondary_asg: Destruction complete after 1m52s
aws_subnet.west_public_subnet_2: Destroying... [id=subnet-05cf33fc2aae38da]
aws_subnet.west_public_subnet_1: Destroying... [id=subnet-0a1940534fc9b413c]
aws_lb_target_group_secondary: Destroying... [id=arn:aws:elasticloadbalancing:us-west-2:134575801911:targetgroup/secondary-target-group/c46ebba4401c8967]
aws_launch_template_secondary_app: Destroying... [id=lt-066730da2d5f4f76a]
aws_launch_template_secondary_app: Destruction complete after 0s
aws_iam_instance_profile_ec2_instance_profile: Destroying... [id=sec2-instance-profile]
aws_security_group.west_app_sg: Destroying... [id=sg-03d04109977fd80d0]
aws_lb_target_group_secondary: Destruction complete after 0s
aws_subnet.west_public_subnet_2: Destruction complete after 0s
aws_subnet.west_public_subnet_1: Destruction complete after 0s
aws_security_group.west_app_sg: Destruction complete after 1s
aws_iam_instance_profile_ec2_instance_profile: Destruction complete after 1s
aws_iam_role.ec2_role: Destroying... [id=ec2-role]
aws_iam_role.ec2_role: Destruction complete after 0s
aws_route53_zone_main: Destruction complete after 1m25s
aws_vpc.west_vpc: Destroying... [id=vpc-0bc8cffa4878b083b]
aws_vpc.main_vpc: Destroying... [id=vpc-0ea415b9040f753dc]
aws_vpc.west_vpc: Destruction complete after 1s
aws_vpc.main_vpc: Destruction complete after 1s

Error: deleting S3 Bucket (aws-tf-backend-vv-bucket): operation error S3: DeleteBucket, https response error StatusCode: 409, RequestID: 809322GX633QD324, HostID: bxuY0GQLK6+fMqUPzoWox8GC02WsylDe6wJutPuff+V1wfo0F5f2oItIISQsujS4E/FggHv5PGQuN+AzeRe9NXKAqPlc, api error BucketNotEmpty: The bucket you tried to delete is not empty

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>terraform destroy -auto-approve
No changes. No objects need to be destroyed.
Either you have not created any objects yet or the existing objects were already deleted outside of Terraform.
Destroy complete! Resources: 0 destroyed.

C:\Users\hp\OneDrive\Desktop\LINKEDIN\Projects\6) AWS TerraOps\Screenshots\Step 1 Initial Setup\Terraform>
```

Conclusion

The AWS TerraOps project demonstrated the deployment of a highly available, scalable, and automated cloud infrastructure using AWS services and Terraform. This initiative was designed to validate practical skills for AWS Cloud Practitioner and HashiCorp Terraform Associate certifications, ensuring a robust understanding of infrastructure automation and cloud scaling. Below is a concise summary of each major step:

Summary of Steps:

1. **Initial Setup:** Installed Terraform and AWS CLI, configured AWS credentials, and set up version control with Git. Established an S3 bucket for storing Terraform state and configured remote backend storage.
2. **VPC and Networking Configuration:** Created a Virtual Private Cloud (VPC) with public and private subnets, set up an Internet Gateway, and associated public route tables to enable communication with the internet.
3. **EC2 and Auto Scaling Setup:** Defined a security group for EC2 instances, set up a launch template with a user data script to install Nginx, and configured an Auto Scaling Group (ASG) with policies for scaling instances based on demand.
4. **IAM and S3 Integration:** Configured IAM roles and policies to allow EC2 instances access to S3, then set up an additional S3 bucket for storing backups.
5. **Multi-Region Deployment and Failover:** Established resources in two AWS regions, created Application Load Balancers (ALBs) in each region, and implemented Route 53 for DNS-based failover. This configuration ensures high availability and resiliency by redirecting traffic to the secondary region in case of primary region failures.
6. **Monitoring and Alerts:** Installed the CloudWatch agent on EC2 instances to monitor CPU, memory, and disk utilization. Configured CloudWatch Alarms and set up SNS notifications to alert based on resource usage.
7. **Parameterization and Validation:** Introduced variables for reusability and validation checks to enforce correct inputs, enhancing flexibility and scalability.
8. **Testing and Validation:**
 - o **ALB Testing:** Verified load balancing across multiple instances within each region.
 - o **Route 53 Failover Testing:** Simulated a failure in the primary region to confirm automatic failover to the secondary region.
 - o **ASG Testing:** Validated scaling policies by monitoring CPU utilization, triggering scale-up and scale-down actions based on load.
 - o **CloudWatch Alarms and Notifications Testing:** Verified alarm configurations by simulating metric thresholds, ensuring alerts are triggered accurately and notifications are sent to designated channels

This project successfully validated the core competencies of deploying and managing AWS infrastructure with Terraform. The automated deployment process and multi-region failover setup illustrated a strong foundation in cloud infrastructure automation, enhancing availability, scalability, and resiliency. This documentation and practical experience add credibility to the theoretical knowledge acquired through AWS and Terraform certifications, providing a complete end-to-end infrastructure solution ideal for real-world applications.

Lessons Learned

- **Significance of Proactive Monitoring and Alerts:** Understood the critical role of proactive monitoring and alert configurations for AWS resources, ensuring any abnormal events are quickly identified and addressed to maintain system reliability and prevent downtime.
- **AWS Security Best Practices:** Gained a deeper understanding of AWS security best practices, including the effective use of IAM policies, Security Groups, encryption, and Multi-Factor Authentication (MFA) to secure cloud infrastructure.
- **Scalability and Fault Tolerance:** Learned how to implement Auto Scaling and failover mechanisms to ensure high availability and resilience, enabling resources to handle varying loads and maintain performance during peak demand.
- **Automation with Terraform:** Acquired hands-on experience in using Terraform for infrastructure automation, streamlining the process of deploying and managing cloud resources while maintaining an organized and efficient environment.
- **Enhanced Understanding of Infrastructure Automation:** Gained in-depth knowledge of how to use Terraform to create, modify, and version control infrastructure efficiently, promoting repeatable and consistent deployments across environments.
- **Efficient Resource Scaling and Load Management:** Realized the importance of automated scaling policies to handle traffic fluctuations without manual intervention, ensuring cost-effective and efficient resource utilization.
- **Cross-Region Failover and Disaster Preparedness:** Learned to plan for and implement cross-region failover strategies with Route 53, enhancing fault tolerance and ensuring high availability even in case of regional disruptions.

Skills Demonstrated

Through the AWS TerraOps Project, the following skills were showcased:

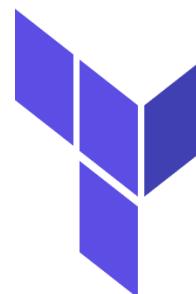
1. **Infrastructure as Code (IaC) Implementation:**
 - Designed and deployed AWS infrastructure using Terraform.
 - Created modular and reusable Terraform code to define cloud resources effectively.
2. **Scalable and Resilient Architecture Design:**
 - Configured Auto Scaling Groups to manage dynamic workloads.
 - Implemented Load Balancing with Application Load Balancer (ALB) for optimized traffic distribution.
3. **Monitoring and Alert Configuration:**
 - Set up monitoring using CloudWatch, capturing metrics for resource health and performance.
 - Configured custom CloudWatch alarms to trigger notifications based on specific metrics and conditions.

4. **Security Hardening:**
 - Applied IAM roles, policies, and security group configurations for secure access control.
 - Enabled encryption at rest and in transit, aligning with best practices for data protection.
5. **Automated Infrastructure Management:**
 - Developed infrastructure automation scripts with Terraform, improving deployment speed and accuracy.
 - Leveraged CloudFormation as needed for additional automation, enhancing deployment flexibility.
6. **Network Architecture and Troubleshooting:**
 - Designed VPC architecture with subnets and route tables to isolate and manage resources effectively.
 - Utilized VPC Flow Logs for network analysis and troubleshooting.
7. **Backup and Recovery Setup:**
 - Configured AWS Backup for critical resources, ensuring a secure recovery strategy.
 - Tested backup and restoration processes to validate disaster recovery readiness.
8. **Cost Management and Resource Optimization:**
 - Practiced cost management by cleaning up resources after testing to prevent unnecessary charges.
 - Configured AWS Cost Explorer and Budgets for ongoing cost tracking and alerting.
9. **Performance and Availability Testing:**
 - Conducted failover testing by simulating region outages, ensuring Route 53 successfully redirected traffic.
 - Monitored EC2 instance health and load balancer performance, identifying and resolving bottlenecks.

And Strong Hands-on Capabilities in

(AWS Infrastructure Management and Automation with Terraform)

Scope of AWS Certified Cloud Practitioner (CLF-C02) and HashiCorp Certified: Terraform Associate (HCTA0-003)



HashiCorp
Terraform